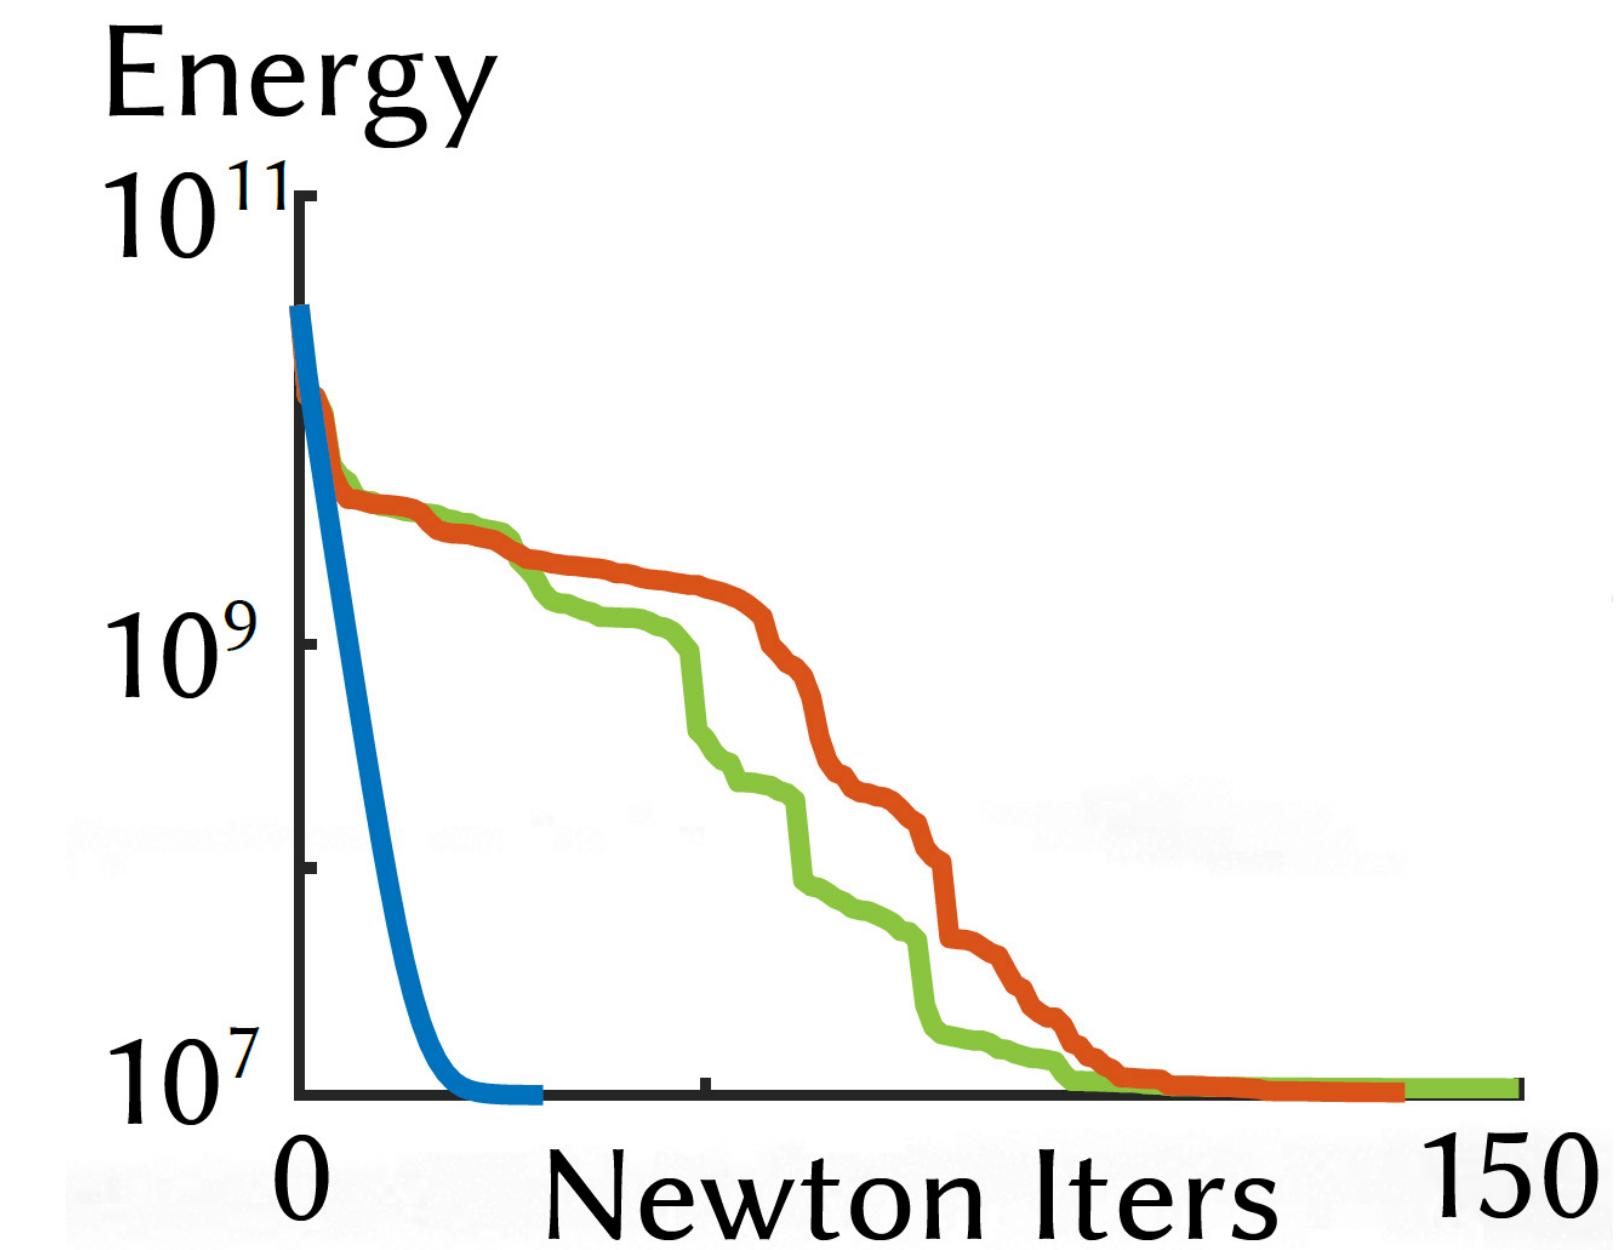


# Introduction to Optimization for Simulation

Honglin Chen, Columbia University



COLUMBIA  
UNIVERSITY



# Quick Intro

- Hi everyone! I'm Honglin Chen.
- Rising 4th year PhD student @ Columbia
  - Physics-based Simulation
  - Geometry Processing
  - Optimization and (a little bit of) ML
- My 2<sup>nd</sup> SCA
  - But my 1<sup>st</sup> time presenting here!



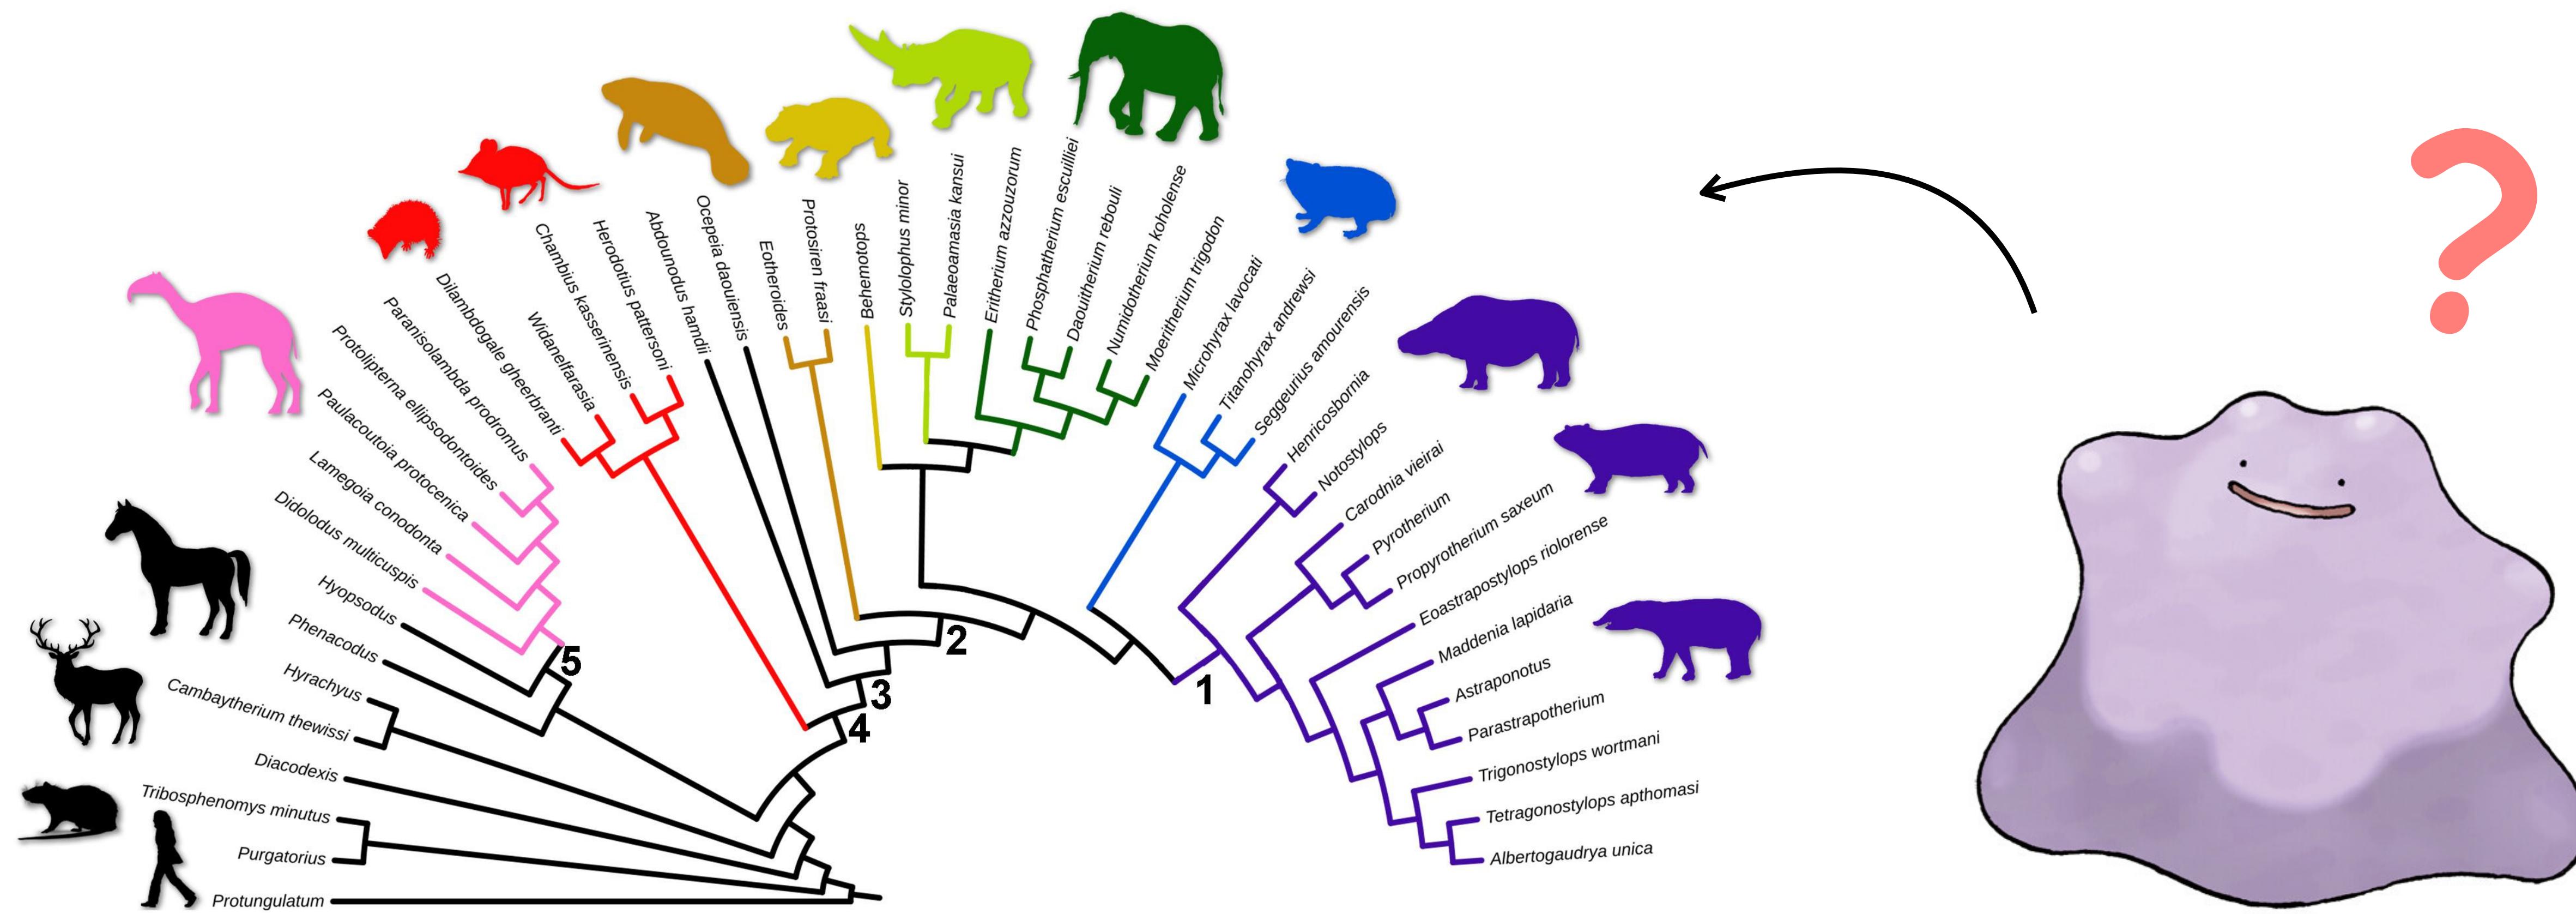
# What's optimization (and why should we study it) ?

**Finding the best feasible solution  
from all the possibilities**  
(potentially subject to some constraints)



# What's an optimization problem?

## Ingredient 1: optimization variables



# What's an optimization problem?

## Ingredient 2: objective function



Optimized for running



Optimized for flying

# What's an optimization problem?

## Ingredient 2: objective function



Optimized for swimming



Optimized for cuteness

# What's an optimization problem?

## Ingredient 2: objective function



What is this optimized for ?? 🤔

# What's an optimization problem?

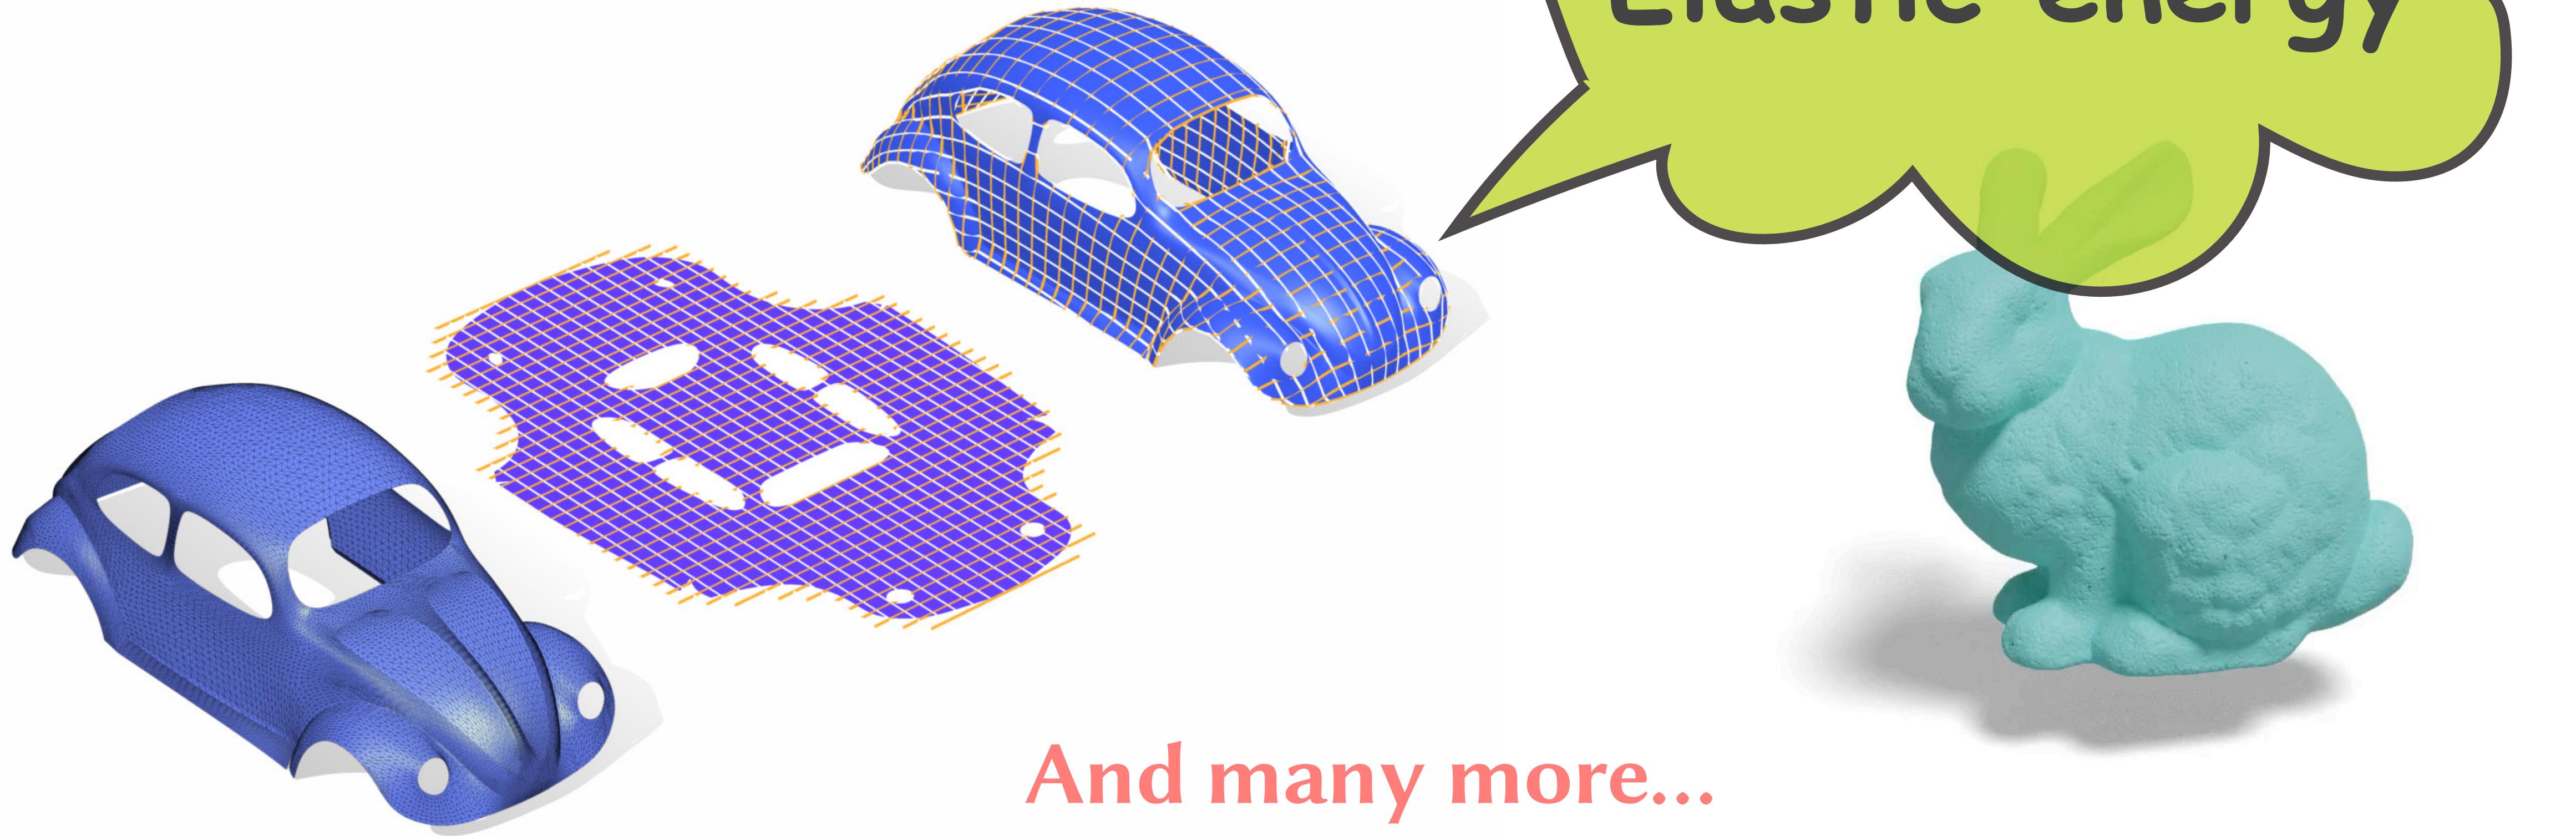
3 components:

- optimization variables that parameterize the space
- an objective function that measures how “good” an arbitrary point in parameter space is
- possibly some constraints



# Why is optimization so important in graphics?

Optimization in Simulation and Geometry



Mesh Parameterization

Elastodynamic Simulation

# Who is this course for?

- New to optimization
- New to elastic simulation
- Want to learn about different techniques for elastic energy minimization



# What will I learn from this course?

- How to minimize an elastic energy (defined on a mesh)?
  - Classical techniques
  - More advanced optimization techniques
- How to make your optimization run better?
  - Common challenges
  - Strategies to tackle these challenges
  - Know where to look if you want to learn more
- Tips for easy implementations
  - Tools and frameworks



# Non-Goals of this Course

This course will not

- Cover how to formulate these energies
  - Eris will talk about that in our next course soon!
- Teach you how to minimize an energy on a neural network
- Dive into details on designing a customized solver for a specific problem



# Agenda

- Elastic Energy
- Optimization Basics & Classical Algorithms
- Optimization Methods for Simulation
  - ✨Challenge 1✨: Handling Nonconvexity
  - ✨Challenge 2✨: Nontrivial Constraints
  - ✨Challenge 3✨: Large-scale Optimization



- Elastic Energy
- Optimization Basics & Classical Algorithms
- Optimization Methods for Simulation
  - Challenge 1 ✨: Handling Nonconvexity
  - Challenge 2 ✨: Nontrivial Constraints
  - Challenge 3 ✨: Large-scale Optimization

# Elastic Energy

---

# Elastic Energy is Everywhere



## Cloth

[Zhang et al. 2022]



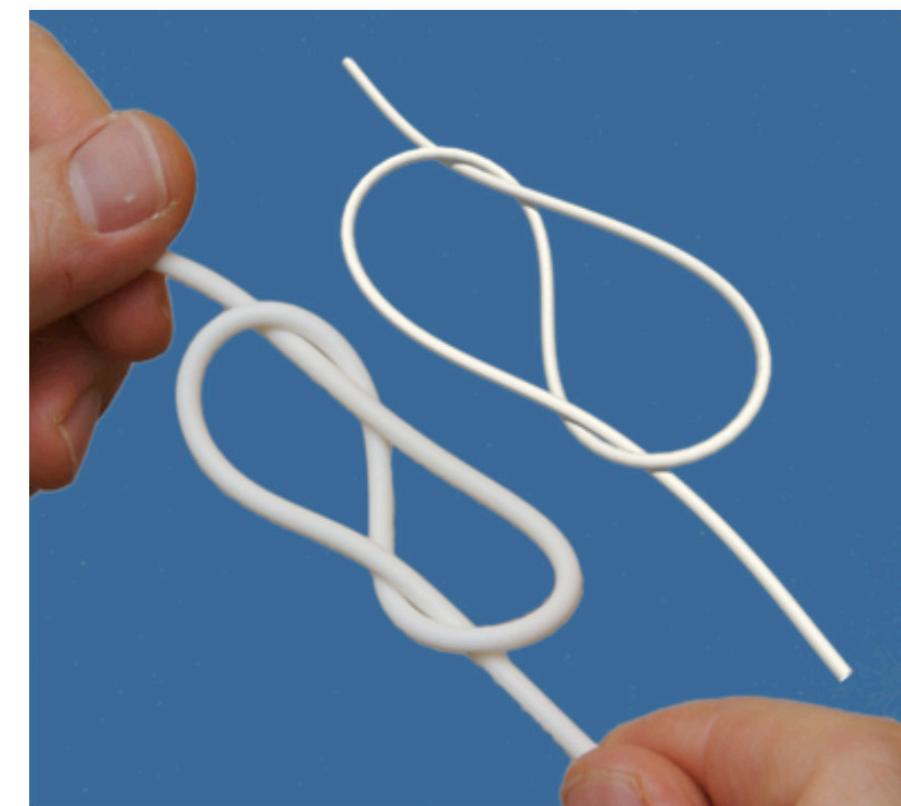
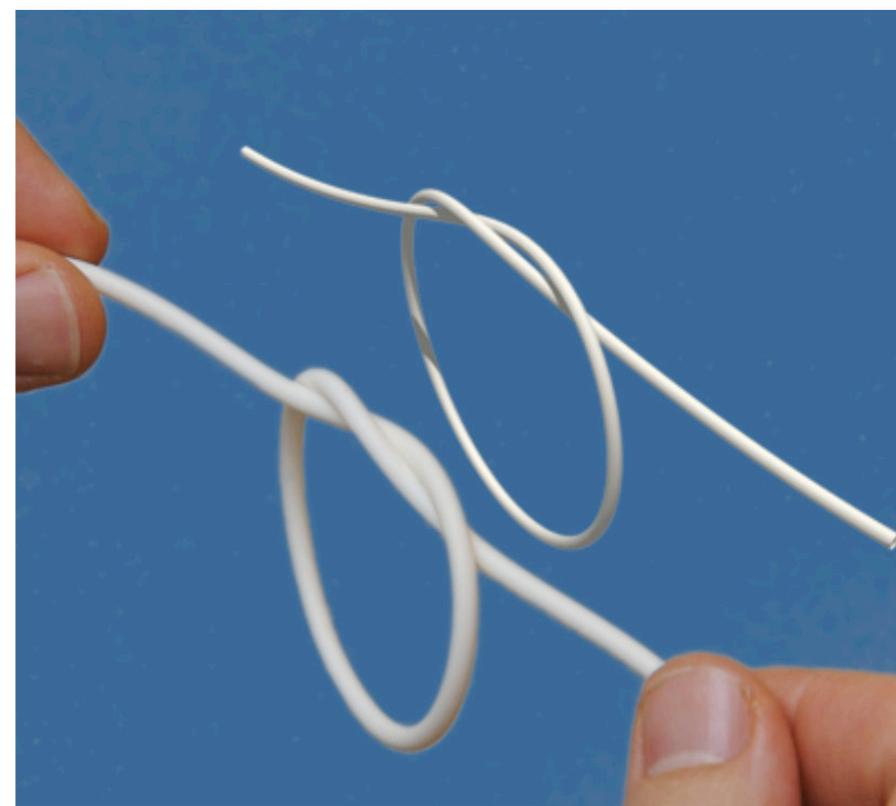
## 3D Deformable Solids

[Smith et al. 2018]



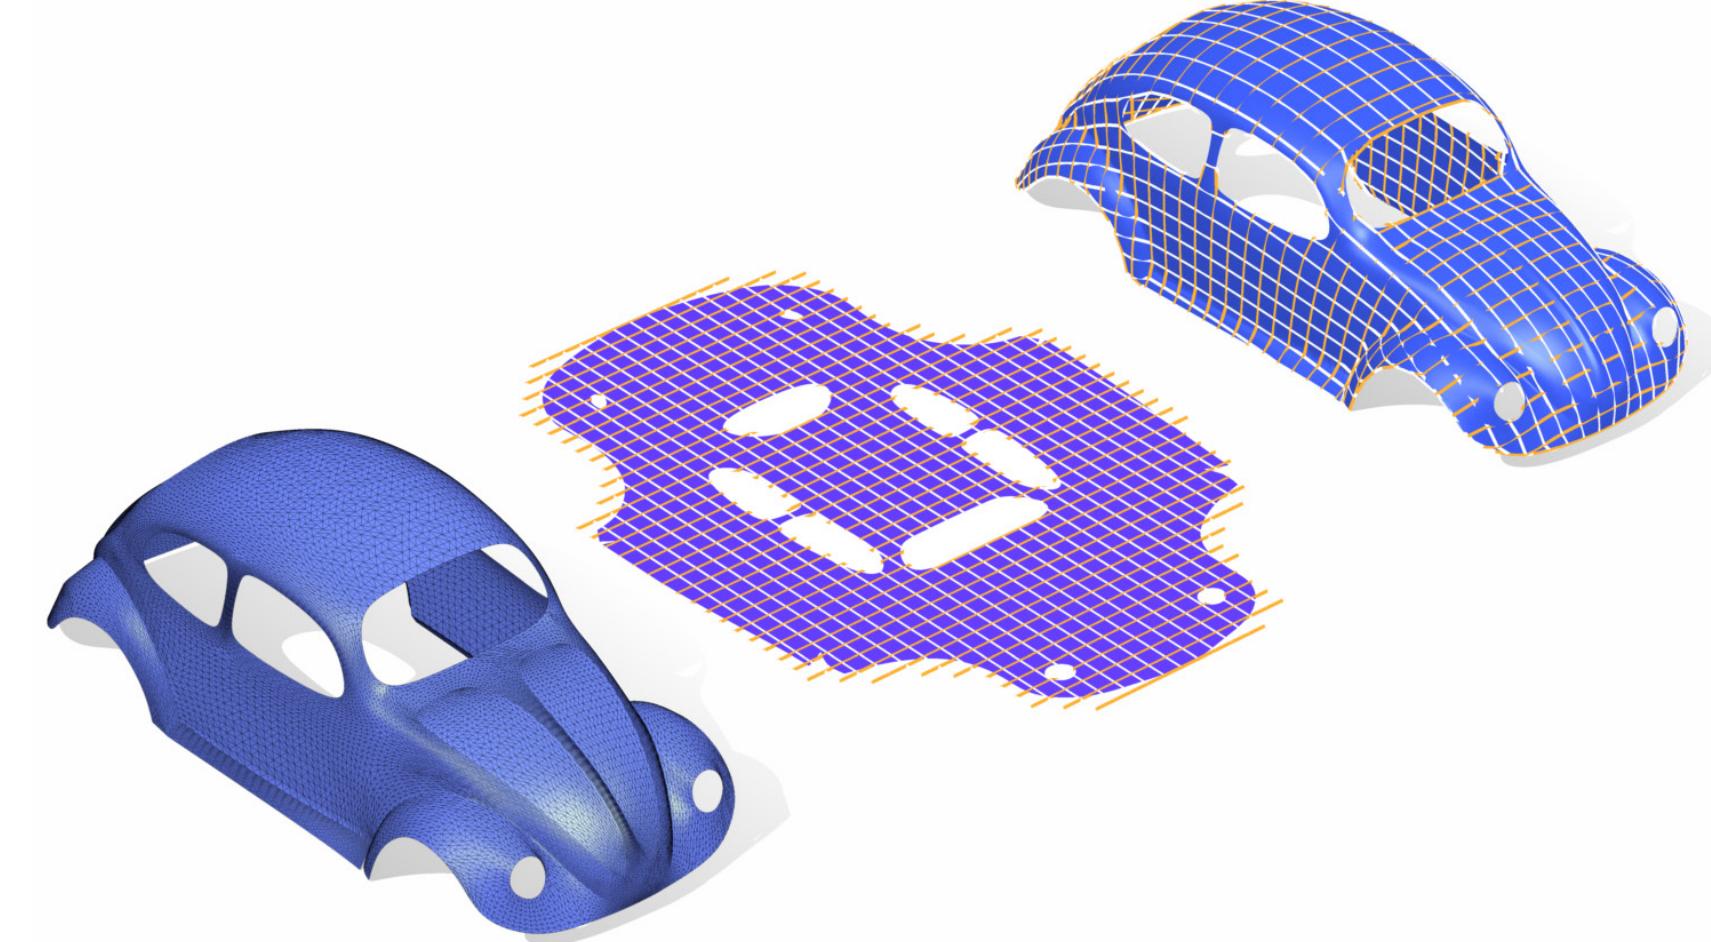
## Thin Shells

[Grinspun et al. 2003]



## Rods

[Bergou et al. 2008]

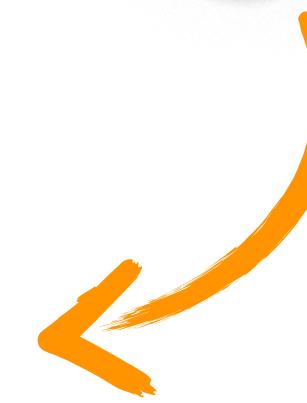
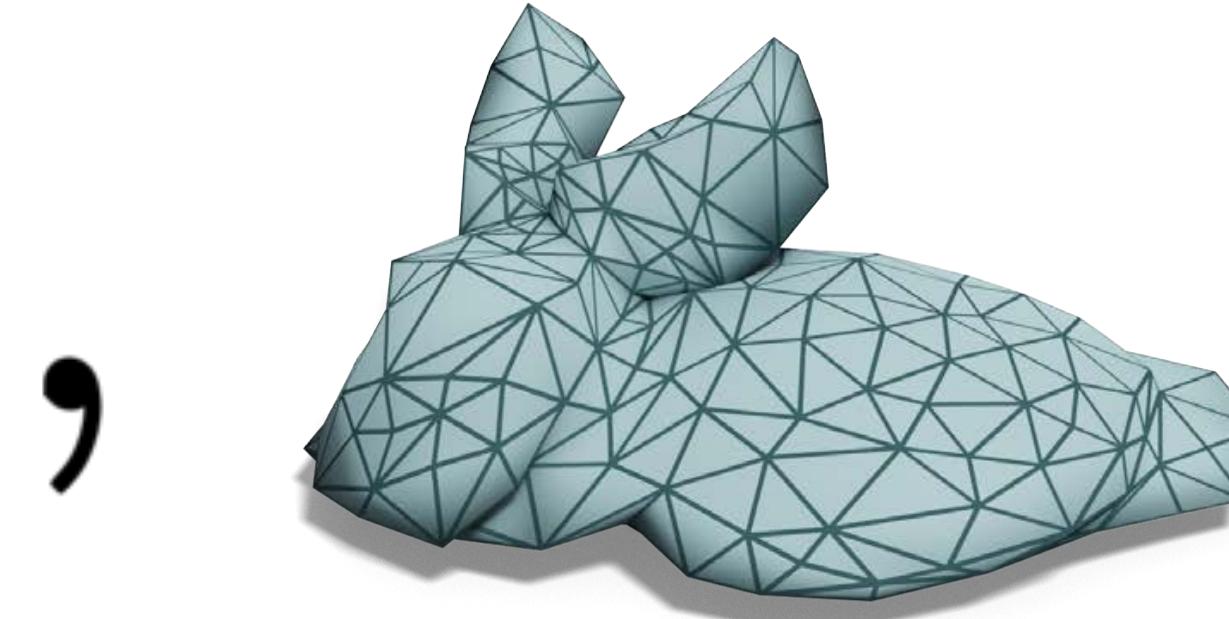
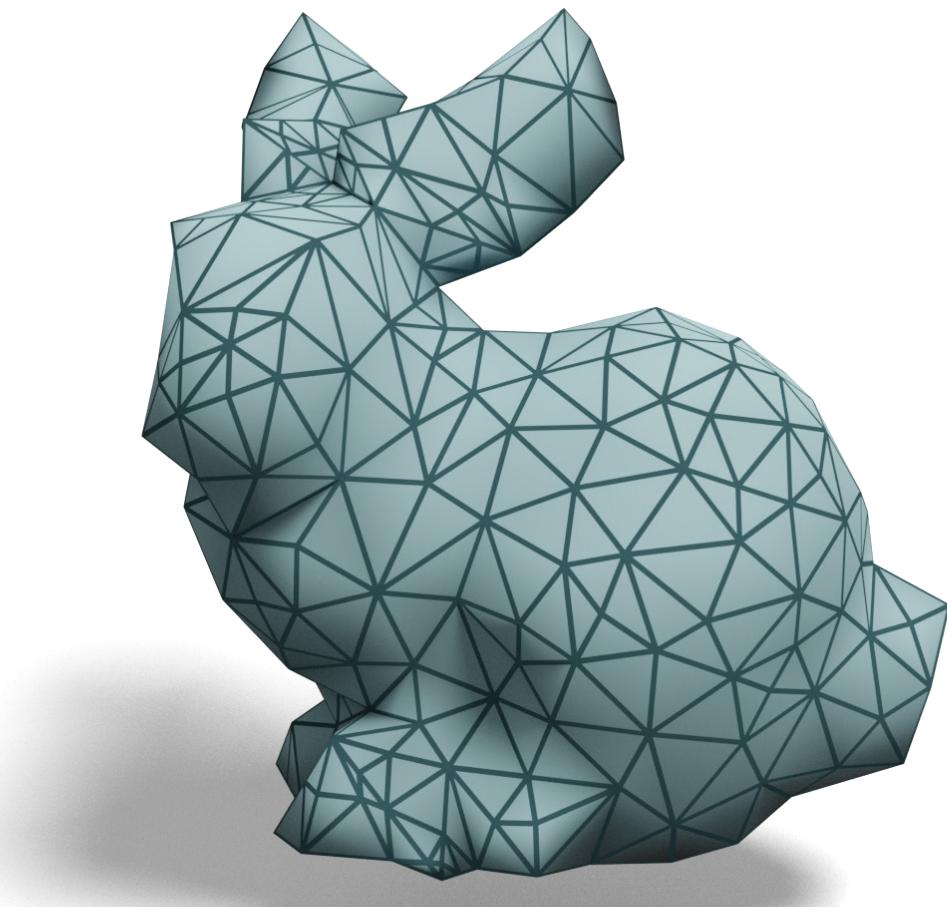


## Surface Parameterization

[Smith et al. 2018]

# Elastic Energy

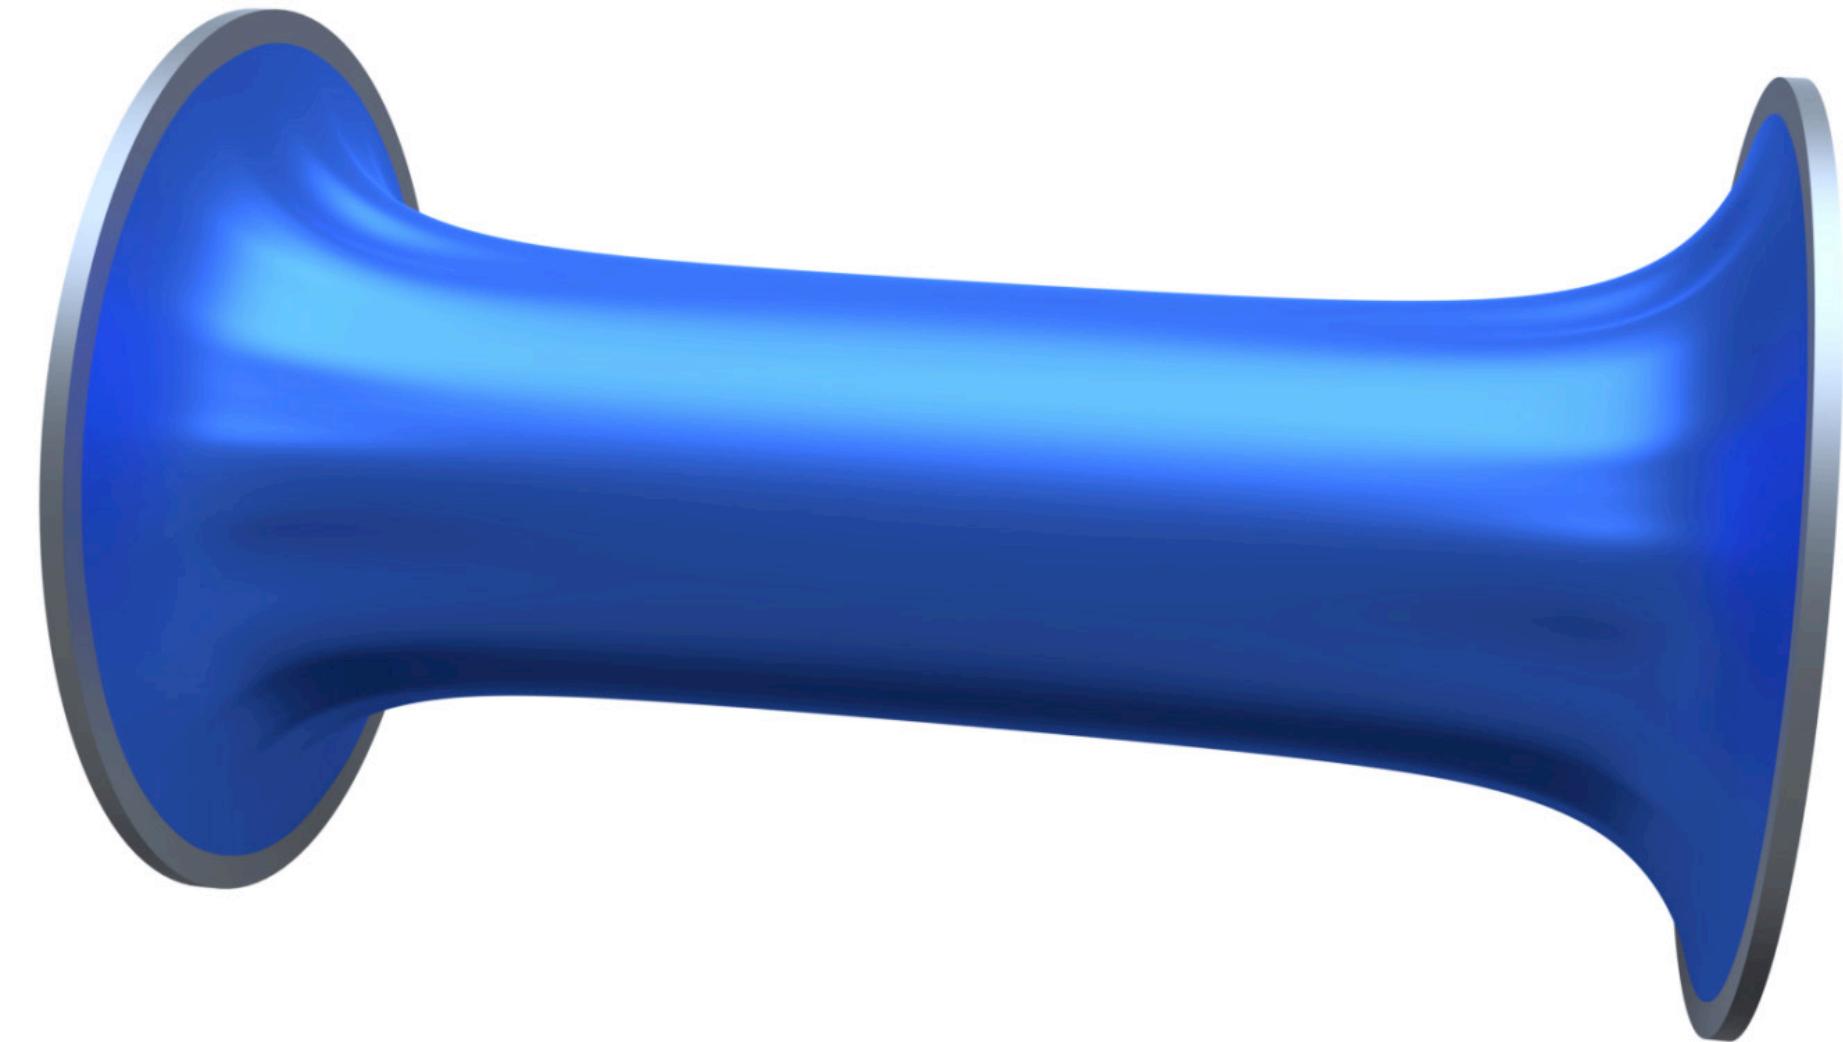
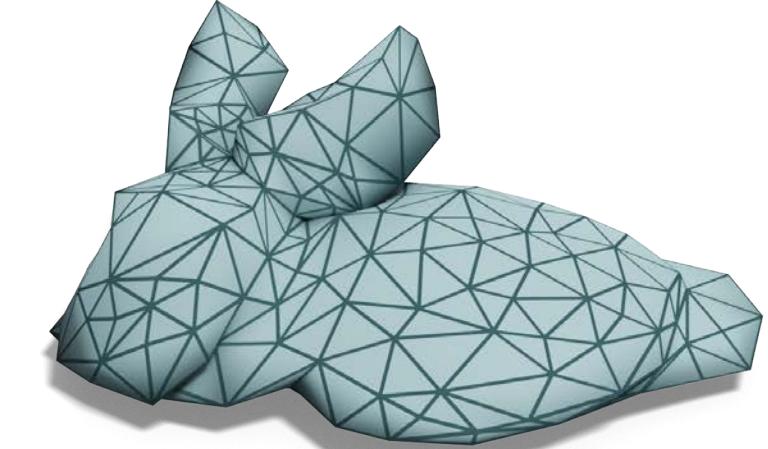
$$E_{\text{elastic}}( \quad , \quad )$$



Target Shape

Current Shape

# Quasistatic Simulation

$$E_{\text{elastic}} ( \quad , \quad )$$


Stable Neo-Hookean Flesh Simulation.

Breannan Smith, Fernando de Goes, and Theodore Kim. ACM Trans. Graph. 2018.

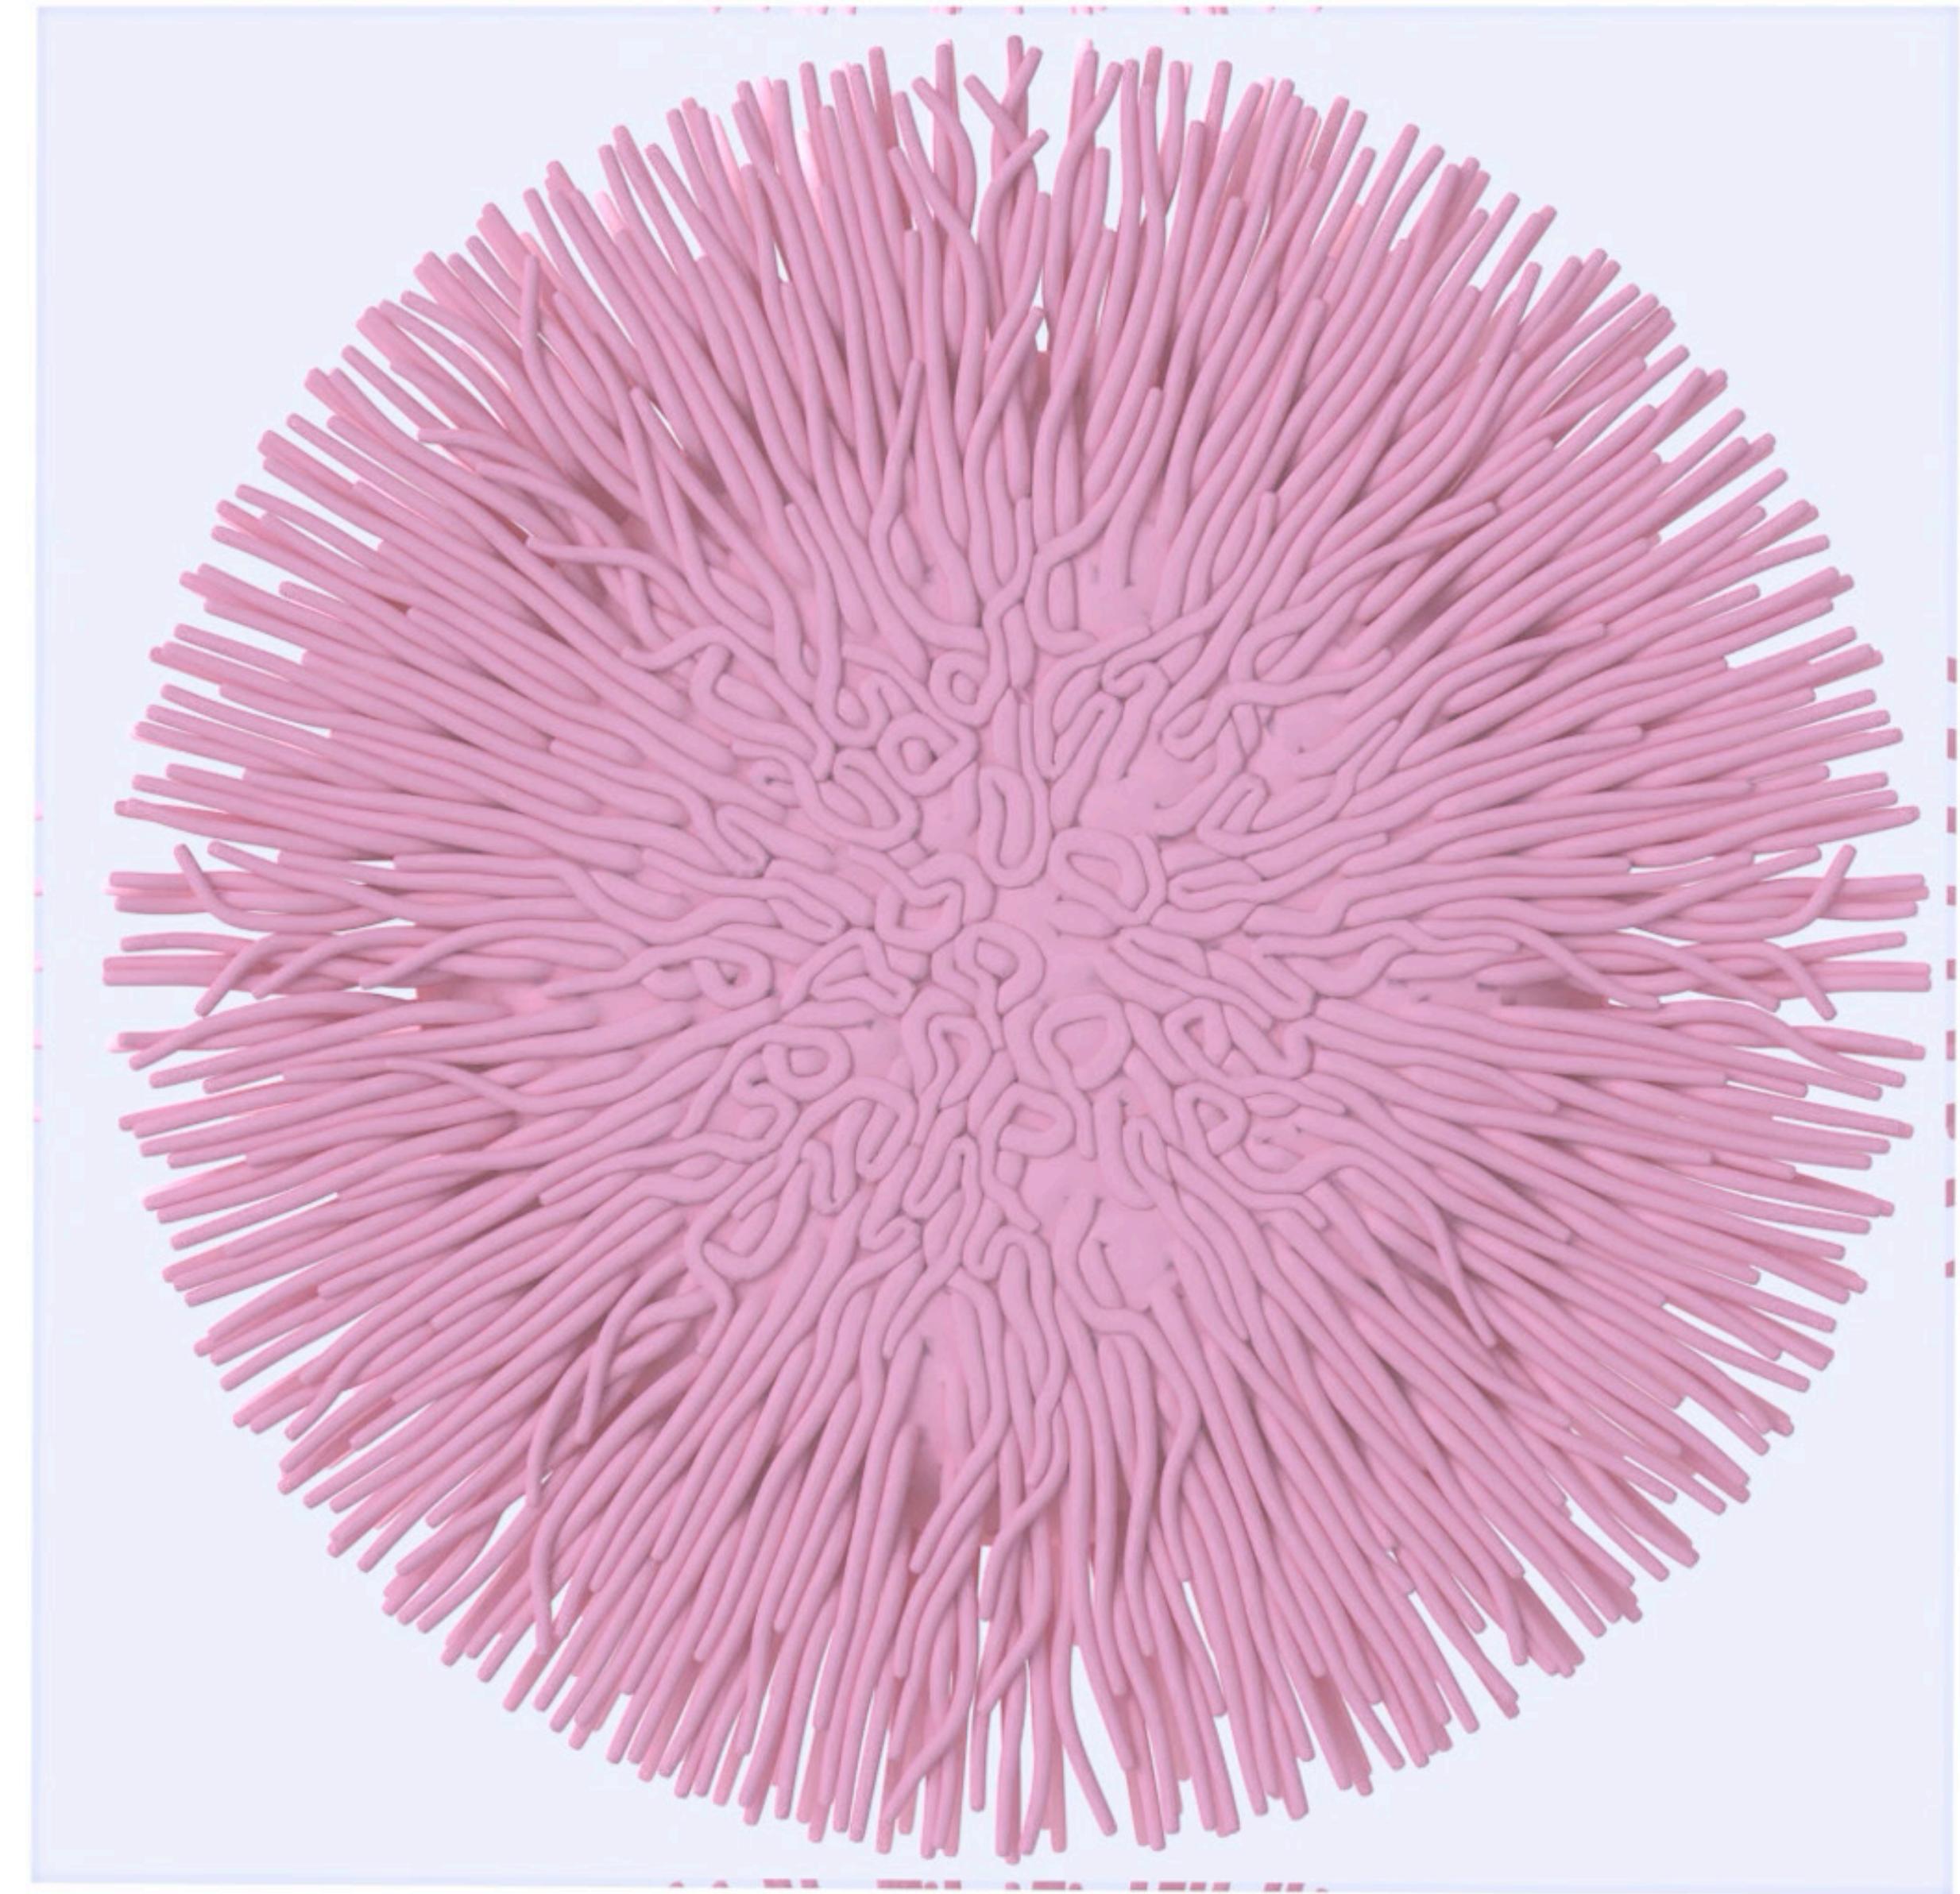
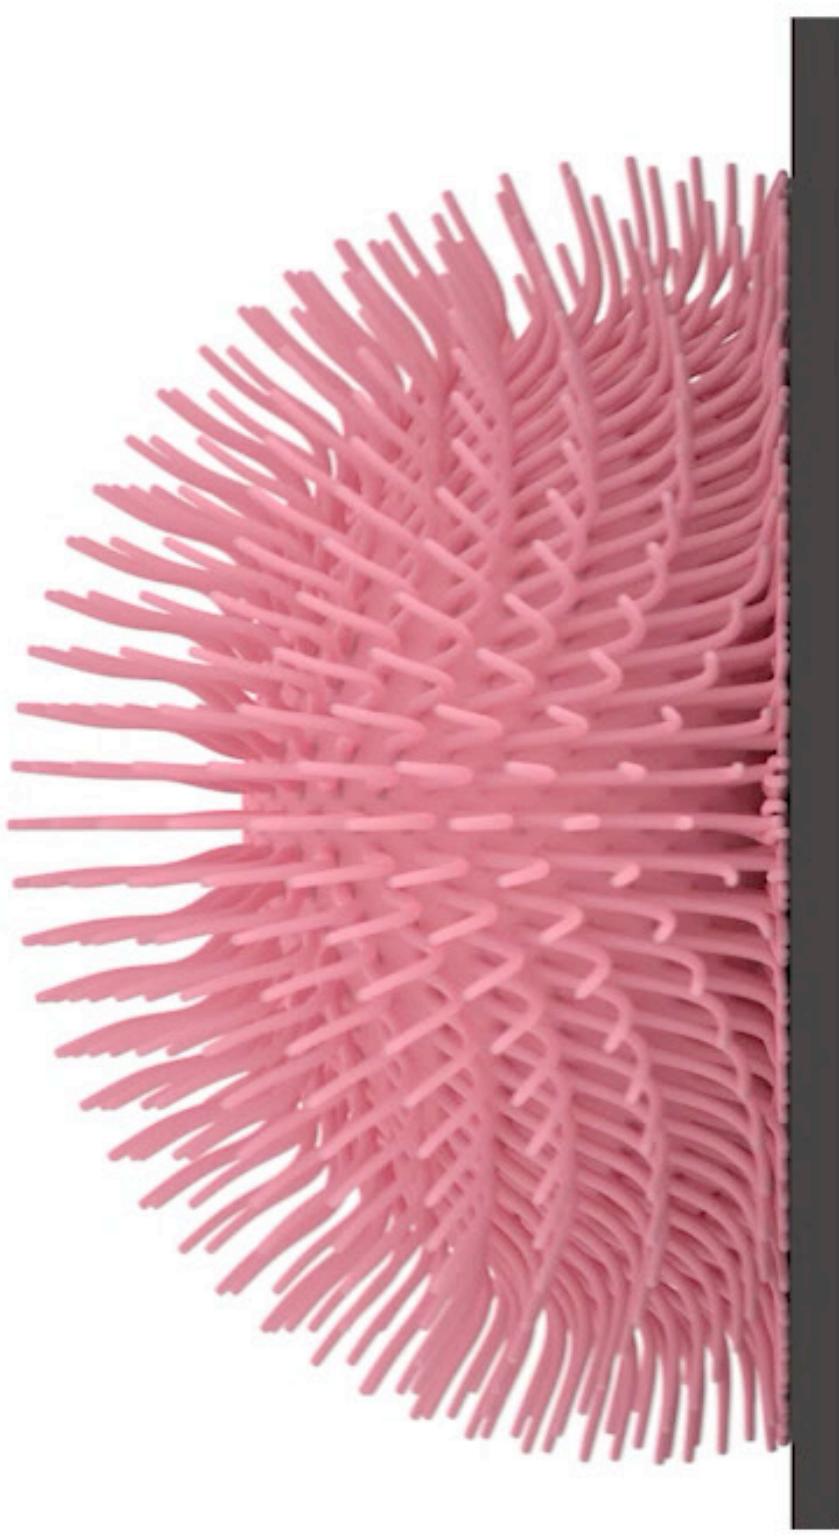
# Dynamic Simulation

$$E_{\text{elastic}} \left( \begin{array}{c} \text{3D Mesh Model} \\ , \end{array} \right) + E_{\text{kinetic}} \\ + E_{\text{external\_force}} \\ + E_{\text{collision}}$$



## Incremental Potential Contact: Intersection- and Inversion-free Large Deformation Dynamics.

Minchen Li, Zachary Ferguson, Teseo Schneider, Timothy Langlois, Denis Zorin, Daniele Panozzo, Chenfanfu Jiang, Danny M. Kaufman.  
ACM Trans. Graph. 2020.

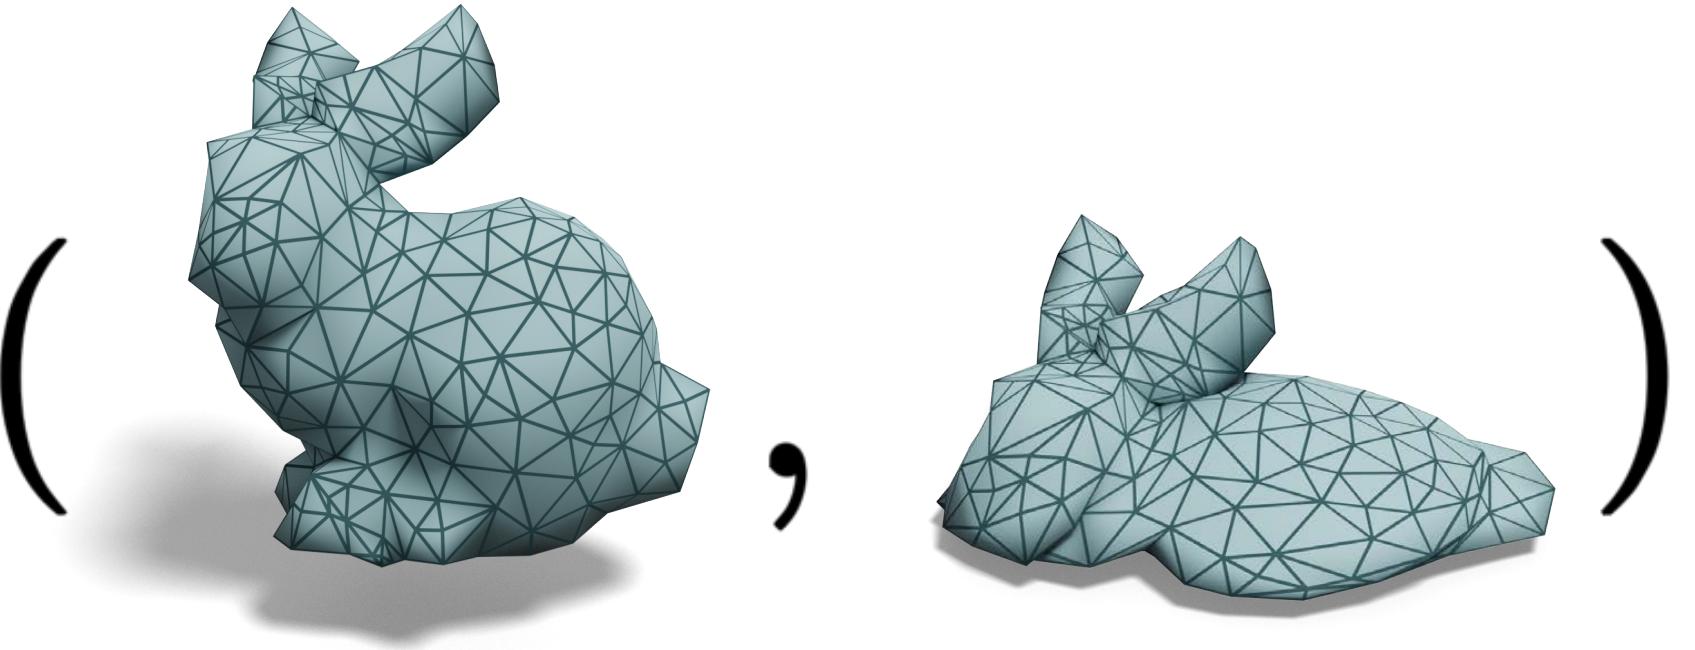


## Incremental Potential Contact: Intersection- and Inversion-free Large Deformation Dynamics.

Minchen Li, Zachary Ferguson, Teseo Schneider, Timothy Langlois, Denis Zorin, Daniele Panozzo, Chenfanfu Jiang, Danny M. Kaufman.  
ACM Trans. Graph. 2020.

# Elastic Energy as Regularization

$$E_{\text{problem\_specific}} + E_{\text{elastic}}( \quad , \quad )$$



## Normal-Driven Spherical Shape Analogies

Hsueh-Ti Derek Liu and Alec Jacobson  
University of Toronto



Figure 1: Our normal-driven spherical shape analogy stylizes an input 3D shape (bottom left) by studying how the surface normal of a style shape (green) relates to the surface normal of a sphere (gray).

### Abstract

This paper introduces a new method to stylize 3D geometry. The key observation is that the surface normal is an effective instrument to capture different geometric styles. Centered around this observation, we cast stylization as a shape analogy problem, where the analogy relationship is defined on the surface normal. This formulation can deform a 3D shape into different styles within a single framework. One can plug-and-play different target styles by providing an exemplar shape or an energy-based style description (e.g., developable surfaces). Our surface stylization methodology enables Normal Captures as a geometric counterpart to material captures (MatCaps) used in rendering, and the prototypical concept of Spherical Shape Analogies as a geometric counterpart to image analogies in image processing.

### 1. Introduction

Analogy of the form  $A : A' :: B : B'$  is a reasoning process that conveys  $A$  is to  $A'$  as  $B$  is to  $B'$ . This formulation has become a core technique for creating artistic 2D digital content, such as image analogies [HIO+01] in Photoshop [Ado21] for image stylization and the Lit Sphere [SMGG01] (a.k.a. MatCap) in ZBrush [PiX20] for non-photorealistic renderings. However, leveraging analogies to stylize 3D geometry is still at a preliminary stage because defining the analogy relationship on surface meshes requires dealing with irregular discretizations, curved metrics, and different topologies.

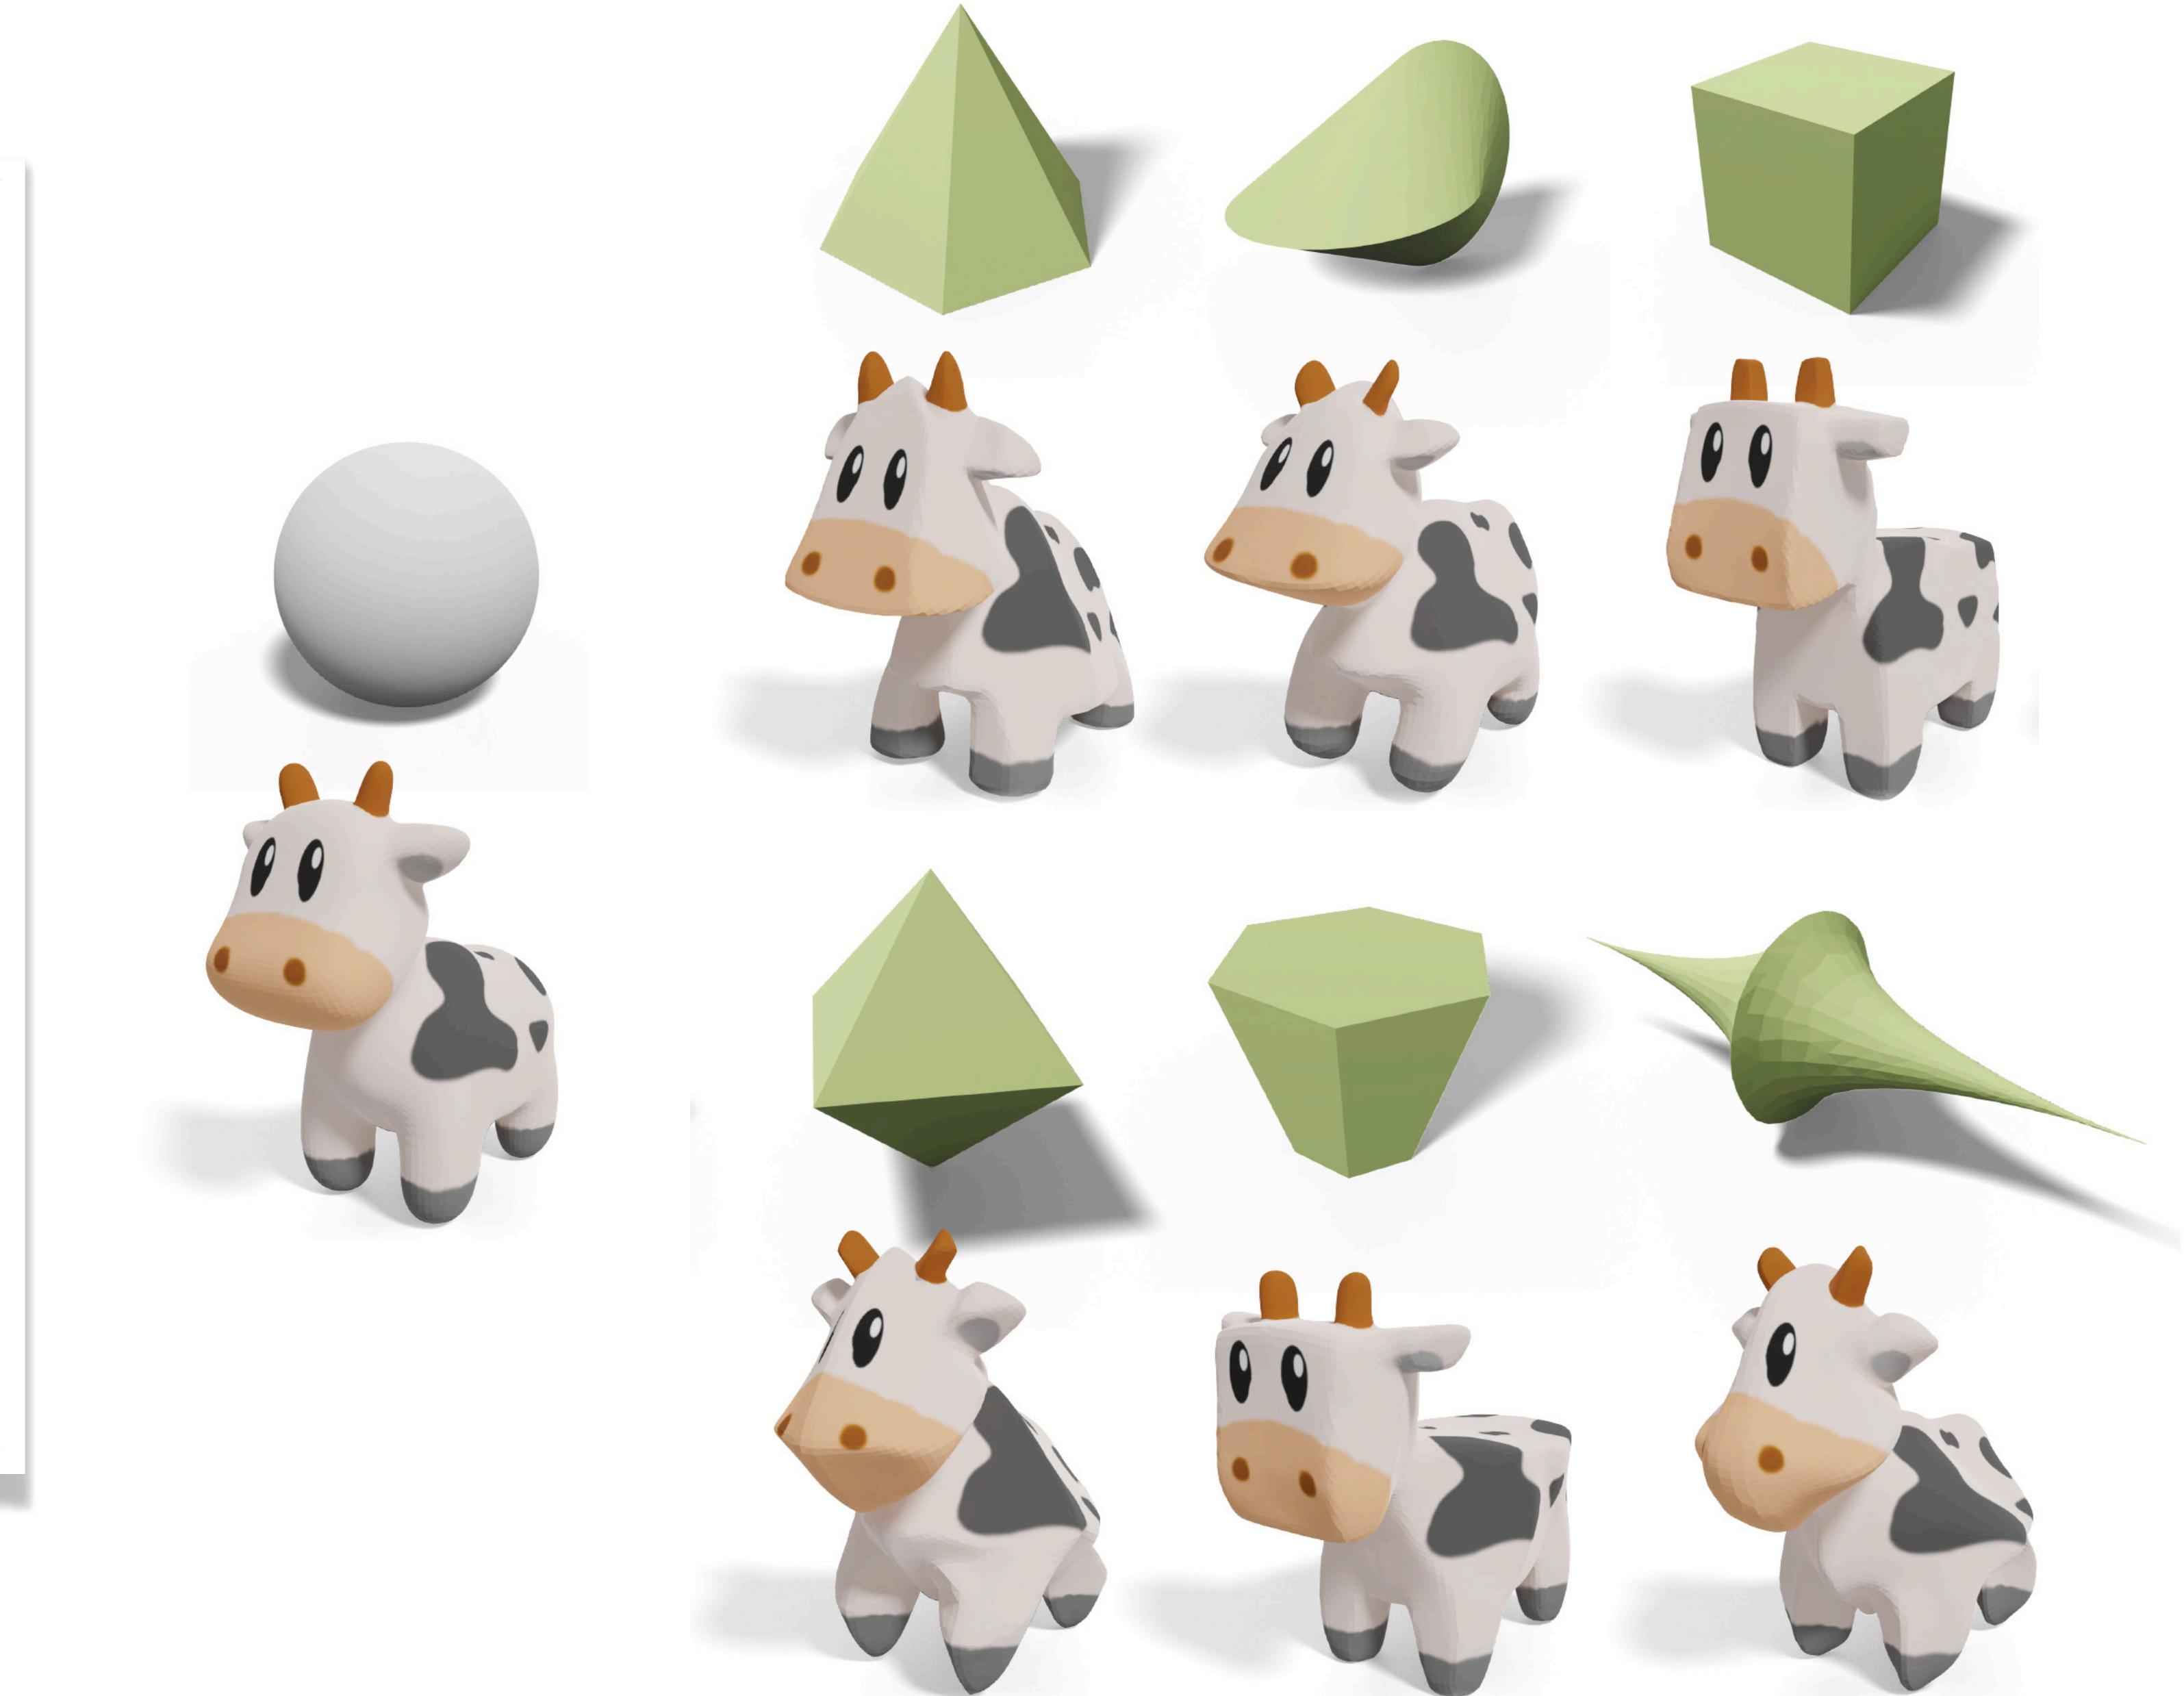
In this paper, we introduce a step towards a more general 3D shape analogies, named *spherical shape analogies*. We consider a specific case where  $A$  is a unit sphere. This restriction enables us to operate on an input mesh  $B$  with arbitrary topologies, boundaries, and geometric complexity. While not fully general, because  $A$  is restricted to be a sphere, we demonstrate that this formulation can immediately achieve different geometric styles within a single

framework. In Fig. 1, we show that by providing different target style shapes  $A'$  to the algorithm, we can turn the input shape  $B$  into different styles. In addition to stylization, our method can encompass many existing applications, such as developable surface approximation and PolyCube deformation.

One key observation in our spherical shape analogies is that the surface normal is an effective instrument to capture geometric styles. Thus, we define the analogy relationship based on normals: we optimize a stylized shape  $B'$  such that the relationship between the surface normals of  $B$  and  $B'$  is the same as the relationship between the surface normals of  $A$  and  $A'$ .

We realize this by casting it as a simple and effective normal-driven shape optimization problem which aims at deforming the input shape towards a set of desired normals. However, such an optimization problem is often difficult due to the nonlinearity of unit normals. We draw inspiration from previous works and apply a change of variables to accelerate the computation: instead of di-

# Shape Stylization

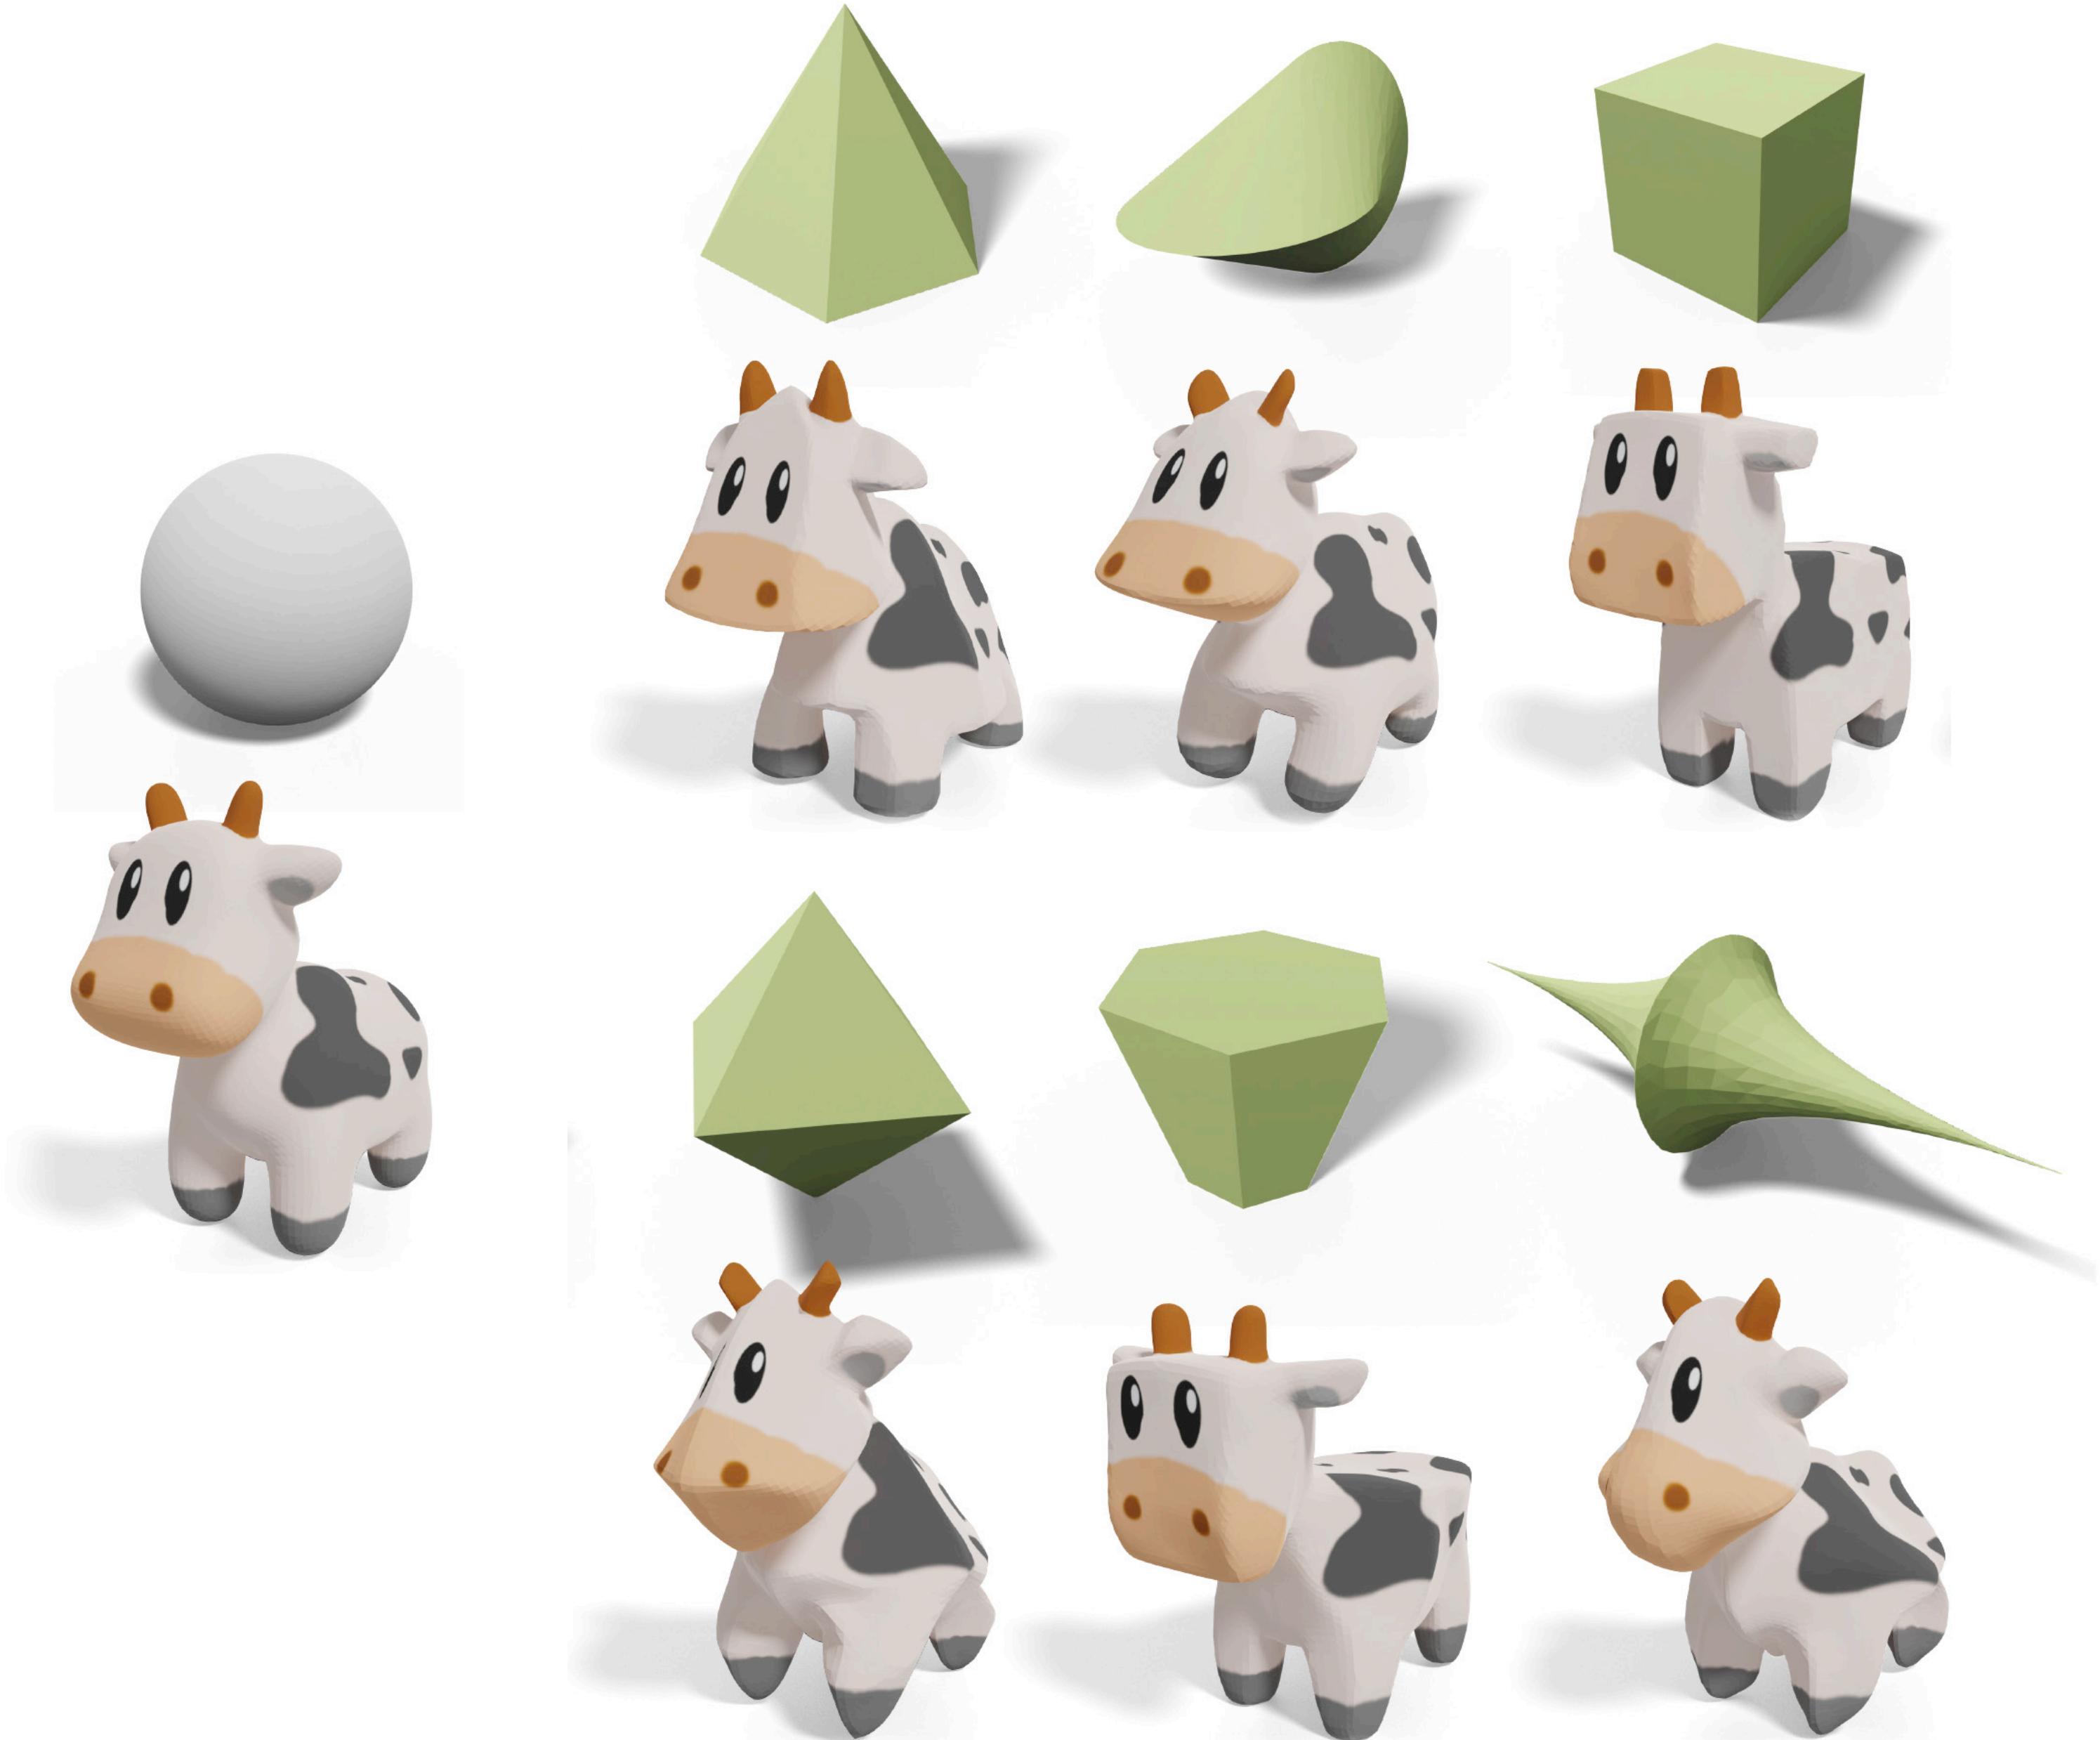


$E_{\text{style}}$

+

$E_{\text{elastic}}$

Shape Stylization



## Local Deformation for Interactive Shape Editing

Honglin Chen  
Columbia University  
New York City, NY, USA  
honglin.chen@columbia.edu

Changxi Zheng  
Columbia University  
New York City, NY, USA  
cxz@cs.columbia.edu

Kevin Wampler  
Adobe Research  
Seattle, WA, USA  
kwampler@adobe.com

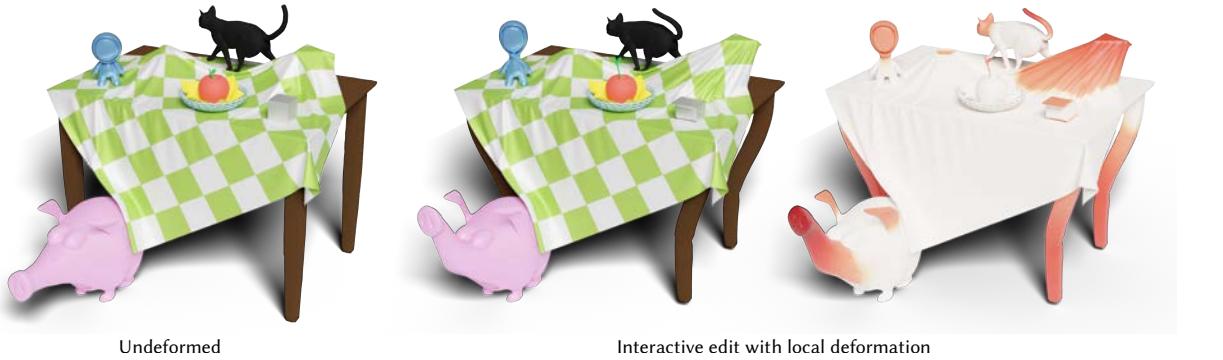


Figure 1: Our method enables the user to edit shapes in an interactive and physically plausible way. The edit is *local*, meaning that the user can focus on one region of the complex scene without worrying about inadvertent changes elsewhere. To visualize the locality, in the rightmost figure we highlight the regions where the vertex displacement is larger than  $10^{-3}$  in red. (Undeformed scene shapes thanks to [Zhang et al. 2022])

### ABSTRACT

We introduce a novel regularization for localizing an elastic-energy-driven deformation to only those regions being manipulated by the user. Our local deformation features a natural region of influence, which is automatically adaptive to the geometry of the shape, the size of the deformation and the elastic energy in use. We further propose a three-block ADMM-based optimization to efficiently minimize the energy and achieve interactive frame rates. Our approach avoids the artifacts of other alternative methods, is simple and easy to implement, does not require tedious control primitive setup and generalizes across different dimensions and elastic energies. We demonstrate the effectiveness and efficiency of our localized deformation tool through a variety of local editing scenarios, including 1D, 2D, 3D elasticity and cloth deformation.

### CCS CONCEPTS

- Computing methodologies → Mesh geometry models.

### KEYWORDS

Local control, shape deformation, elasticity, sparsity, ADMM.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyright held by the author(s). All rights reserved. Copying or distribution by others than the author(s) may be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.  
SIGGRAPH '23 Conference Proceedings, August 6–10, 2023, Los Angeles, CA, USA  
© 2023 Copyright held by the owner(s). Publication rights licensed to ACM.  
ACM ISBN 978-1-4503-9308-515.00  
<https://doi.org/10.1145/3588432.3591501>

### ACM Reference Format:

Honglin Chen, Changxi Zheng, and Kevin Wampler. 2023. Local Deformation for Interactive Shape Editing. In *Special Interest Group on Computer Graphics and Interactive Techniques Conference Proceedings (SIGGRAPH '23 Conference Proceedings)*, August 6–10, 2023, Los Angeles, CA, USA. ACM, New York, NY, USA, 15 pages. <https://doi.org/10.1145/3588432.3591501>

### 1 INTRODUCTION

Local deformation is a core component in modeling and animation. In a localized deformation, only the parts of the shape near where the user is currently manipulating move—everything else stays still, ensuring that the user can focus entirely on one region of the shape without worrying about inadvertent changes elsewhere. However, existing localized deformation tools tend to have practical impediments for interactive design: they are either too slow to run, unaware of the geometry, introduce artifacts, or require a careful control point setup.

When global deformation is acceptable, a widely useful approach is to solve for the deformation by minimizing an elastic energy defined over the shape, subject to positional constraints derived from the user's input. This paradigm has many advantages. The deformation accounts for the geometry of the shape, generalizes well to 2D, 3D, and cloth, and the elastic energy can be used to model a wide range of both real-world and stylized materials. Unfortunately elastic energy minimization is by its nature global, and jointly solves for all the degrees of freedom in the shape. This necessitates a rigging step to “pin down” certain aspects of the deformation lest the optimizer move them. This limits the applicability of such methods to situations where a suitable rig is available, or the region of influence for a deformation is known in advance.



# Local Deformation

$E_{\text{local}}$ 

+

 $E_{\text{elastic}}$ 

Local Deformation

## Large Steps in Inverse Rendering of Geometry

BAPTISTE NICOLET, École Polytechnique Fédérale de Lausanne (EPFL), Switzerland  
ALEC JACOBSON, University of Toronto, Canada  
WENZEL JAKOB, École Polytechnique Fédérale de Lausanne (EPFL), Switzerland

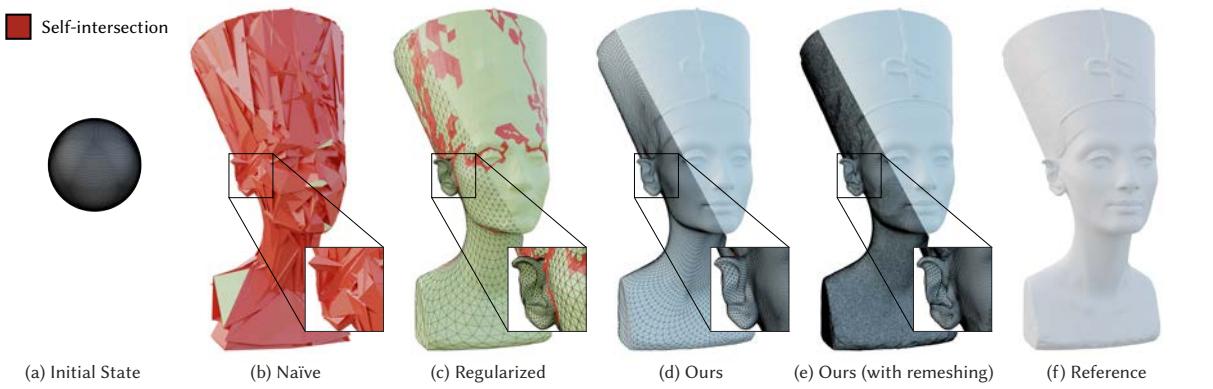


Fig. 1. (a) Inverse reconstruction of the NEFERTITI bust from a spherical starting guess with 25 rendered views (1 shown). (b) Naive application of a differentiable renderer produces an unusable tangled mesh when gradient steps pull on the silhouette without regard for distortion or self-intersections. (c) Regularization can alleviate such problems by making the optimization aware of mesh quality. On the flipside, this penalizes non-smooth parts of the geometry and causes unsatisfactory convergence by gradient-based optimizers. While the final mesh undeniably looks better, a closer inspection of the wireframe rendering reveals countless self-intersections. (d) Our method addresses both problems and converges to a high-quality mesh. (e) Combined with an isotropic remeshing step, our reconstruction captures fine details of the reference. (f) The hyper-parameters of each method were optimized to obtain the best convergence at equal time. Self-intersections are shown in red.

Inverse reconstruction from images is a central problem in many scientific and engineering disciplines. Recent progress on differentiable rendering has led to methods that can efficiently differentiate the full process of image formation with respect to millions of parameters to solve such problems via gradient-based optimization.

At the same time, the availability of cheap derivatives does not necessarily make an inverse problem easy to solve. Mesh-based representations remain a particular source of irritation: an adverse gradient step involving vertex positions could turn parts of the mesh inside-out, introduce numerous local self-intersections, or lead to inadequate usage of the vertex budget due to distortion. These types of issues are often irrecoverable in the sense that subsequent optimization steps will further exacerbate them. In other words, the optimization lacks robustness due to an objective function with substantial non-convexity.

Such robustness issues are commonly mitigated by imposing additional regularization, typically in the form of Laplacian energies that quantify and improve the smoothness of the current iterate. However, regularization introduces its own set of problems: solutions must now compromise between solving the problem and being smooth. Furthermore, gradient steps involving

a Laplacian energy resemble Jacobi's iterative method for solving linear equations that is known for its exceptionally slow convergence.

We propose a simple and practical alternative that casts differentiable rendering into the framework of preconditioned gradient descent. Our preconditioner biases gradient steps towards smooth solutions without requiring the final solution to be smooth. In contrast to Jacobi-style iteration, each gradient step propagates information among all variables, enabling convergence using fewer and larger steps.

Our method is not restricted to meshes and can also accelerate the reconstruction of other representations, where smooth solutions are generally expected. We demonstrate its superior performance in the context of geometric optimization and texture reconstruction.

CCS Concepts: • Computing methodologies → Rendering; Shape modeling.

Additional Key Words and Phrases: differentiable rendering, geometry reconstruction, Laplacian mesh processing

ACM Reference Format:  
Baptiste Nicolet, Alec Jacobson, and Wenzel Jakob. 2021. Large Steps in Inverse Rendering of Geometry. *ACM Trans. Graph.*, 40, 6, Article 248 (December 2021), 13 pages. <https://doi.org/10.1145/3478513.3480501>

© 2021 Copyright held by the owner/author(s). Publication rights licensed to ACM. This is the author's version of the work. It is posted here for your personal use. Not for redistribution. The definitive Version of Record was published in *ACM Transactions on Graphics*, <https://doi.org/10.1145/3478513.3480501>.

ACM Trans. Graph., Vol. 40, No. 6, Article 248. Publication date: December 2021.

# Inverse Rendering

Target:



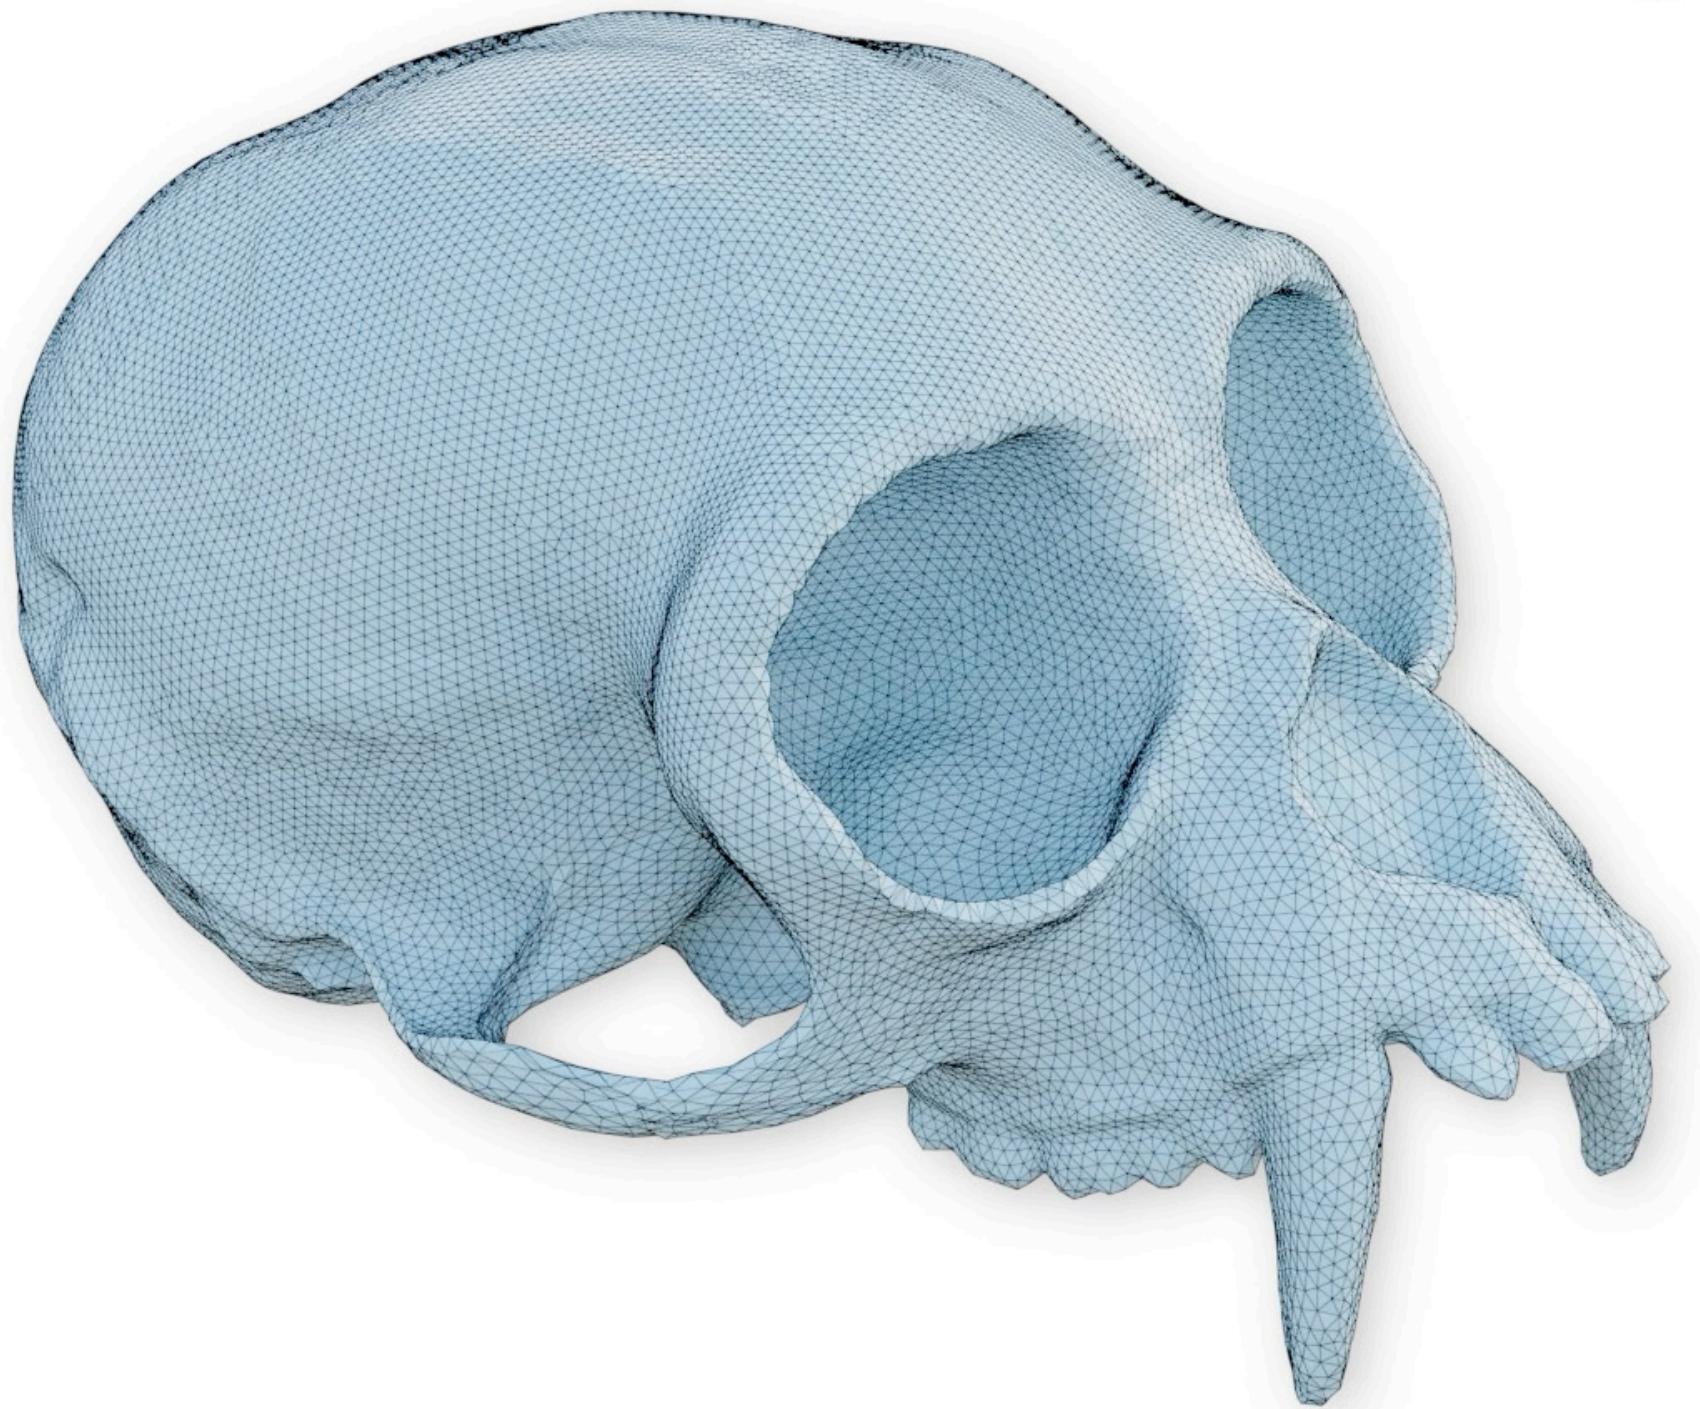
$E_{\text{render}}$

+

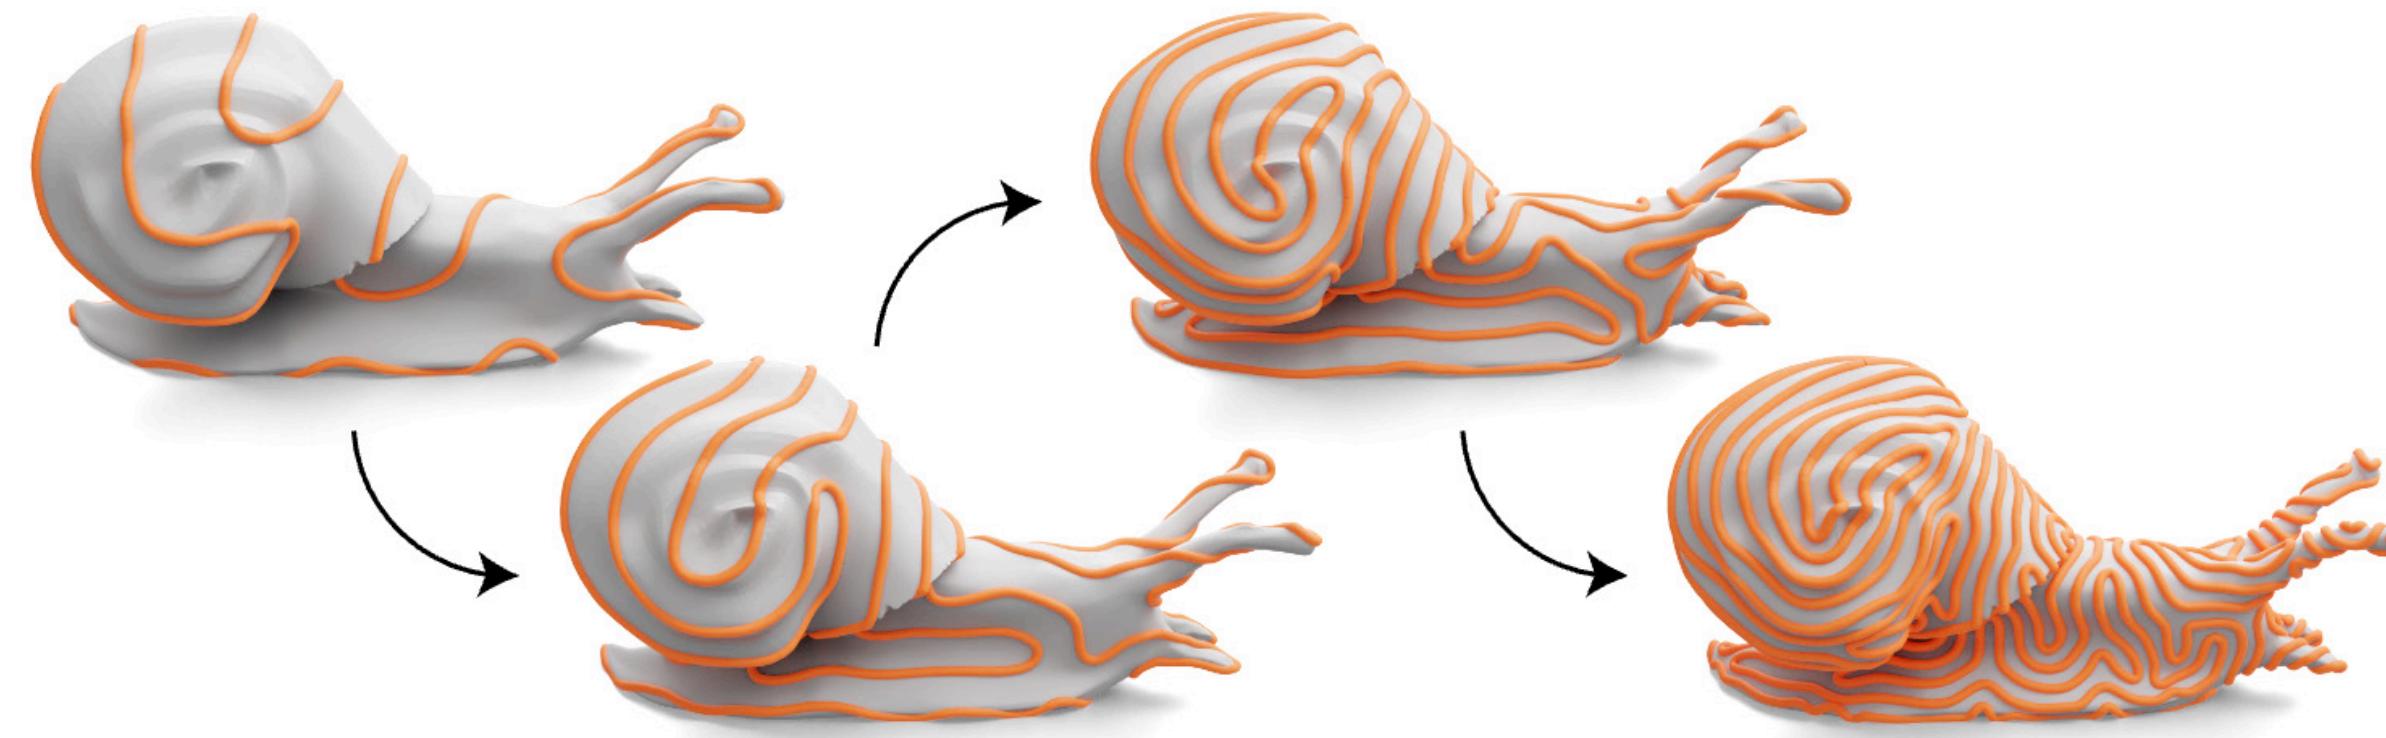
$E_{\text{elastic}}$

Inverse Rendering

Target:



# More than just Elastic Energy



Surface-Filling Curves

[Noma et al. 2024]



Quad Mesh Planarization

[Liu et al. 2006]



Developable Surfaces

[Sellán et al. 2020]



Magnetic Simulation

[Chen et al. 2022]

And many more...!

- Elastic Energy
- Optimization Basics & Classical Algorithms
- Optimization Methods for Simulation
  - Challenge 1 ✨: Handling Nonconvexity
  - Challenge 2 ✨: Nontrivial Constraints
  - Challenge 3 ✨: Large-scale Optimization

# Optimization Basics

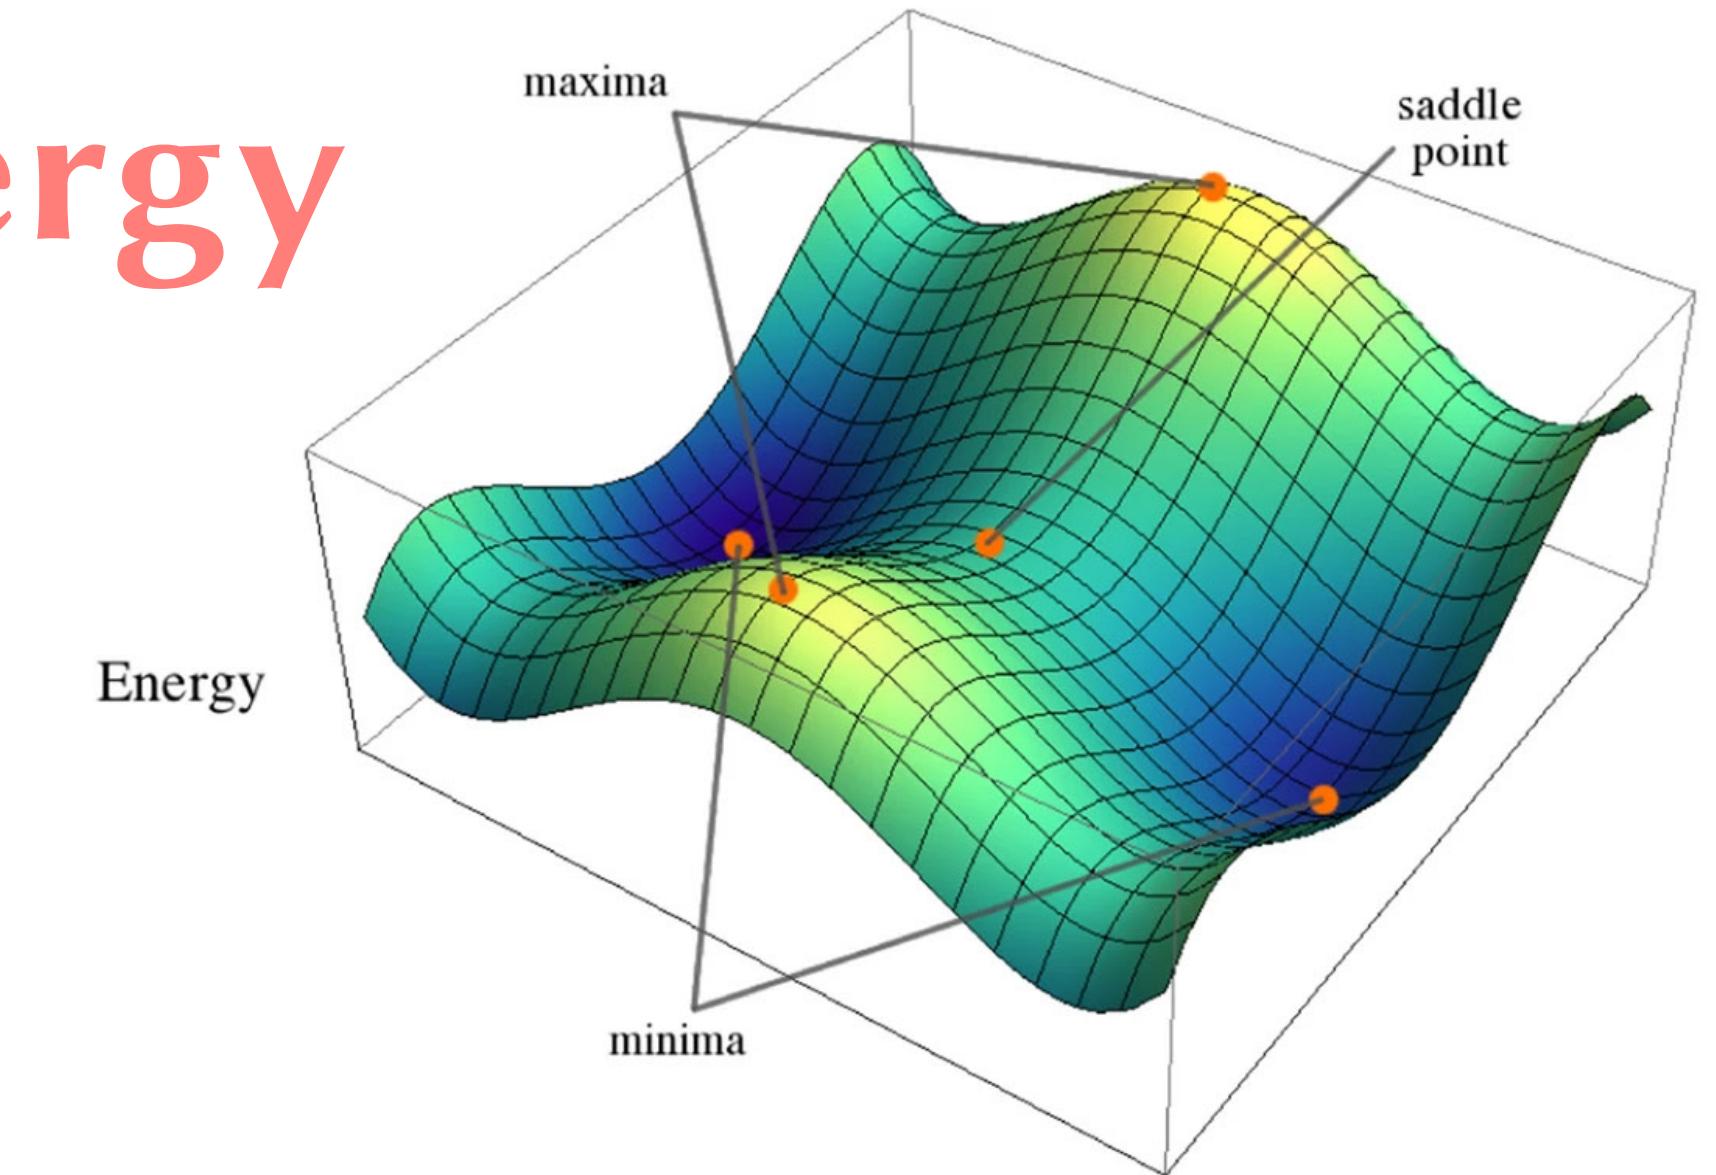
---

# What is Energy Minimization?



Optimization  
Variables

$$\min_x f(x)$$





**High Energy**  
Undesired State

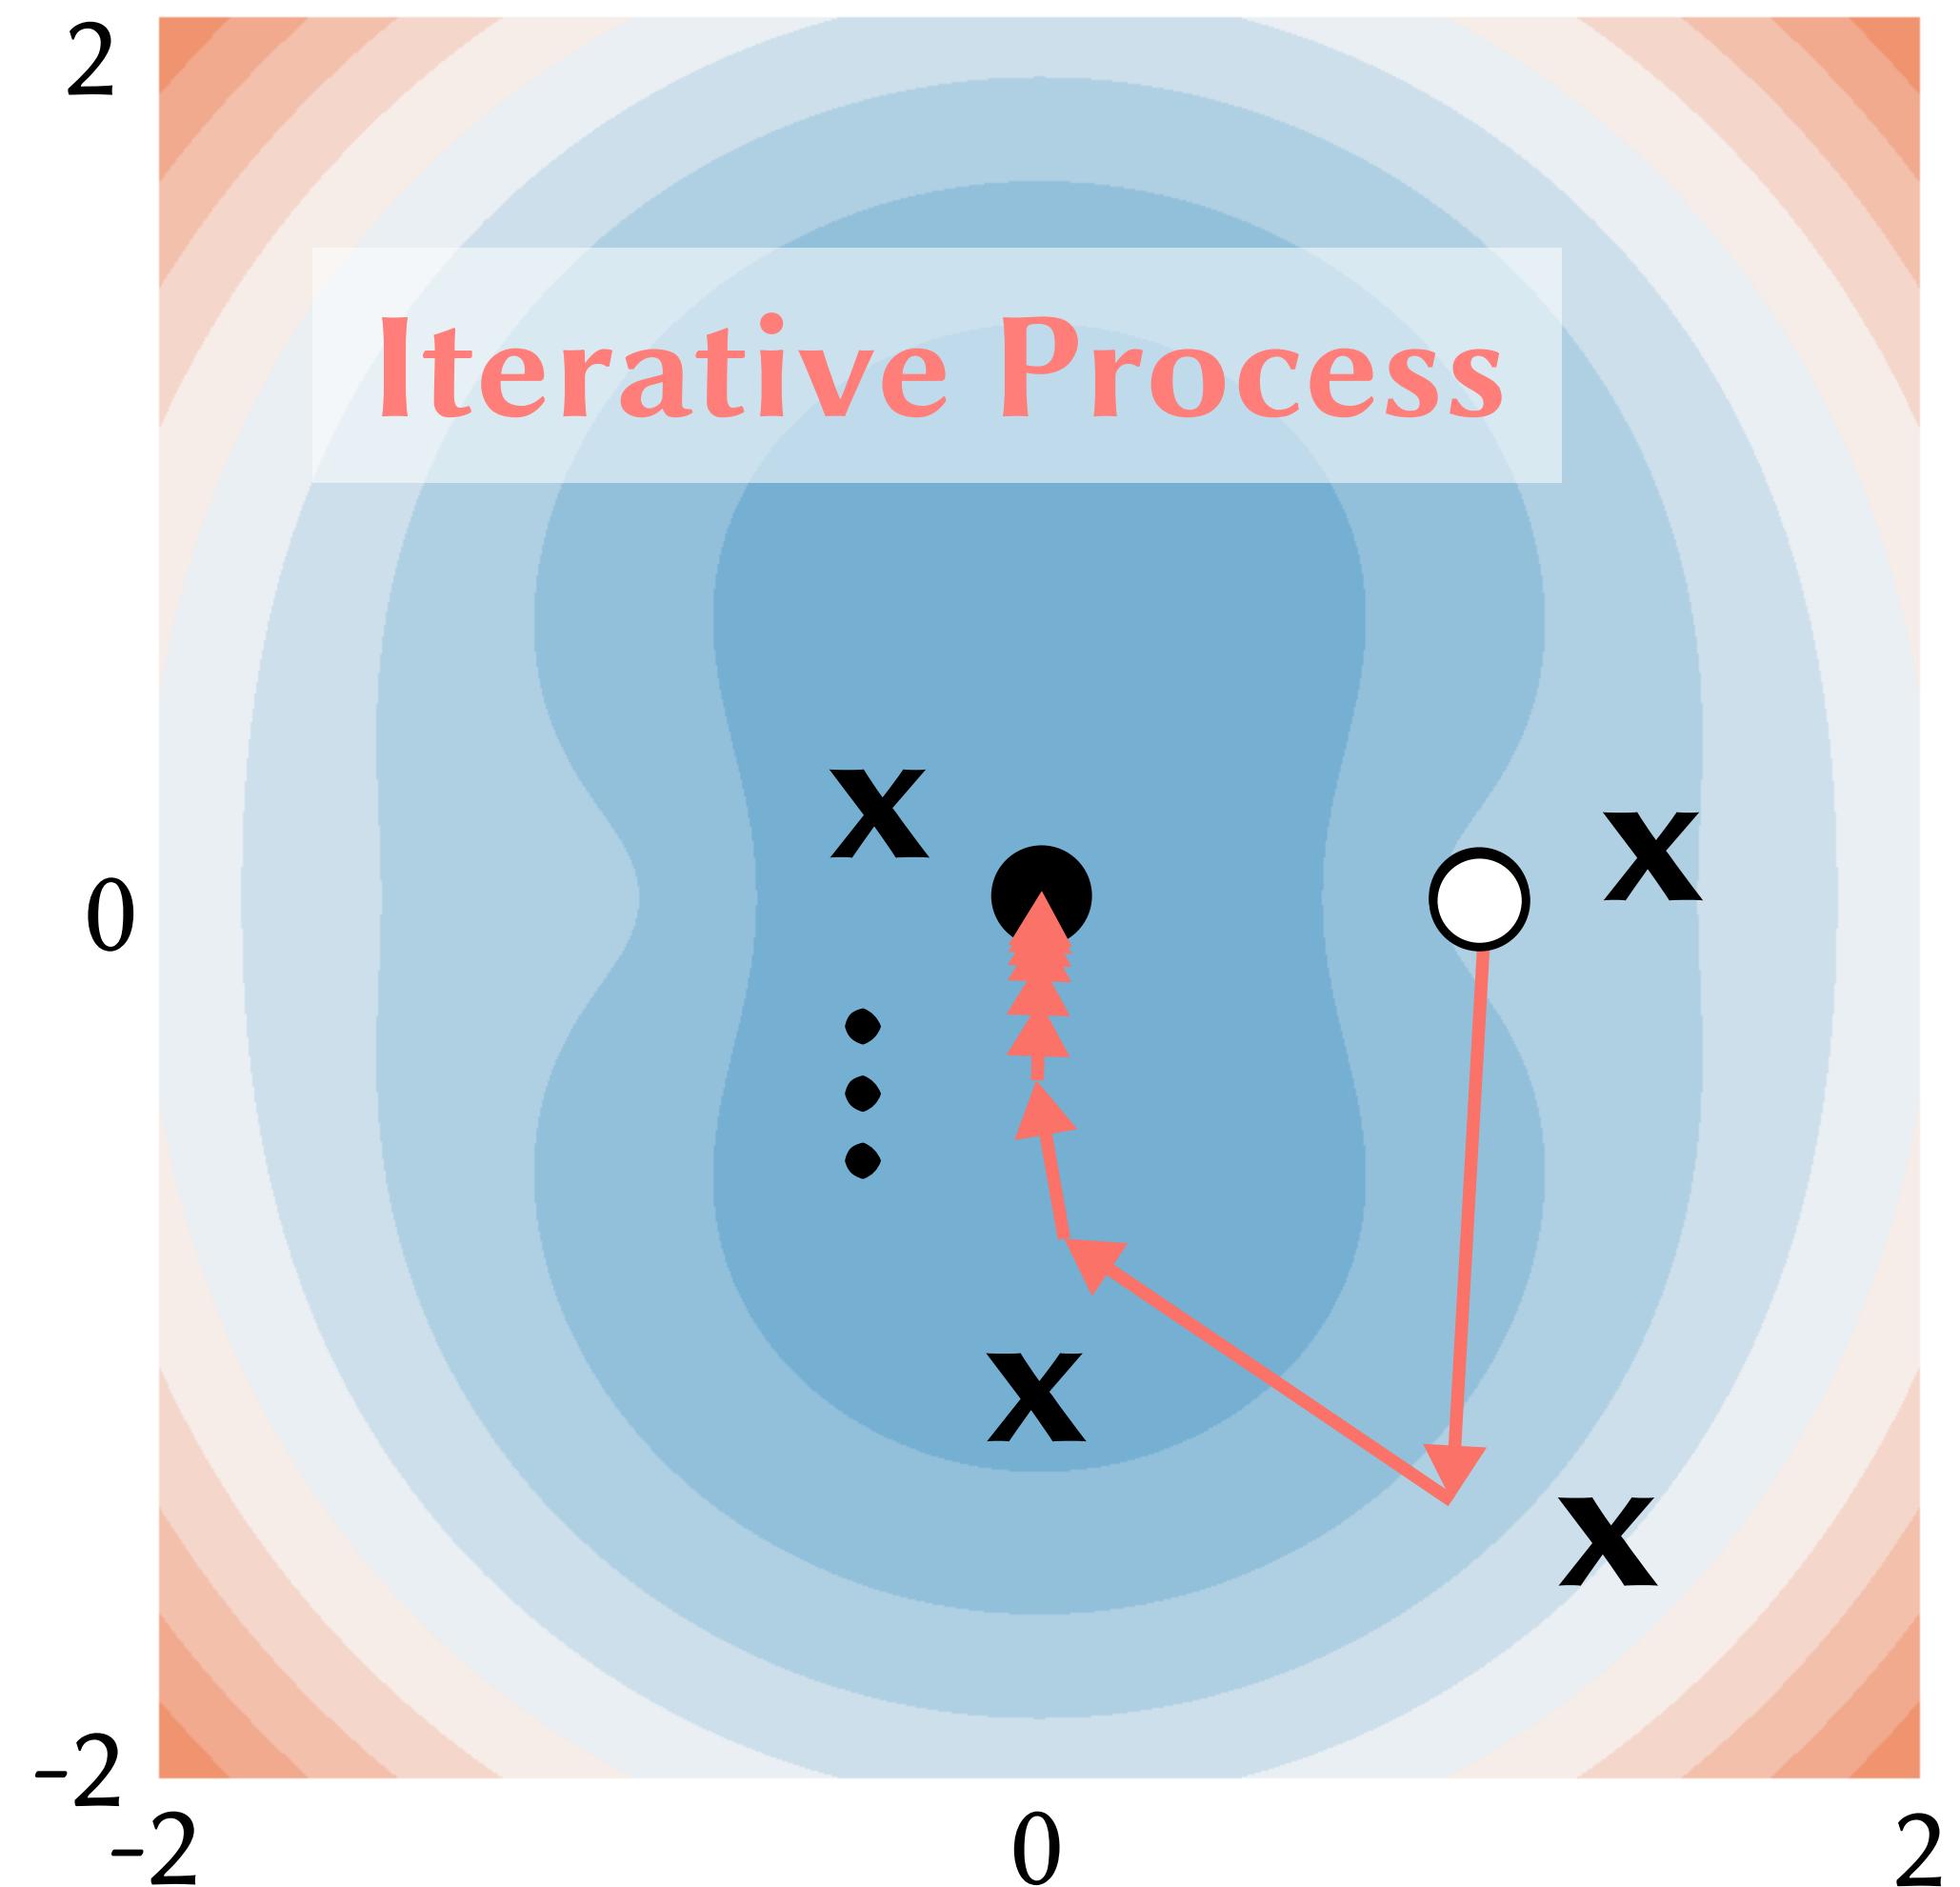
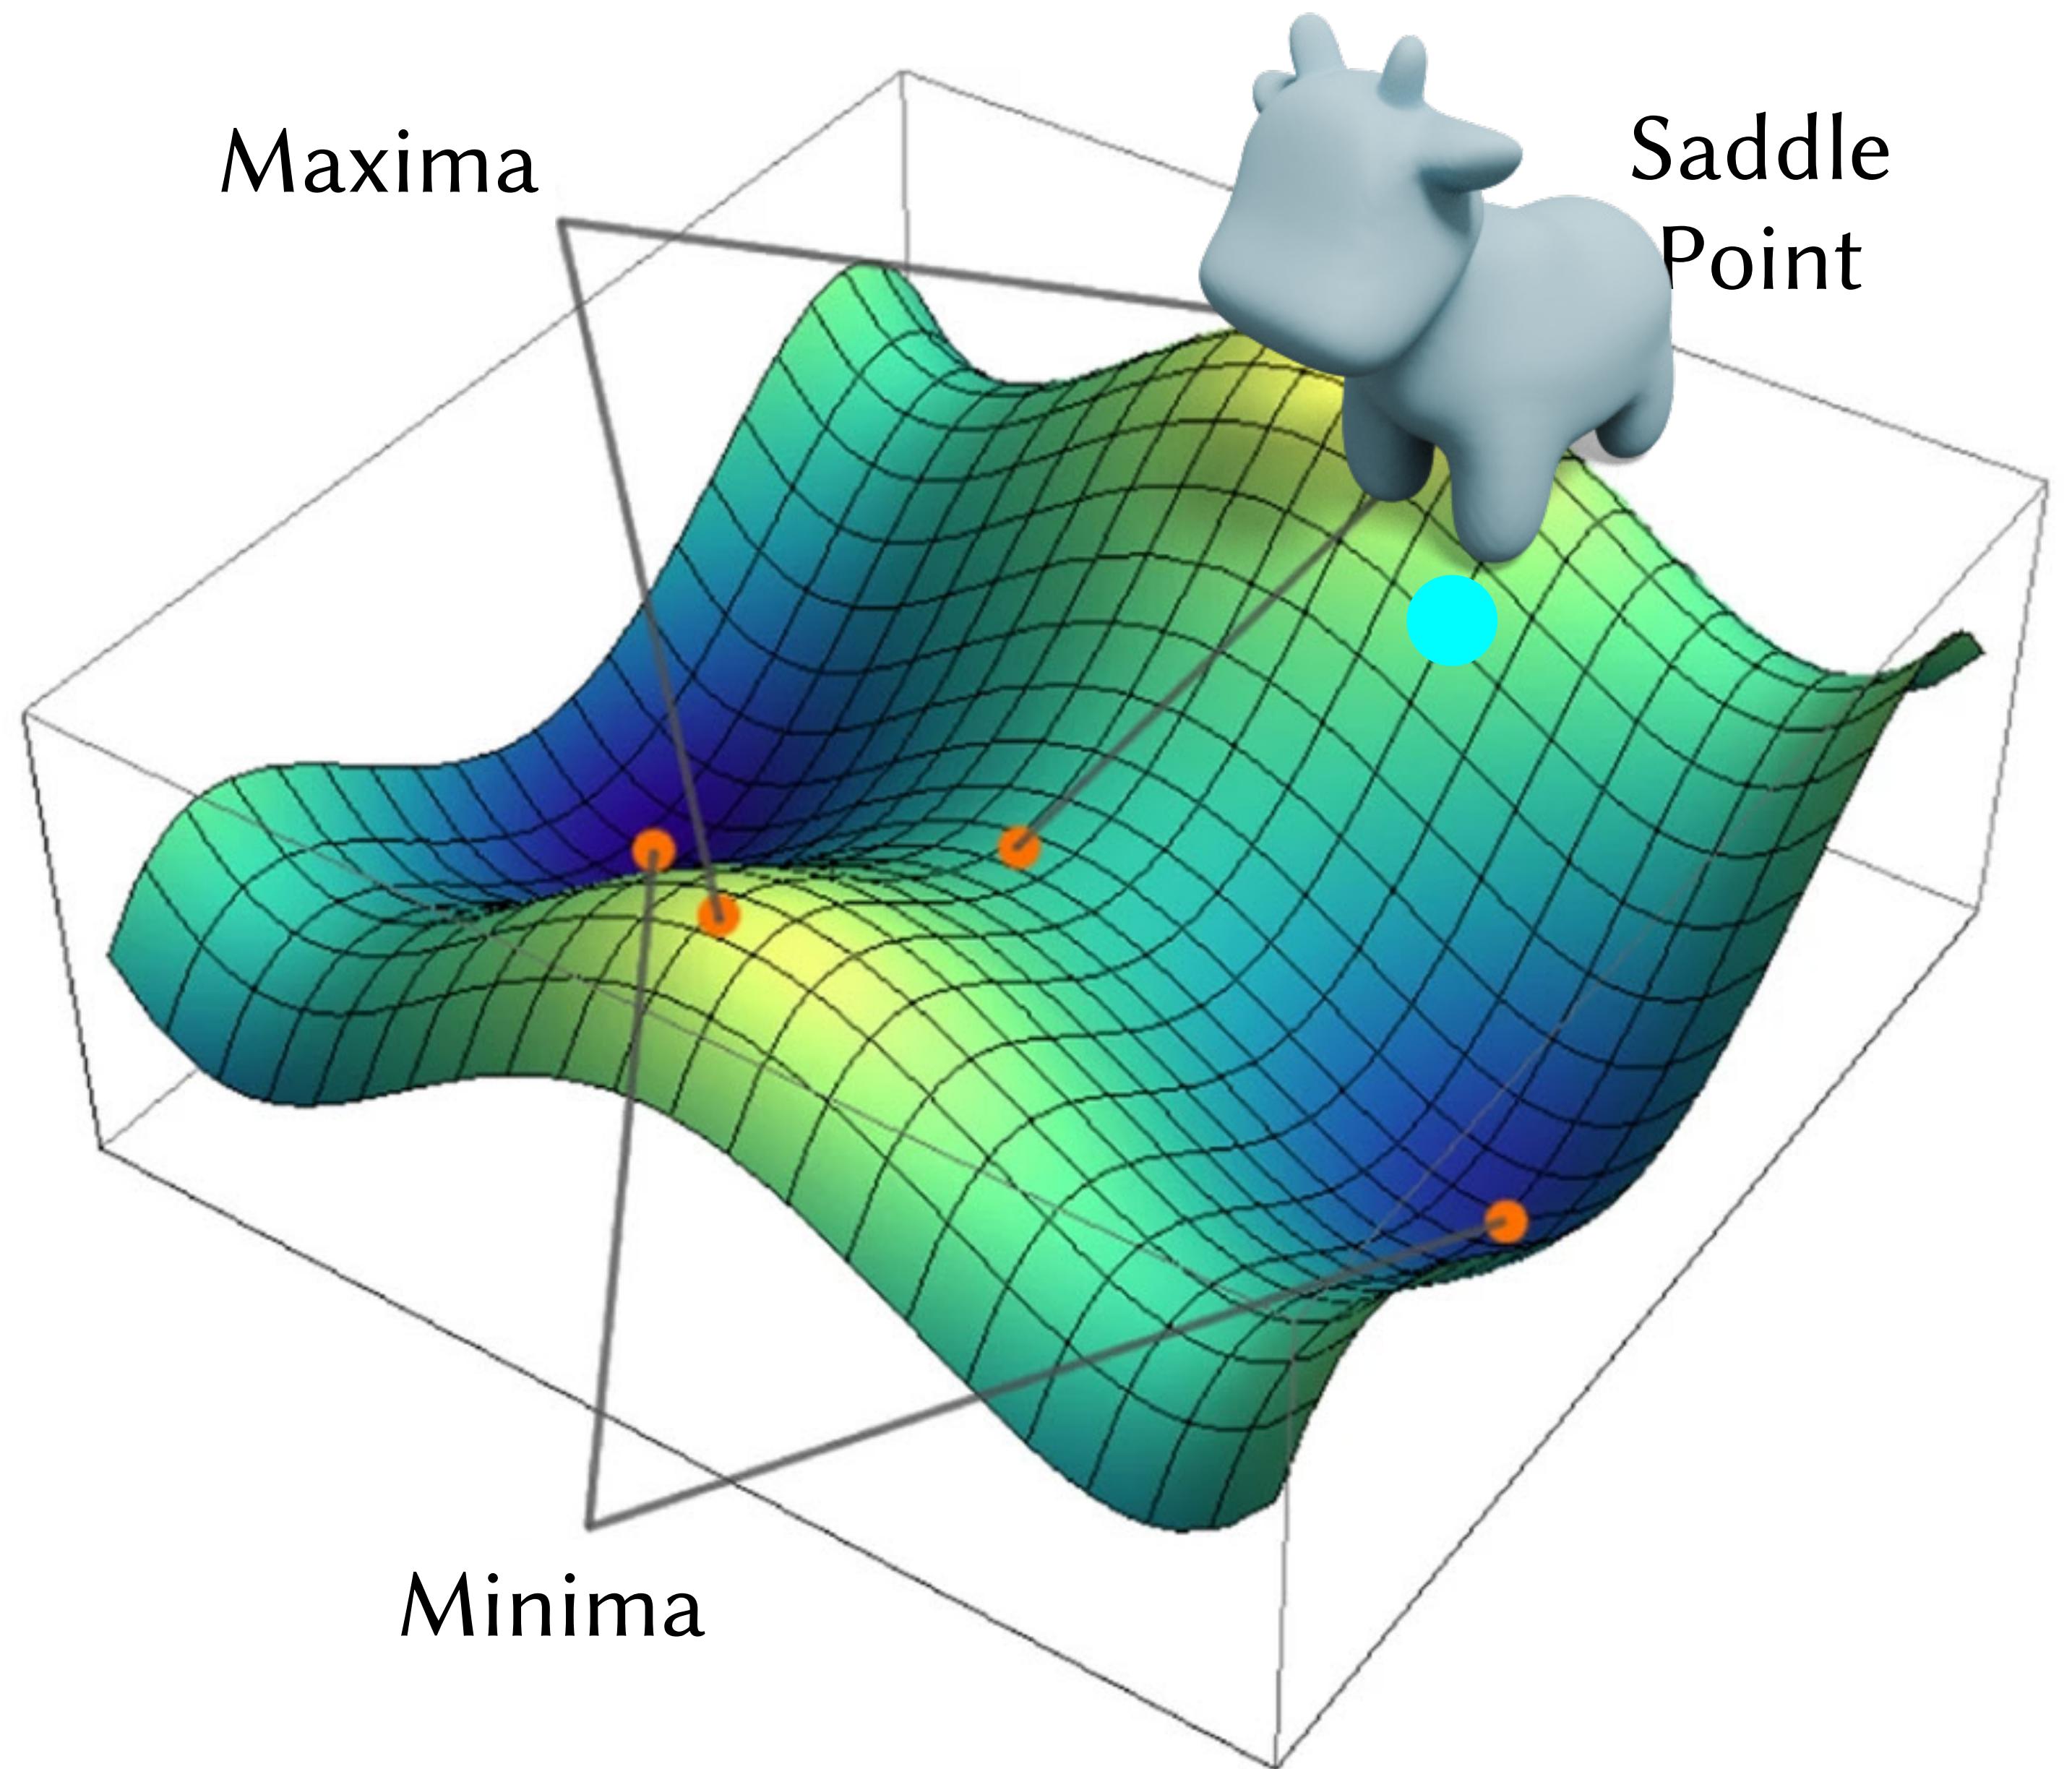
**Low Energy**  
Desired State

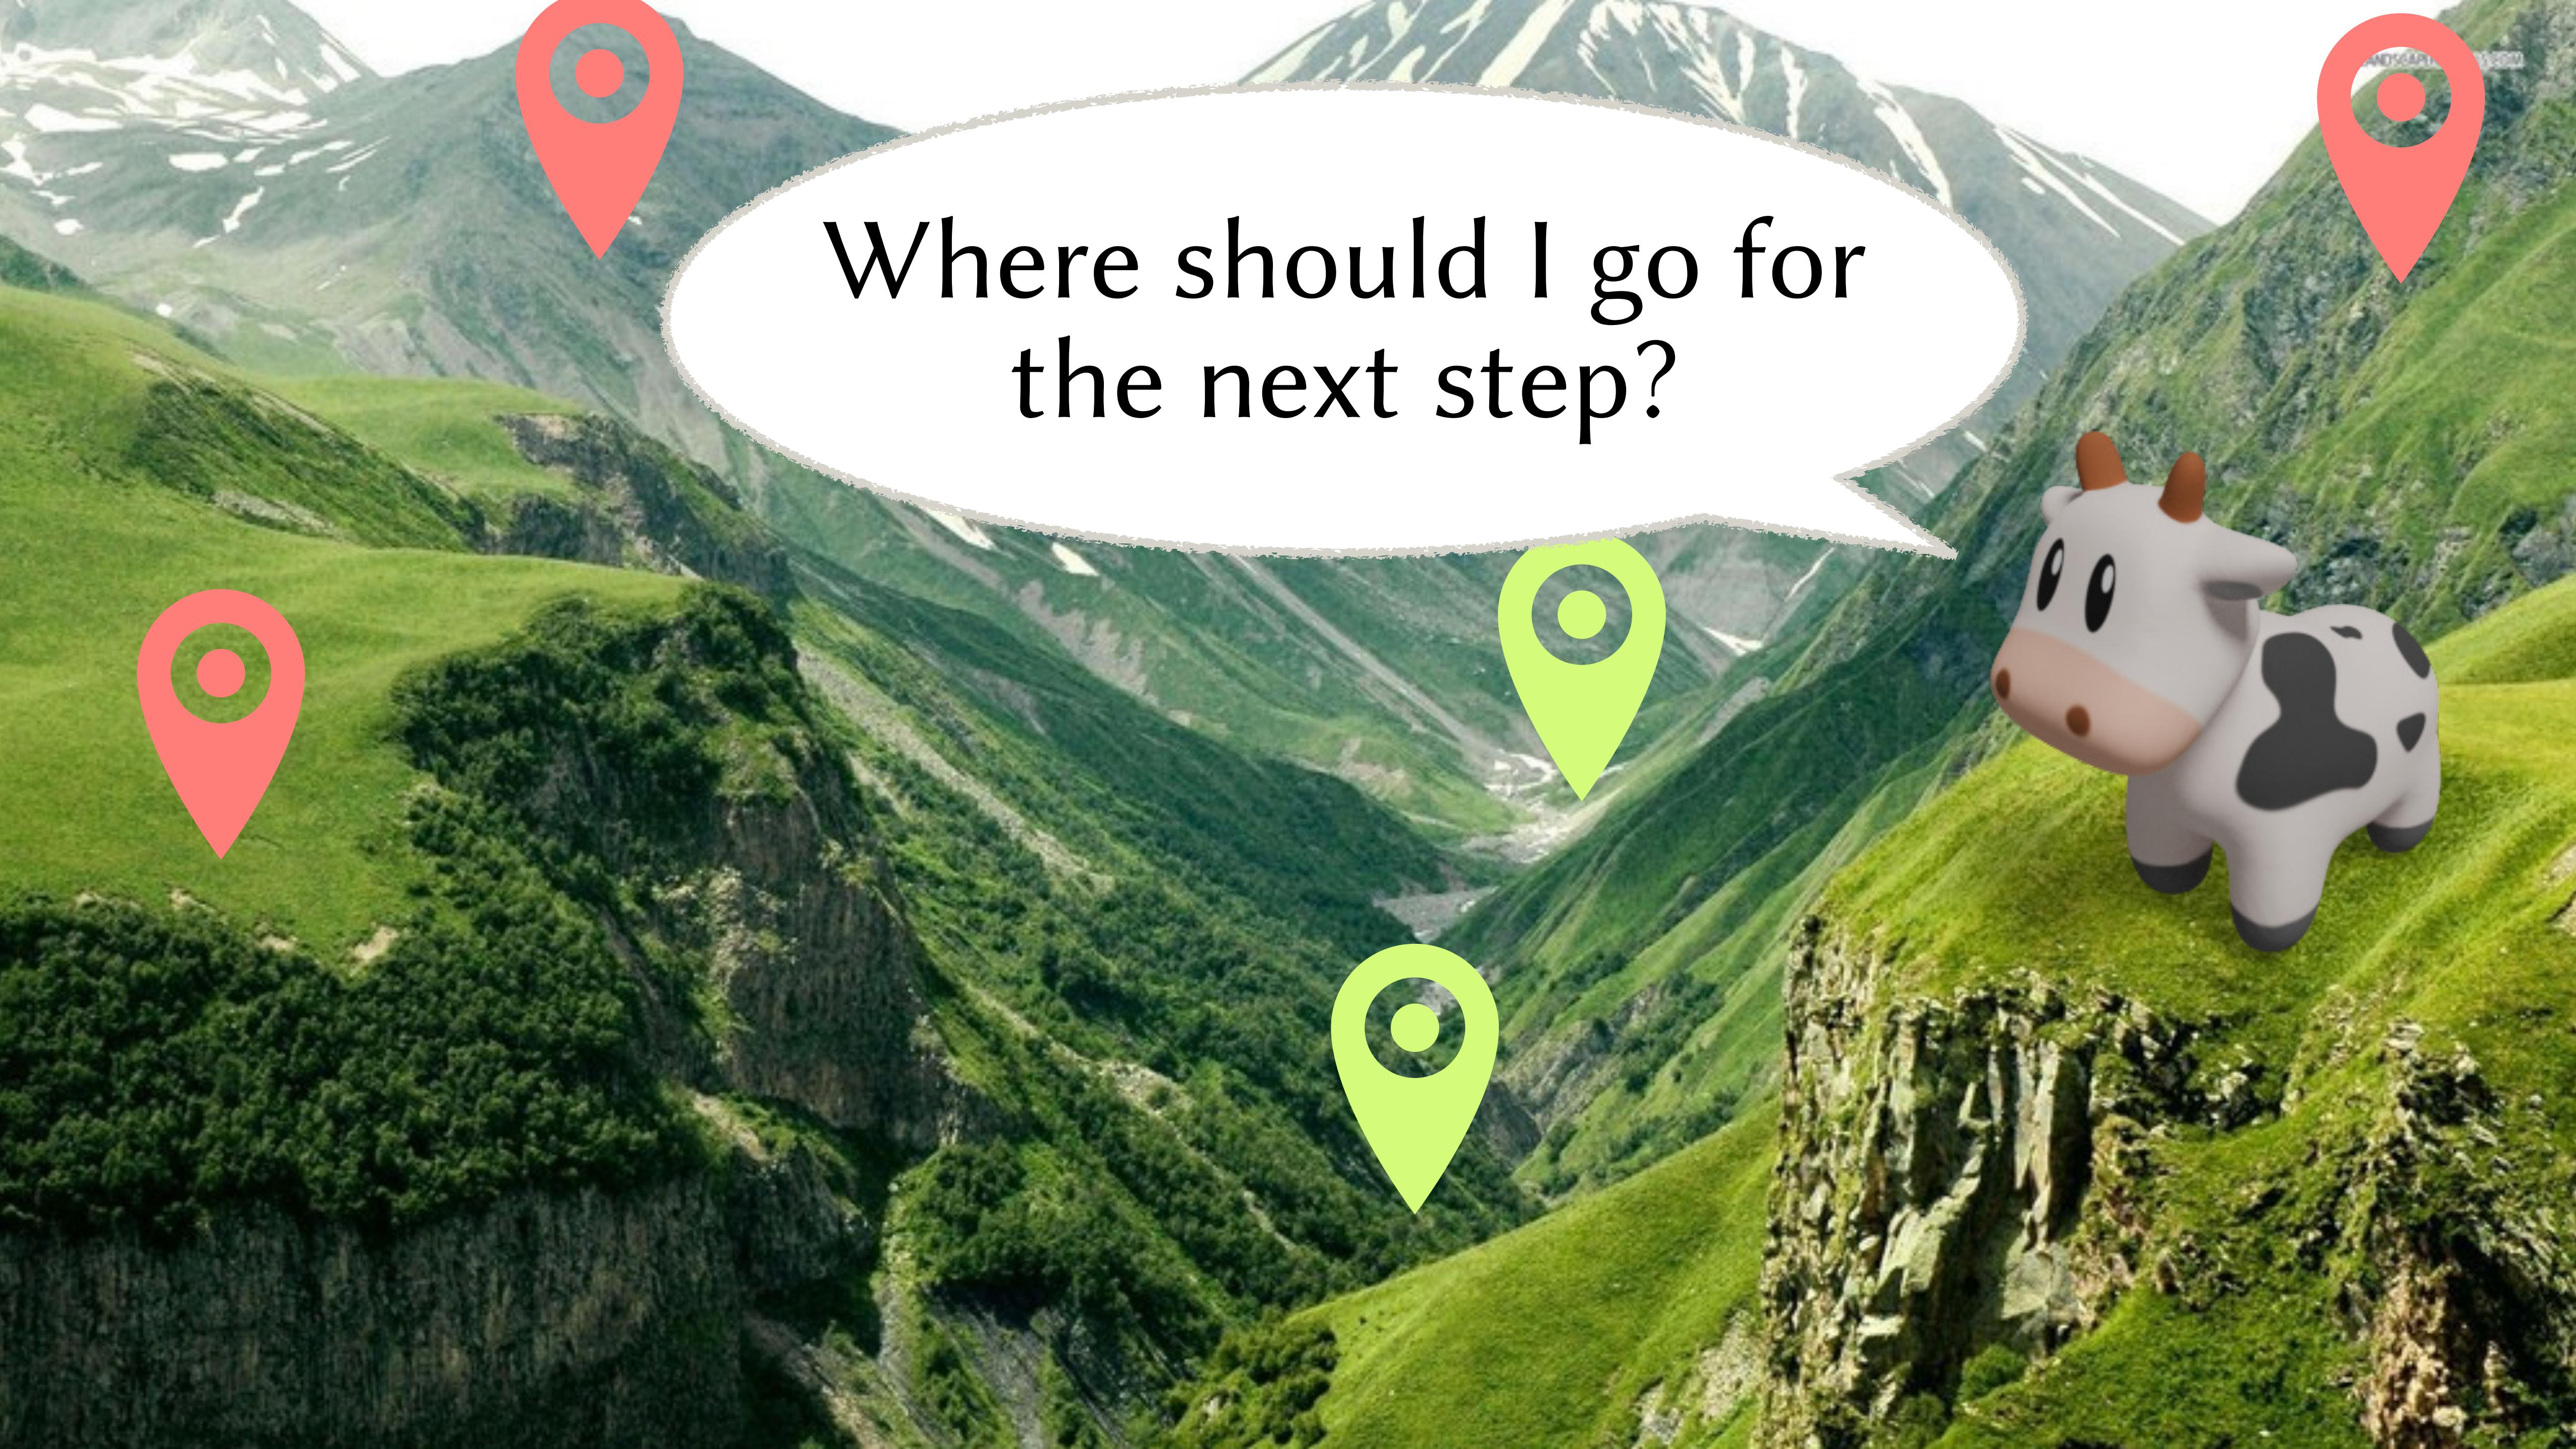
Spot

How can I go to the lowest point of  
the energy landscape?

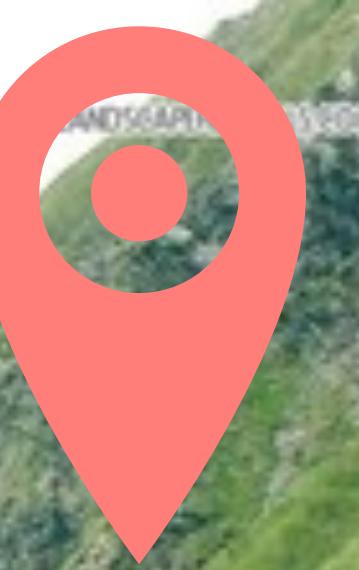
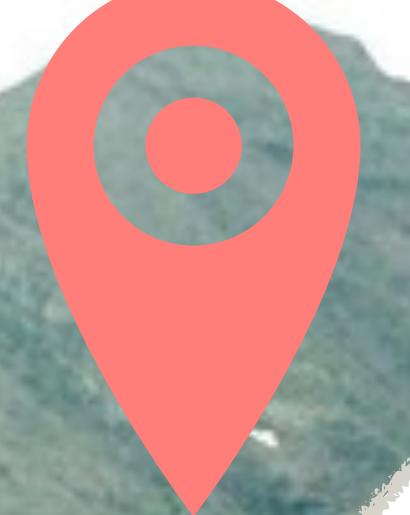
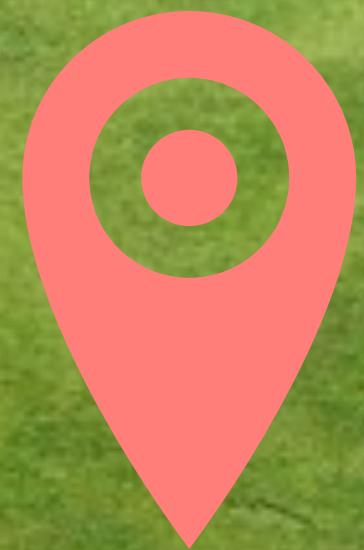


# Iterative Optimization





Where should I go for  
the next step?

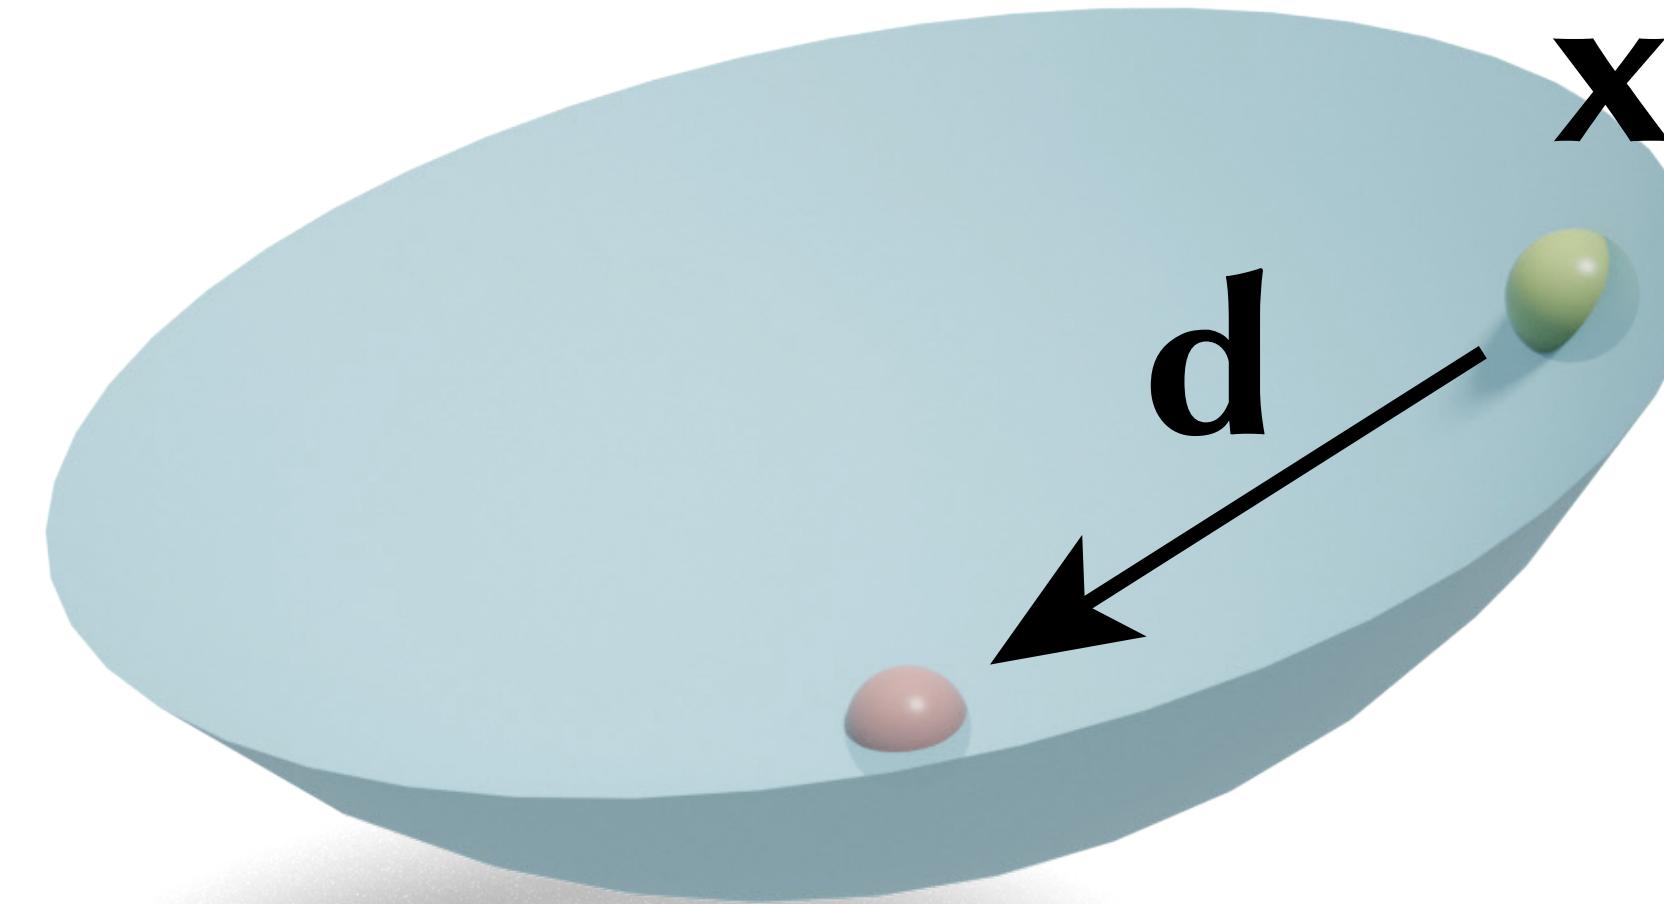


1. Which direction?
2. How far should I go?



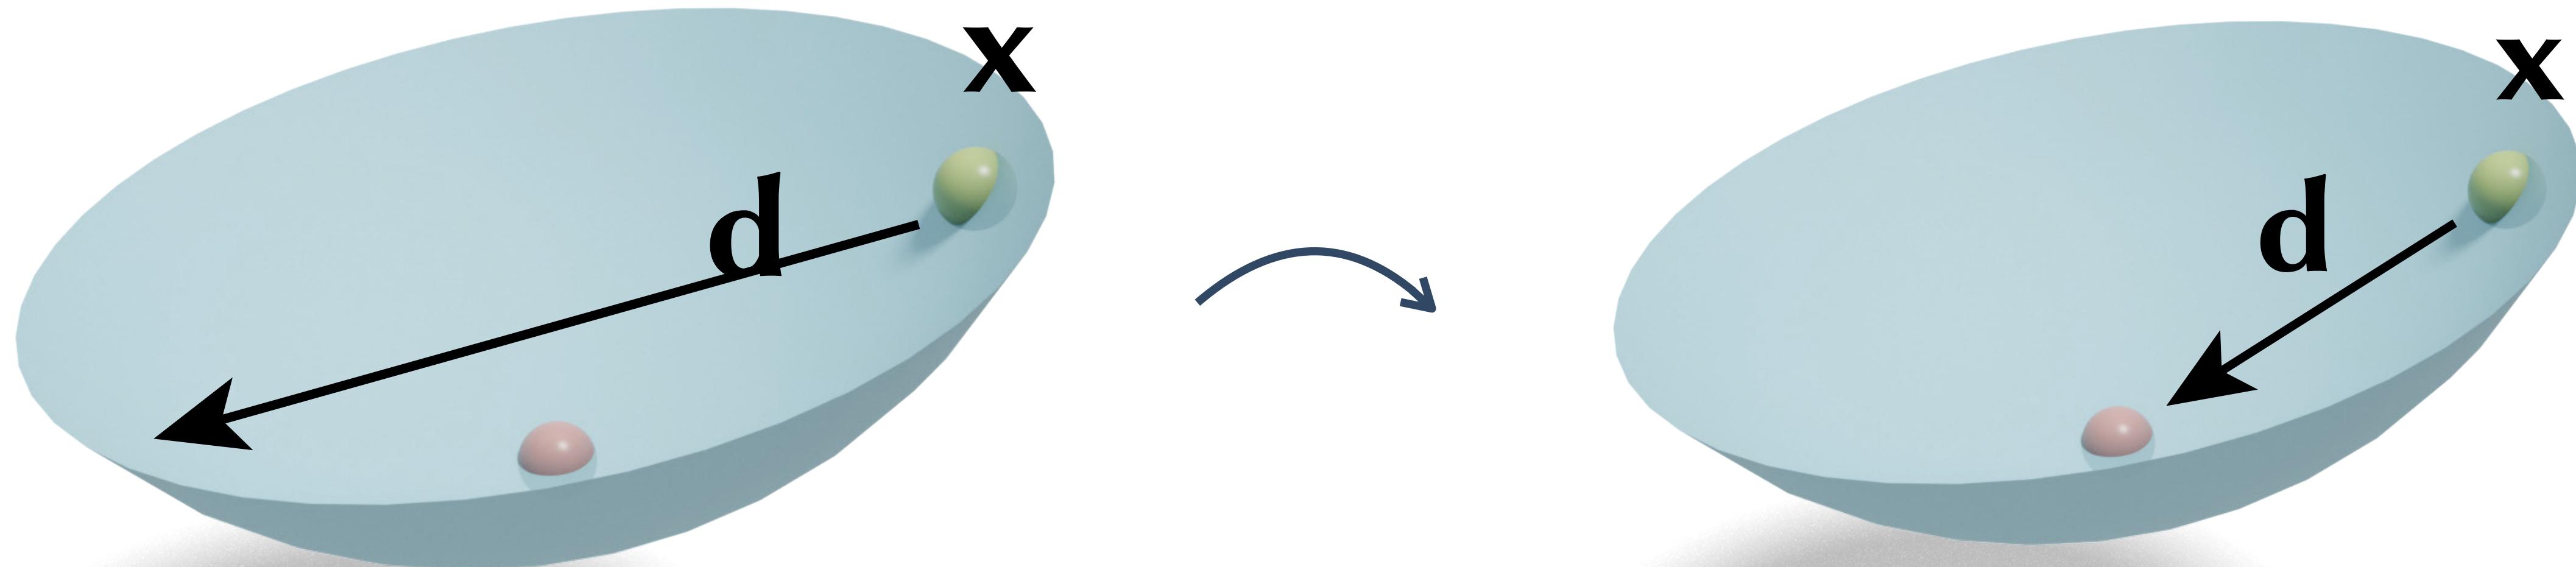
# Search Direction

A descent direction is a vector that points towards a local minimum of an objective function.

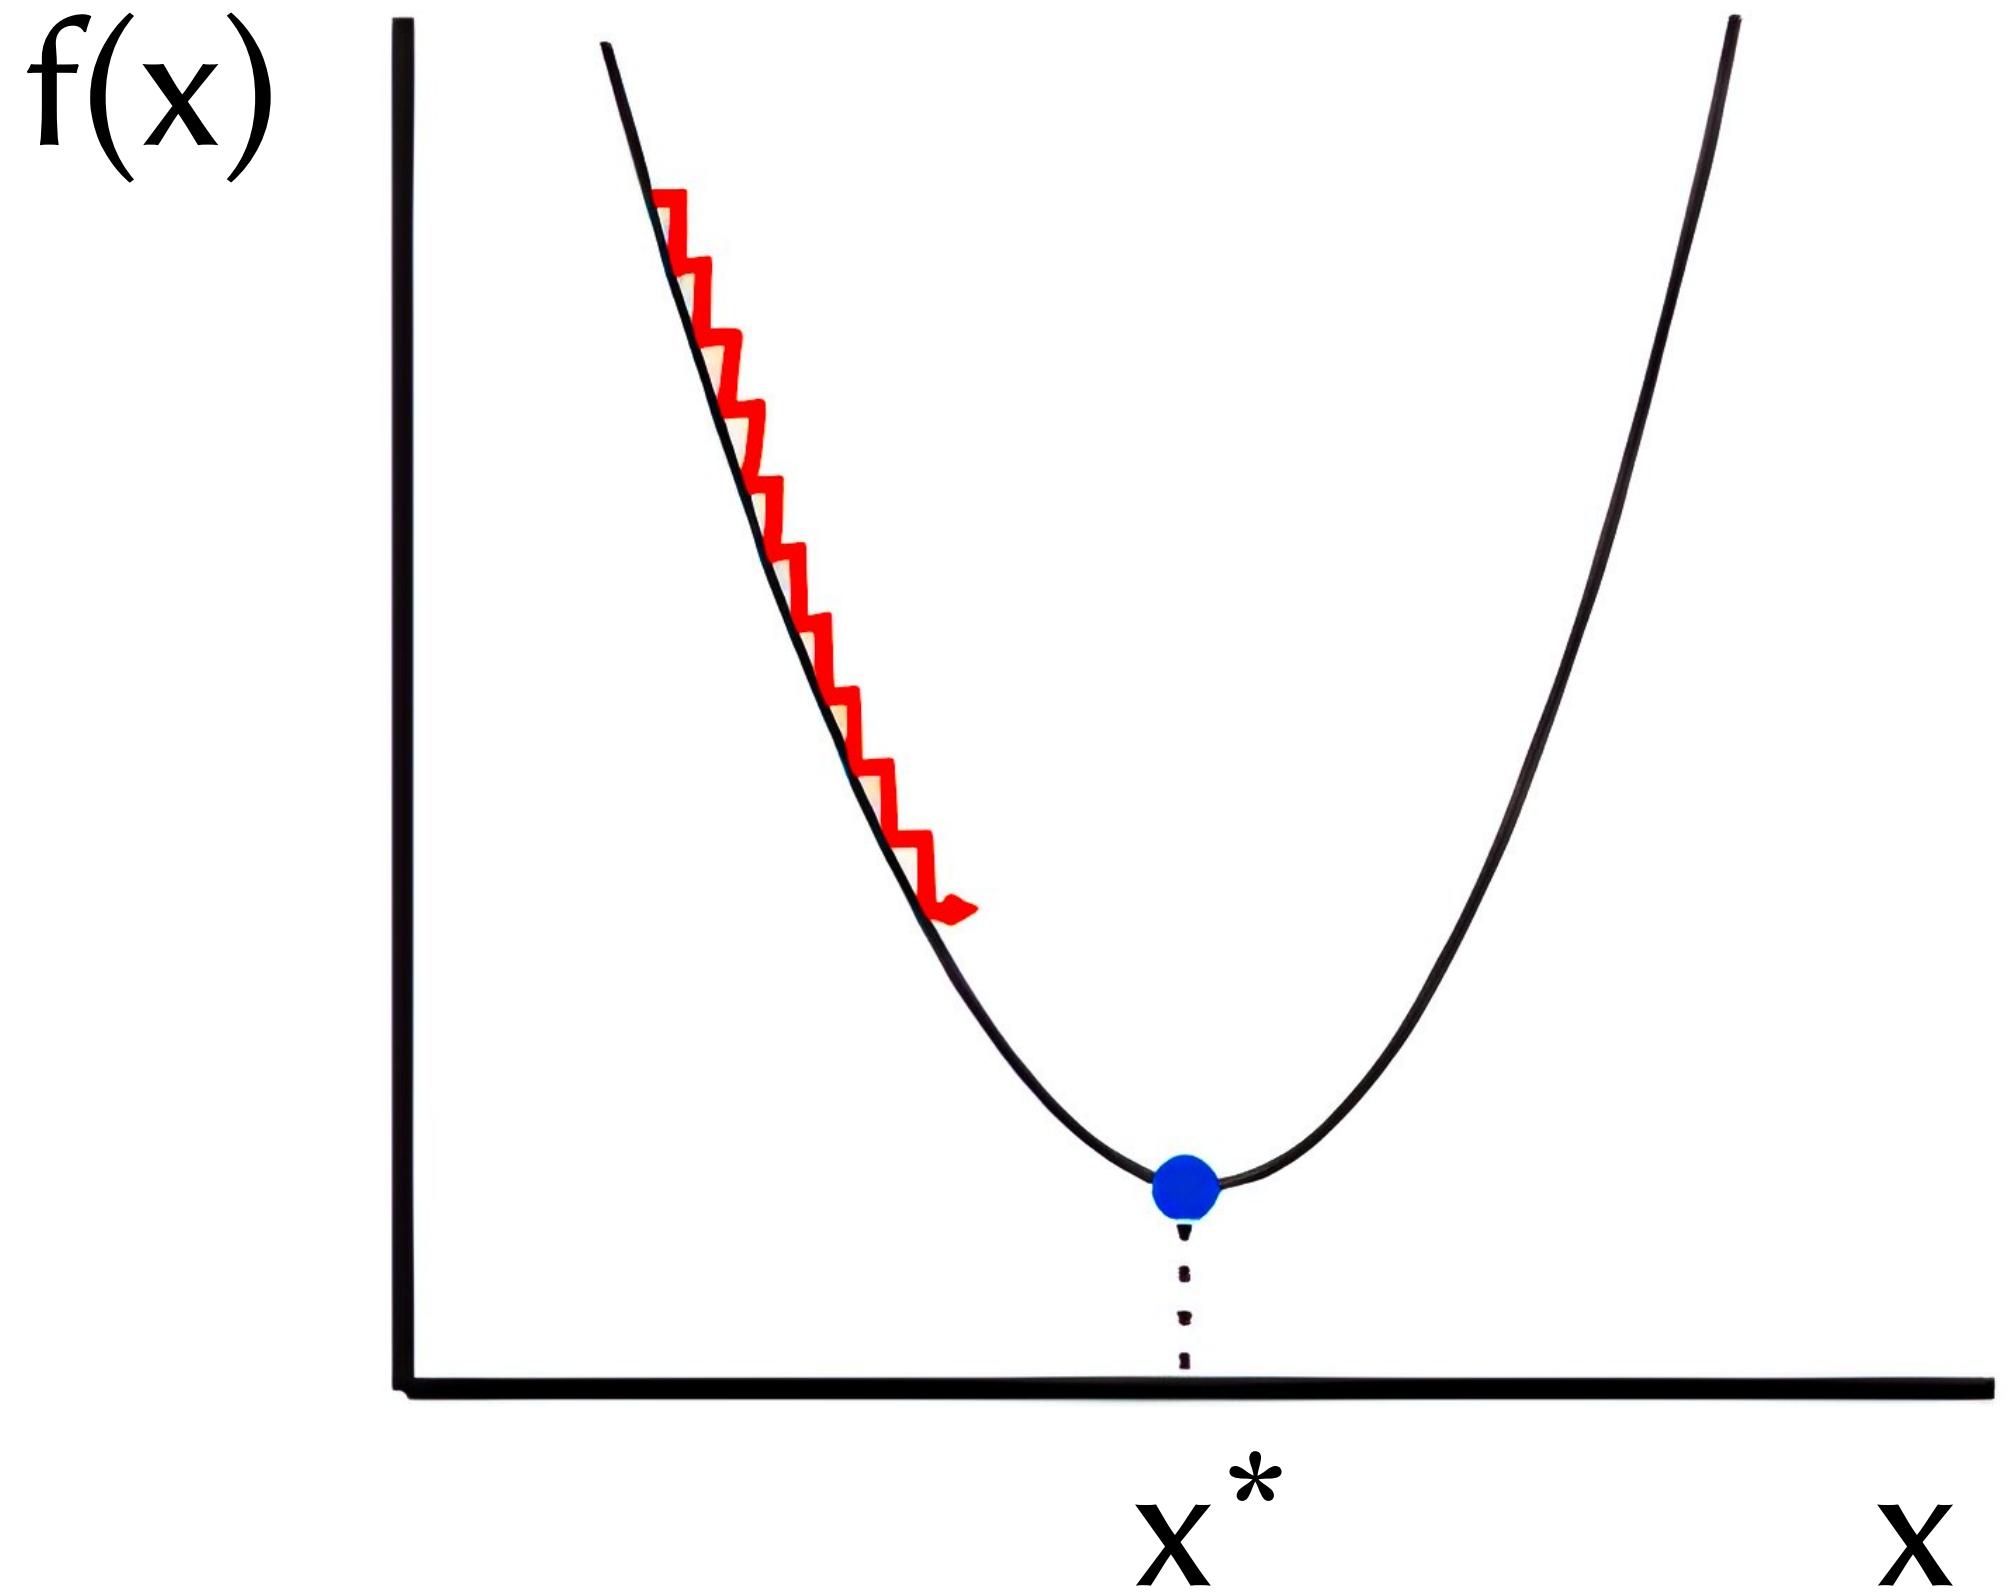


# Step Size

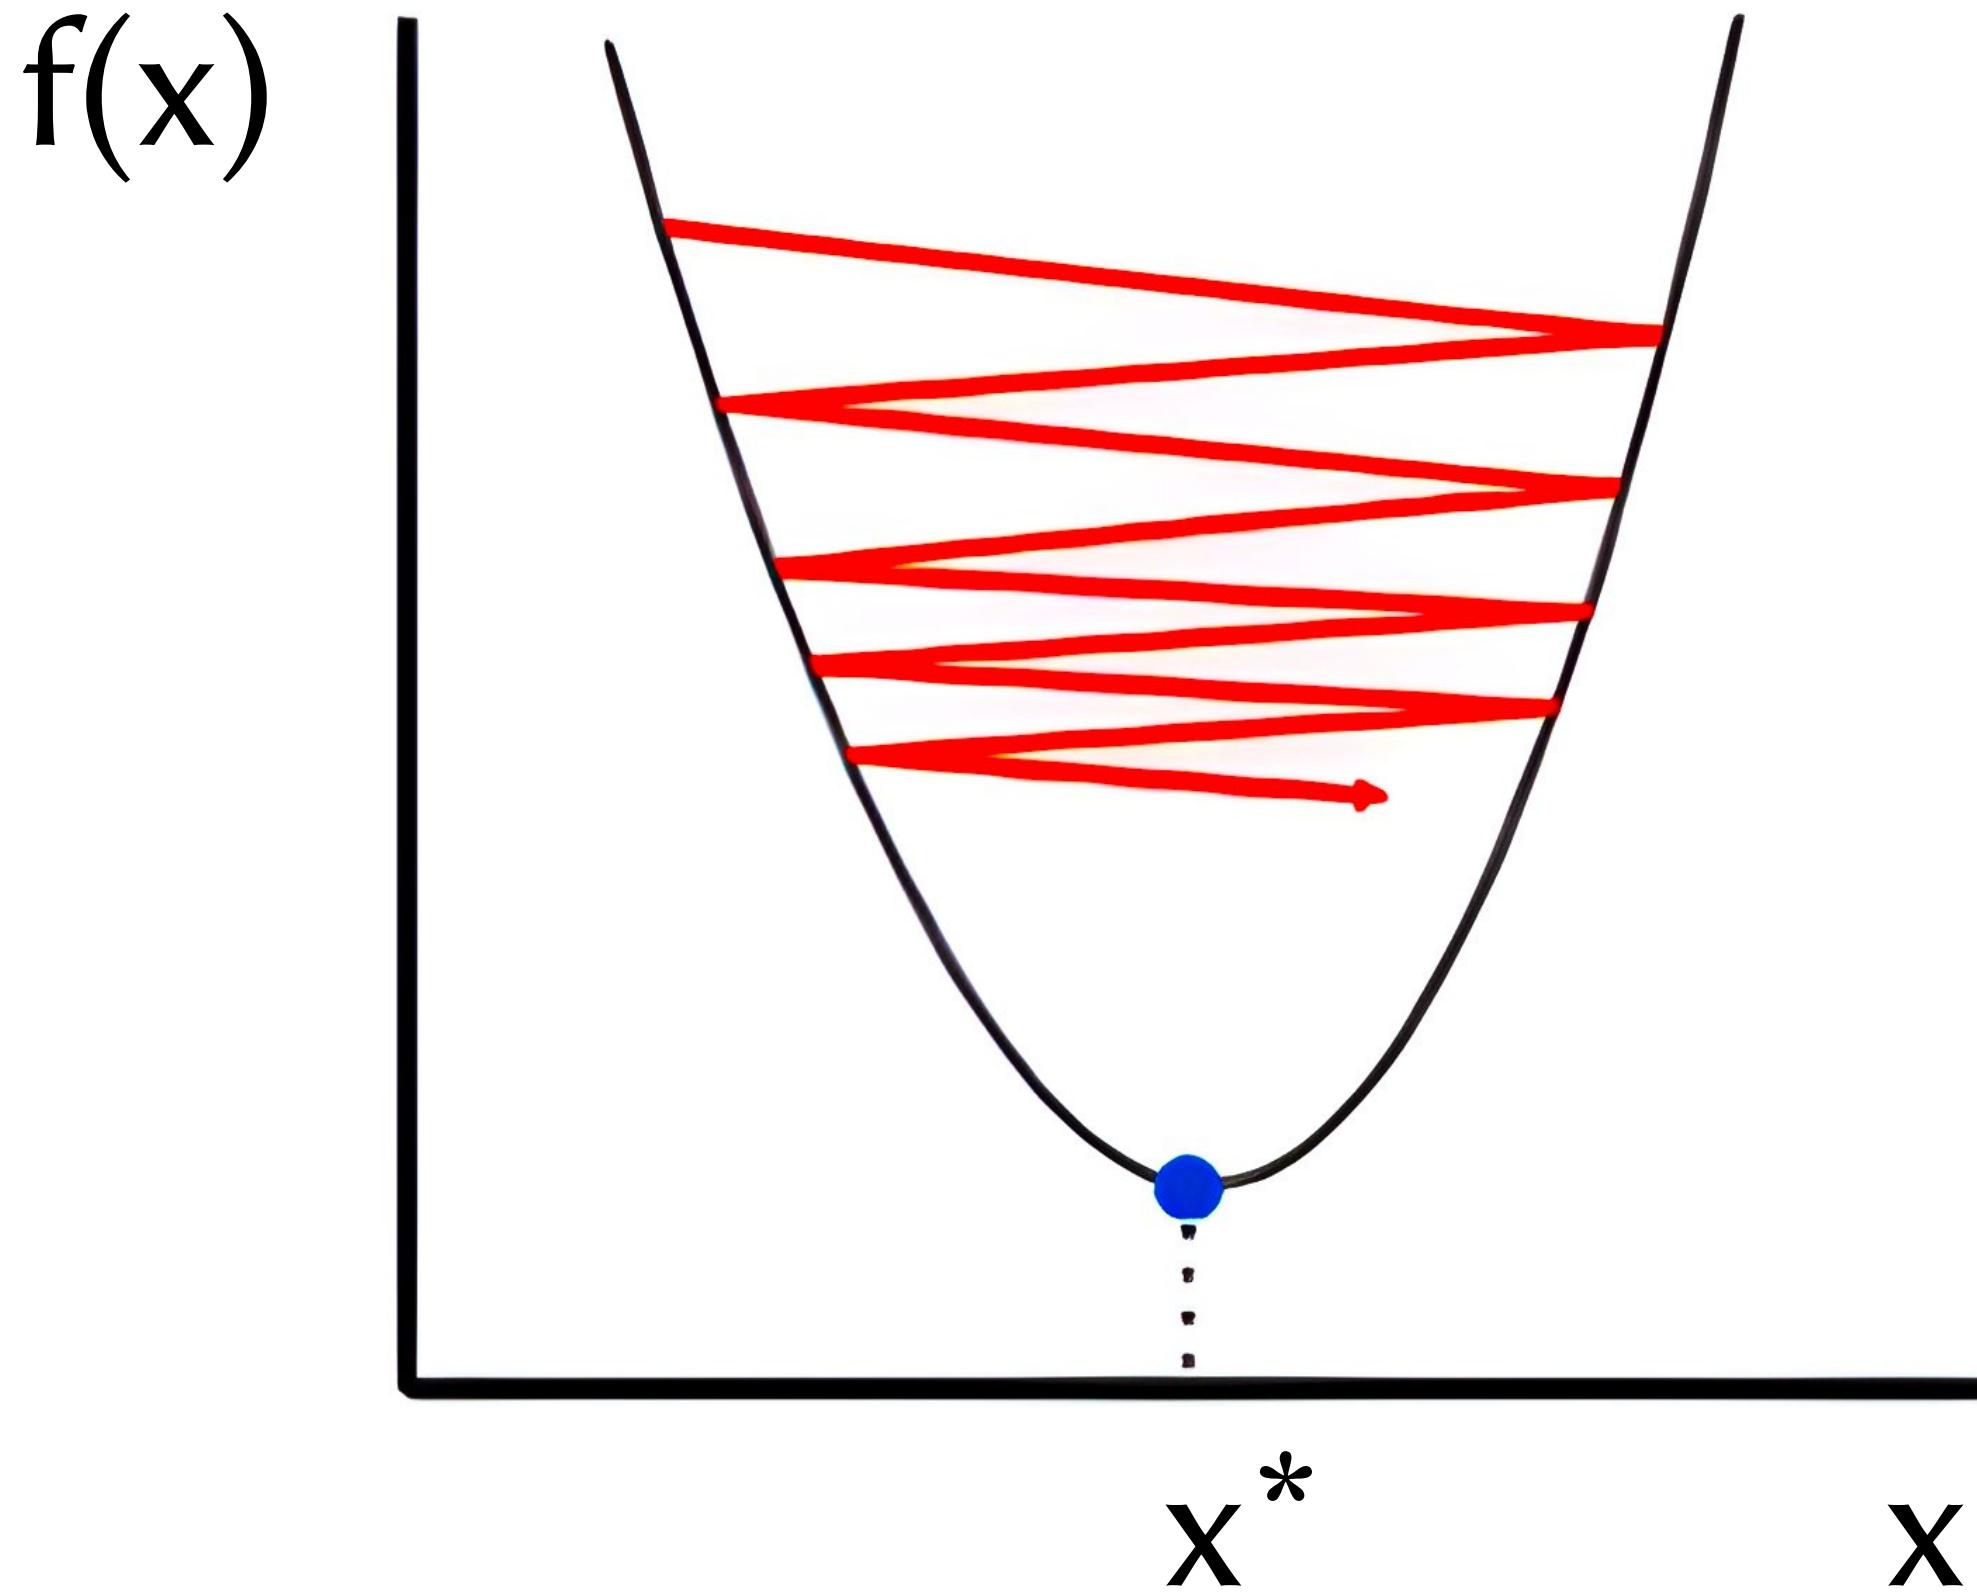
The step size determines how far we should move along the search direction.



# Picking the Right Step Size



Too small: converge  
vert slowly



Too big: overshoot  
and even diverge

# Backtracking Line Search

---

**Result:**  $\alpha$

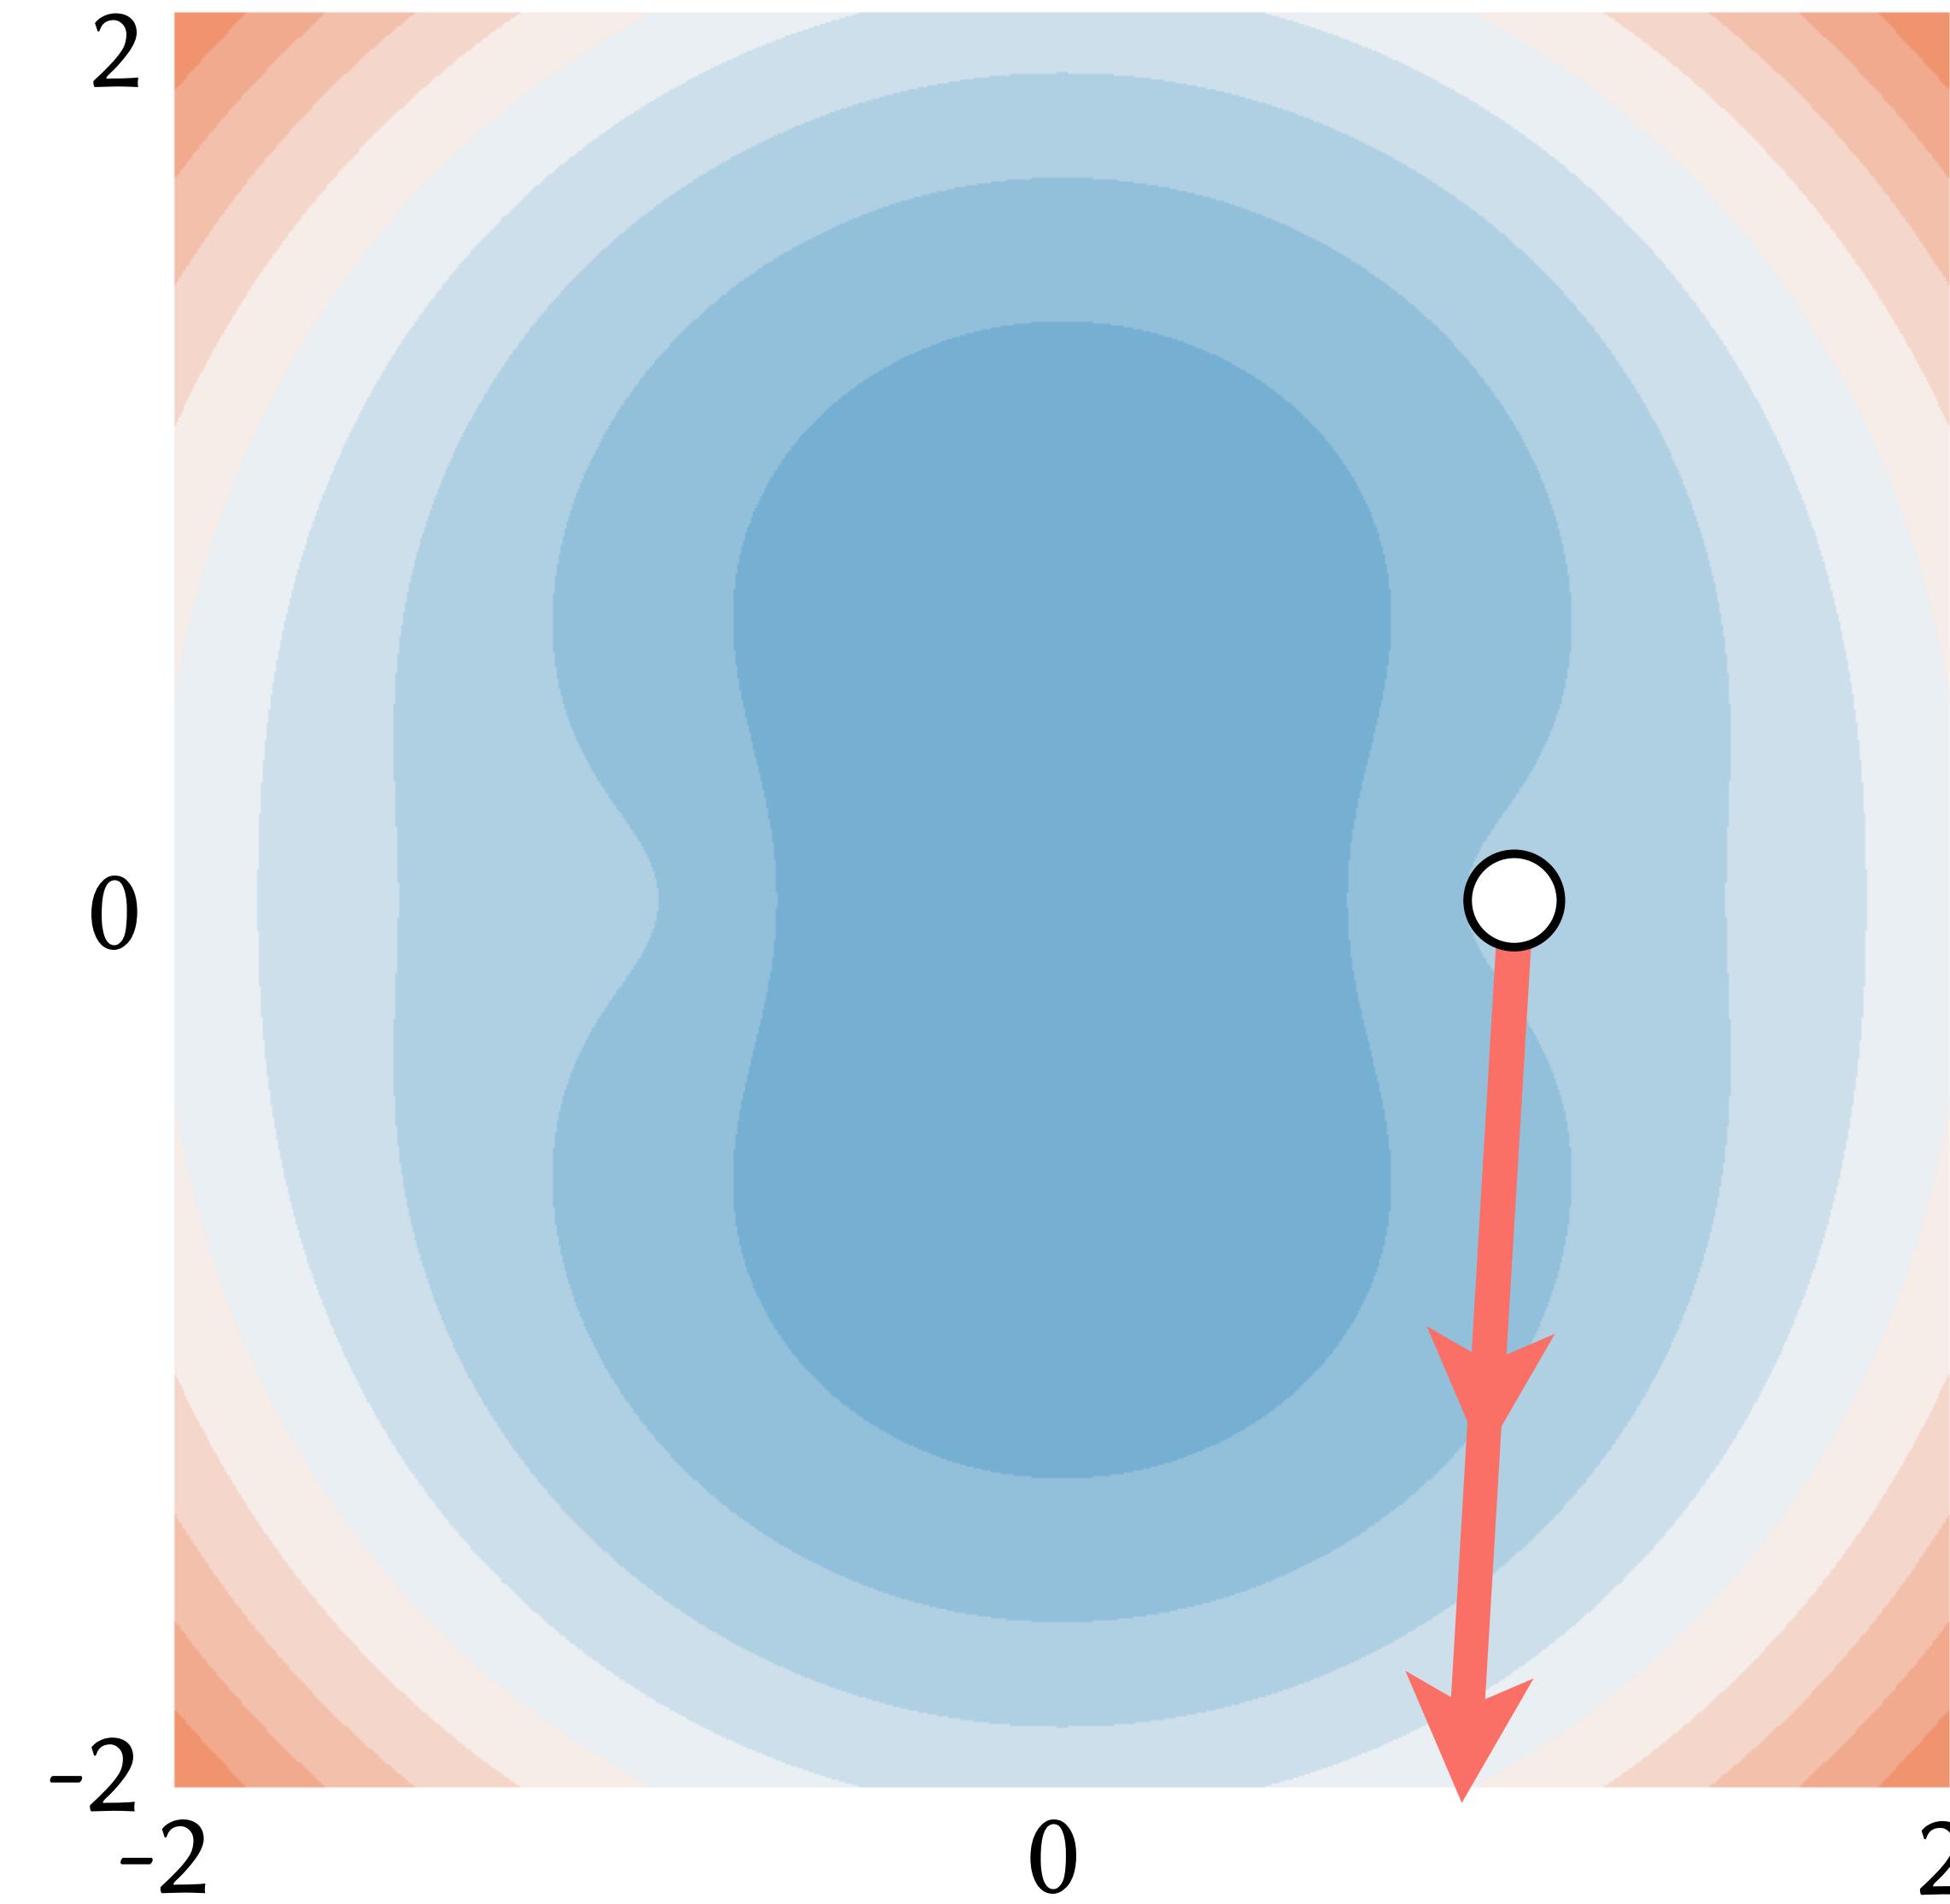
$\alpha \leftarrow 1;$

**while**  $E(x^i + \alpha p) > E(x^i)$  **do**

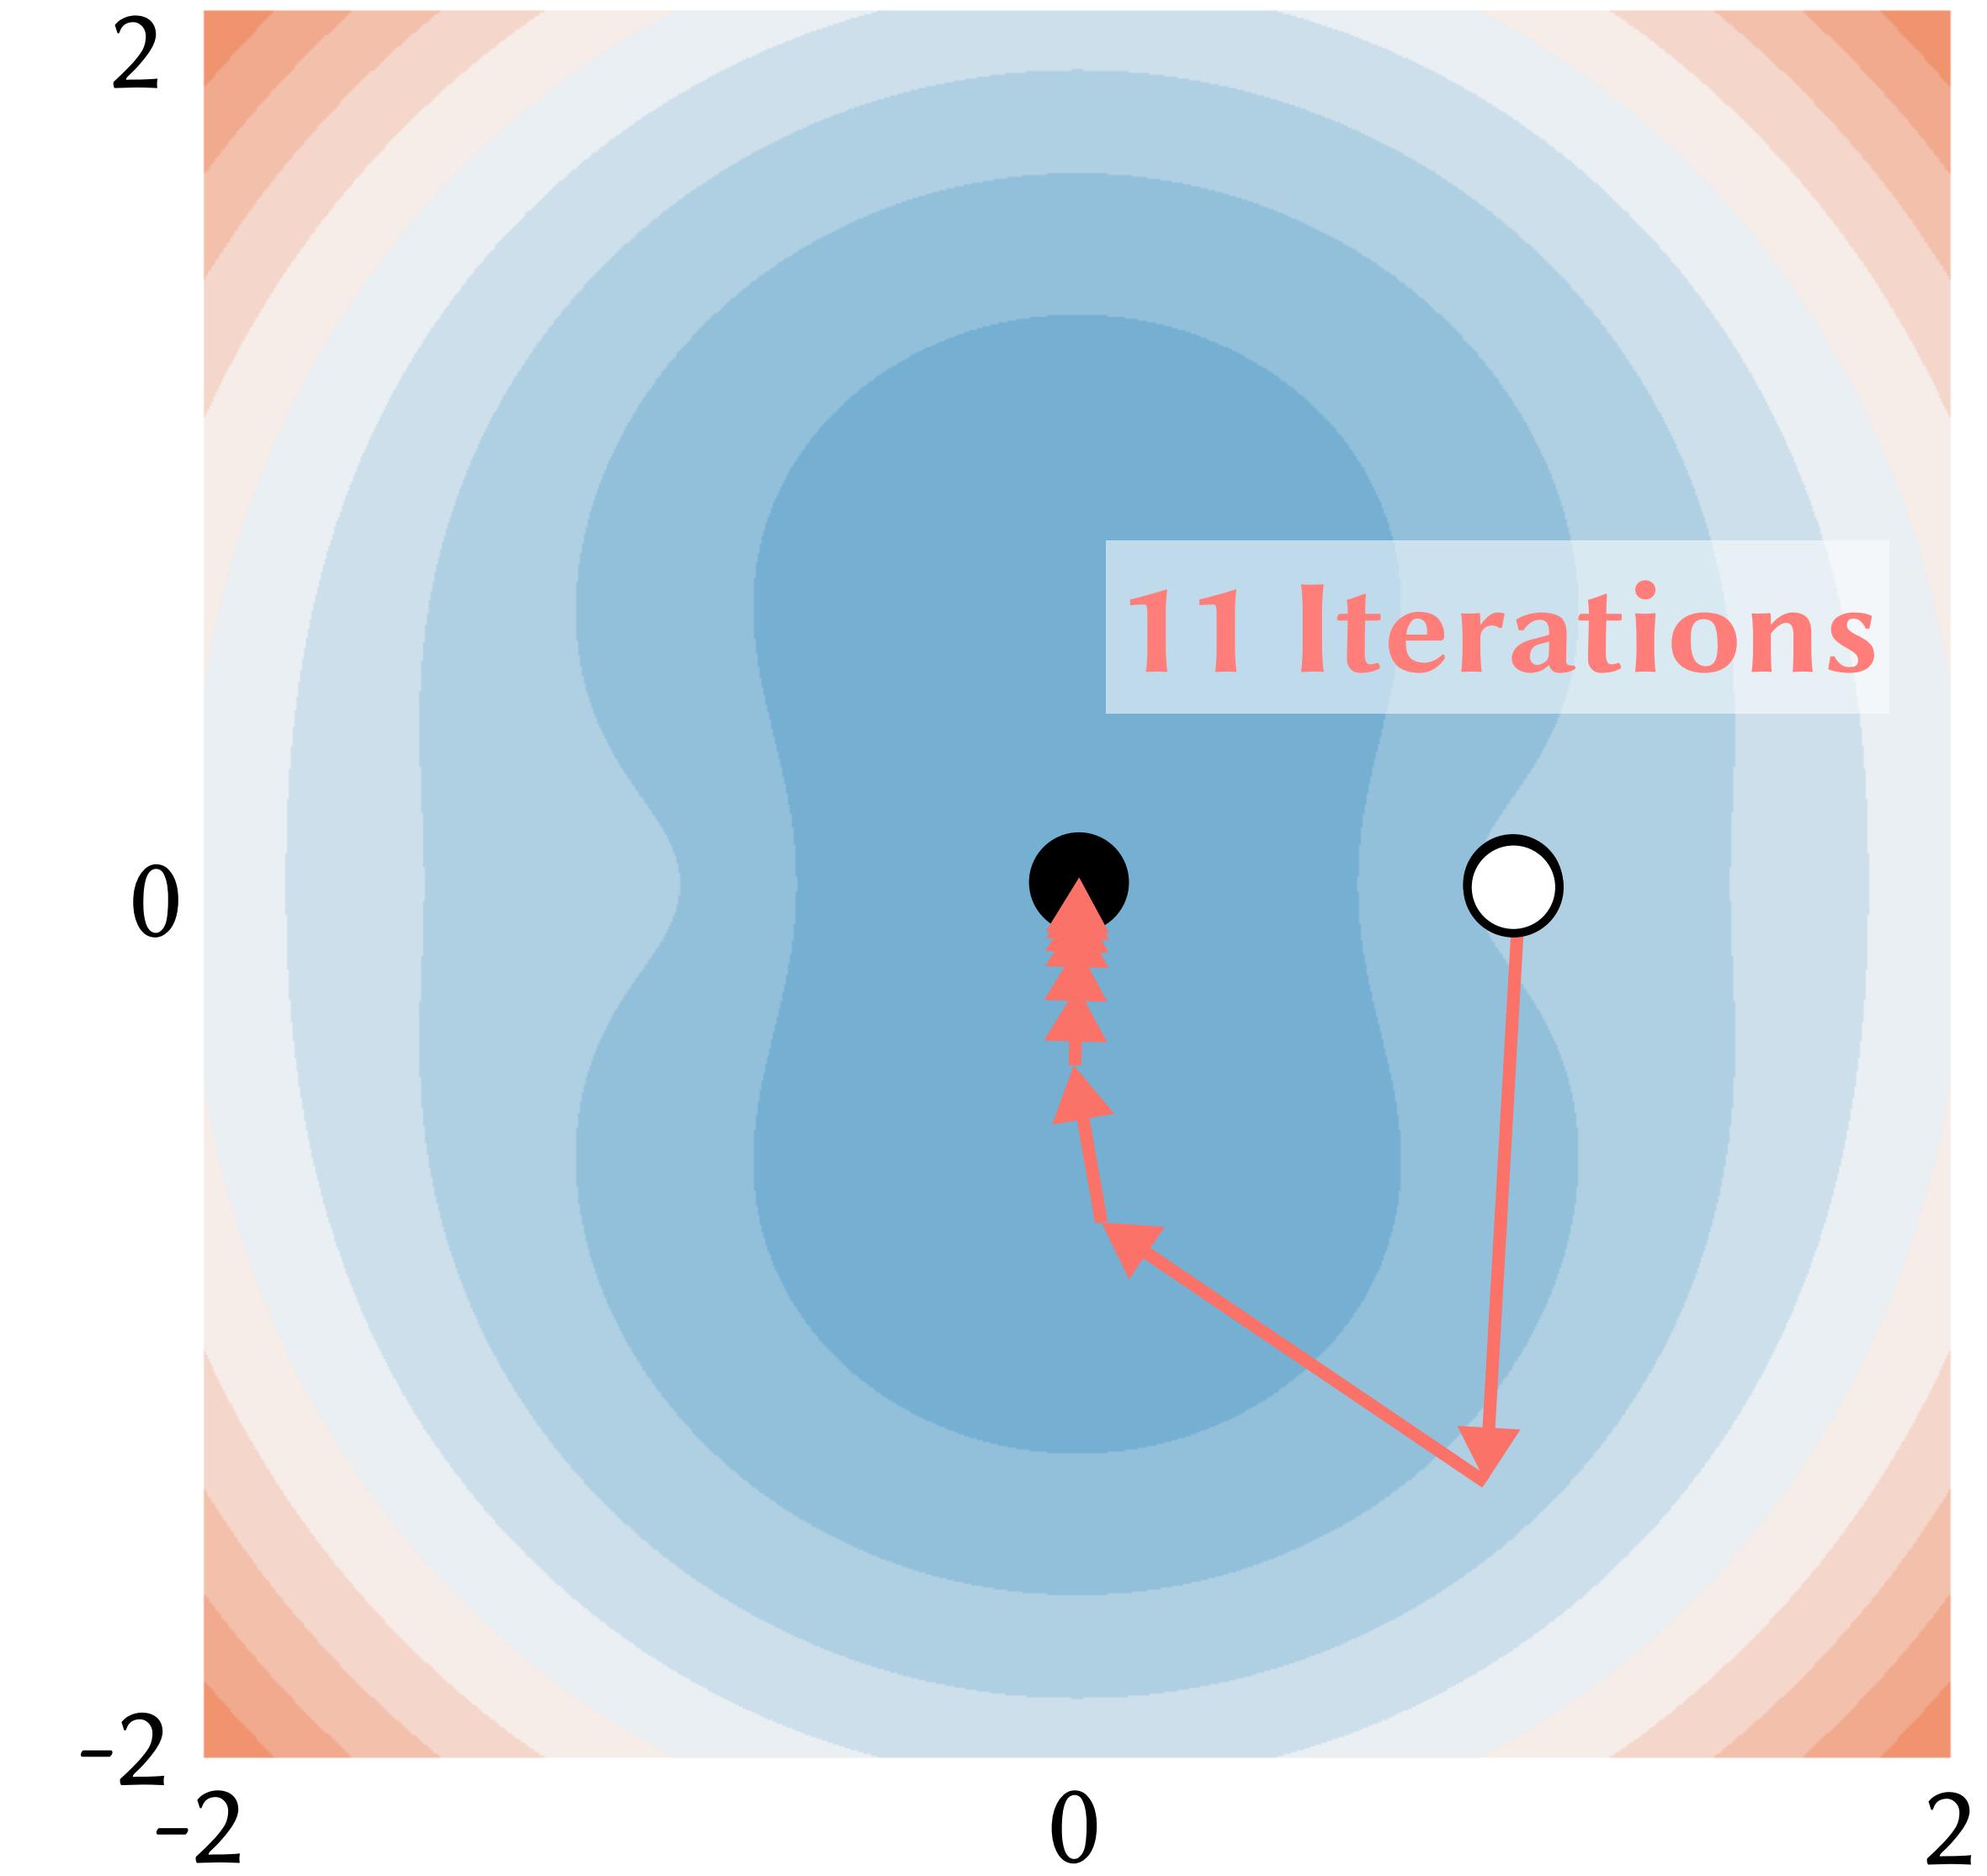
$\alpha \leftarrow \alpha/2;$

---

## Descent Direction



## Energy Minimization



- Elastic Energy
- Optimization Basics & Classical Algorithms
- Optimization Methods for Simulation
  - Challenge 1 ✨: Handling Nonconvexity
  - Challenge 2 ✨: Nontrivial Constraints
  - Challenge 3 ✨: Large-scale Optimization

# Classical Algorithms

---



# How can you minimize a function without knowing much about it?

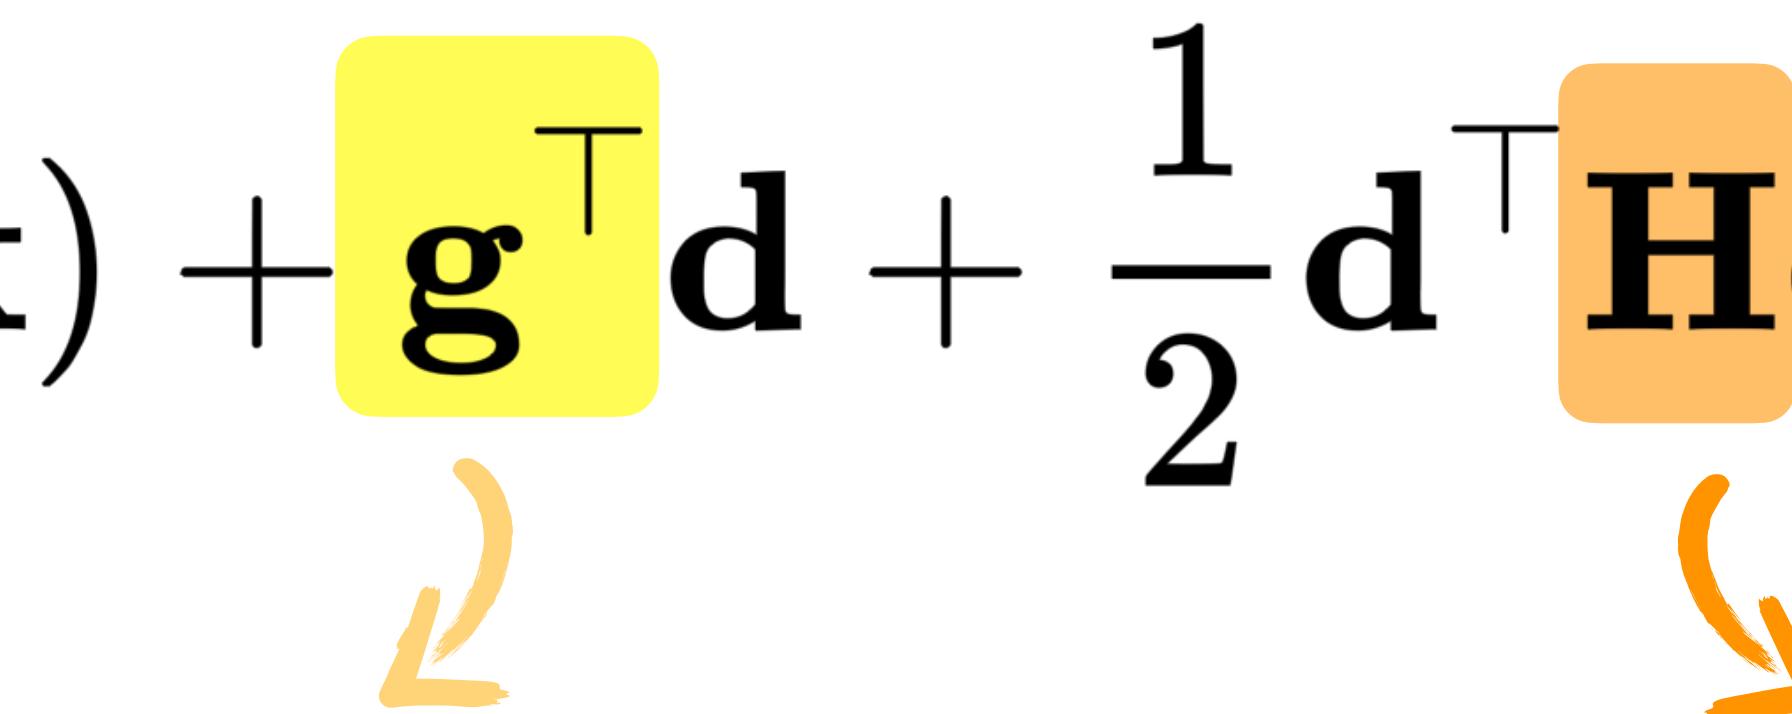
Assume it is much simpler  
than it really is



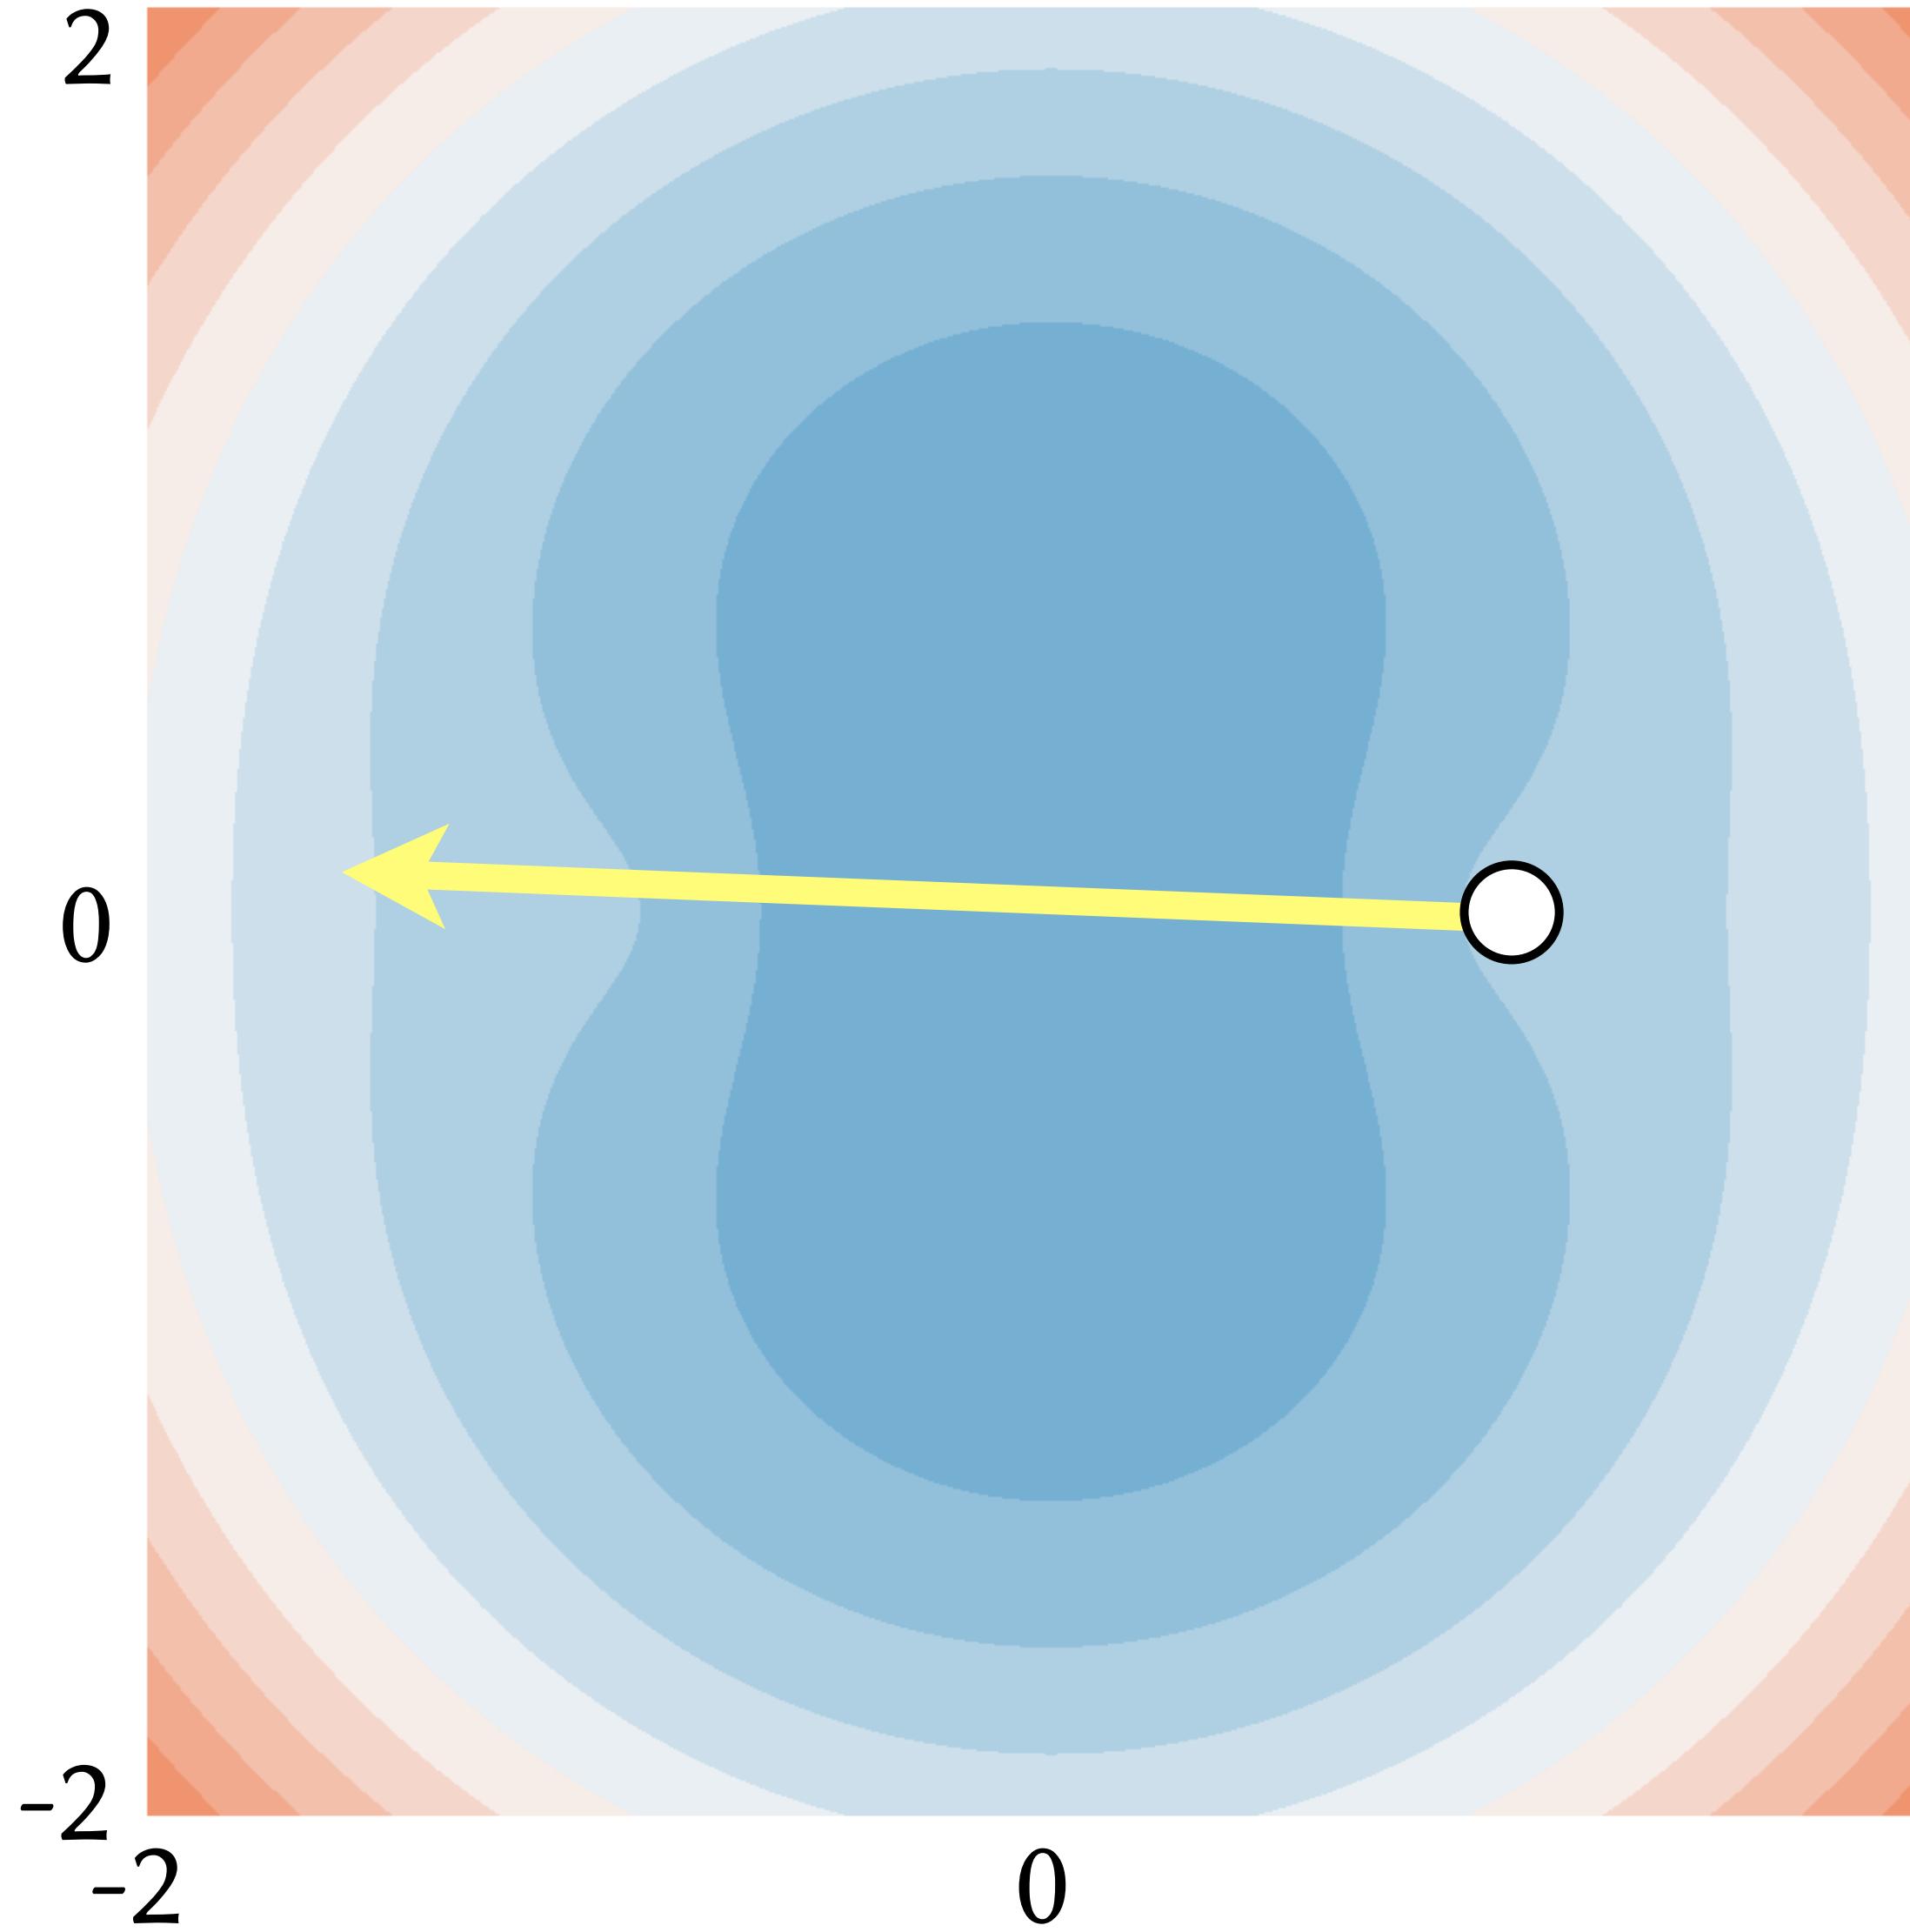
# Taylor Expansion for Approximation

If  $f$  is continuously twice differentiable, then

$$f(\mathbf{x} + \mathbf{d}) \approx f(\mathbf{x}) + \mathbf{g}^\top \mathbf{d} + \frac{1}{2} \mathbf{d}^\top \mathbf{H} \mathbf{d} + O(\|\mathbf{d}\|^3)$$

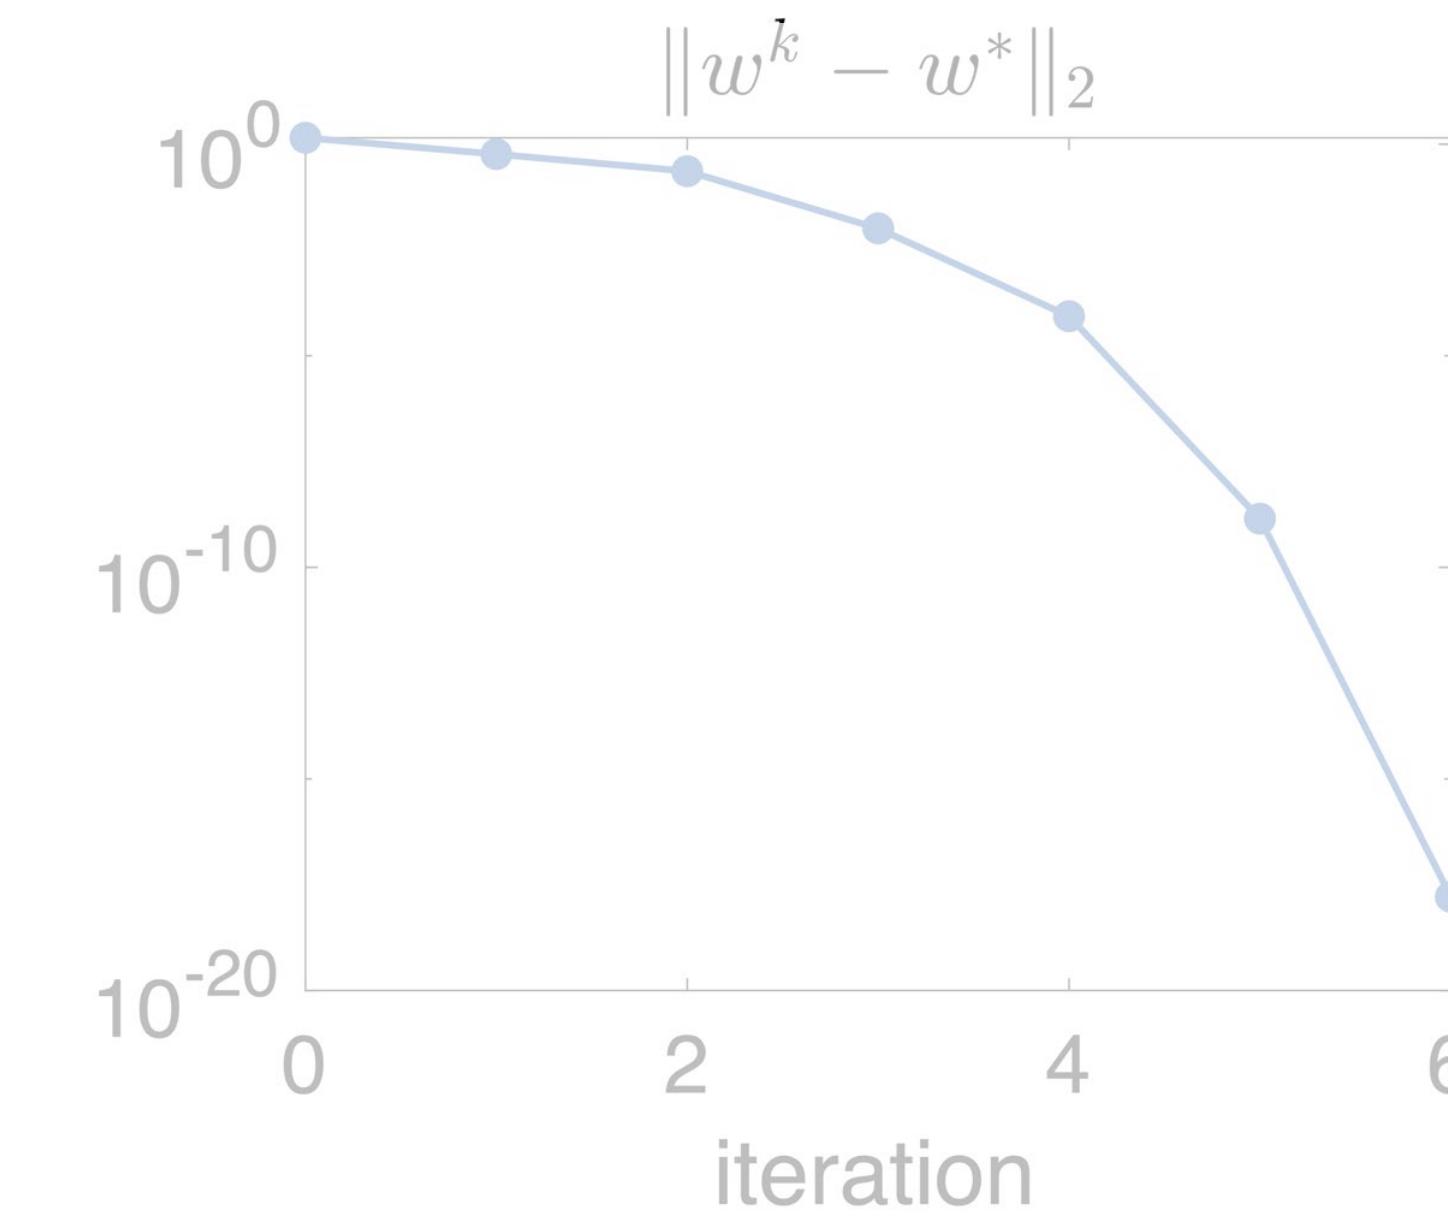
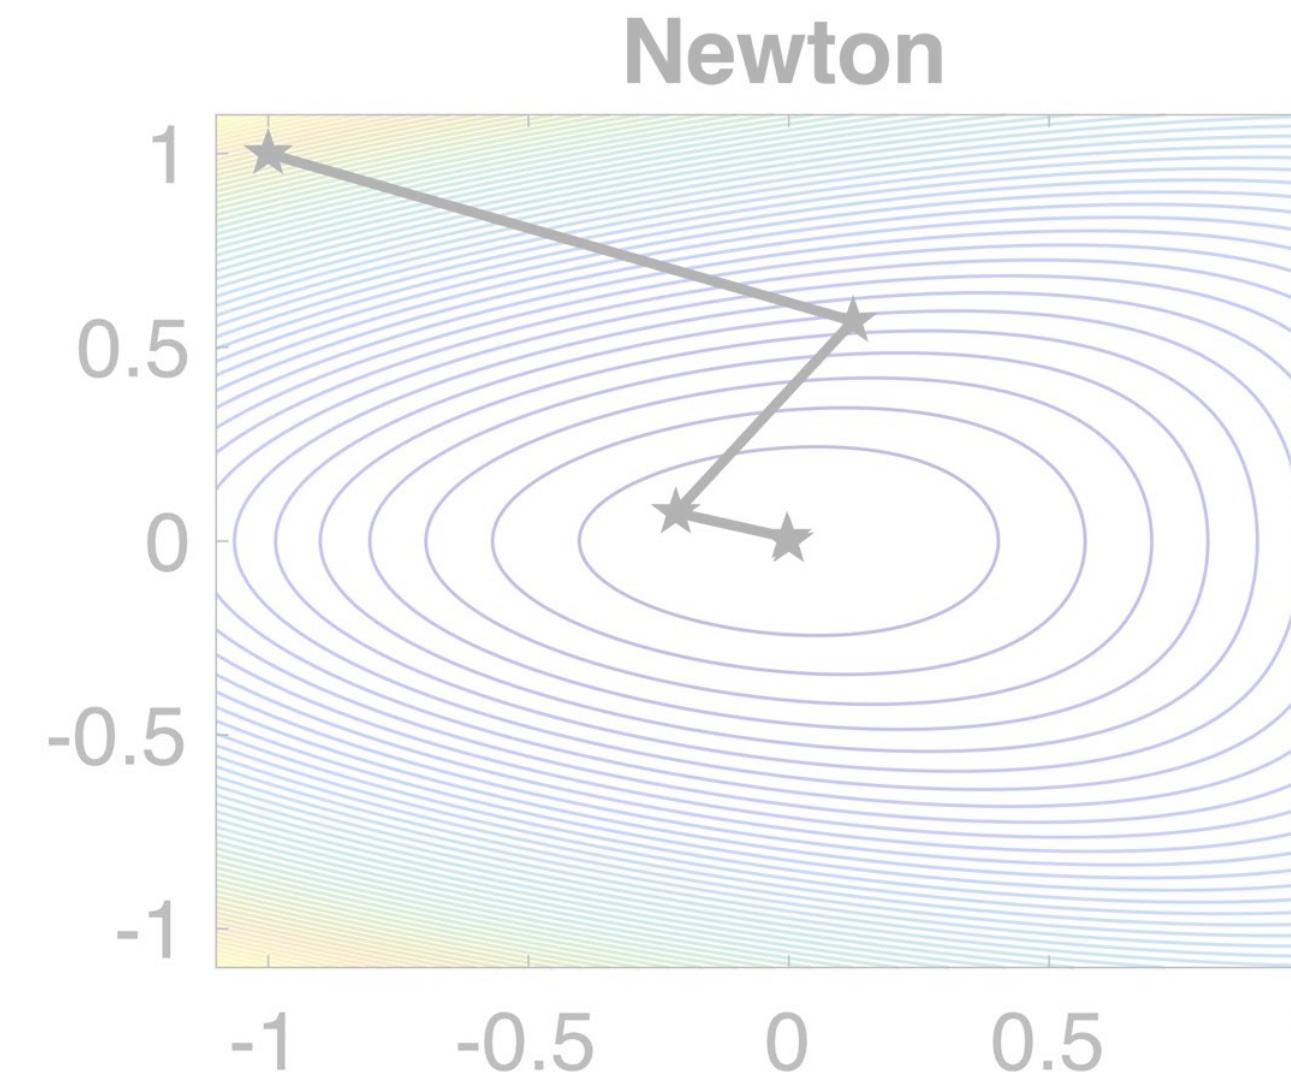
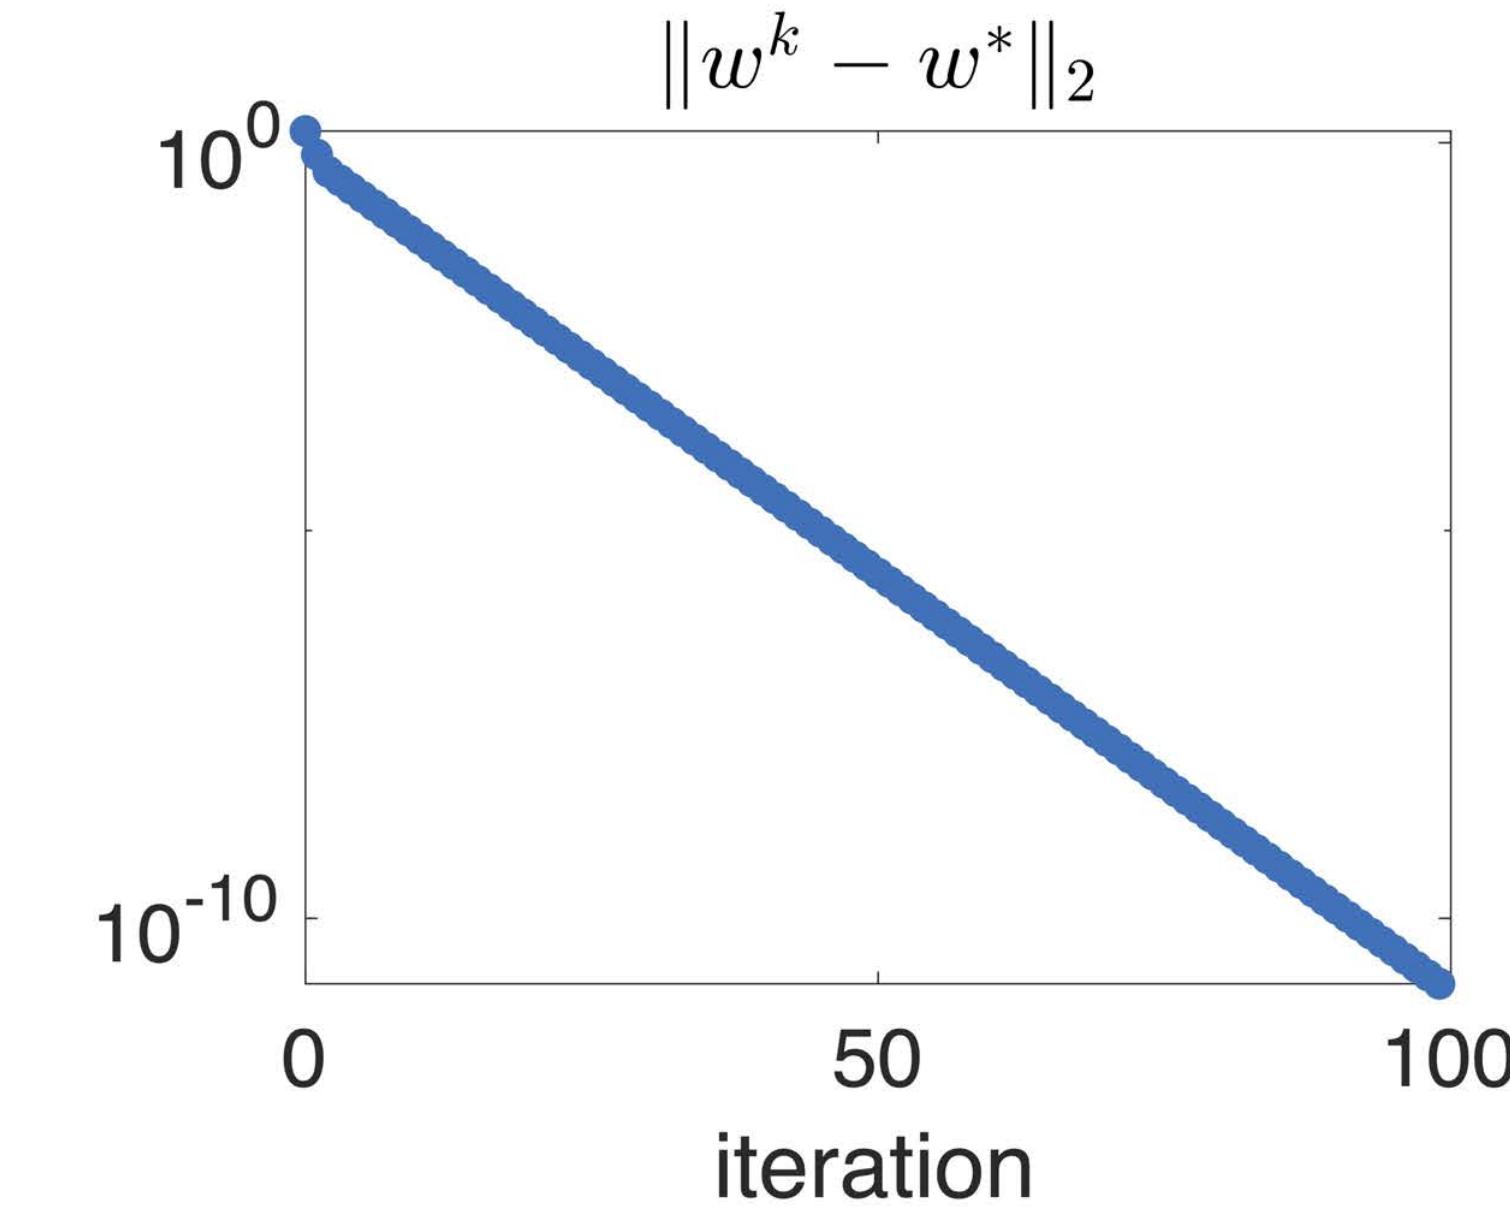
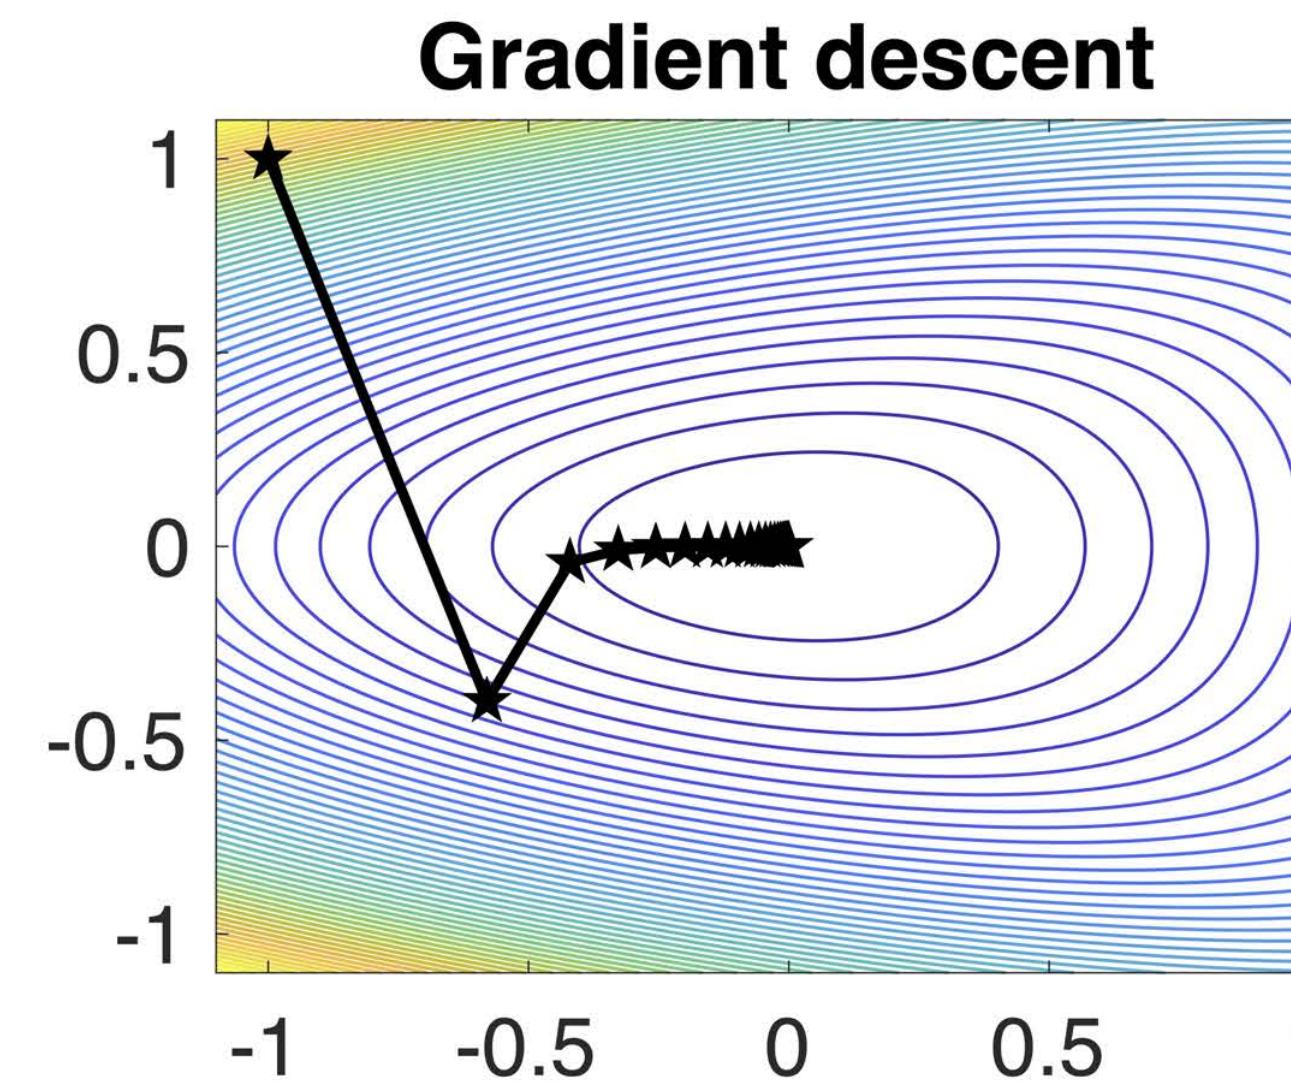
  
Gradient      Hessian

# Gradient Descent: Use First-order Information



$$\mathbf{d} = -\mathbf{g}$$

# Gradient Descent Can be Slow

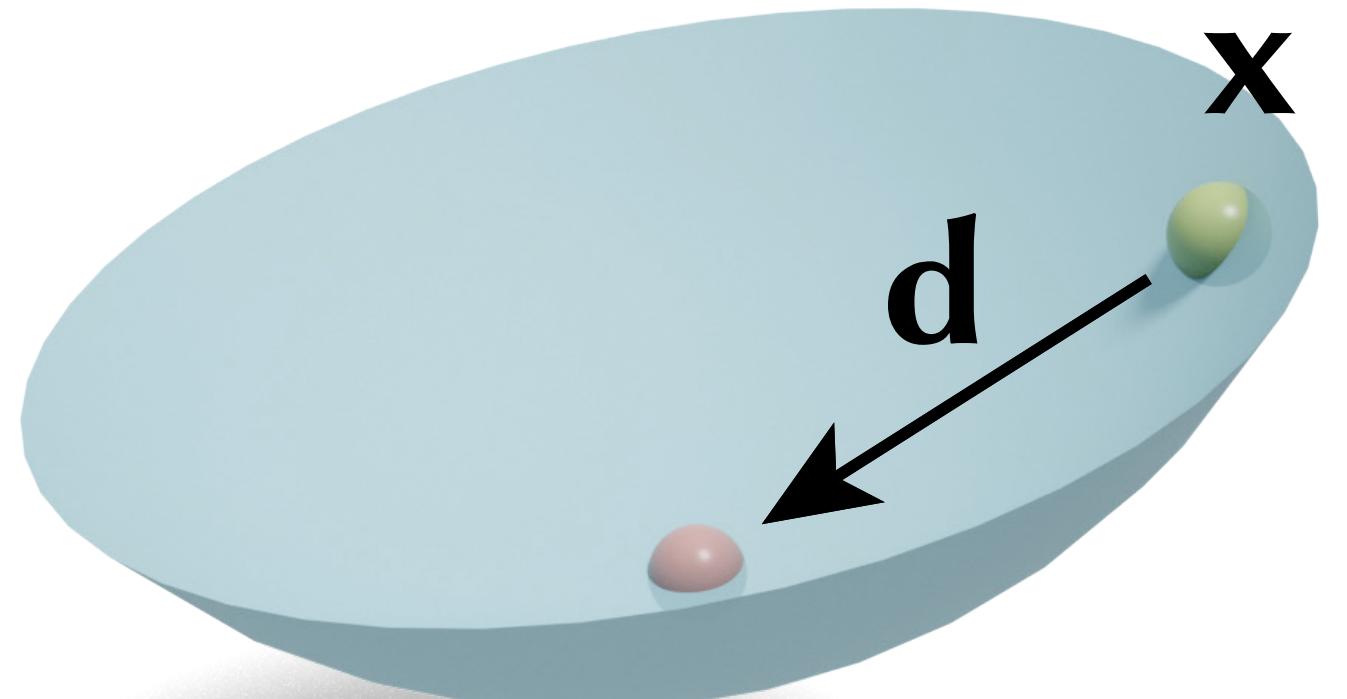


# Newton's Method: Use Second-order Information

“Pretend” the function is quadratic:

$$\min_{\mathbf{d}} f(\mathbf{x}) + \mathbf{g}^\top \mathbf{d} + \frac{1}{2} \mathbf{d}^\top \mathbf{H} \mathbf{d}$$

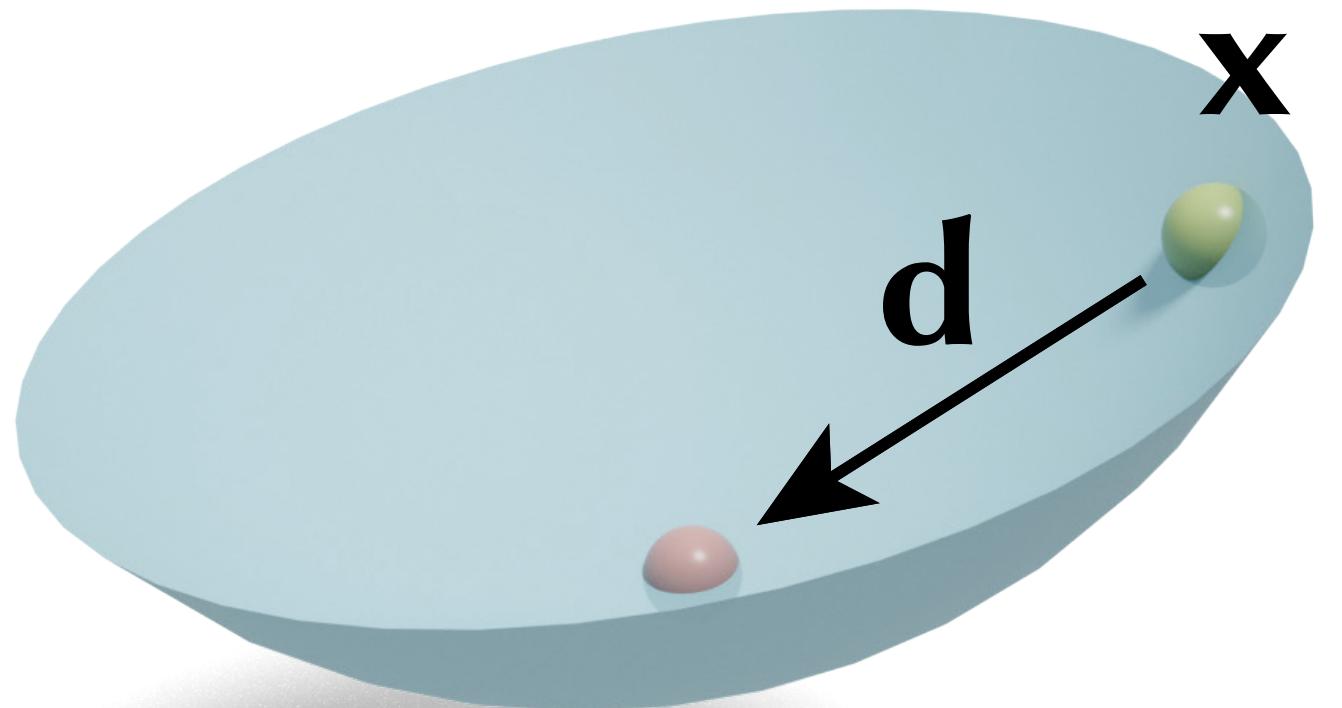
Set the gradient w.r.t.  $\mathbf{d}$  to be 0



# Newton's Method

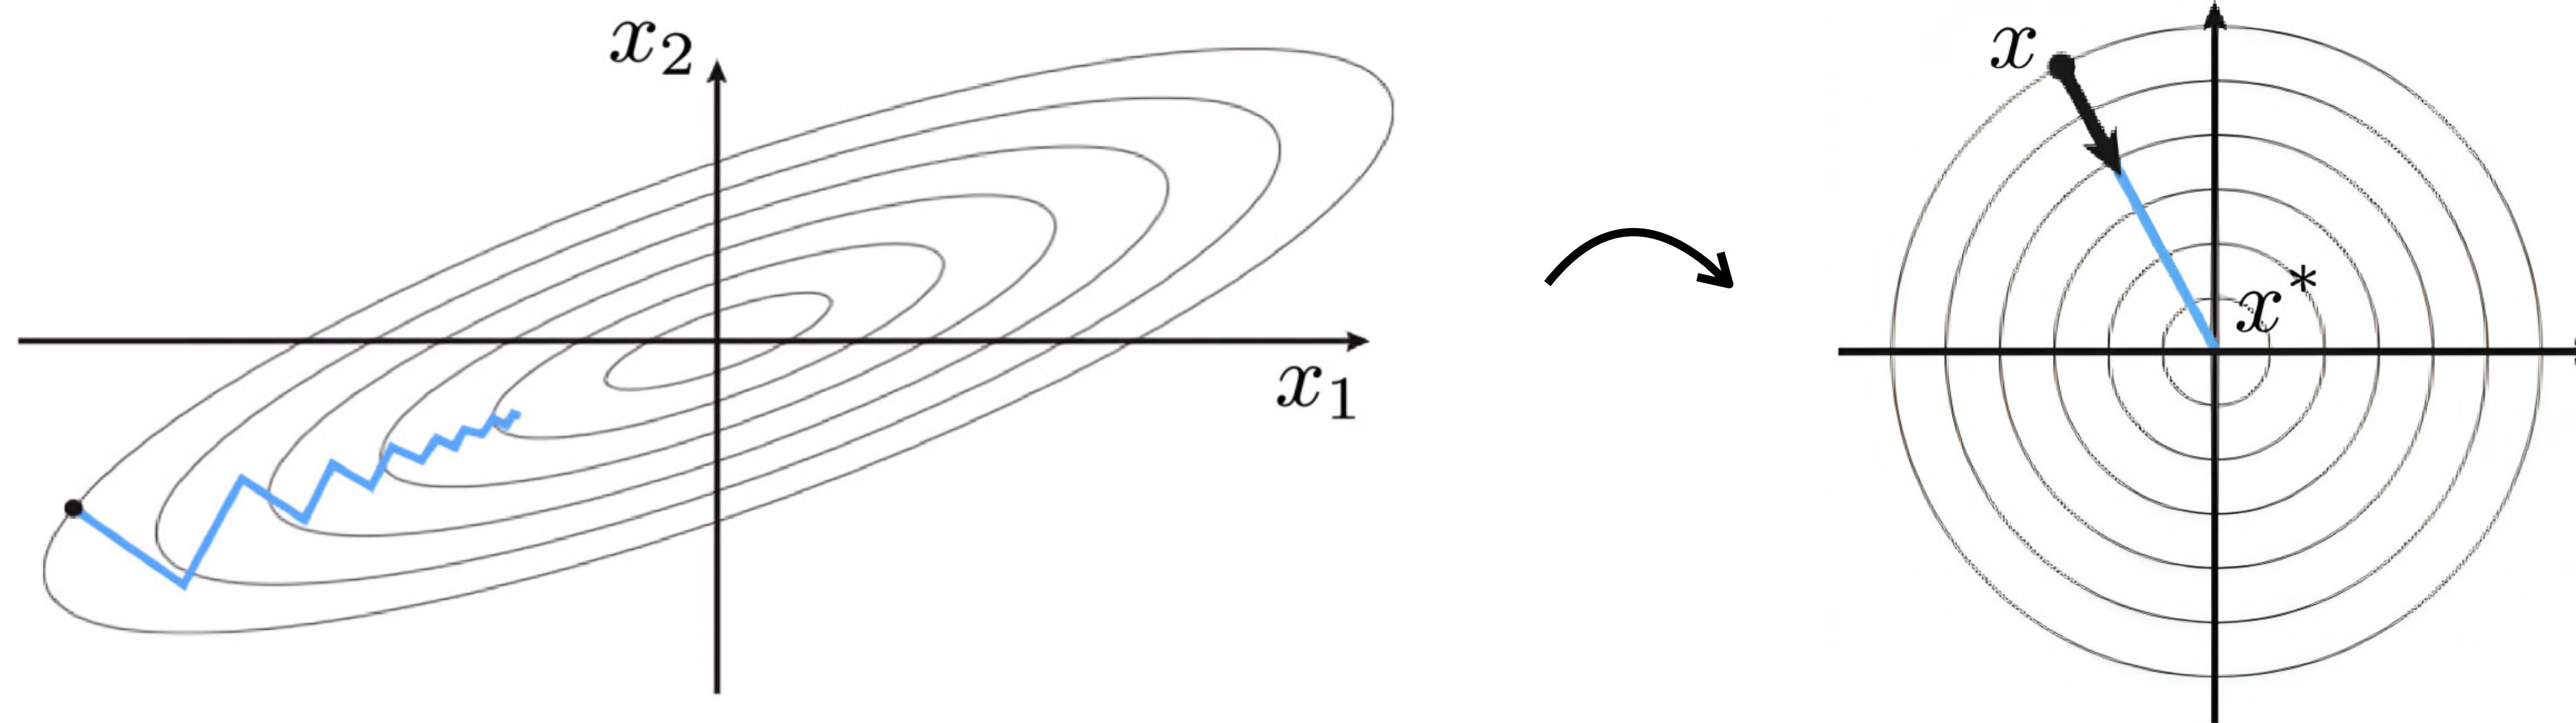
If the function  $f$  is convex,

$$\mathbf{d} = -\mathbf{H}^{-1}\mathbf{g}$$

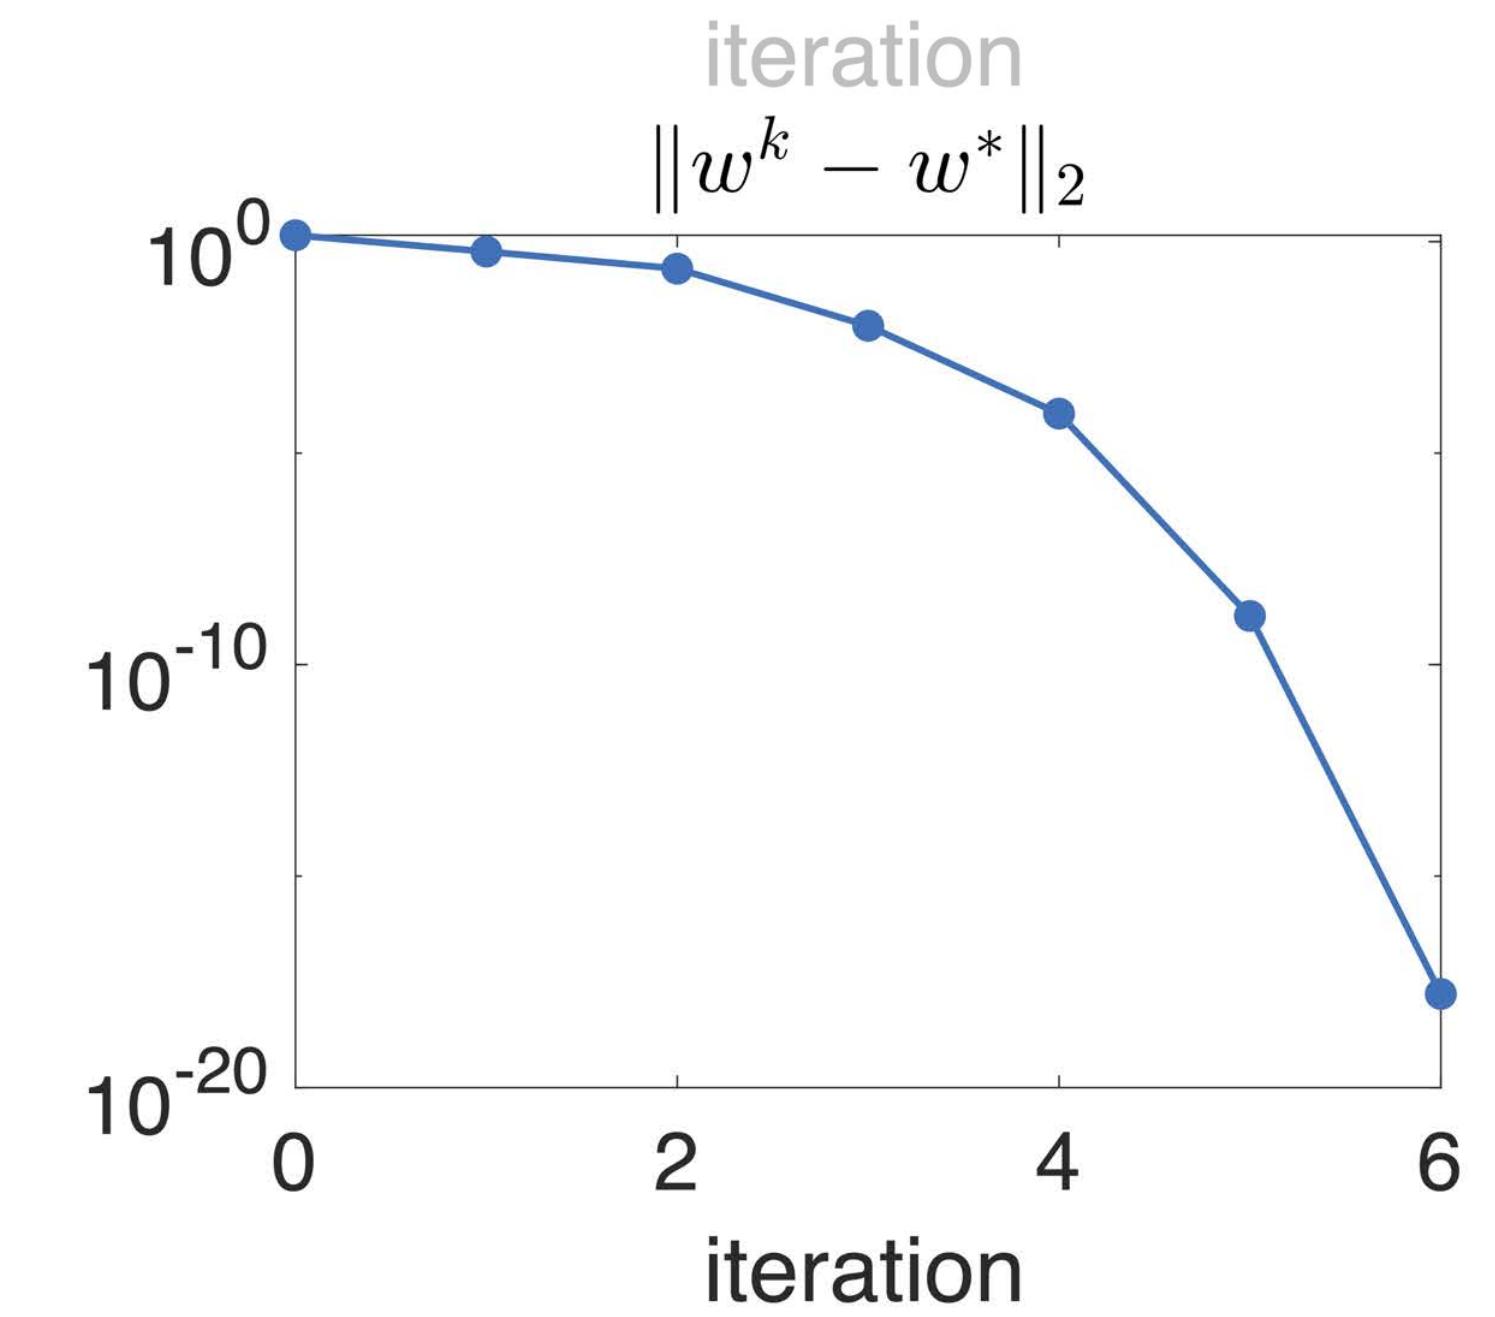
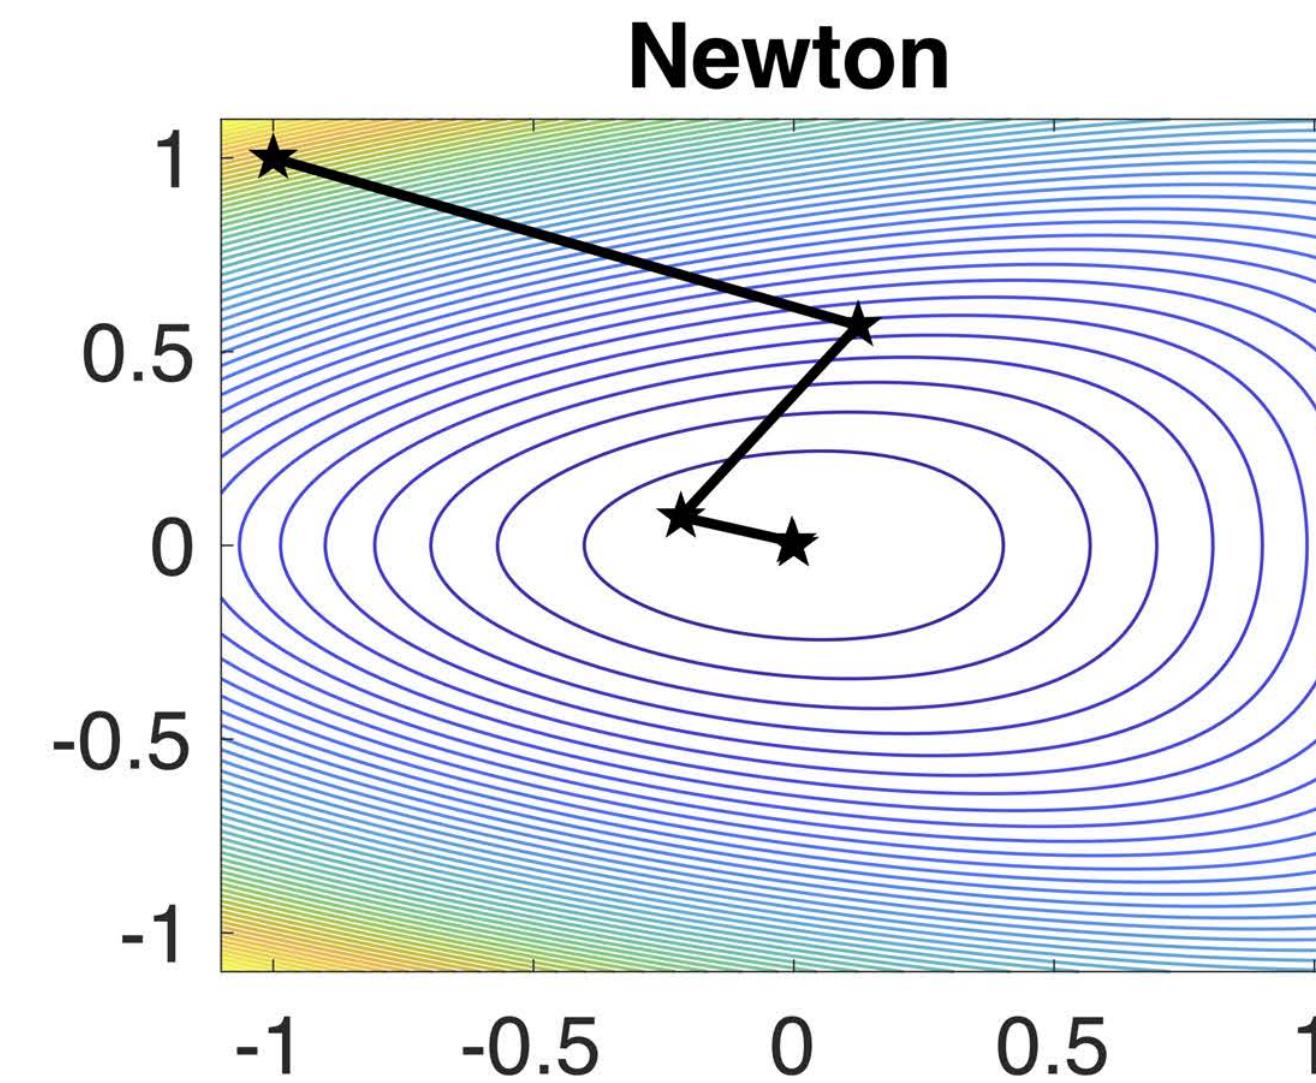
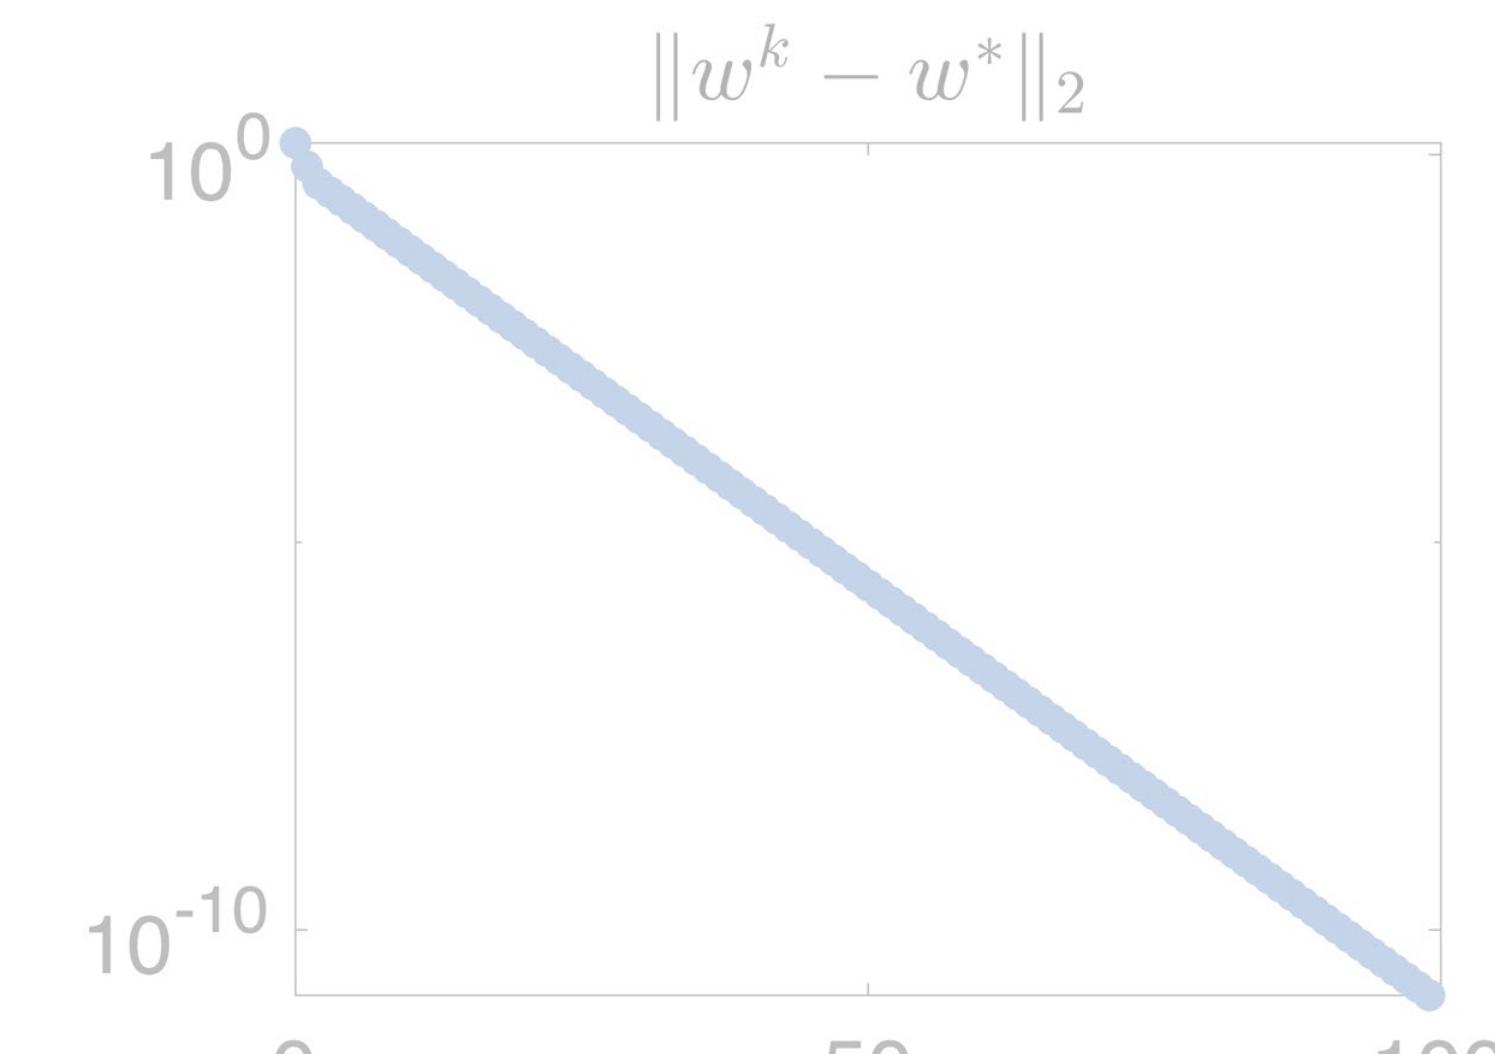
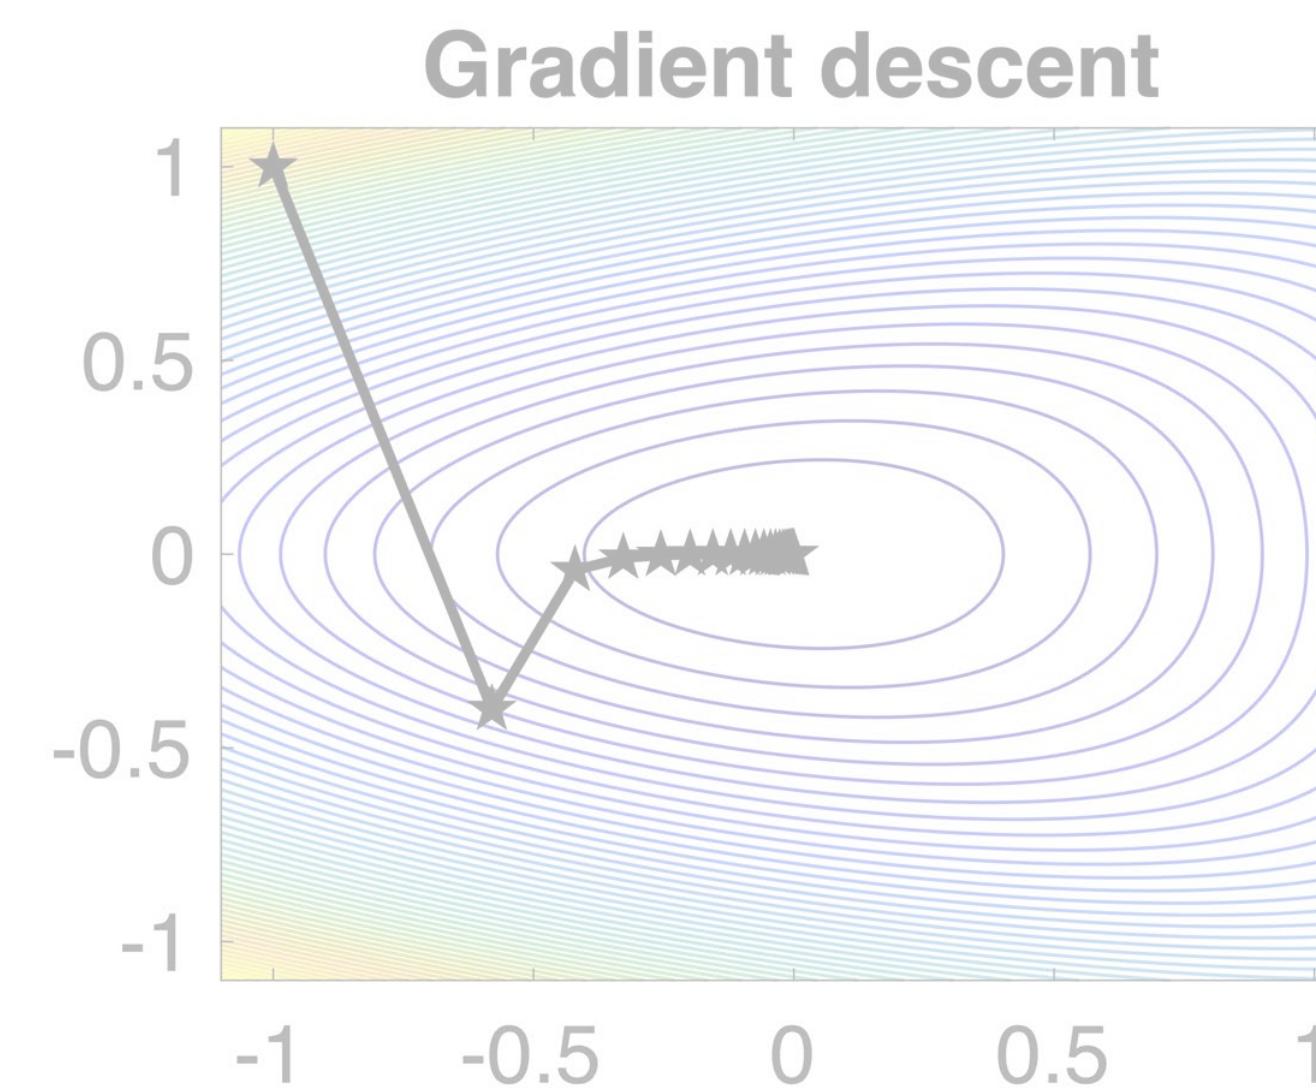


# Newton's Method

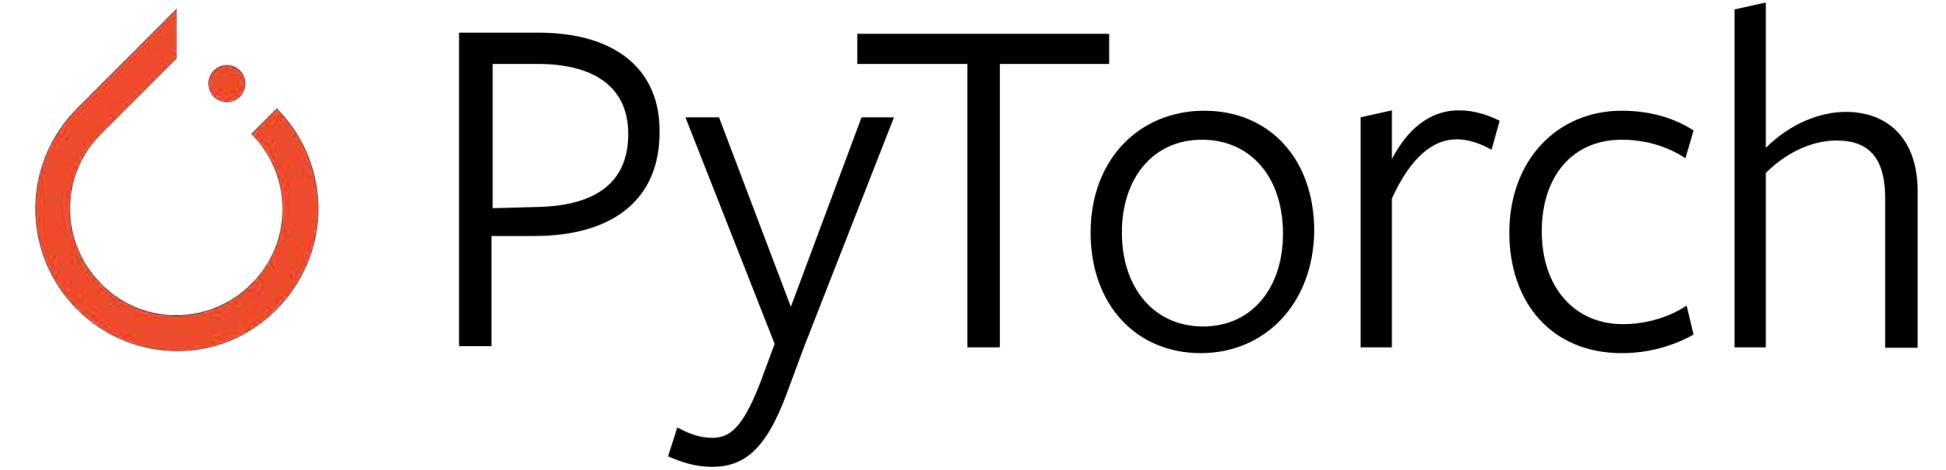
- Idea: apply a coordinate transformation so that the local energy landscape looks more like a “round bowl”
  - The gradient now directly points toward the nearby minimizer



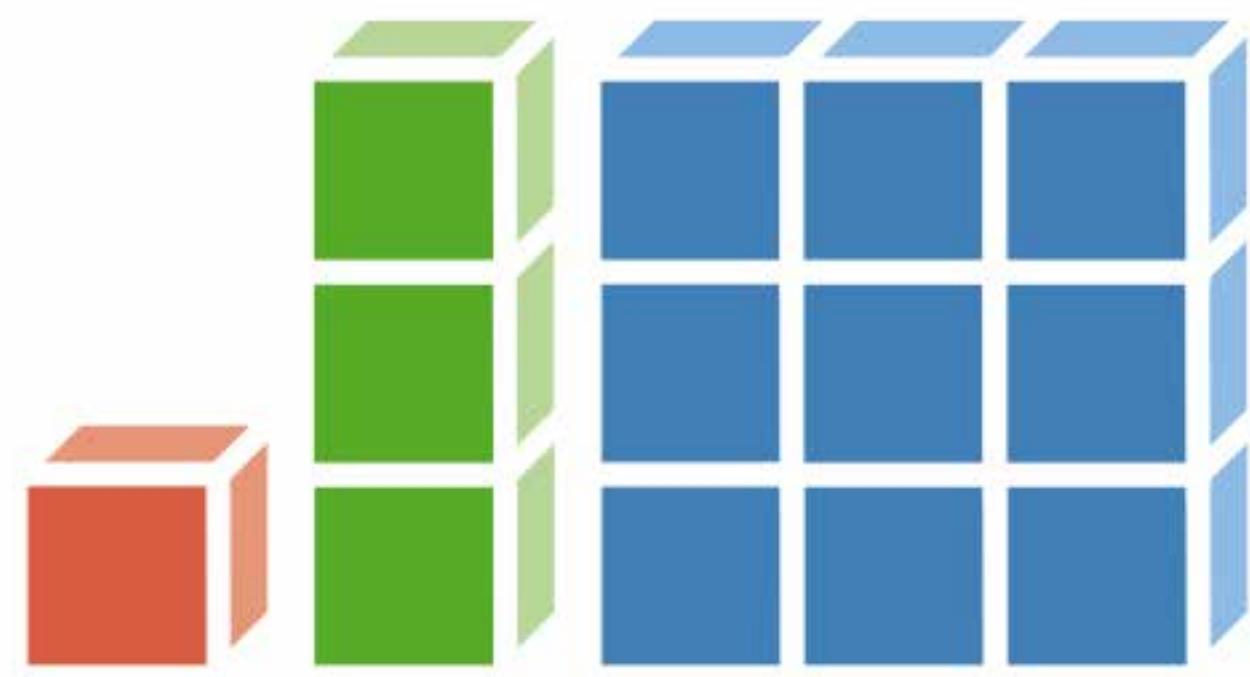
# Fast Convergence of Newton's Method



# Auto-differentiation Libraries



## TinyAD



Easy to compute the gradient  
and the (sparse) Hessian :)

# In C++

# TinyAD



```
// Compute symmetric Dirichlet energy
Eigen::Matrix2<T> J = M * Mr.inverse();
return A * (J.squaredNorm() + J.inverse().squaredNorm());
```

```
#include <TinyAD/ScalarFunction.hh>

// Set up a function with 2D vertex positions as variables
auto func = TinyAD::scalar_function<2>(mesh.vertices());

// Add an objective term per triangle. Each connecting 3 vertices
func.add_elements<3>(mesh.faces(), [&] (auto& element)
{
    // Element is evaluated with either double or TinyAD::Double<6>
    using T = TINYAD_SCALAR_TYPE(element);

    // Get variable 2D vertex positions of triangle t
    OpenMesh::SmartFaceHandle t = element.handle;
    Eigen::Vector2<T> a = element.variables(t.halfedge().to());
    Eigen::Vector2<T> b = element.variables(t.halfedge().next().to());
    Eigen::Vector2<T> c = element.variables(t.halfedge().from());

    return ...
});

// Evaluate the function using any of these methods:
double f = func.eval(x);
auto [f, g] = func.eval_with_gradient(x);
auto [f, g, H] = func.eval_with_derivatives(x);
auto [f, g, H_proj] = func.eval_with_hessian_proj(x);
...
```

Just write out  
the energy

- Elastic Energy
- Optimization Basics & Classical Algorithms

## ● Optimization Methods for Simulation

- ⚡Challenge 1⚡: Handling Nonconvexity
- ⚡Challenge 2⚡: Nontrivial Constraints
- ⚡Challenge 3⚡: Large-scale Optimization

# Optimization Methods for Simulation

---

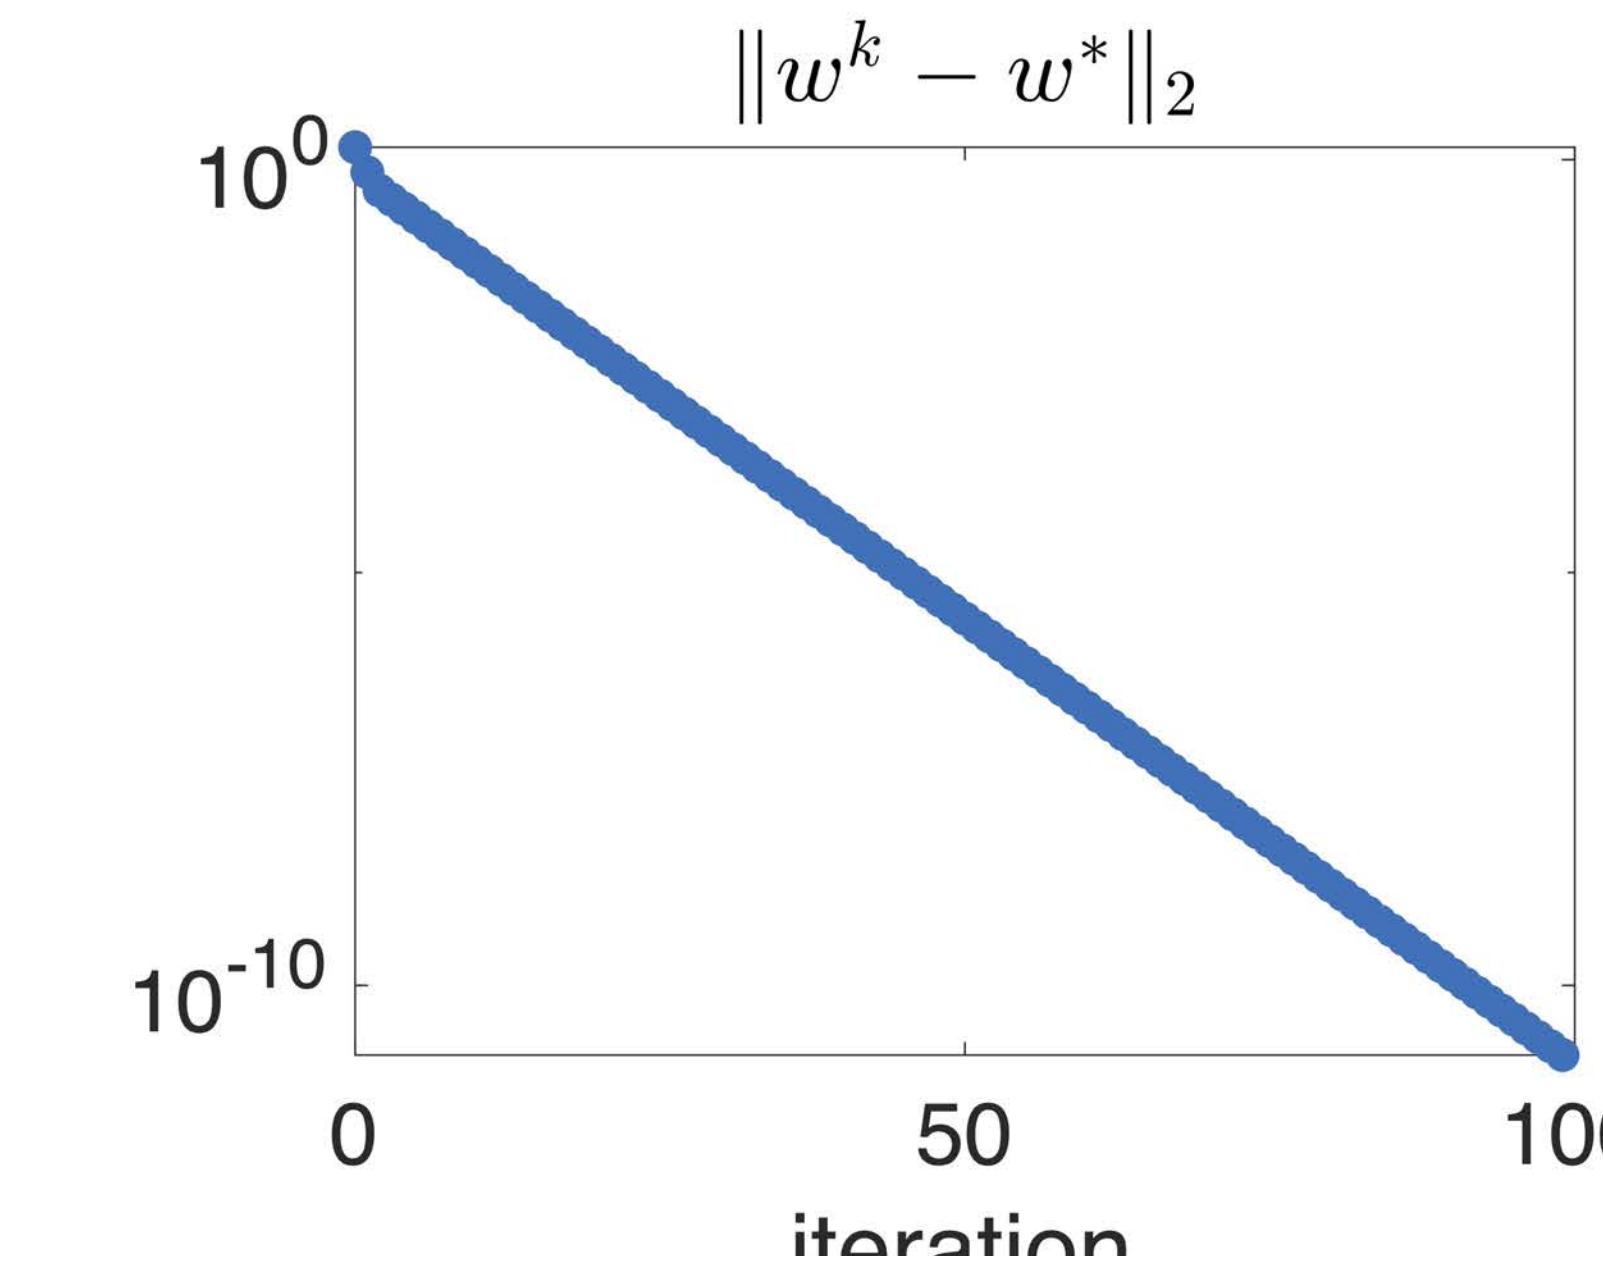
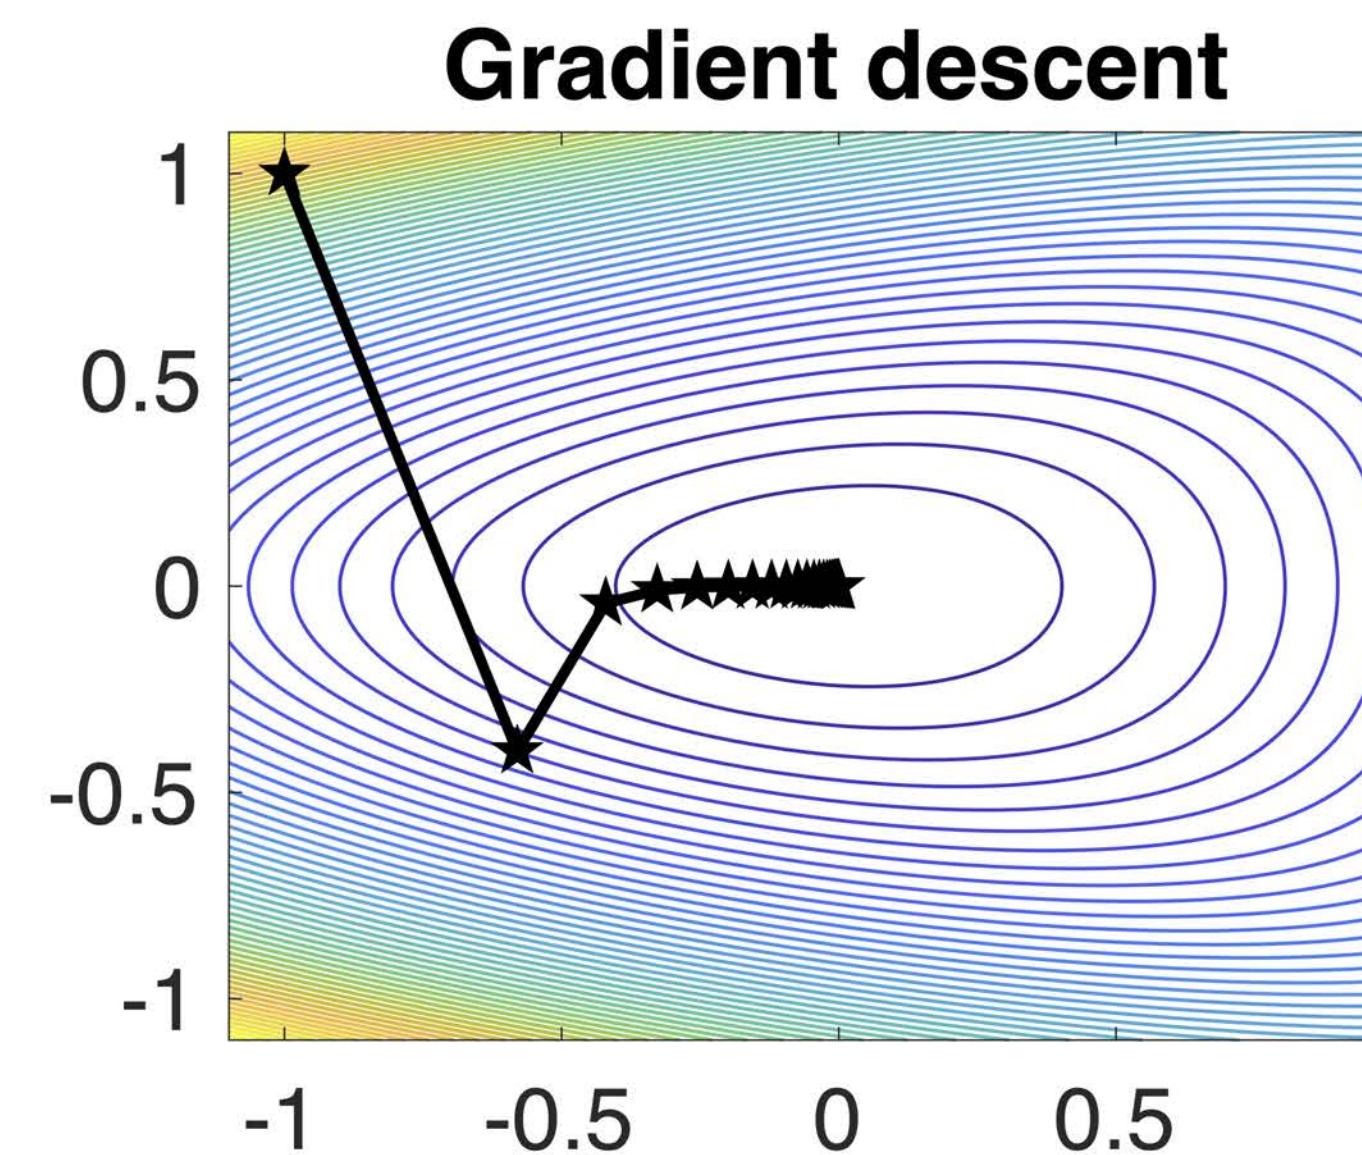
What are the challenges of applying those classical optimization methods to simulation?



# Why not just use gradient descent?

$$\mathbf{d} = -\mathbf{g}$$

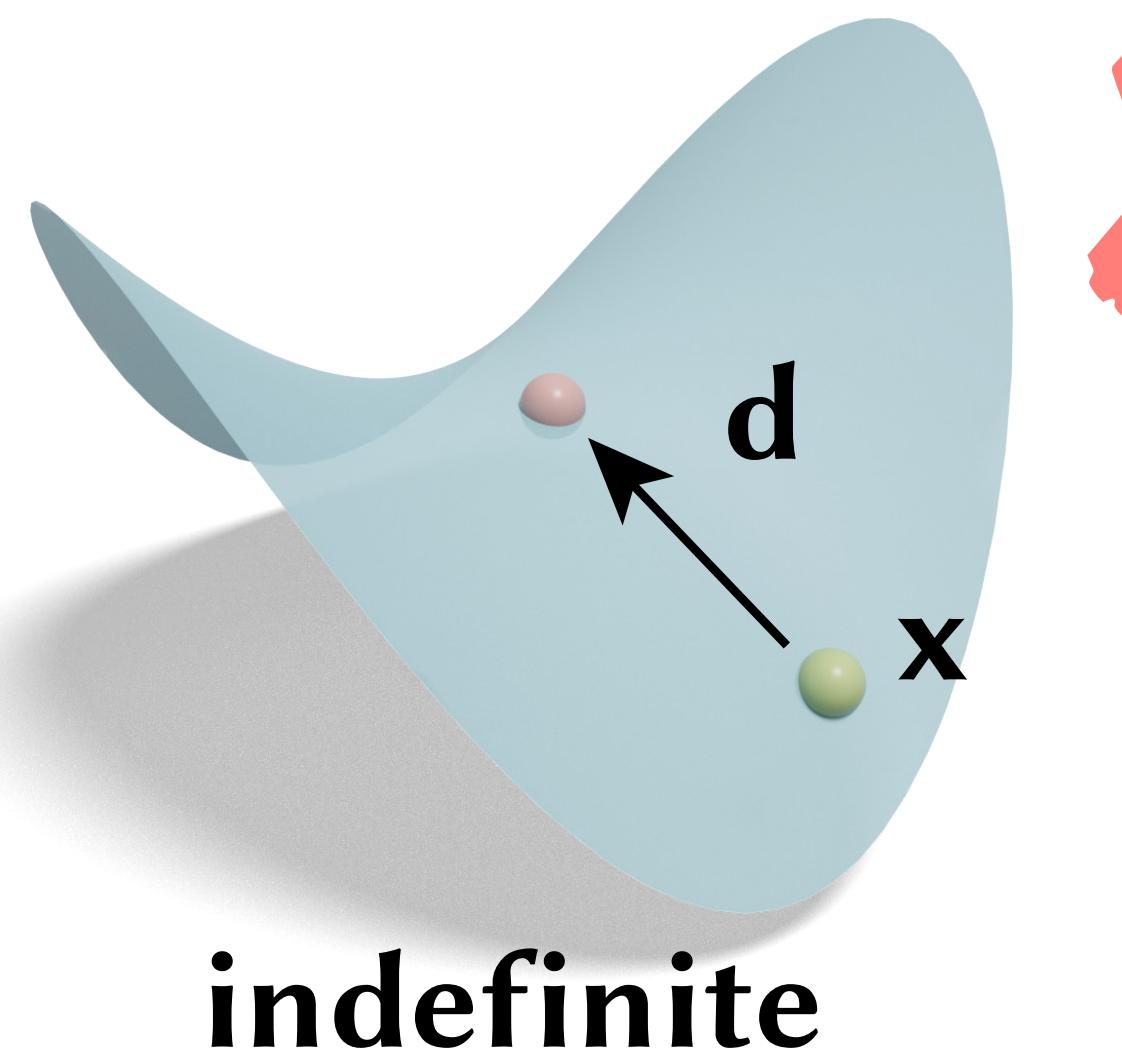
- The convergence of vanilla gradient descent can be very slow in higher dimensions



# Why not just directly use Newton's method?

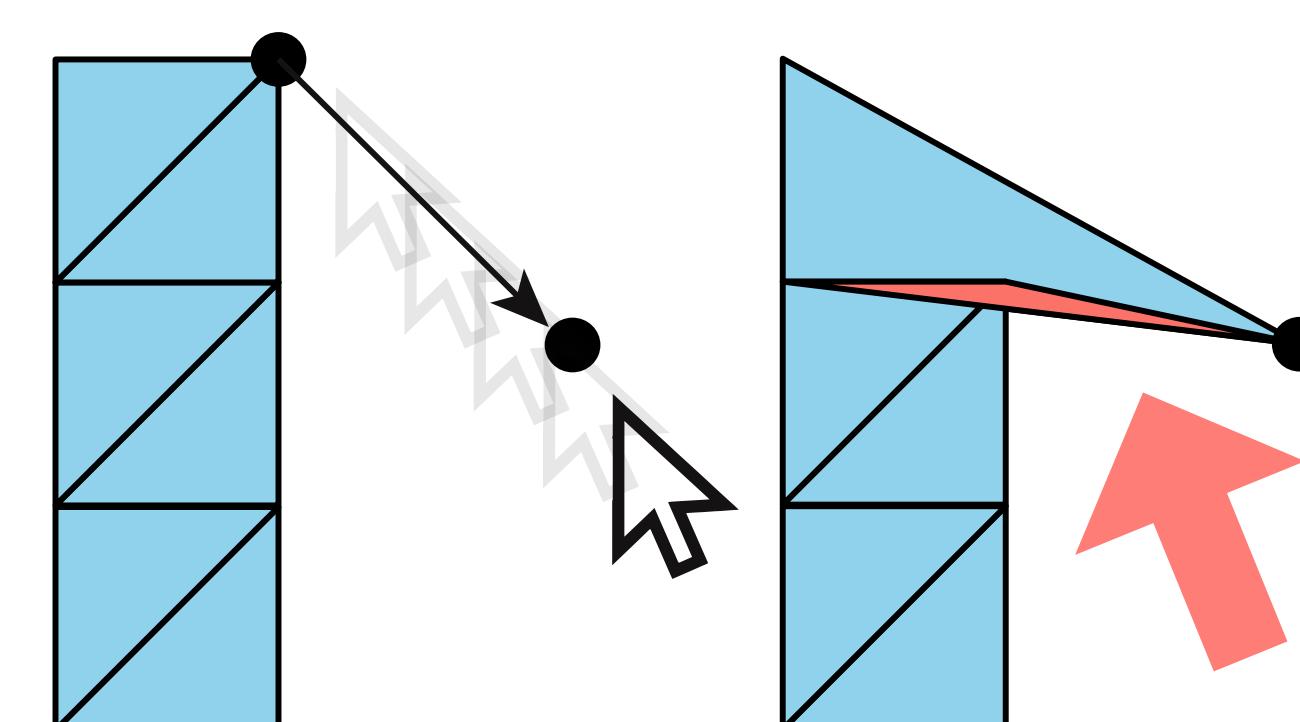
$$\mathbf{d} = -\mathbf{H}^{-1}\mathbf{g}$$

- Nonconvexity

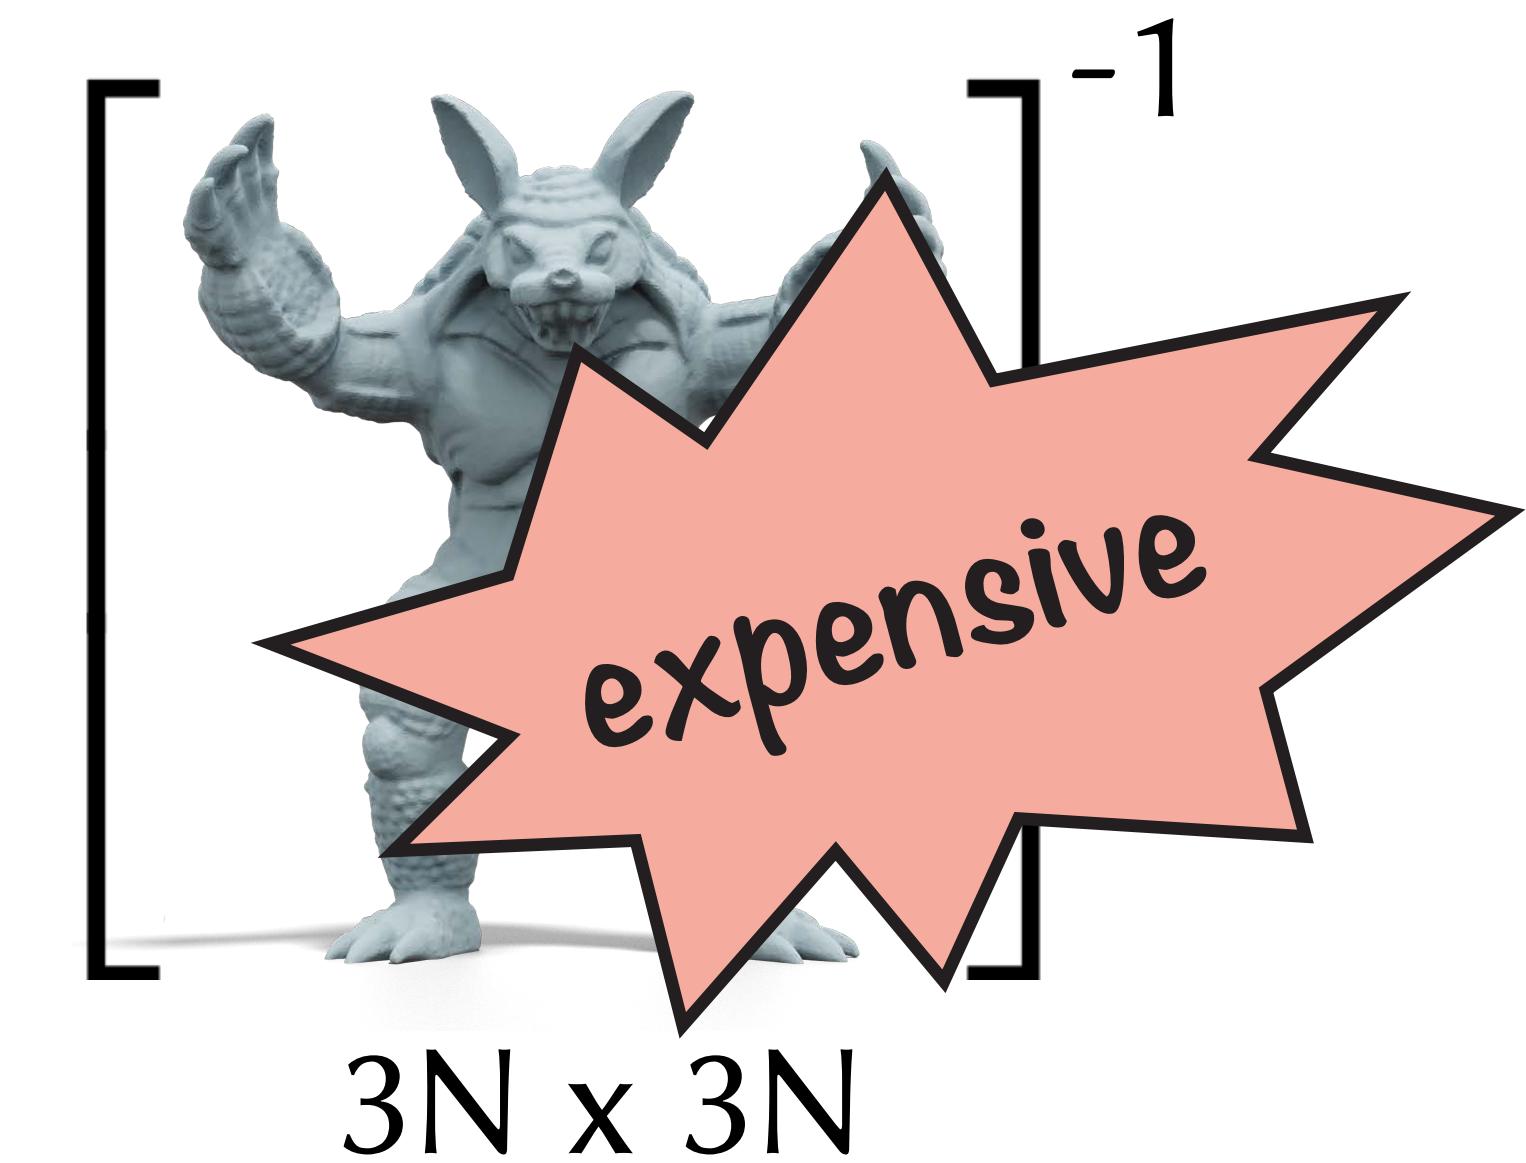


indefinite

- Infeasible points



- Large-scale

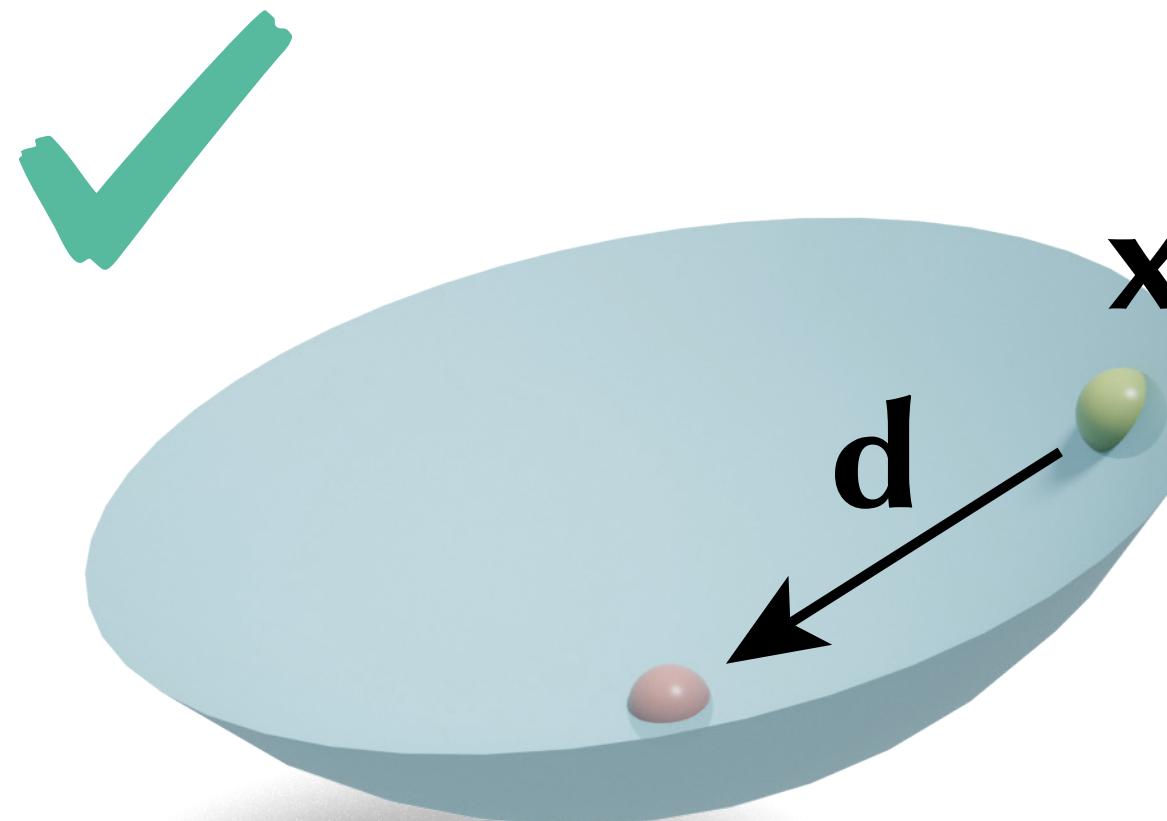


- Elastic Energy
- Optimization Basics & Classical Algorithms
- Optimization Methods for Simulation
  - ⚡Challenge 1⚡: Handling Nonconvexity
  - ⚡Challenge 2⚡: Nontrivial Constraints
  - ⚡Challenge 3⚡: Large-scale Optimization

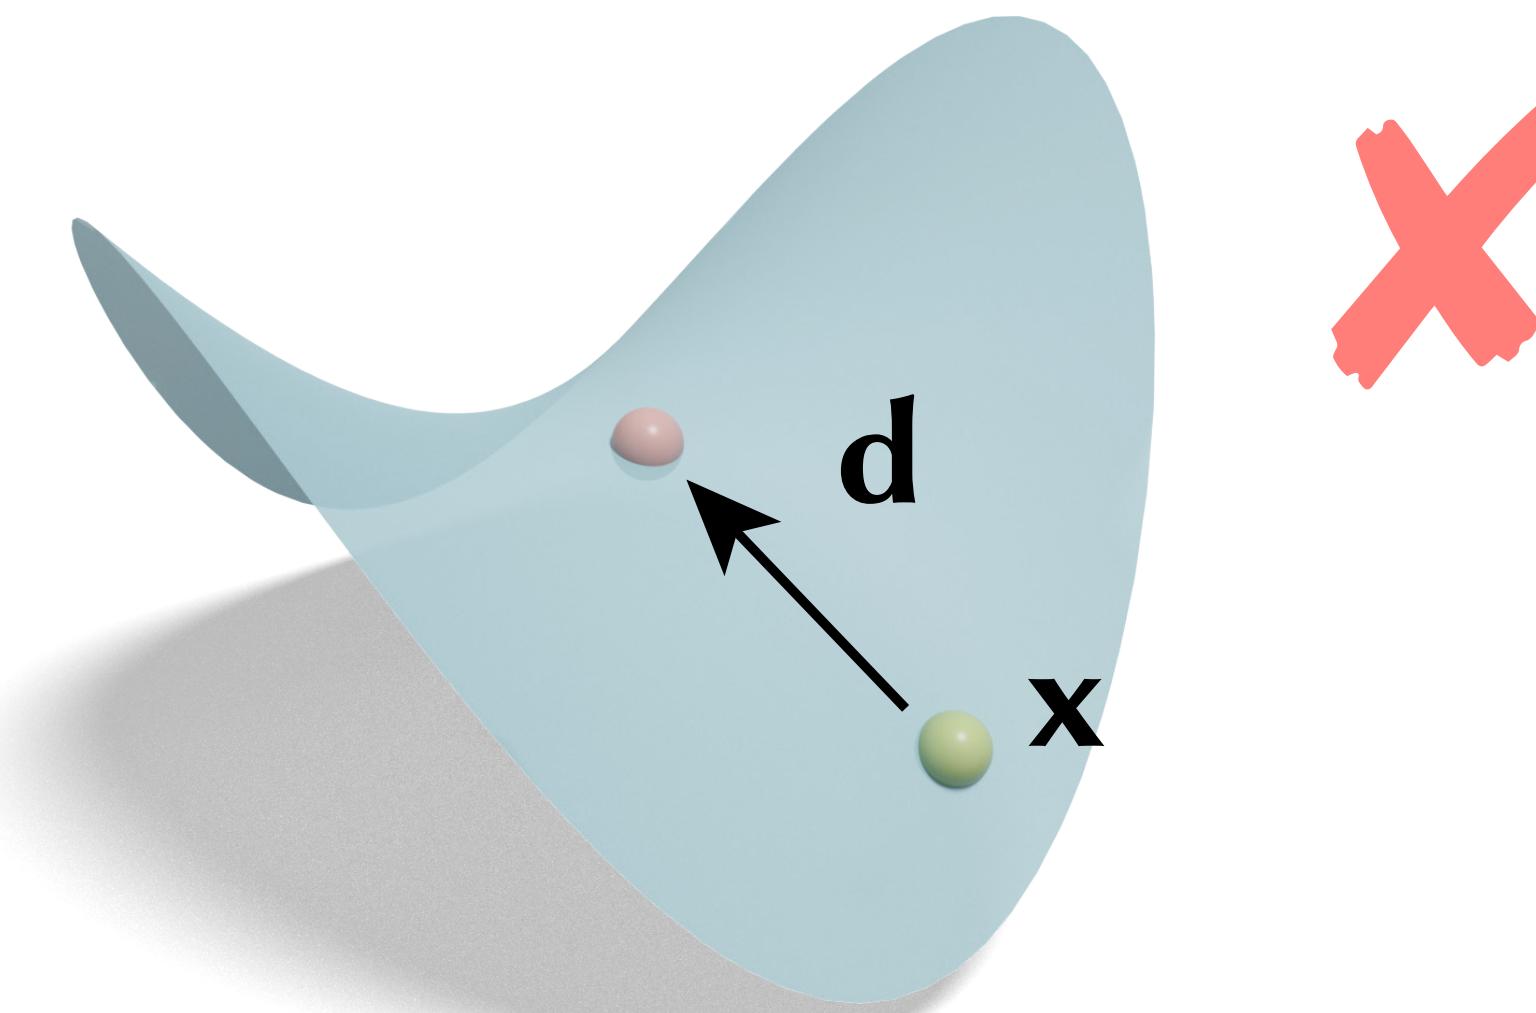
# Challenge 1: Handling Nonconvexity

# Challenge 1: Handling Nonconvexity

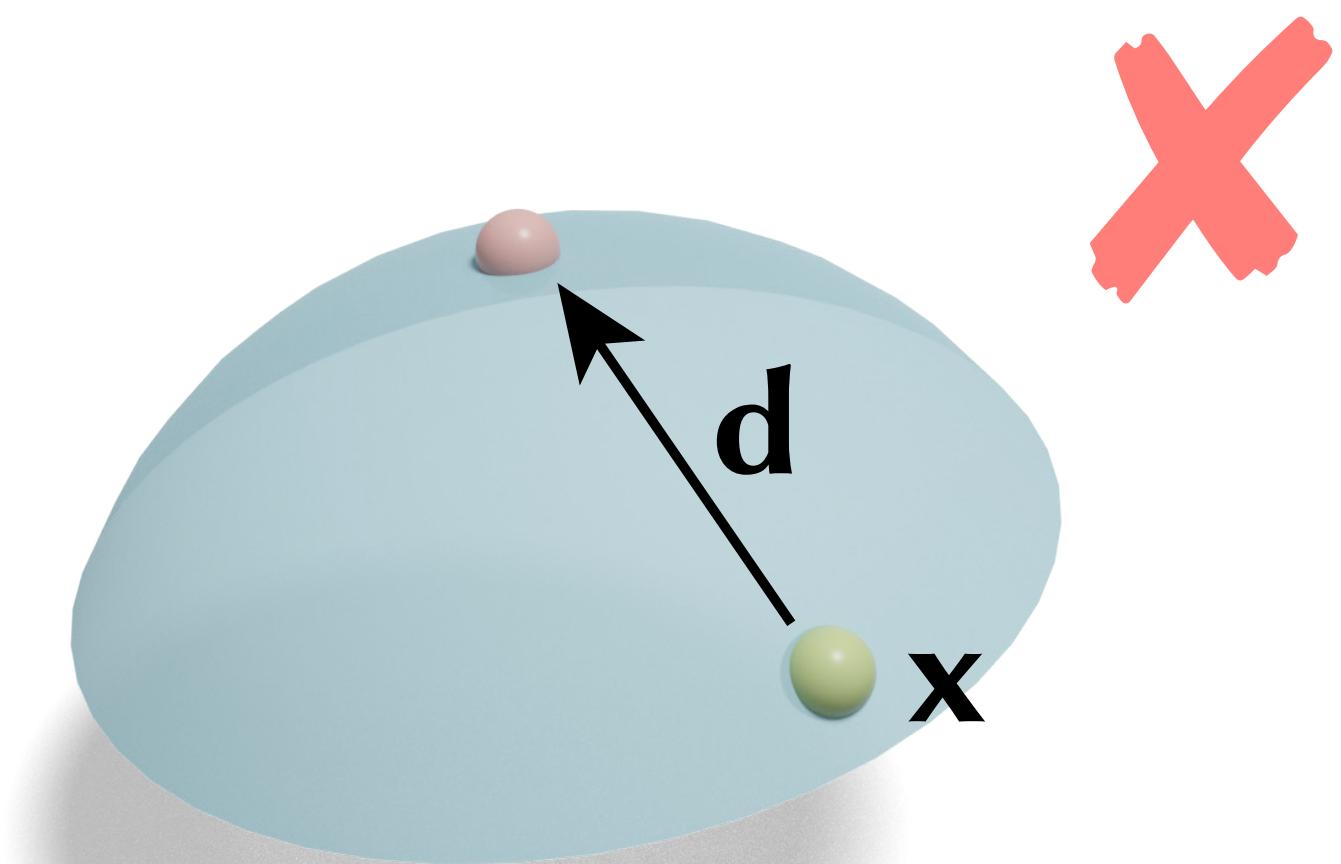
$$\mathbf{d} = -\mathbf{H}^{-1}\mathbf{g}$$



**positive definite**

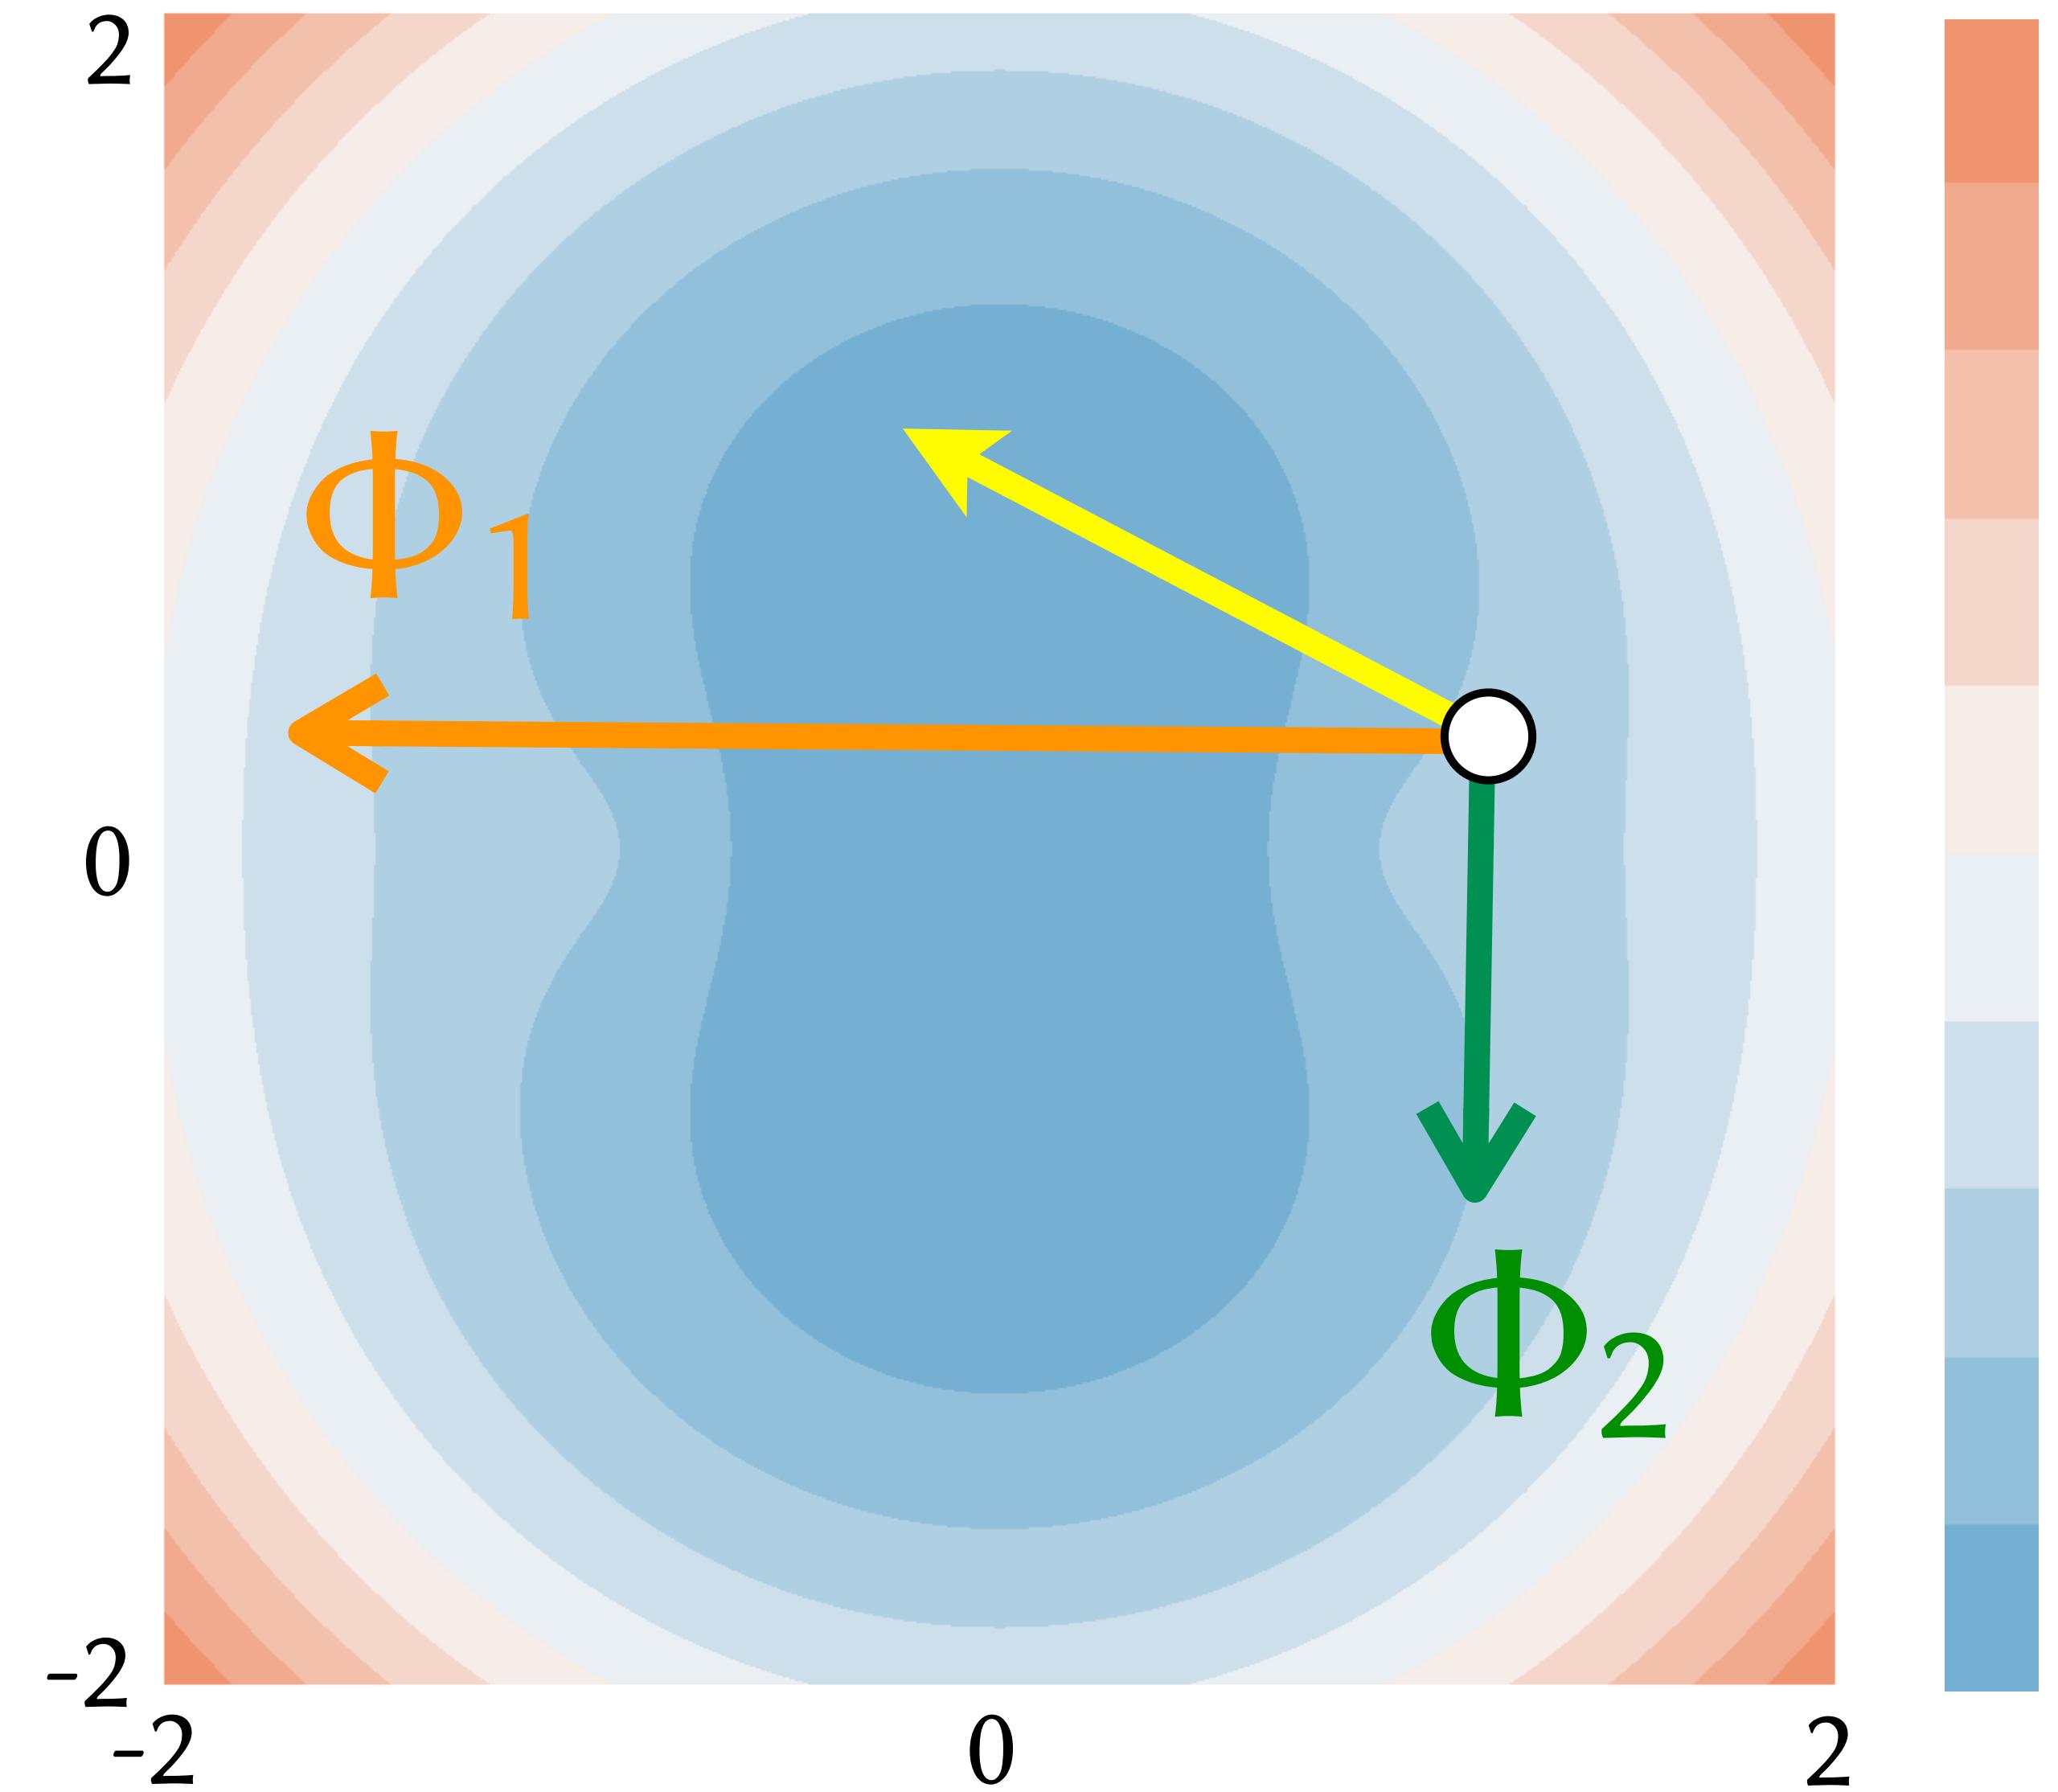


**indefinite**



**negative definite**

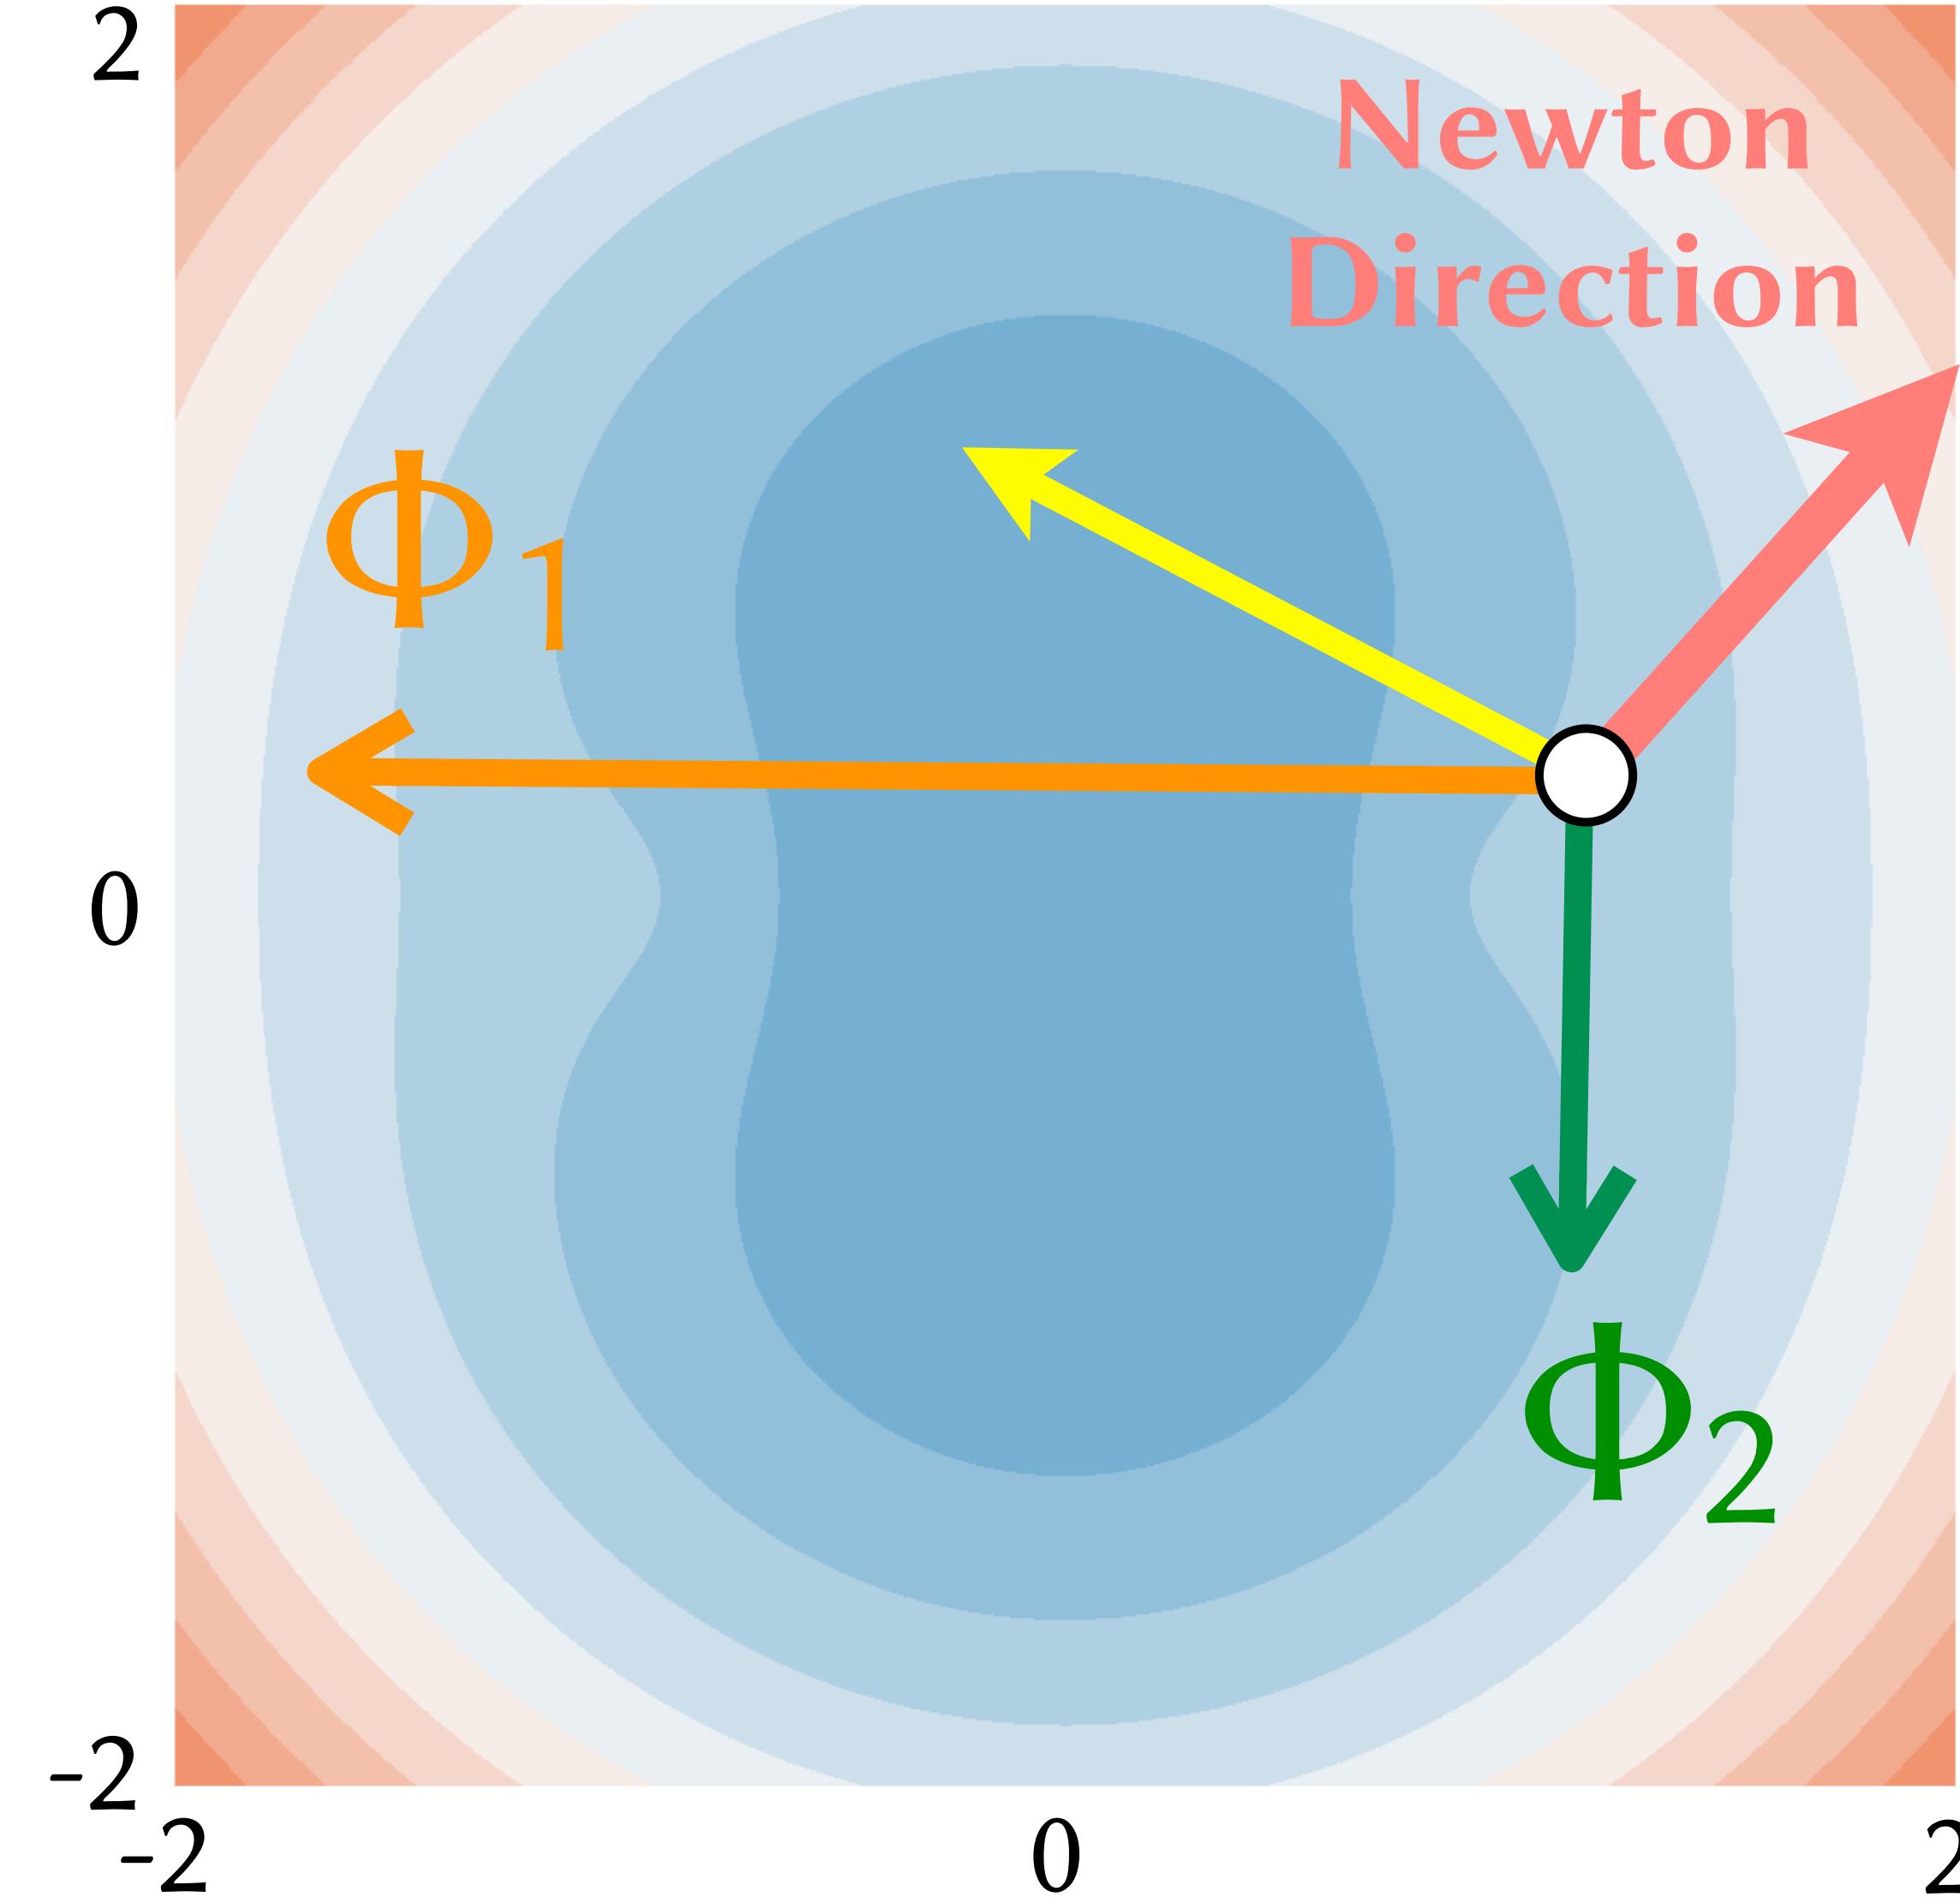
# Why is Nonconvexity Problematic?



$$\mathbf{H} = \Phi \begin{bmatrix} -6 & 0 \\ 0 & 3 \end{bmatrix} \Phi^\top$$

$$f(x, y) = \left( \sqrt{(x+1)^2 + y^2} - 1 \right)^2 + \left( \sqrt{(x-1)^2 + y^2} - 1 \right)^2$$

Choose  $(x, y) = (1, 0.2)$



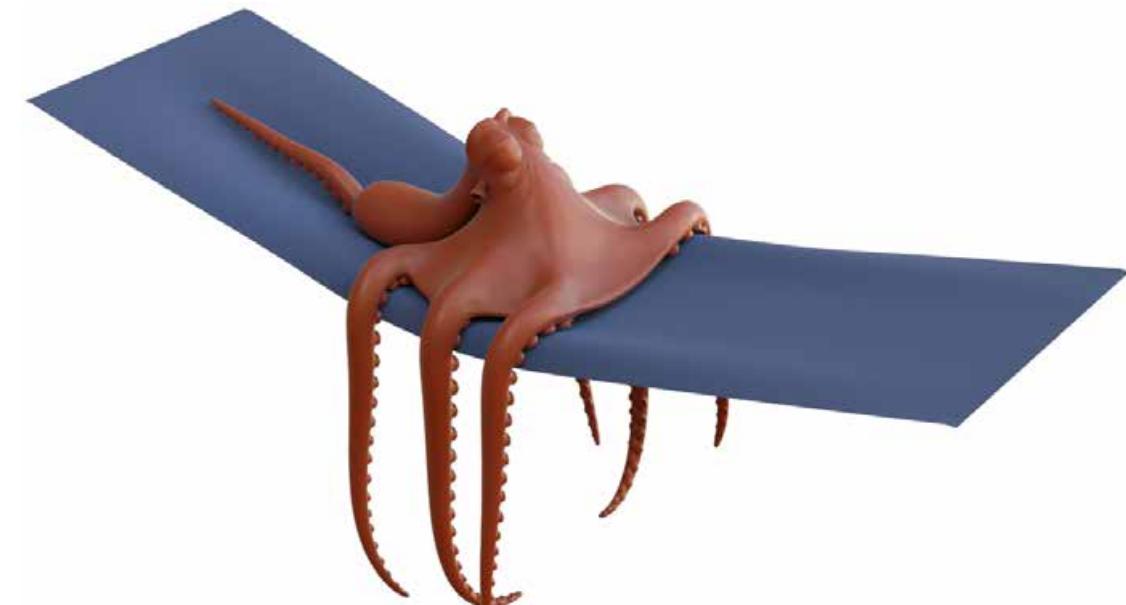
$$\begin{aligned}
 -\mathbf{H}^{-1} \nabla f &= -\left( \Phi \begin{bmatrix} -6 \\ 0 \\ 3 \end{bmatrix} \Phi^\top \right) \nabla f \\
 &= -\left( \Phi \begin{bmatrix} -\frac{1}{6} \\ 0 \\ \frac{1}{3} \end{bmatrix} \Phi^\top \right) \nabla f
 \end{aligned}$$

# Where does this non-convexity come from?

COURSE, SIGGRAPH 2022

**Dynamic Deformables:**  
**Implementation and Production**  
**Practicalities (Now With Code!)**

**Instructors:**  
Theodore Kim, Yale University  
David Eberle, Pixar Animation Studios



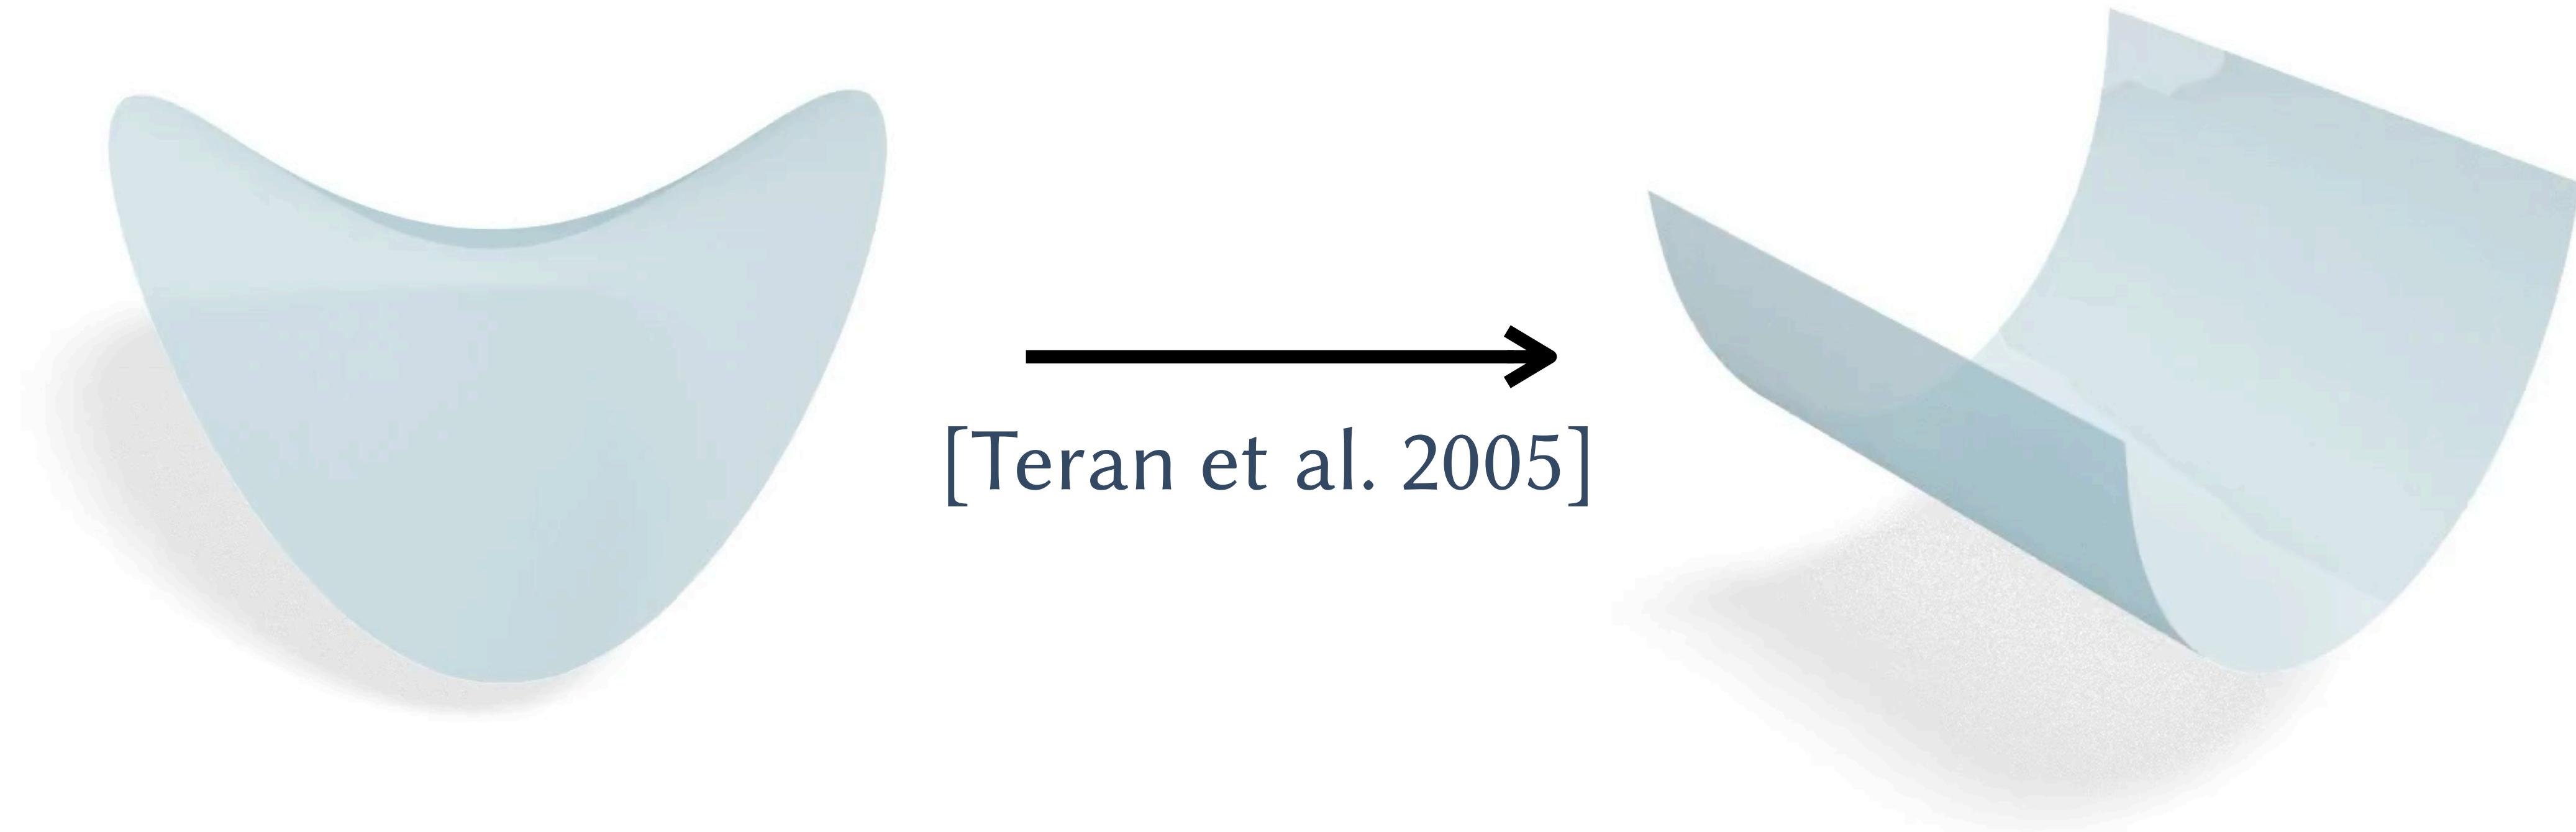
Built on: April 25, 2024

# Projected Newton

$$\mathbf{d} = -(\mathbf{H}^+)^{-1} \mathbf{g}$$

Positive Definite

# Eigenvalue Clamping

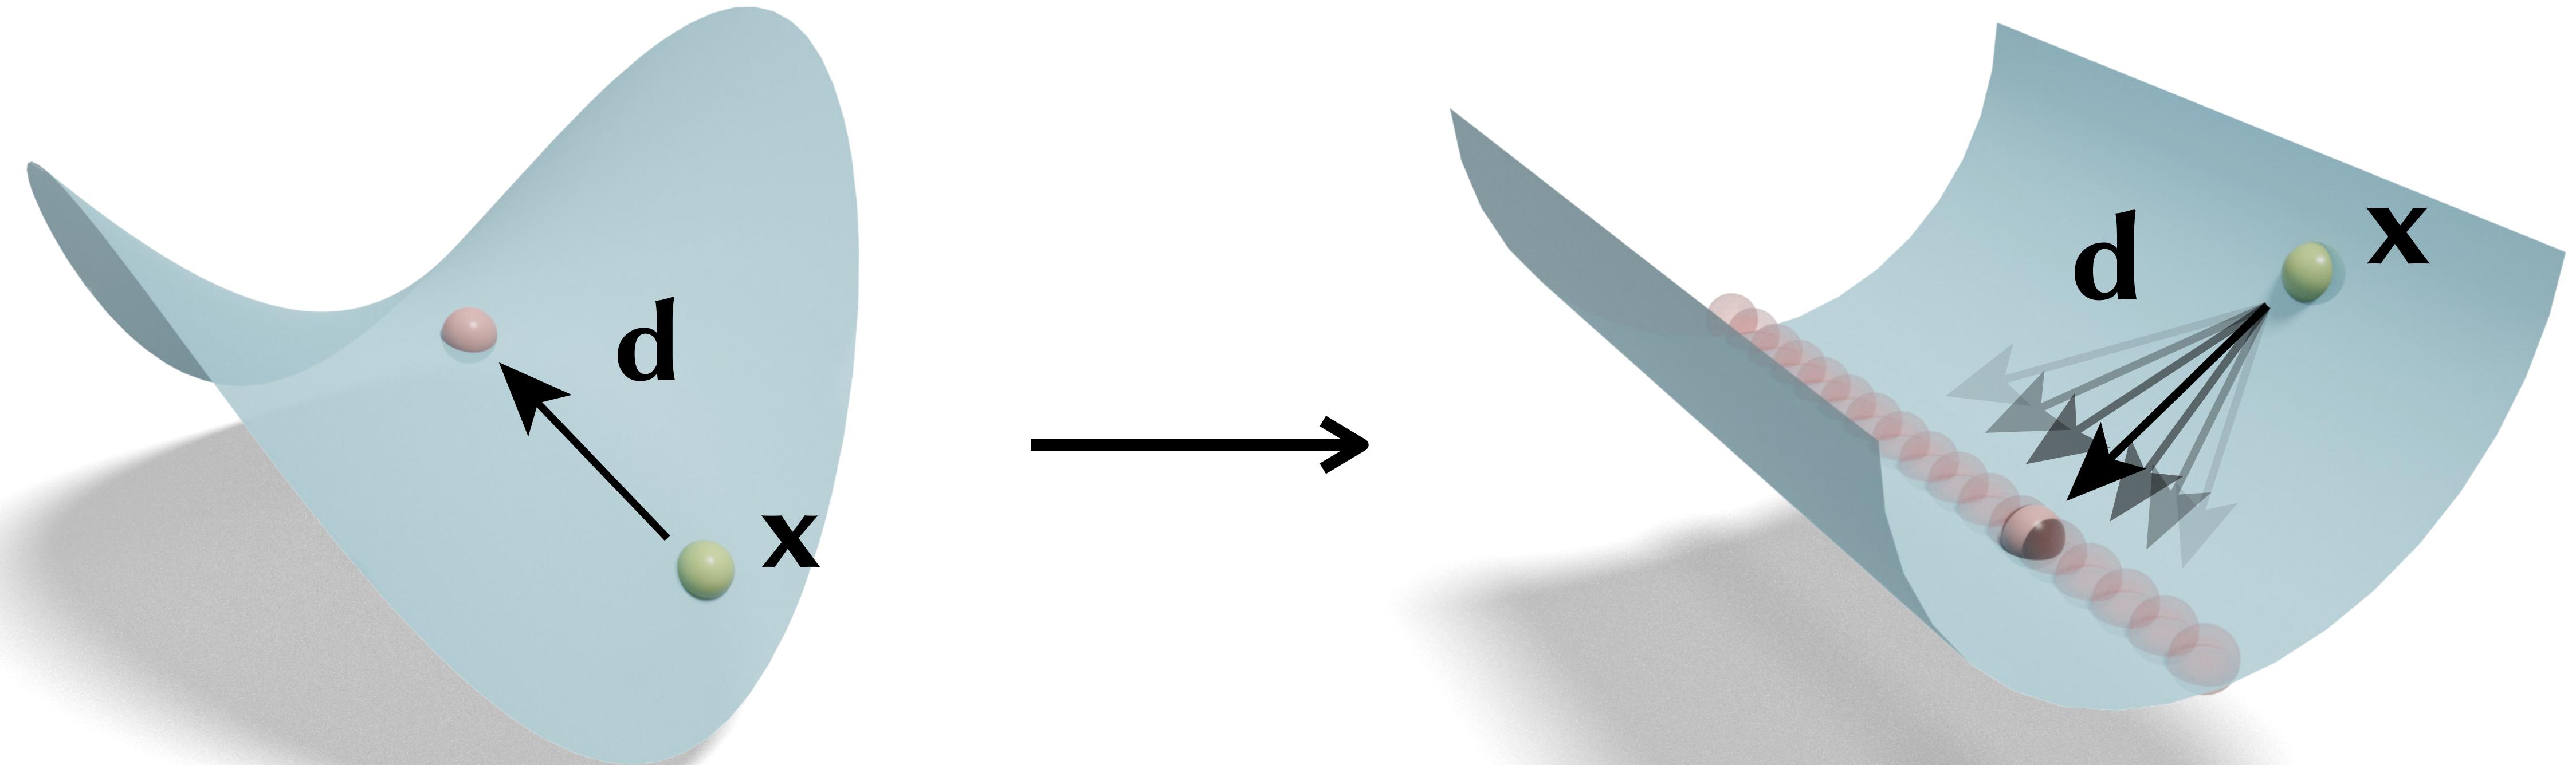


$$\mathbf{H} = \sum_i \mathbf{P}_i^\top \mathbf{H}_i \mathbf{P}_i$$

$$\mathbf{H}_i = \mathbf{U} \Lambda \mathbf{U}^T$$

$$\Lambda_k^+ = \max(\Lambda_k, \epsilon)$$

# Eigenvalue Clamping



indefinite

$\lambda_{\max} > 0, \lambda_{\min} = 0$

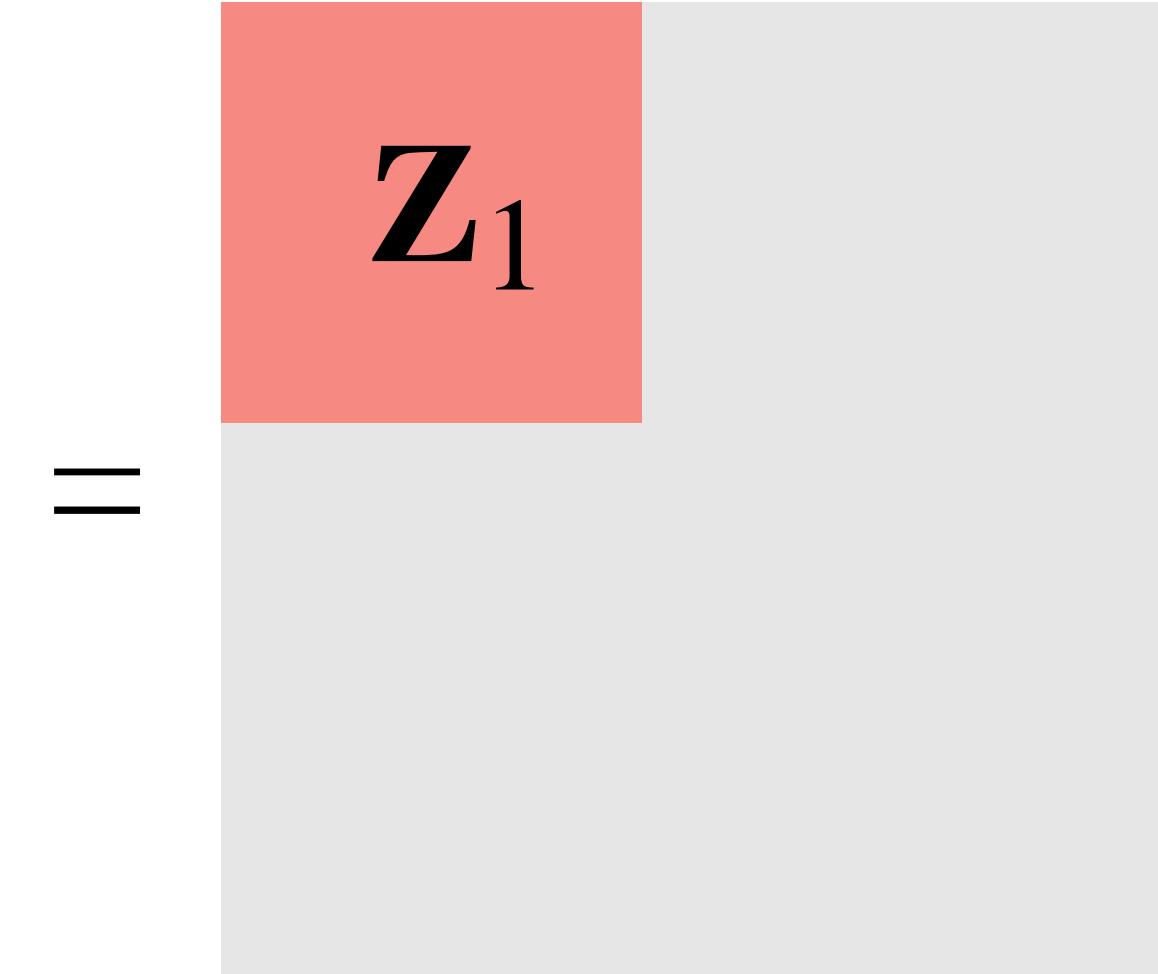
Positive Semidefinite

# Additive Contribution of Neighboring Elements

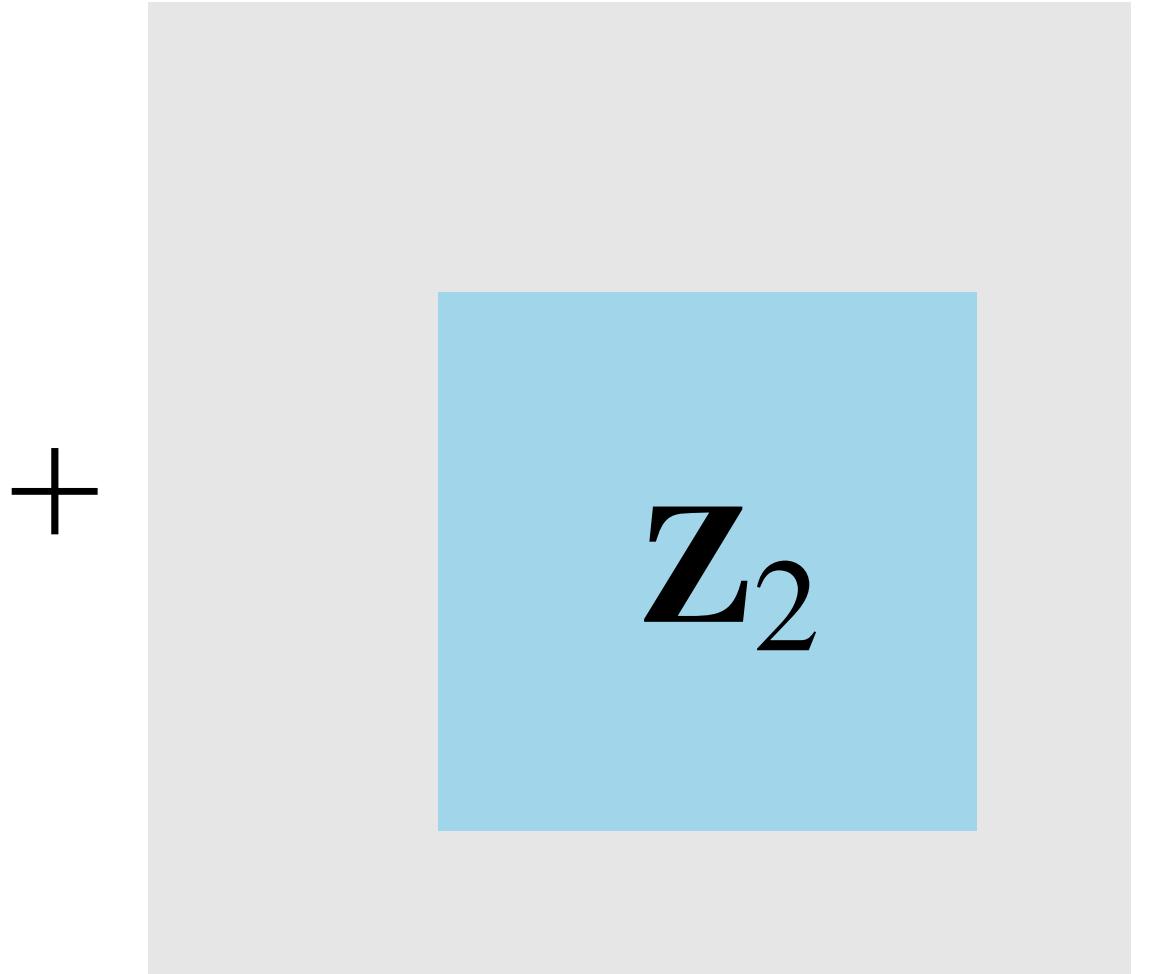
+ Also boundary conditions



$$X \succeq 0$$



$$P_1^T Z_1 P_1$$



$$P_2^T Z_2 P_2$$



$$P_3^T Z_3 P_3$$

$$Z_1 \succeq 0$$

$$Z_2 \succeq 0$$

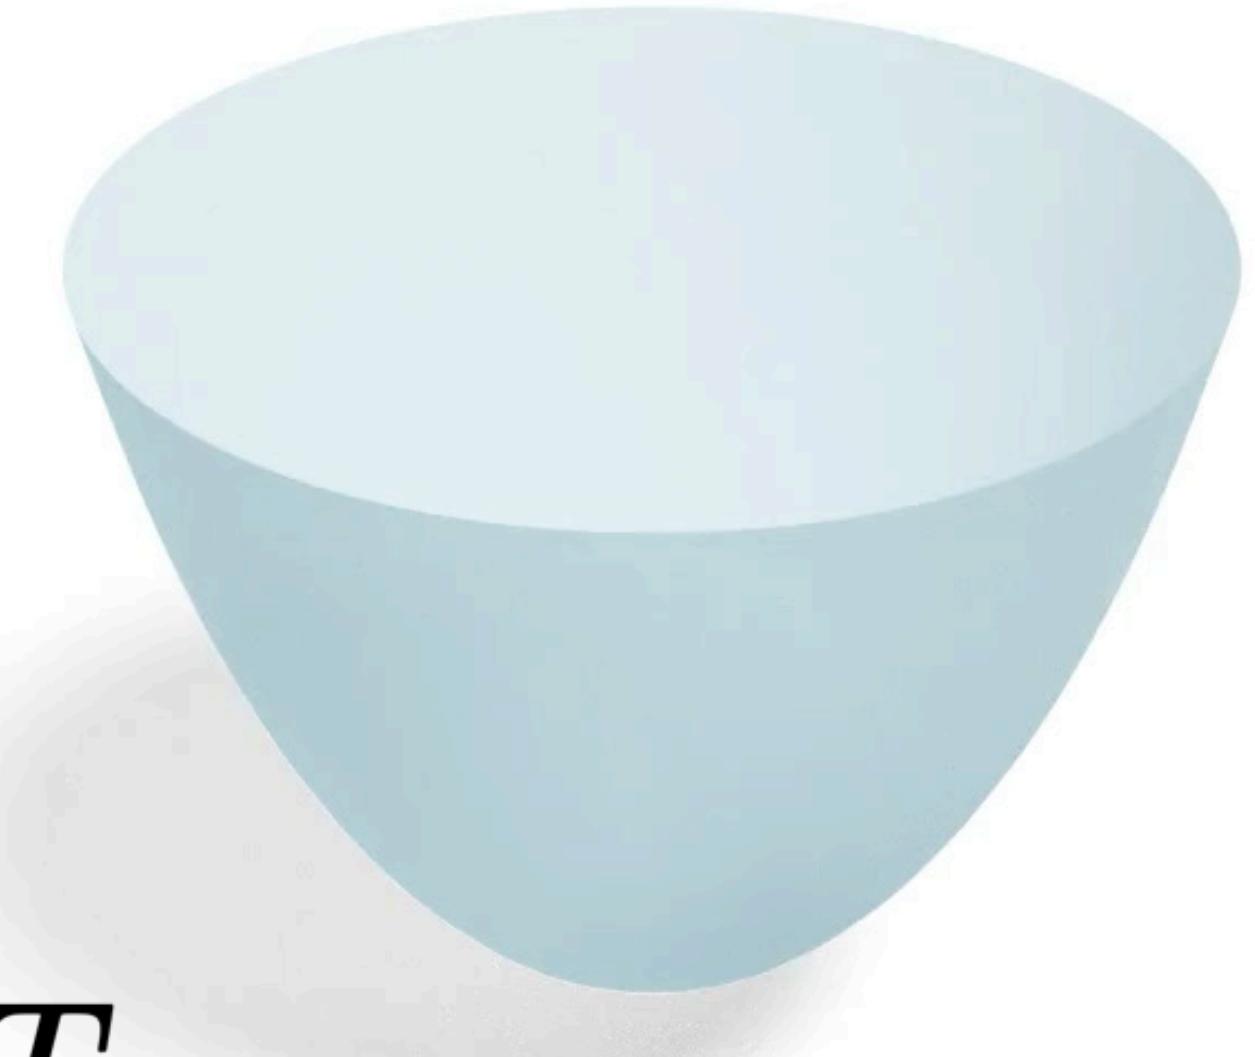
$$Z_3 \succeq 0$$

The global Hessian can often be positive definite  
 $(X > 0)$

# Add a multiple of Identity Matrix



→  
[Tikhonov 1943]



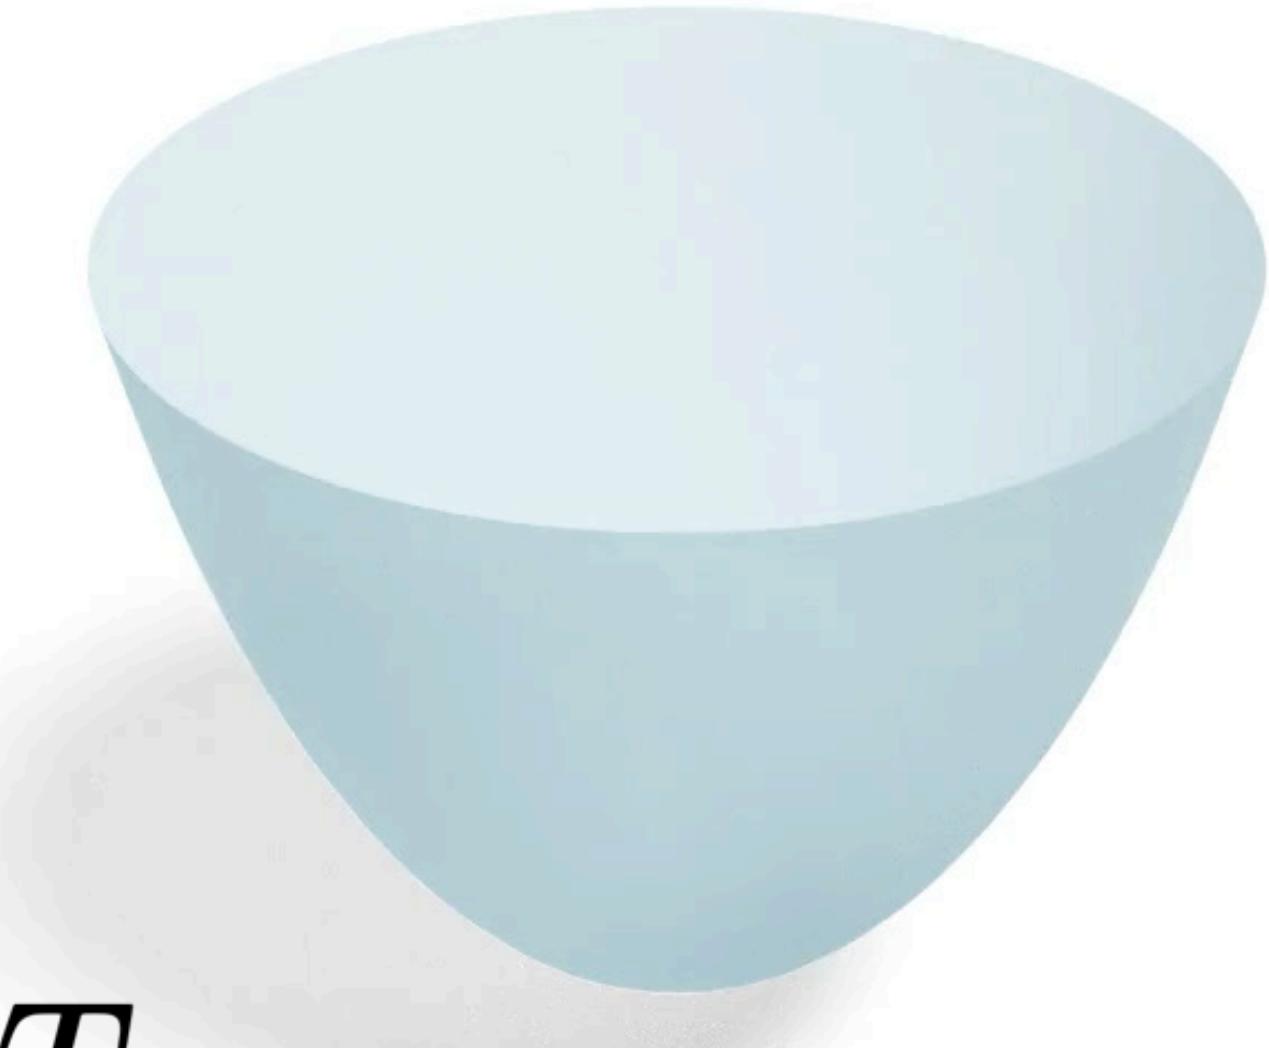
$$\mathbf{H}_i = \mathbf{U}\Lambda\mathbf{U}^T$$

$$\Lambda_k^+ = \Lambda_k + m\mathbf{I}$$

# Absolute Eigenvalue Projection



→  
[Chen et al. 2024]



$$\mathbf{H}_i = \mathbf{U} \Lambda \mathbf{U}^T$$

$$\Lambda_k^+ = |\Lambda_k|$$

(for highly non-convex scenarios)

# Recap: Projected Newton

$$\mathbf{d} = -(\mathbf{H}^+)^{-1} \mathbf{g}$$

Positive Definite

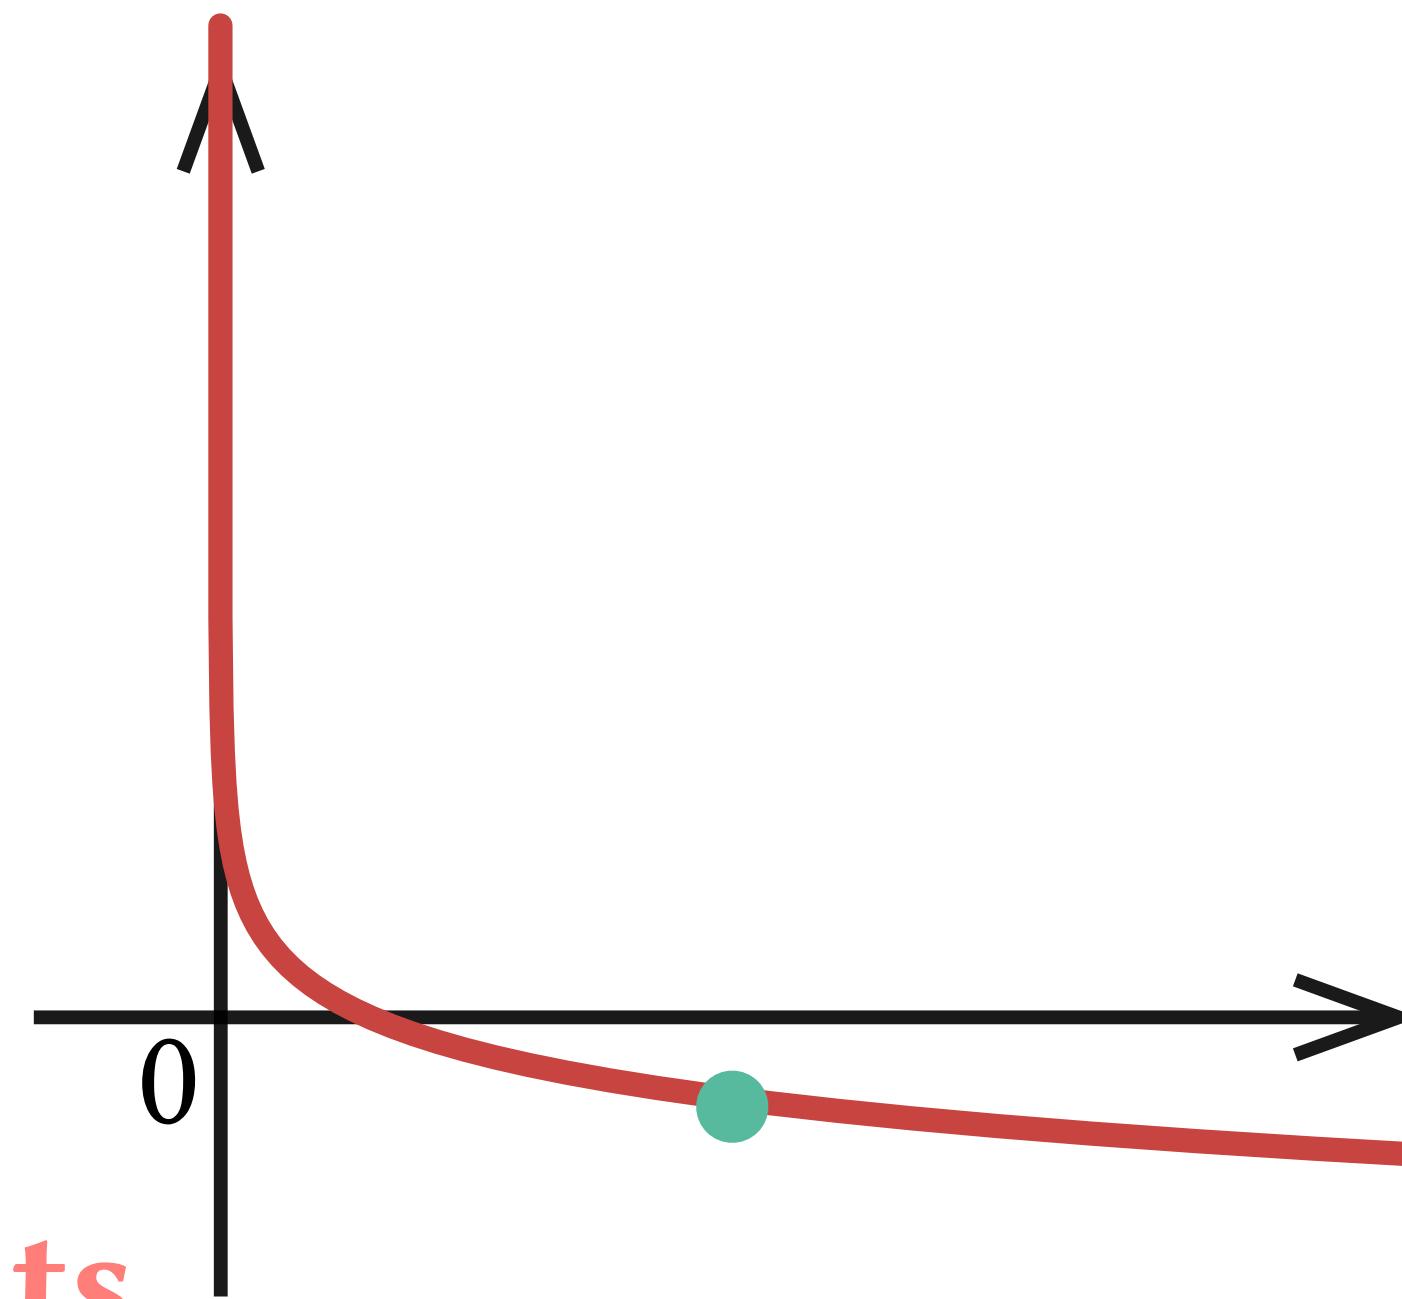
- Elastic Energy
- Optimization Basics & Classical Algorithms
- Optimization Methods for Simulation
  - Challenge 1: Handling Nonconvexity
  - Challenge 2: Nontrivial Constraints
  - Challenge 3: Large-scale Optimization

# Challenge 2: Nontrivial Constraints

---

# Nontrivial Constraints as a Barrier Function

Undefined  
when  $x < 0$



Infeasible Points

$$f(x) = -\log(x)$$

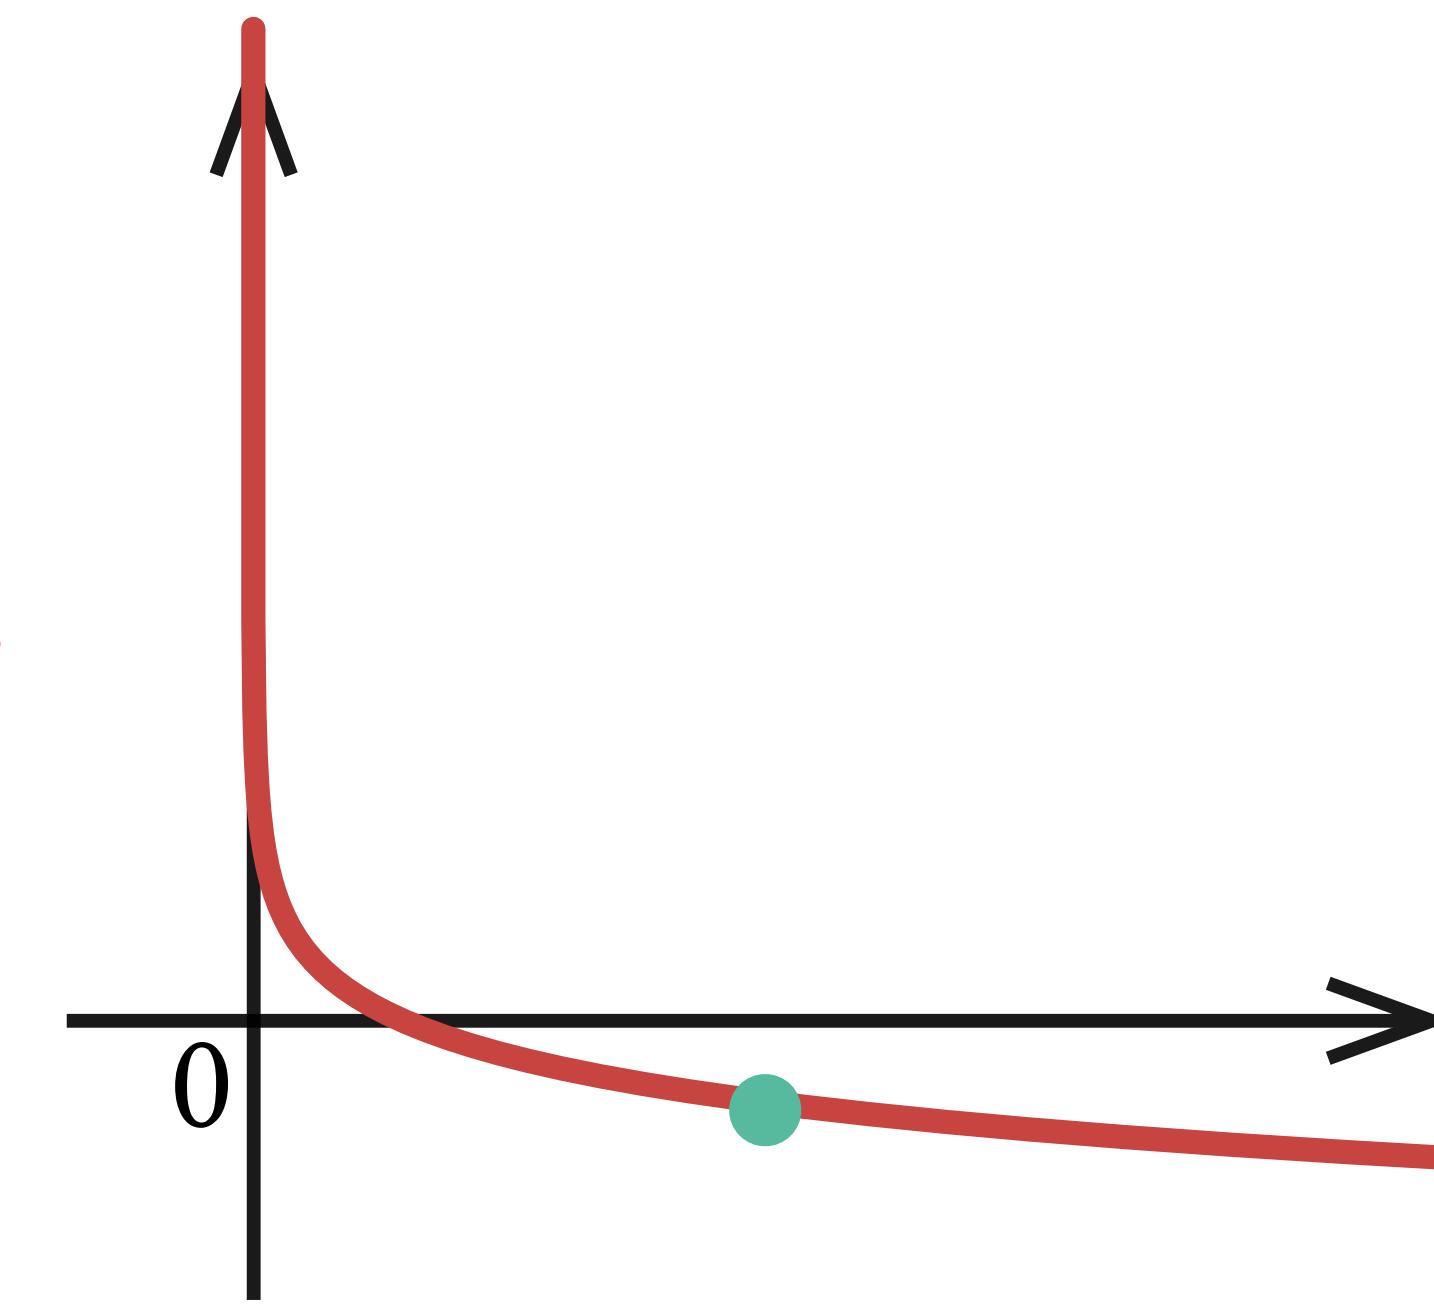
Barrier term

# Example 1: Neo-Hookean Energy

$$E_{\text{stretch}}(\text{triangle}, \text{diamond}) + E_{\text{vol}}(\text{triangle}, \text{square})$$

$$-\mu \log(\text{triangle}, \text{inverted triangle})$$

(inversion)



# Strategy 1: Remove the infeasibility in the energy

$$E_{\text{stretch}} \left( \begin{array}{c} \text{green tetrahedron} \\ , \\ \text{green parallelepiped} \end{array} \right) + E_{\text{vol}} \left( \begin{array}{c} \text{green tetrahedron} \\ , \\ \text{green cube} \end{array} \right)$$

[Smith et al. 2018]

$$\cancel{-\mu \log \left( \begin{array}{c} \text{green tetrahedron} \\ , \\ \text{orange triangle} \end{array} \right)}$$

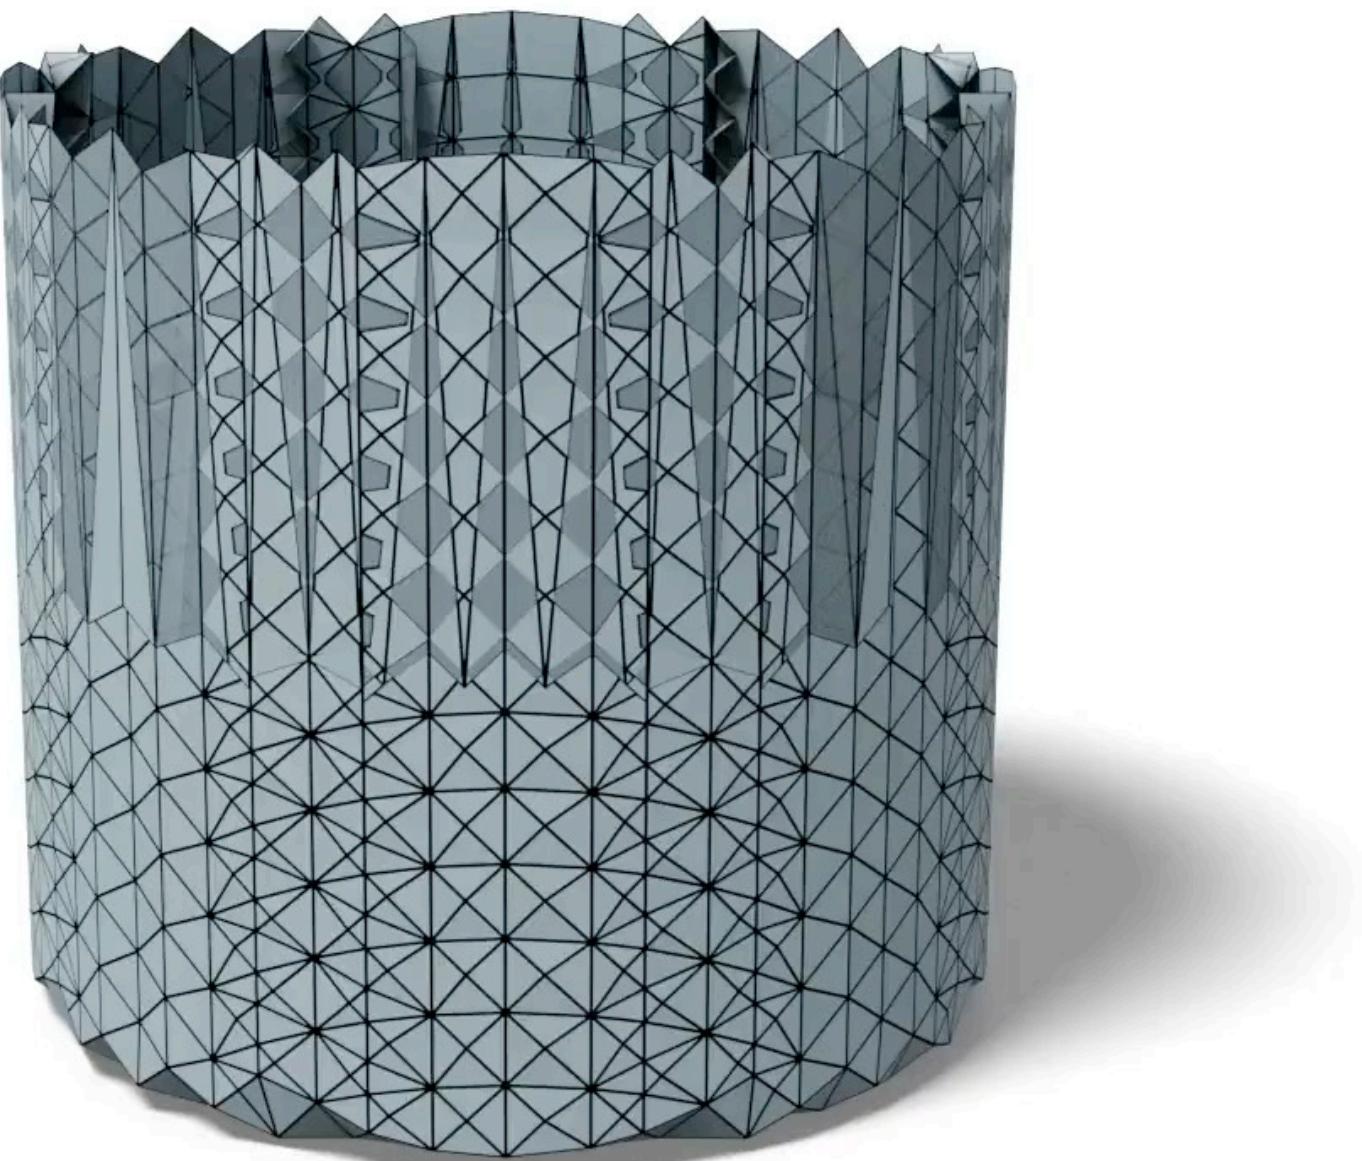


Stable Neo-Hookean Flesh Simulation.

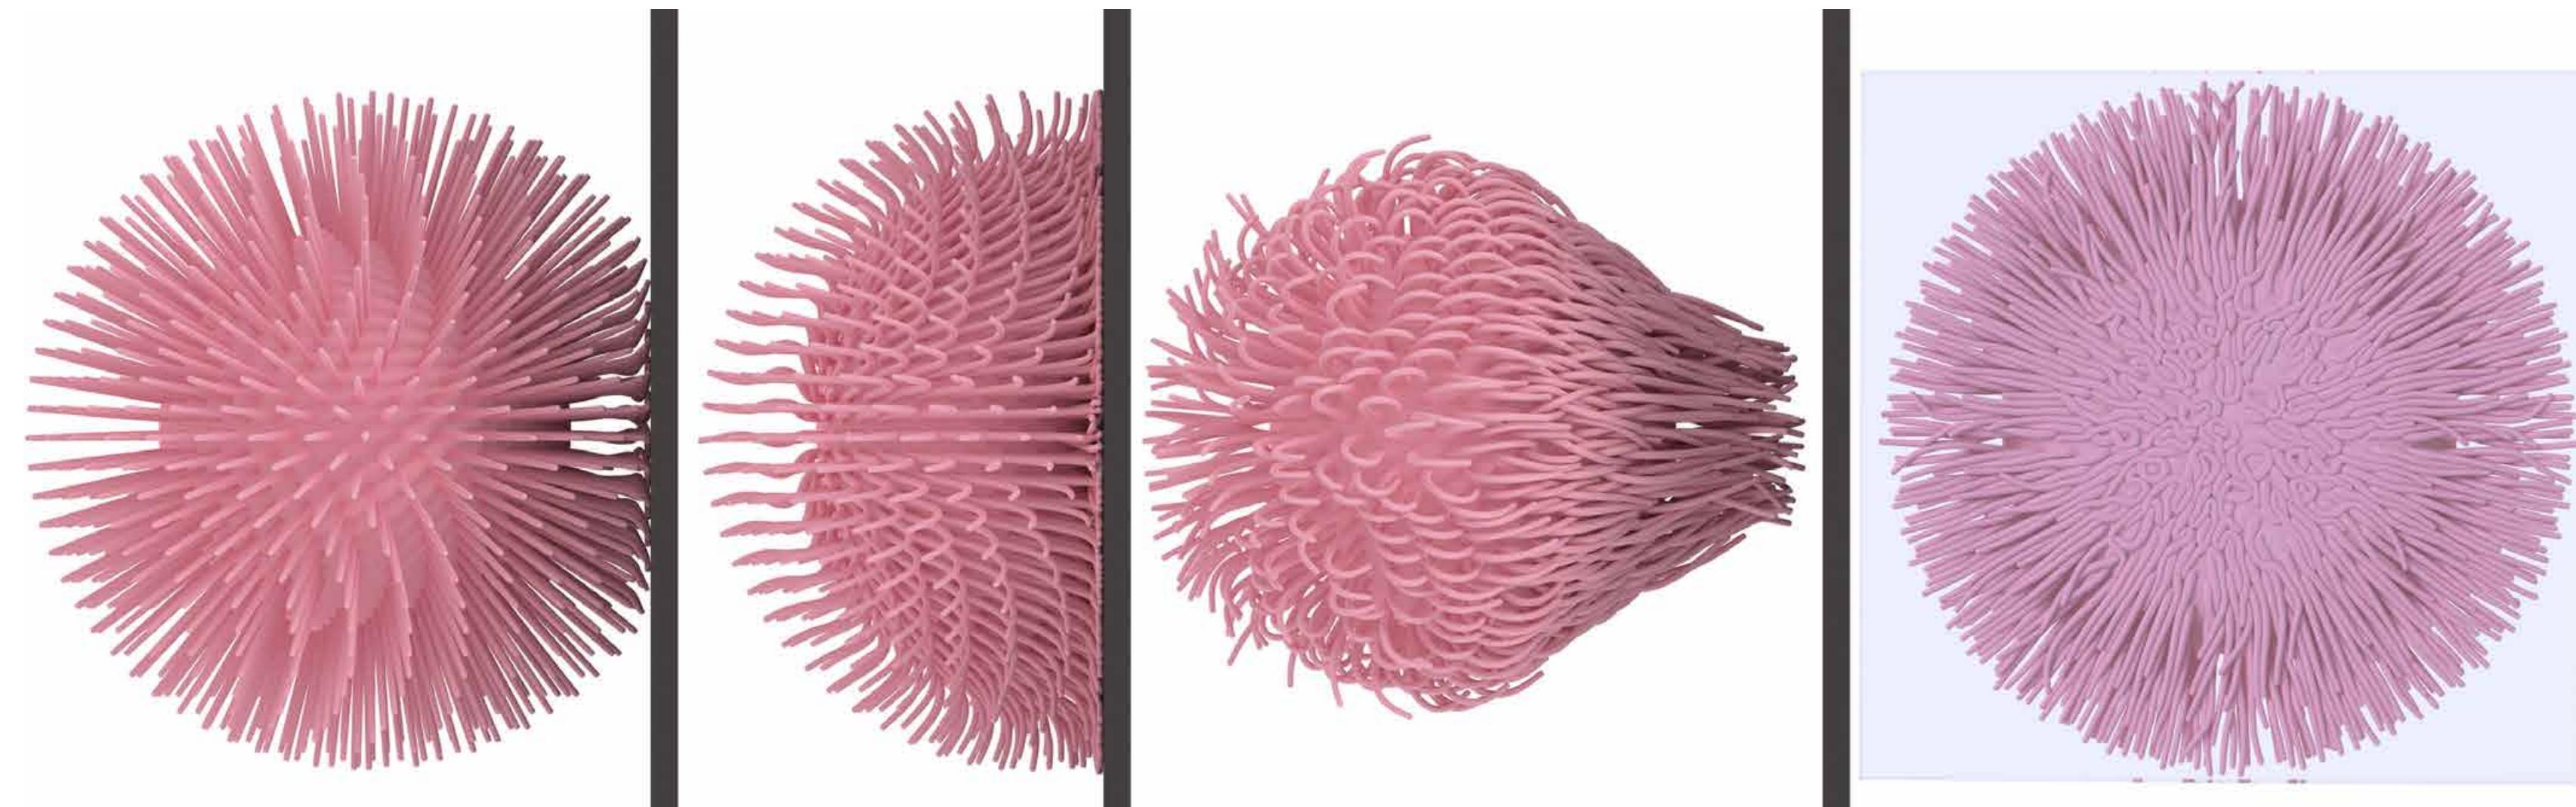
Breannan Smith, Fernando de Goes, and Theodore Kim. ACM Trans. Graph. 2018.

# Strategy 1: Remove the infeasibility in the energy

- ✓ Stable even when inversion exists before or during the optimization
- ✗ No guarantee to be inversion-free

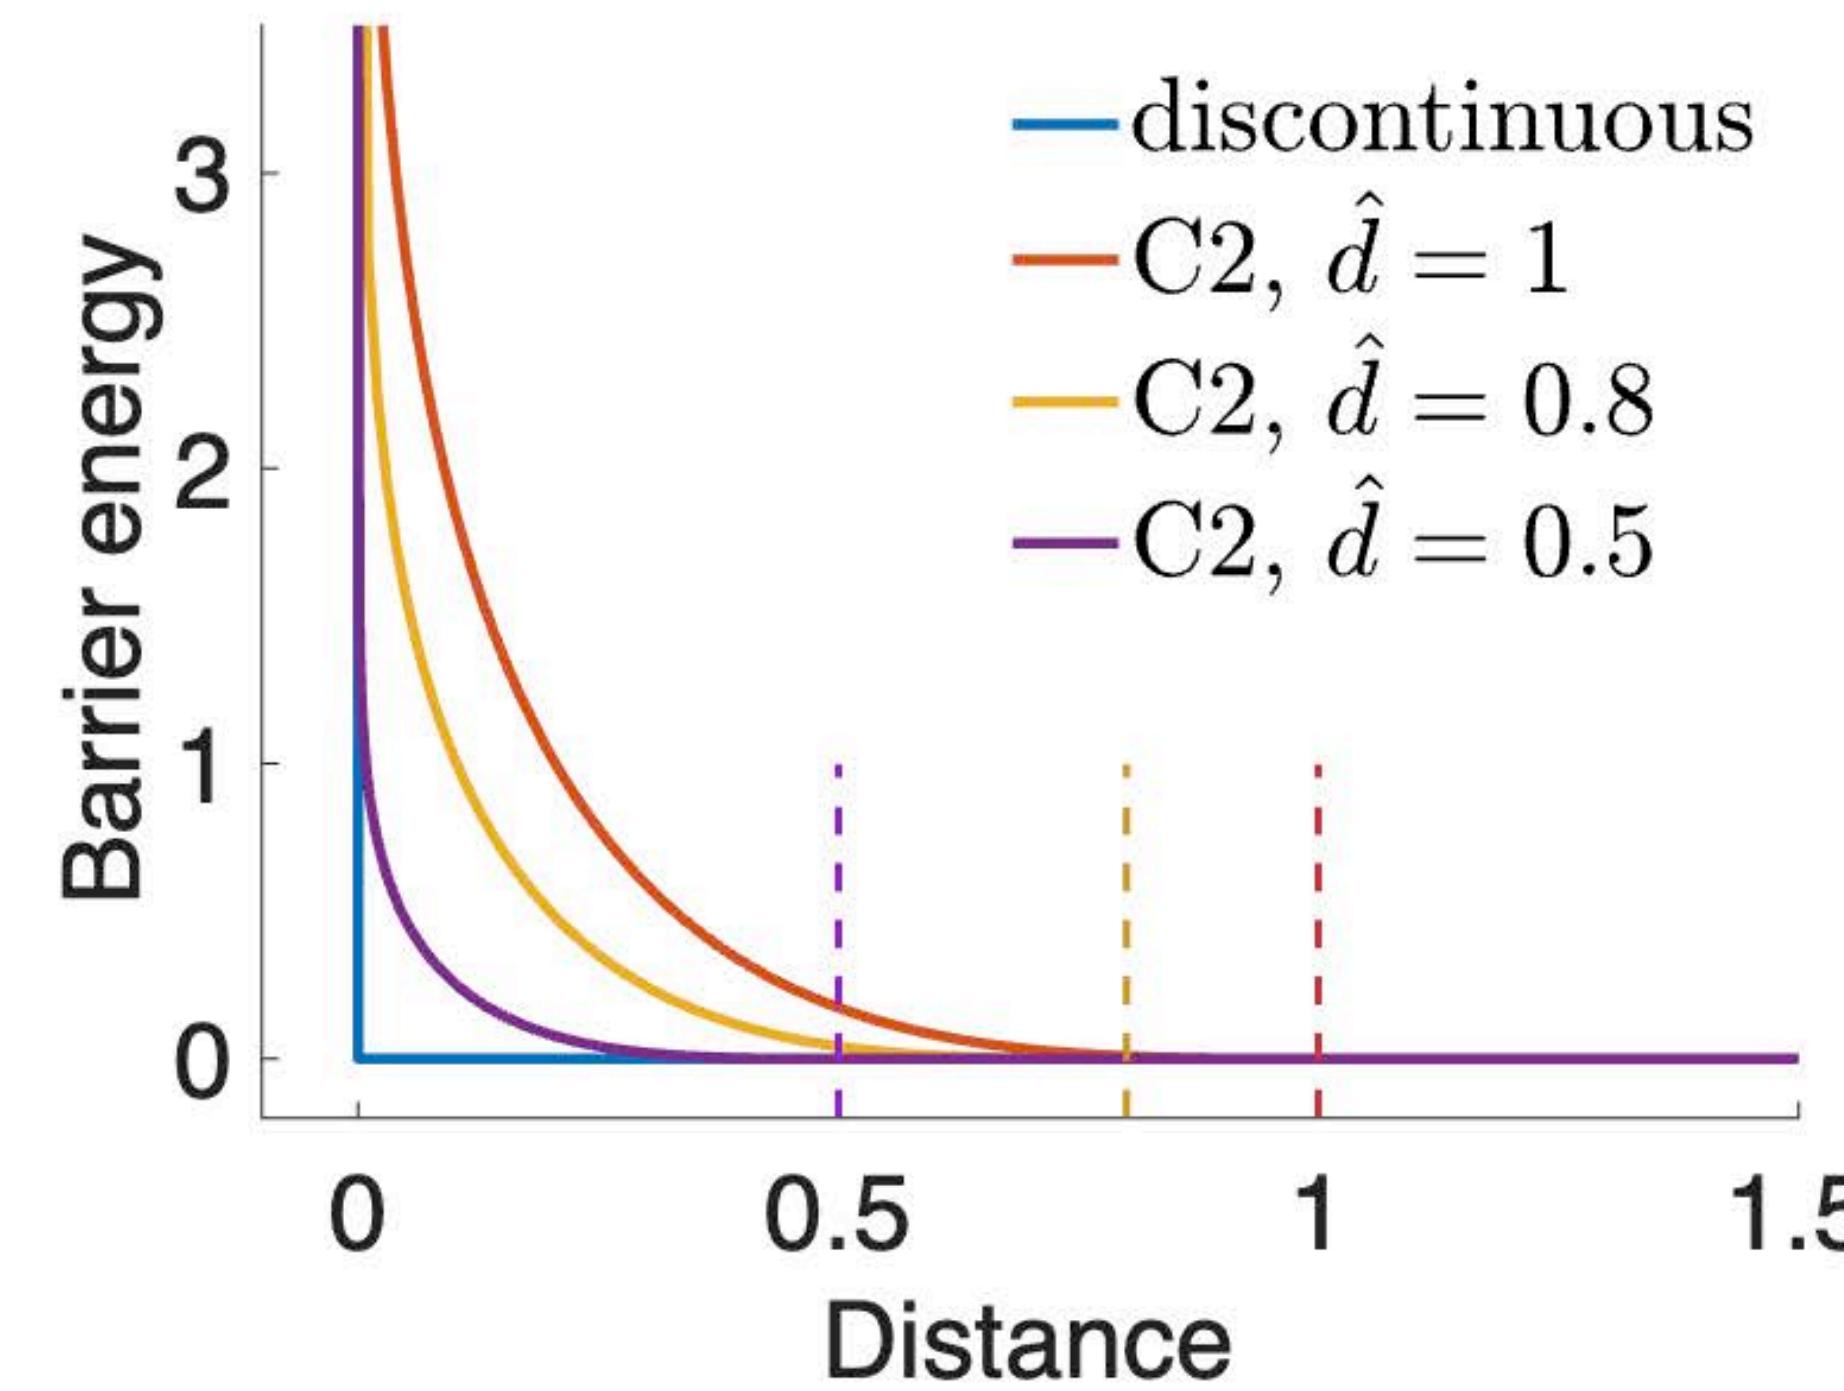


# Example 2: Incremental Potential Contact (IPC)

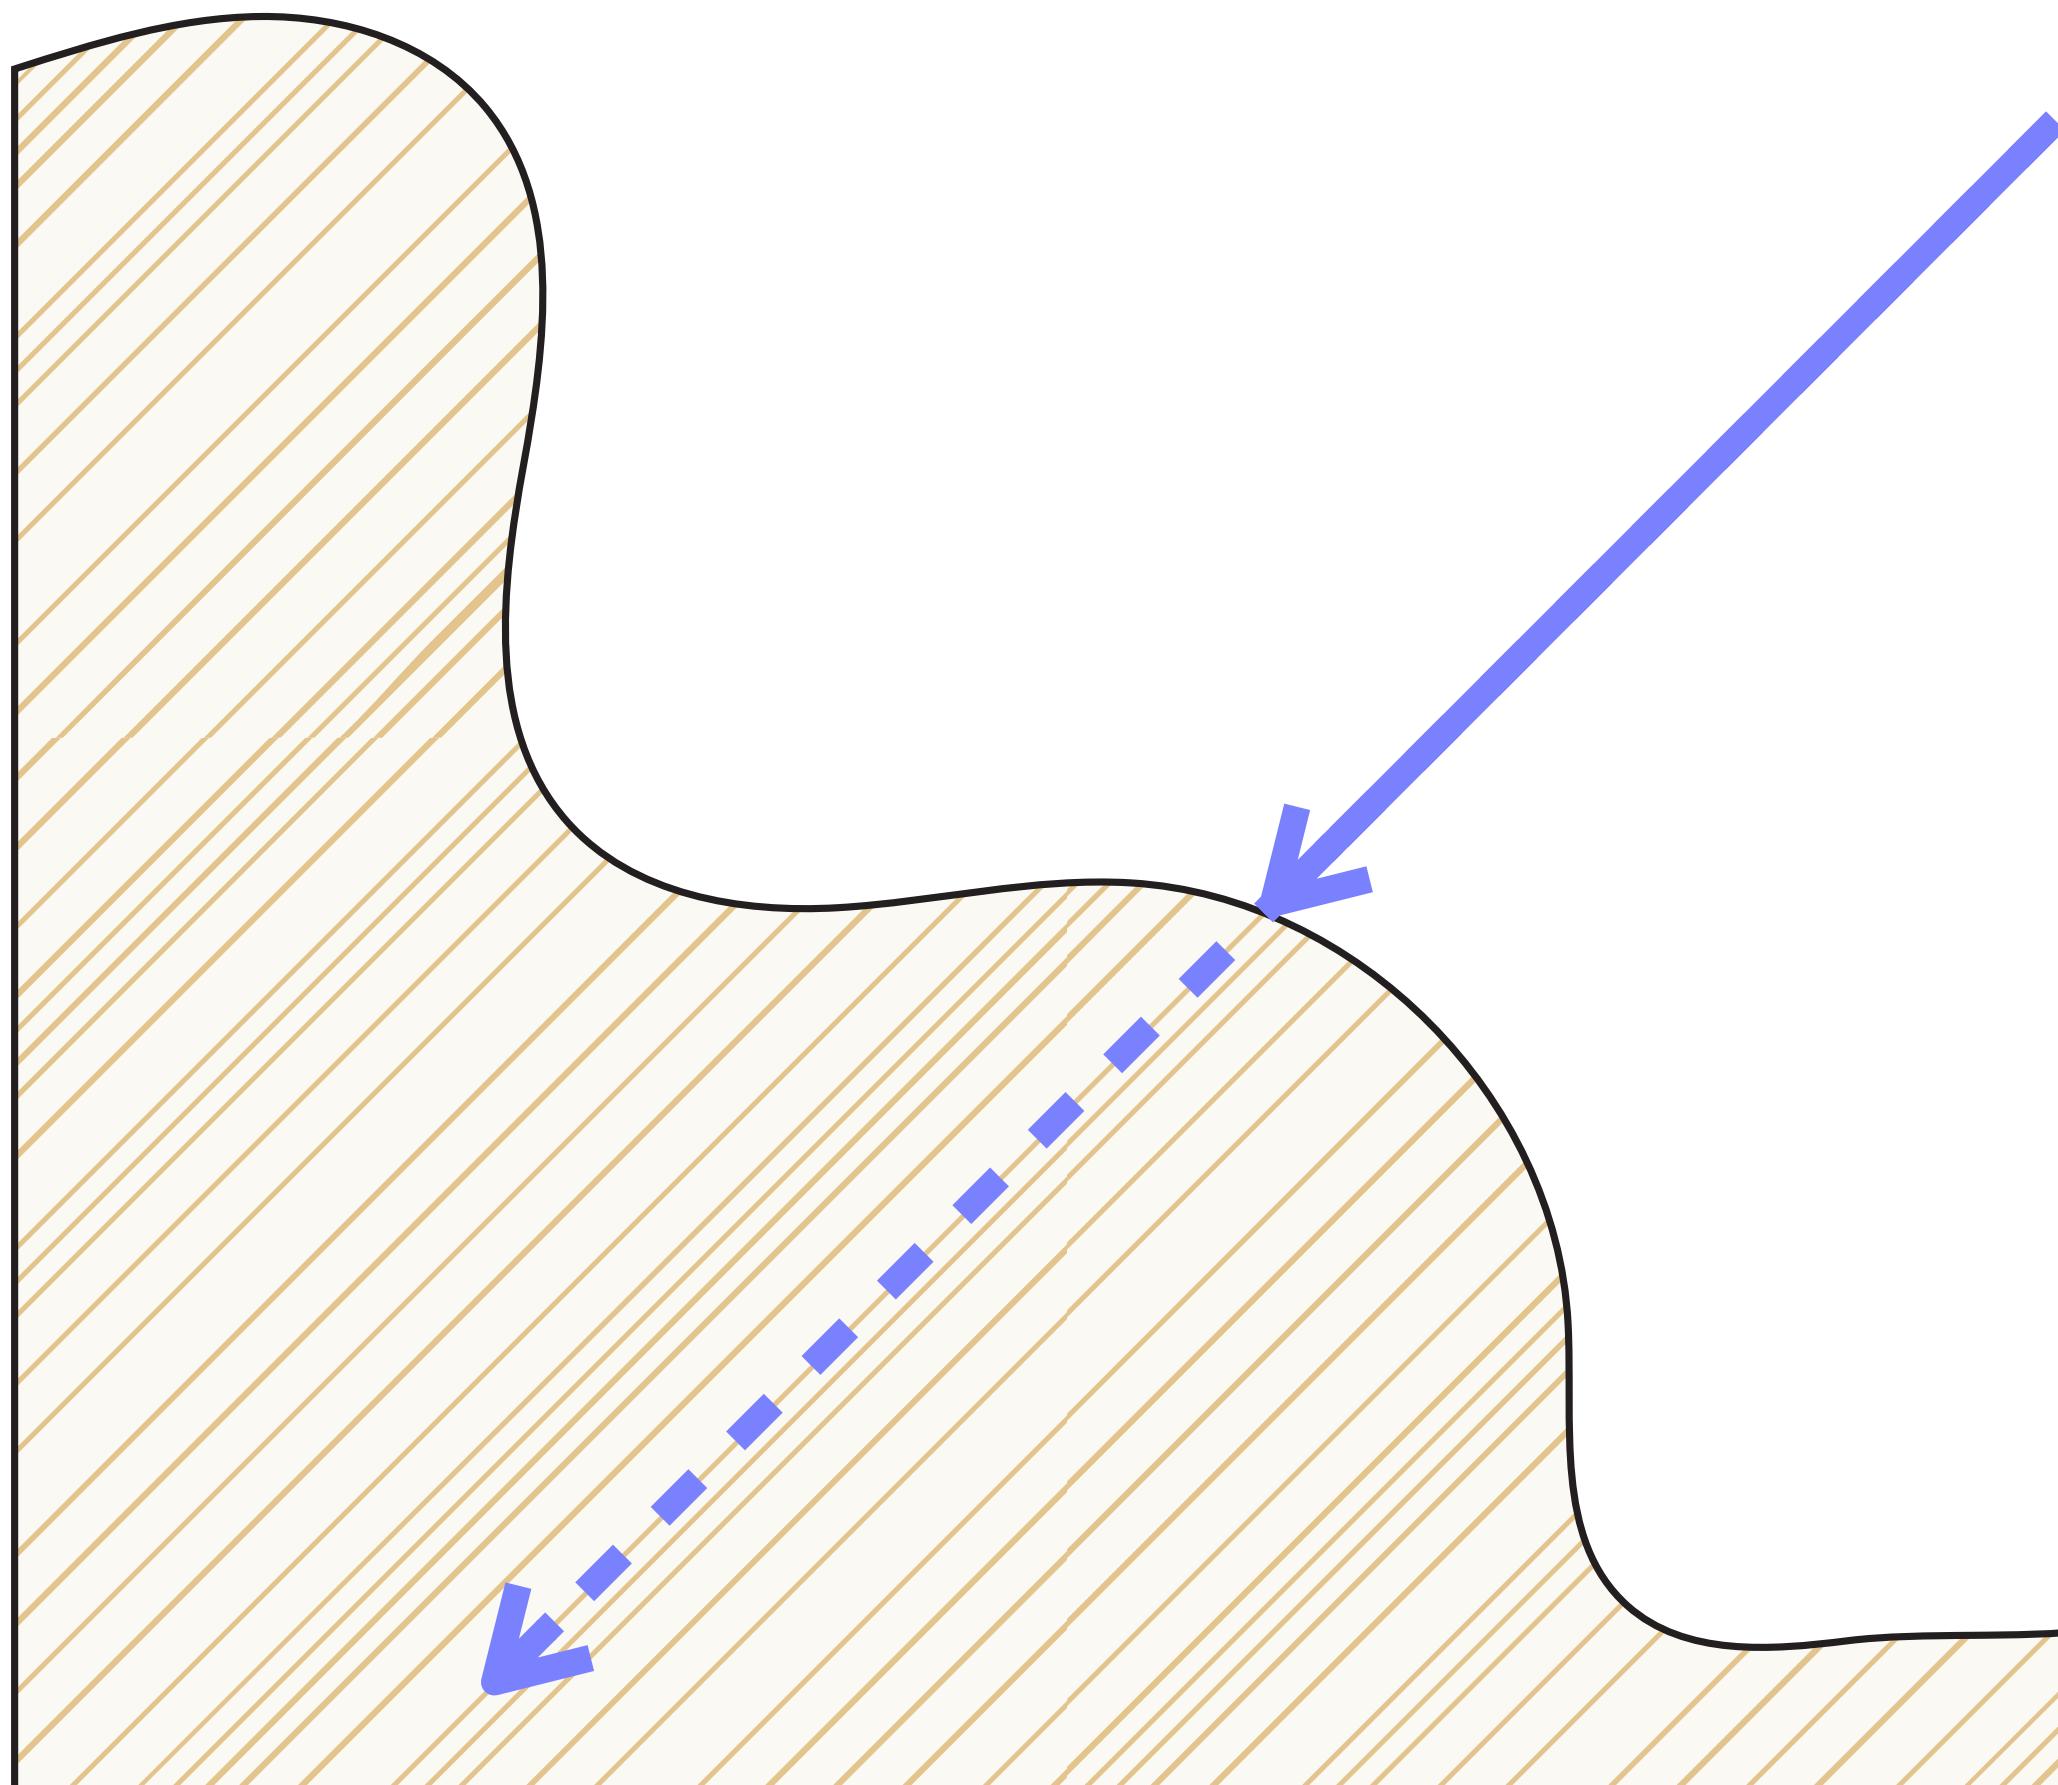


# Incremental Potential Contact

$$b(d, \hat{d}) = \begin{cases} -(d - \hat{d})^2 \ln\left(\frac{d}{\hat{d}}\right), & 0 < d < \hat{d} \\ 0 & d \geq \hat{d} \end{cases}$$



# Strategy 2: Avoid going to infeasible regions



**feasibility-aware line search**

- ✓ **Guarantee to satisfy the constraints**
- ✗ **Line search is more complex and may clamp down the step size too much**

- Elastic Energy
- Optimization Basics & Classical Algorithms
- Optimization Methods for Simulation
  - Challenge 1 ✨: Handling Nonconvexity
  - Challenge 2 ✨: Nontrivial Constraints
  - Challenge 3 ✨: Large-scale Optimization

# Challenge 3: Large-scale Optimization

---

# Large-Scale Optimization

$$\mathbf{d} = -\mathbf{H}^{-1}\mathbf{g}$$



$3N \times 3N$

# Total Computational Cost

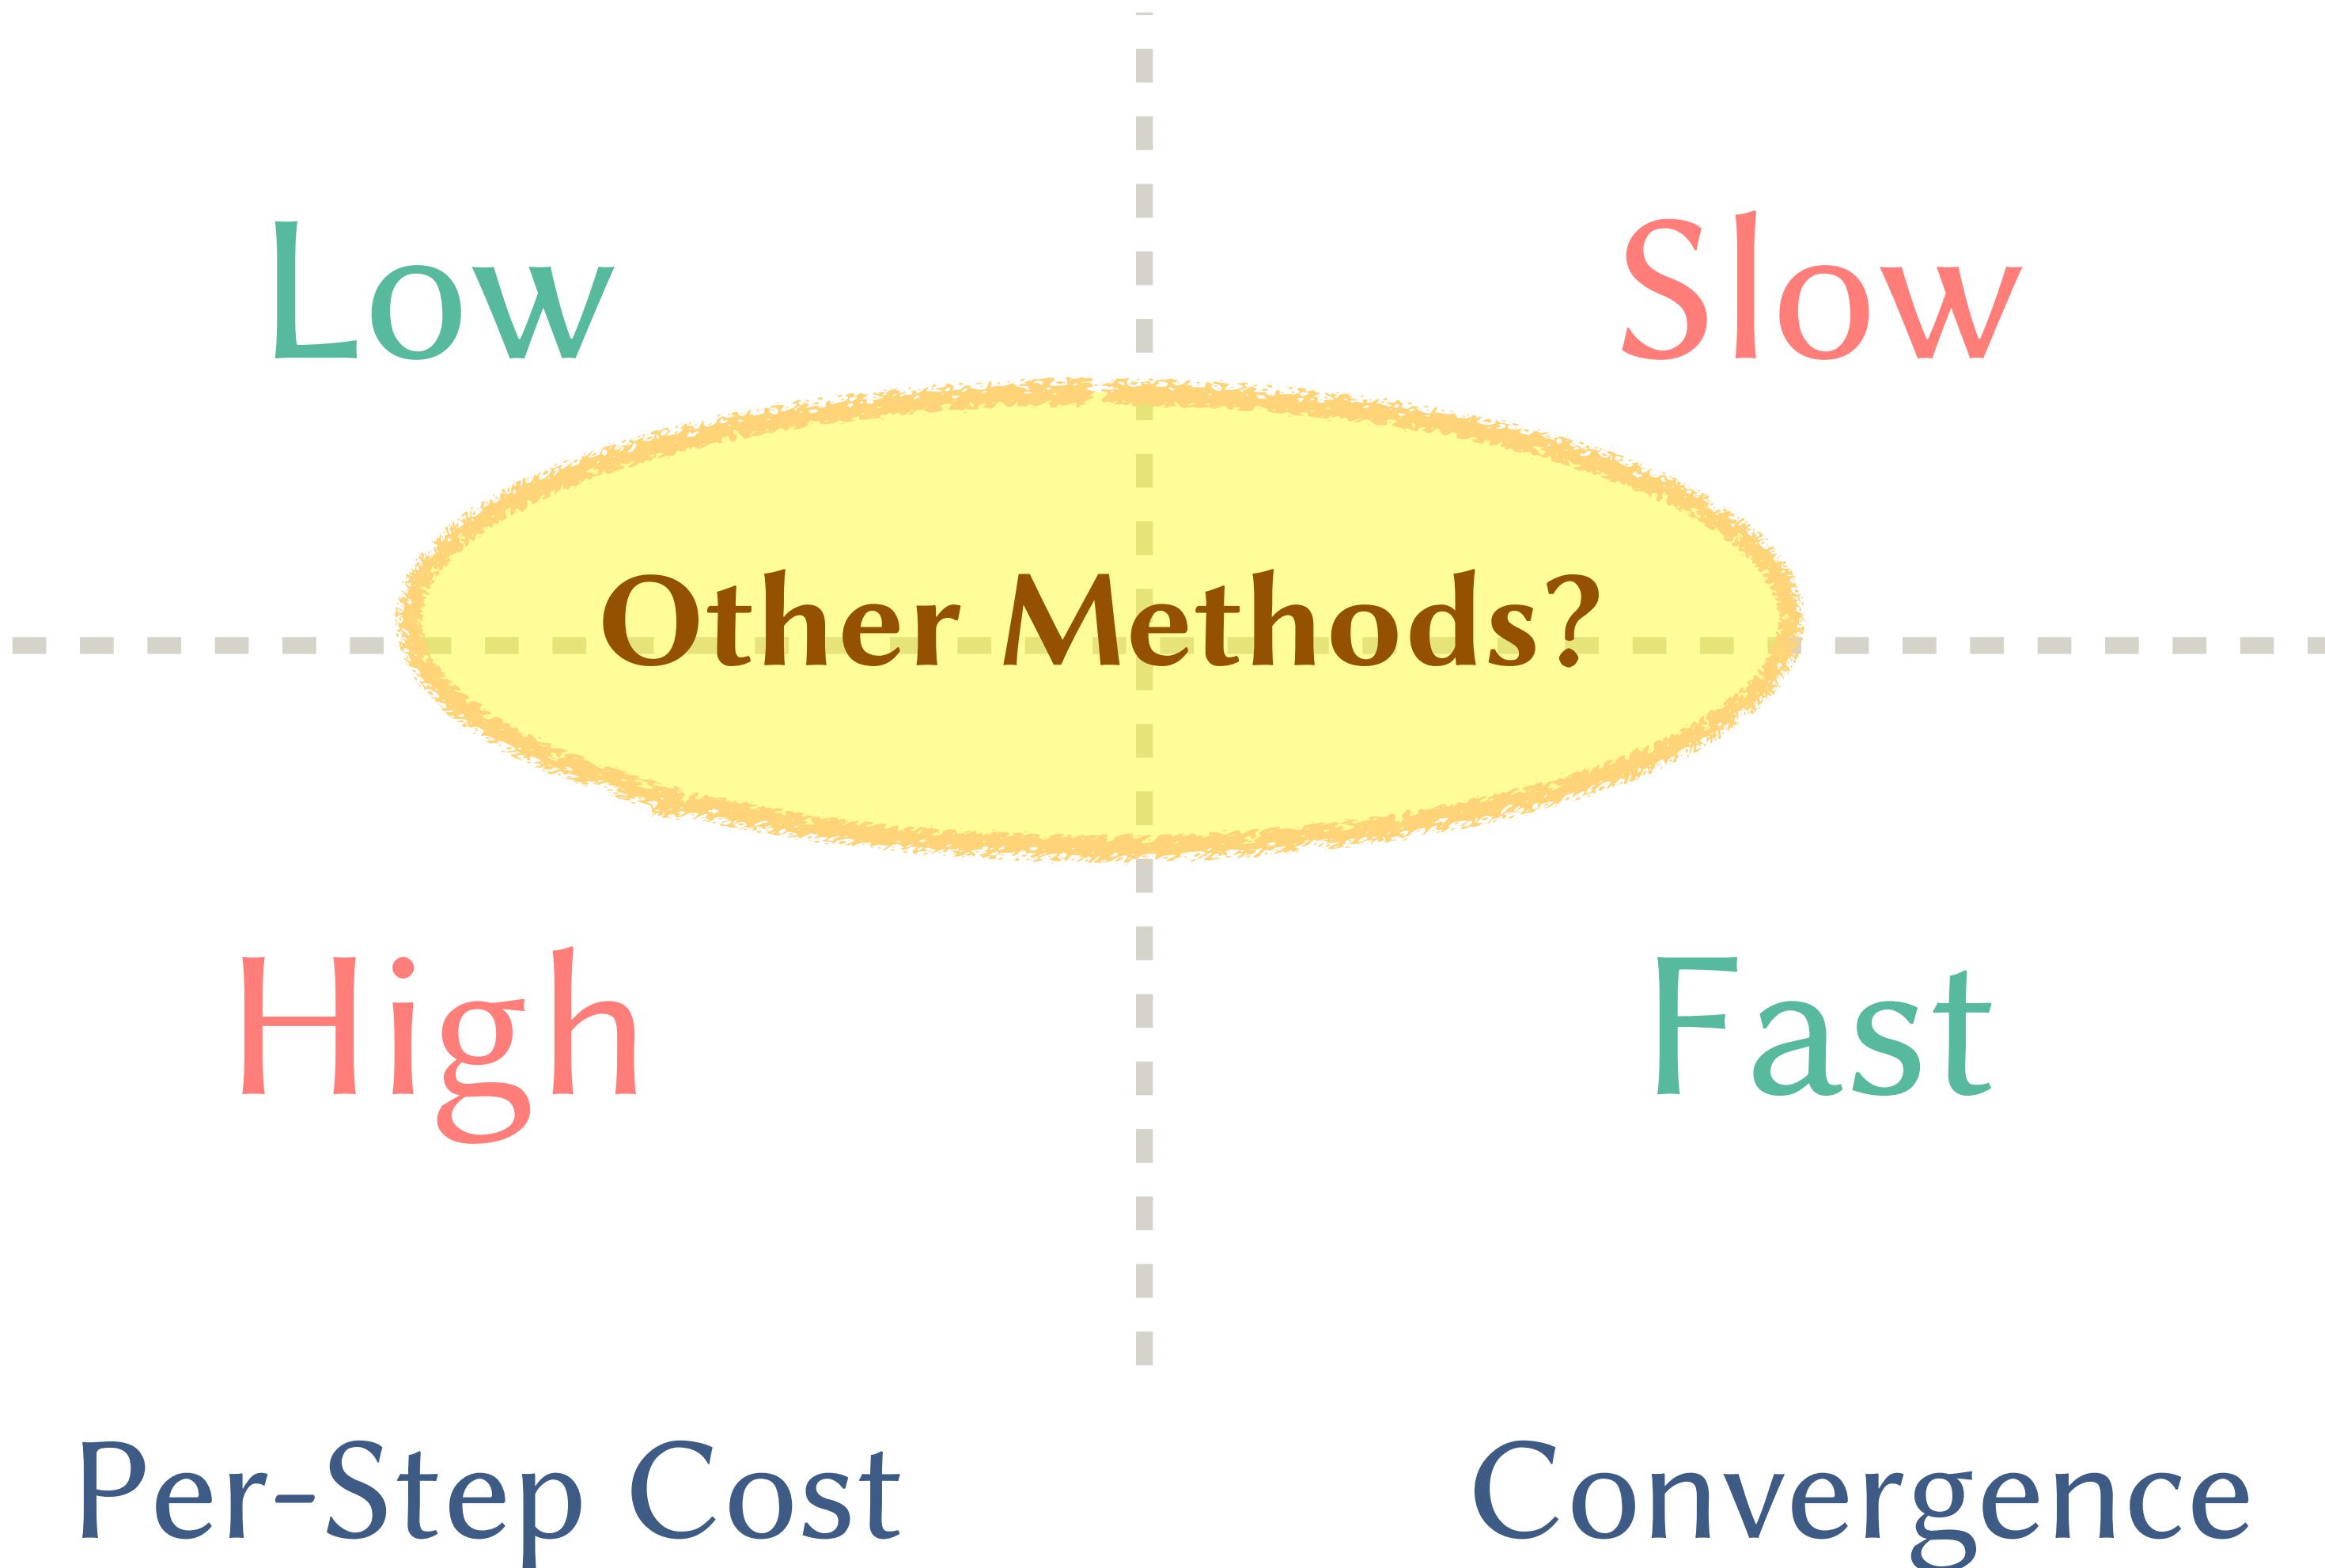
Total Cost =

Per-step Cost  $\times$  Iterations

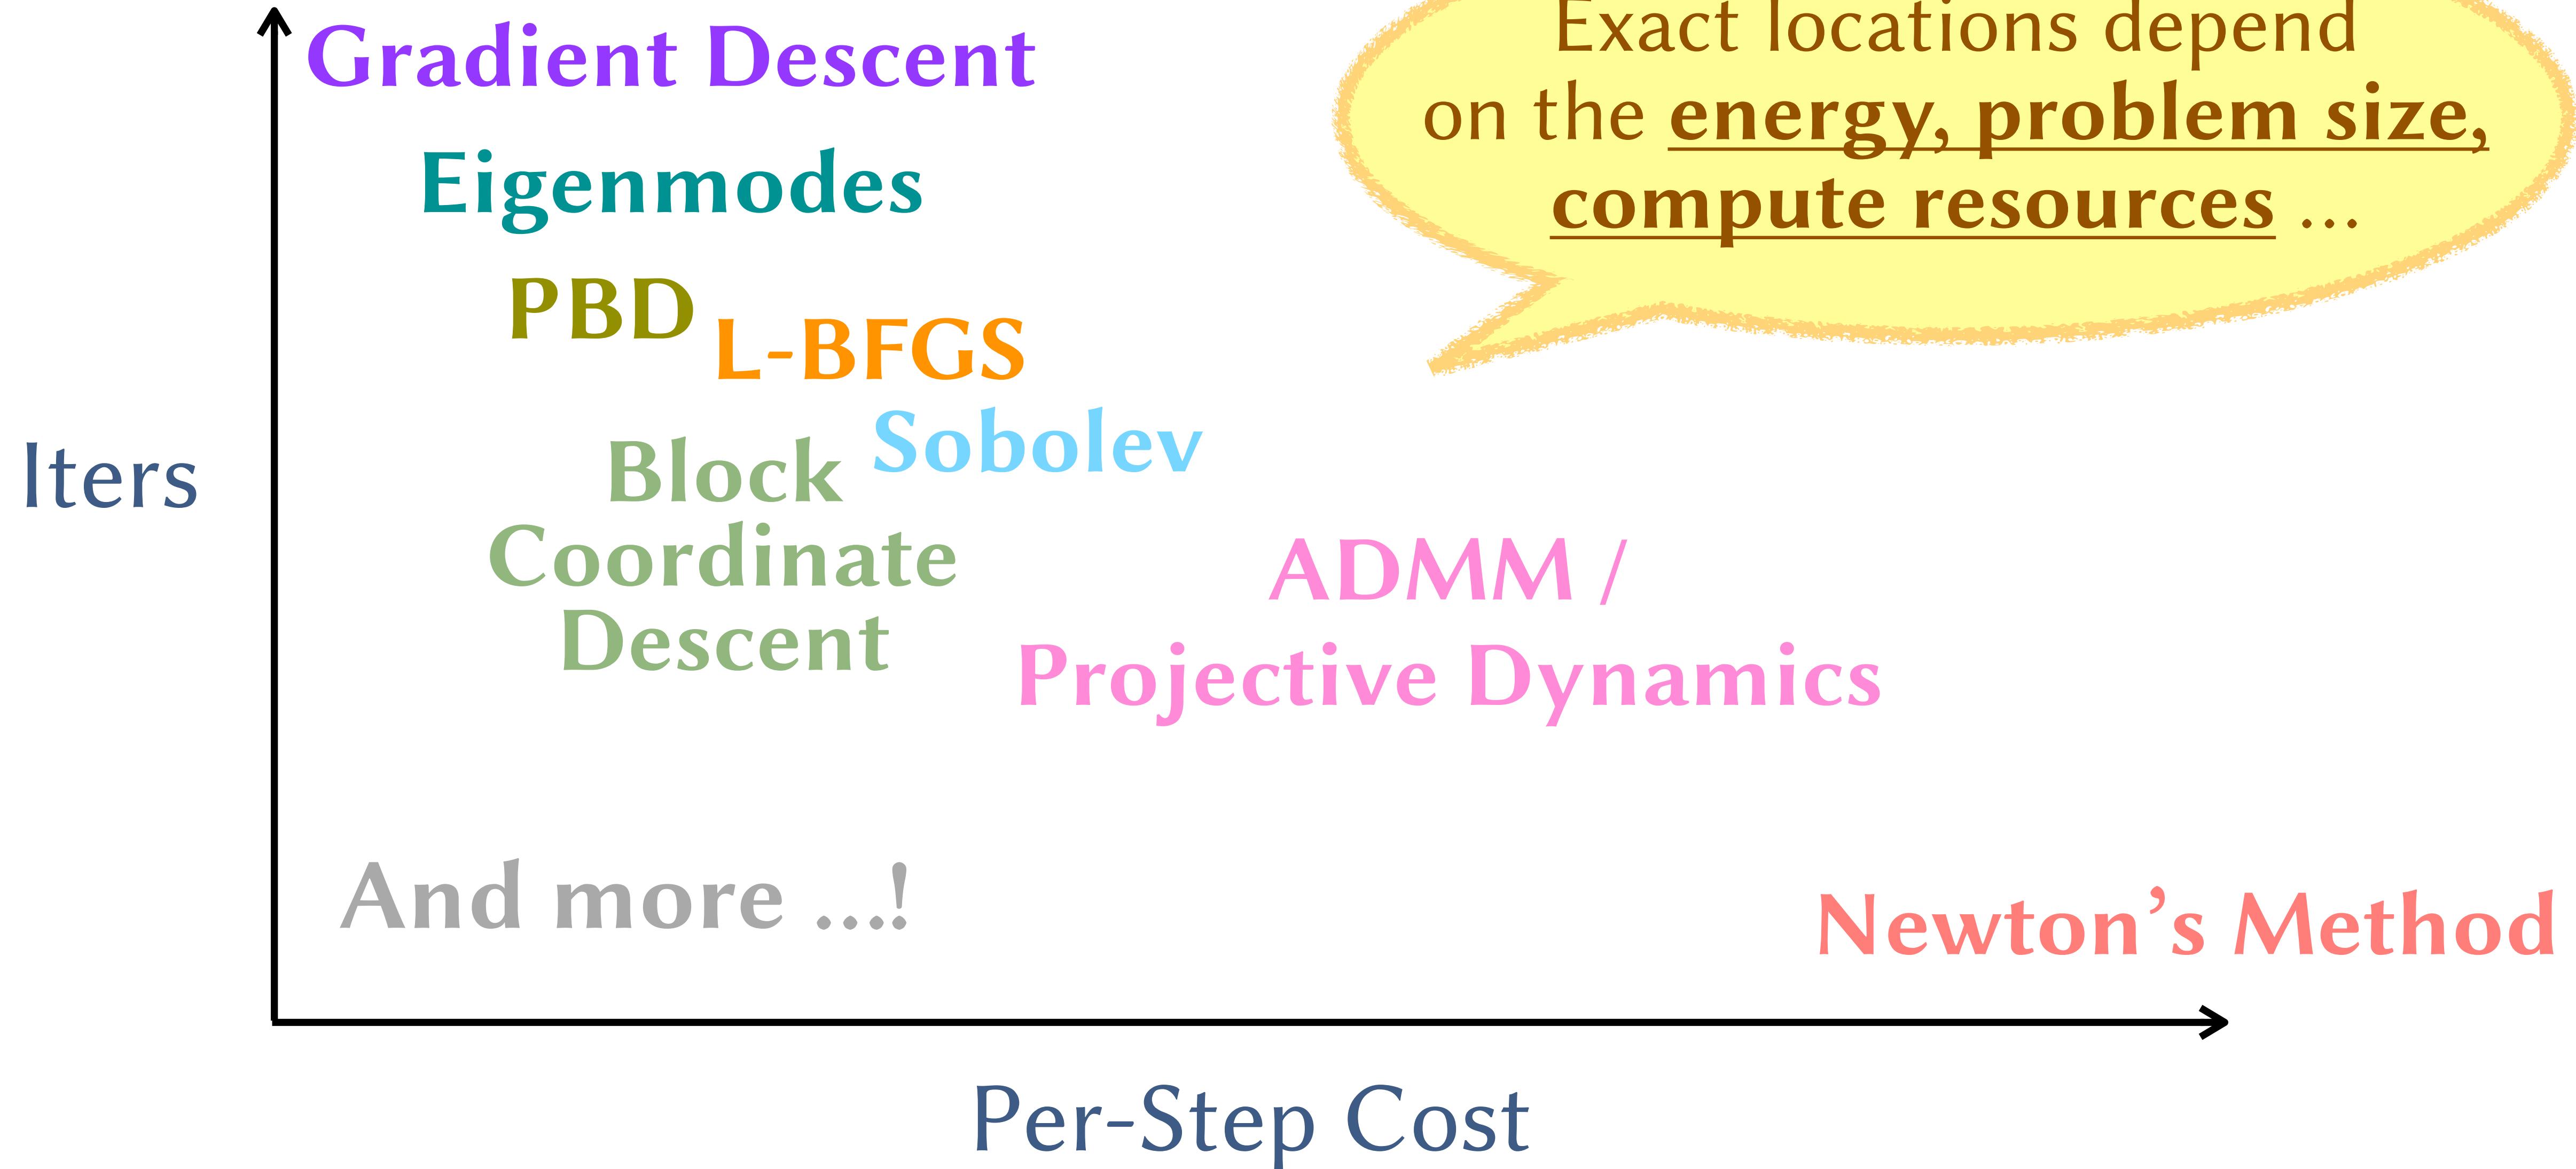
$$\text{Total Cost} = \text{Per-step Cost} \times \text{Iterations}$$

Gradient Descent  
(First-order)

Newton's Method  
(Second-order)



# An Inaccurate Map



# Overall Goal

$$\mathbf{d} = -\mathbf{H}^{-1}\mathbf{g}$$

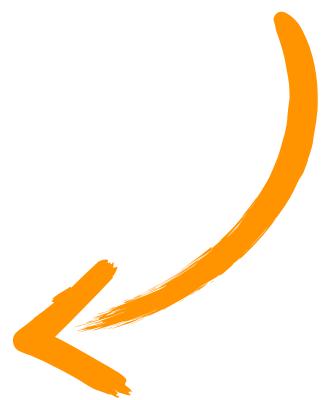


Simpler Hessian

Better gradient

# Strategy 1: Simplify H

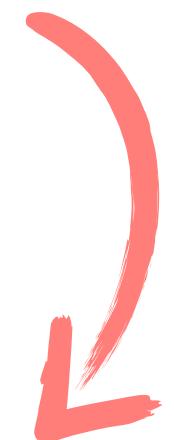
$$\mathbf{d} = -\mathbf{H}^{-1}\mathbf{g}$$



Simpler Hessian

# Steepest Descent

$$\mathbf{d} = -\boxed{\mathbf{H}^{-1}} \mathbf{g}$$



$$\mathbf{d} = -\boxed{\mathbf{I}^{-1}} \mathbf{g}$$

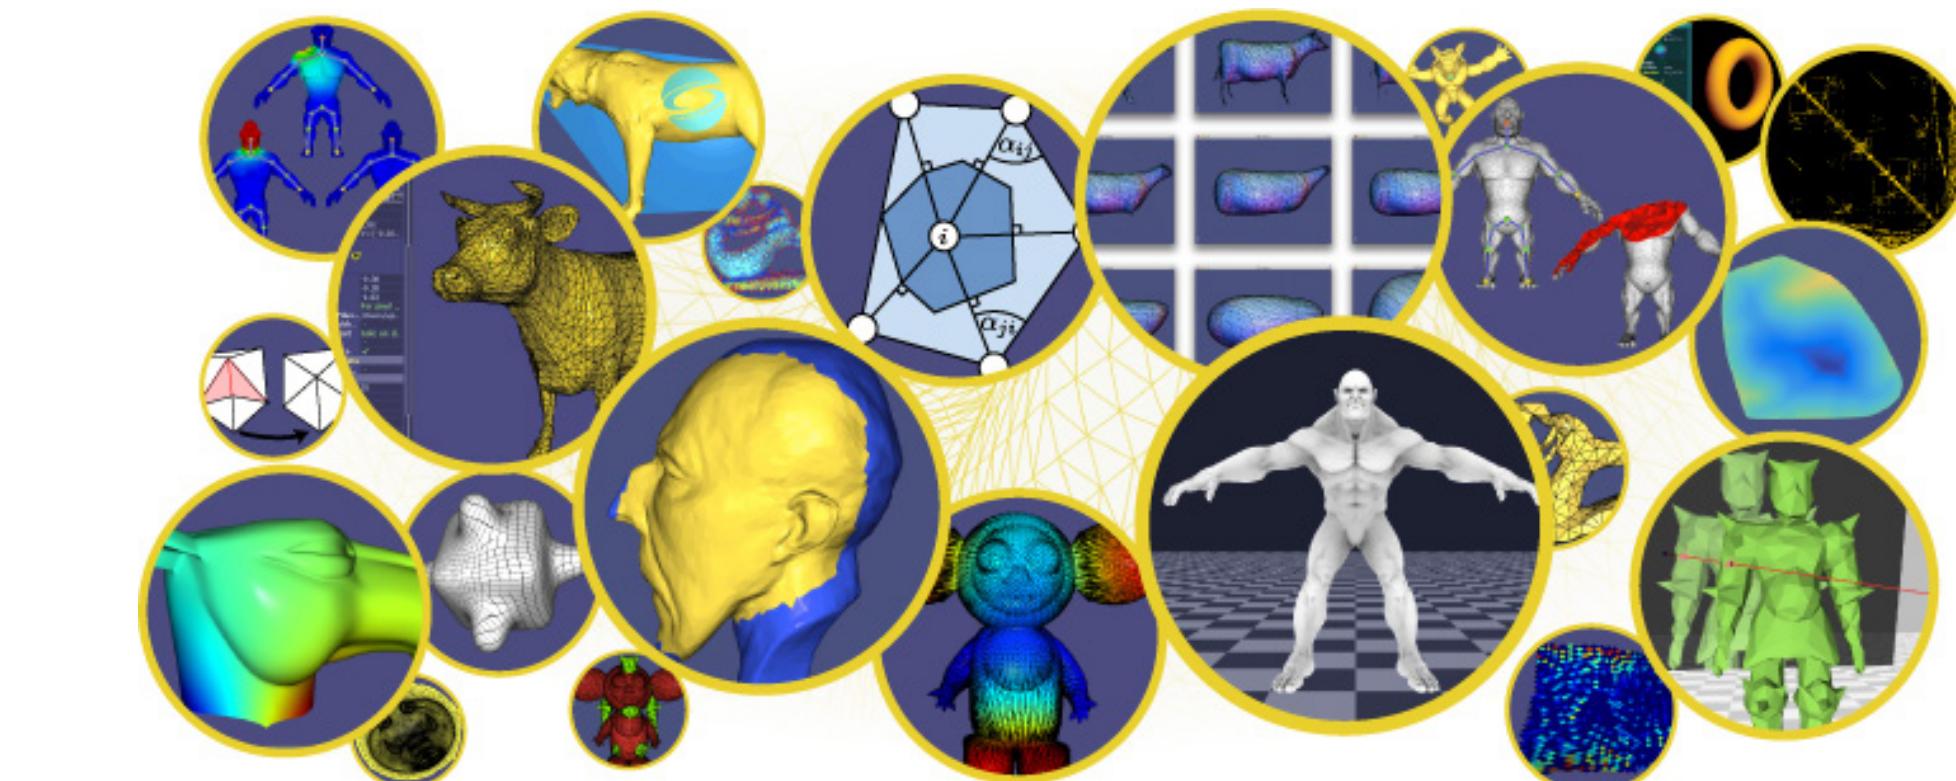
(Gradient Descent)

# Sobolev-Preconditioned Gradient Descent

$$\mathbf{d} = -\mathbf{L}^{-1}\mathbf{g}$$



Gpytoolbox (python)

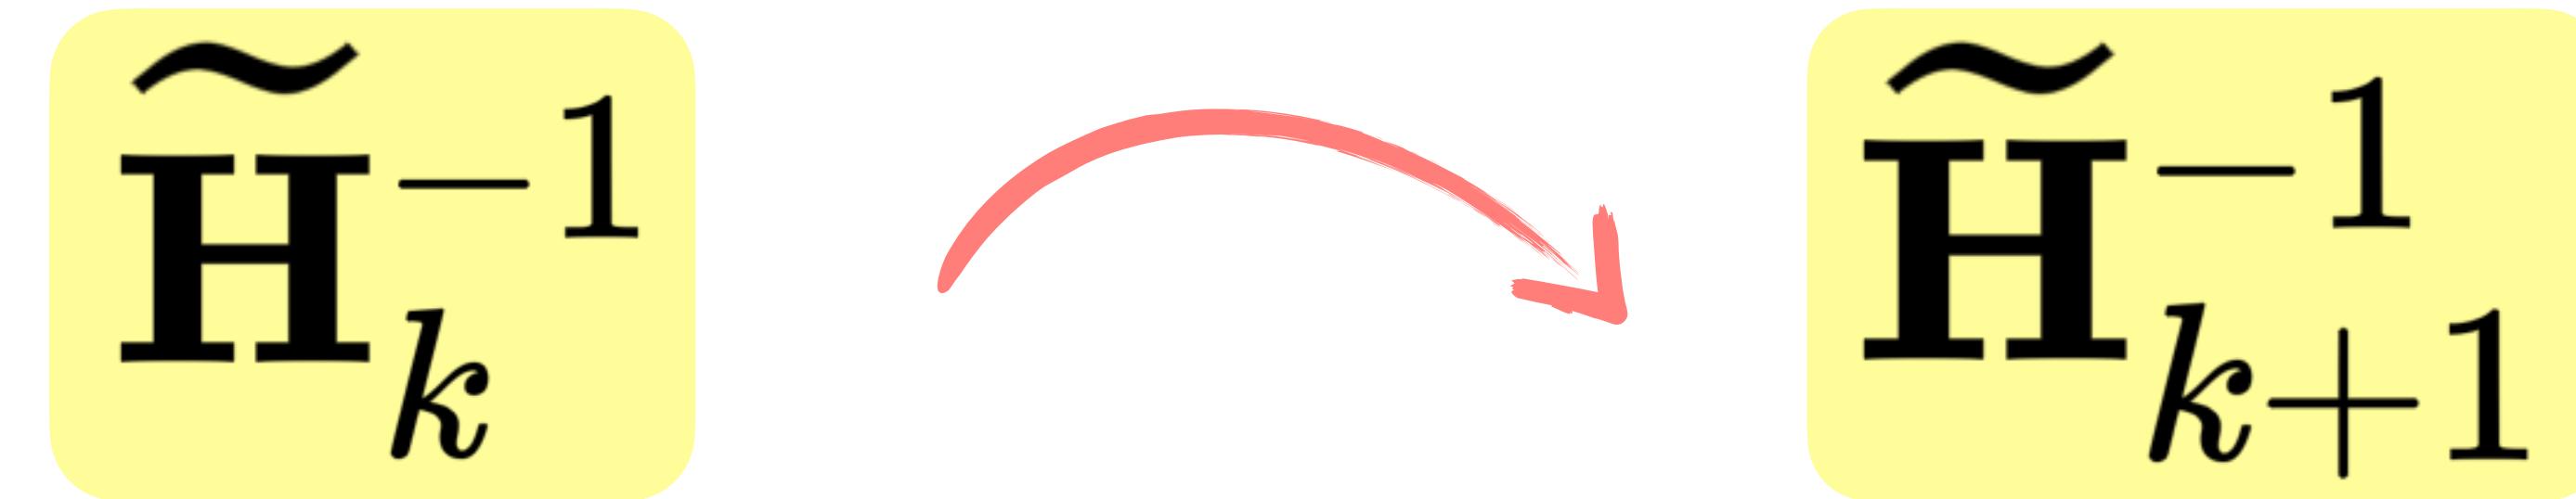


libigl (C++)

# Limited-memory BFGS (L-BFGS)

$$\mathbf{d} = -\mathbf{H}^{-1}\mathbf{g}$$

Low-rank Update



Previous Approximation  
(Known)

# BFGS

(Rank-2 Update)

$$\mathbf{H}_{k+1}^{-1} = \left( \mathbf{I} - \frac{\mathbf{s}_k \mathbf{y}_k^\top}{\mathbf{y}_k^\top \mathbf{s}_k} \right) \mathbf{H}_k^{-1} \left( \mathbf{I} - \frac{\mathbf{y}_k \mathbf{s}_k^\top}{\mathbf{y}_k^\top \mathbf{s}_k} \right) + \frac{\mathbf{s}_k \mathbf{s}_k^\top}{\mathbf{y}_k^\top \mathbf{s}_k}$$

$$\mathbf{y}_k = \mathbf{g}_{k+1} - \mathbf{g}_k$$

$$\mathbf{s}_k = \mathbf{x}_{k+1} - \mathbf{x}_k$$

# L-BFGS

Use only the last  $m$  pairs of  $\{\mathbf{y}_k, \mathbf{s}_k\}$  and  $\mathbf{H}_0^{-1}$



# Sobolev Gradient 🤝 L-BFGS

**Blended Cured Quasi-Newton for Distortion Optimization**

YUFENG ZHU, University of British Columbia & Adobe Research  
ROBERT BRIDSON, Autodesk & University of British Columbia  
DANNY M. KAUFMAN, Adobe Research

Fig. 1. **Twisting.** A stress-test 3D deformation problem. Left: we initialize a 1.5M element tetrahedral mesh bar with a straight rest shape into a tightly twisted coil, constraining both ends to stay fixed. Right: minimizing the ISO deformation energy to find a constrained equilibrium with (top to bottom) Projected Newton (PN), Accelerated Quadratic Proxy (AQP) and our Blended Cured quasi-Newton (BCQN) method, we show intermediate shapes at reported wall-clock time (seconds) and iteration counts at those times (BCQN/AQP/PN). AQP, much slower than BCQN, requires many more iterations to converge while PN, despite requiring fewer iterations, is well over 25X slower due to per-iteration costs dominated by factorization.

Optimizing distortion energies over a mesh, in two or three dimensions, is a common and critical problem in physical simulation and geometry processing. We present three new improvements to the state of the art: a barrier-aware line-search filter that cures blocked descent steps due to element barrier terms and so enables rapid progress; an energy proxy model that adaptively blends the Sobolev (inverse-Laplacian-processed) gradient and L-BFGS descent to gain the advantages of both, while avoiding L-BFGS's current limitations in distortion optimization tasks; and a characteristic gradient norm providing a robust and largely mesh- and energy-independent convergence criterion that avoids wrongful termination when algorithms temporarily slow their progress. Together these improvements form the basis for Blended Cured Quasi-Newton (BCQN), a new distortion optimization algorithm. Over a wide range of problems over all scales we show that BCQN is generally the fastest and most robust method available, making some previously intractable problems practical while offering up to an order of magnitude improvement in others.

CCS Concepts: • Mathematics of computing → Continuous optimization; • Computing methodologies → Shape modeling;

Additional Key Words and Phrases: numerical optimization, geometry optimization, distortion, deformation, elasticity, simulation, preconditioning, fast solvers

Authors' addresses: Yufeng Zhu, University of British Columbia & Adobe Research; Robert Bridson, Autodesk & University of British Columbia; Danny M. Kaufman, Adobe Research.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](mailto:permissions@acm.org).  
© 2018 Association for Computing Machinery.  
0730-0301/2018/8-ART40 \$15.00  
<https://doi.org/10.1145/3197517.3201359>

ACM Reference Format:  
Yufeng Zhu, Robert Bridson, and Danny M. Kaufman. 2018. Blended Cured Quasi-Newton for Distortion Optimization. *ACM Trans. Graph.* 37, 4, Article 40 (August 2018), 14 pages. <https://doi.org/10.1145/3197517.3201359>

ACM Trans. Graph., Vol. 37, No. 4, Article 40. Publication date: August 2018.

## Low-rank update

$$\{\mathbf{z}_k, \mathbf{s}_k\} \text{ and }$$

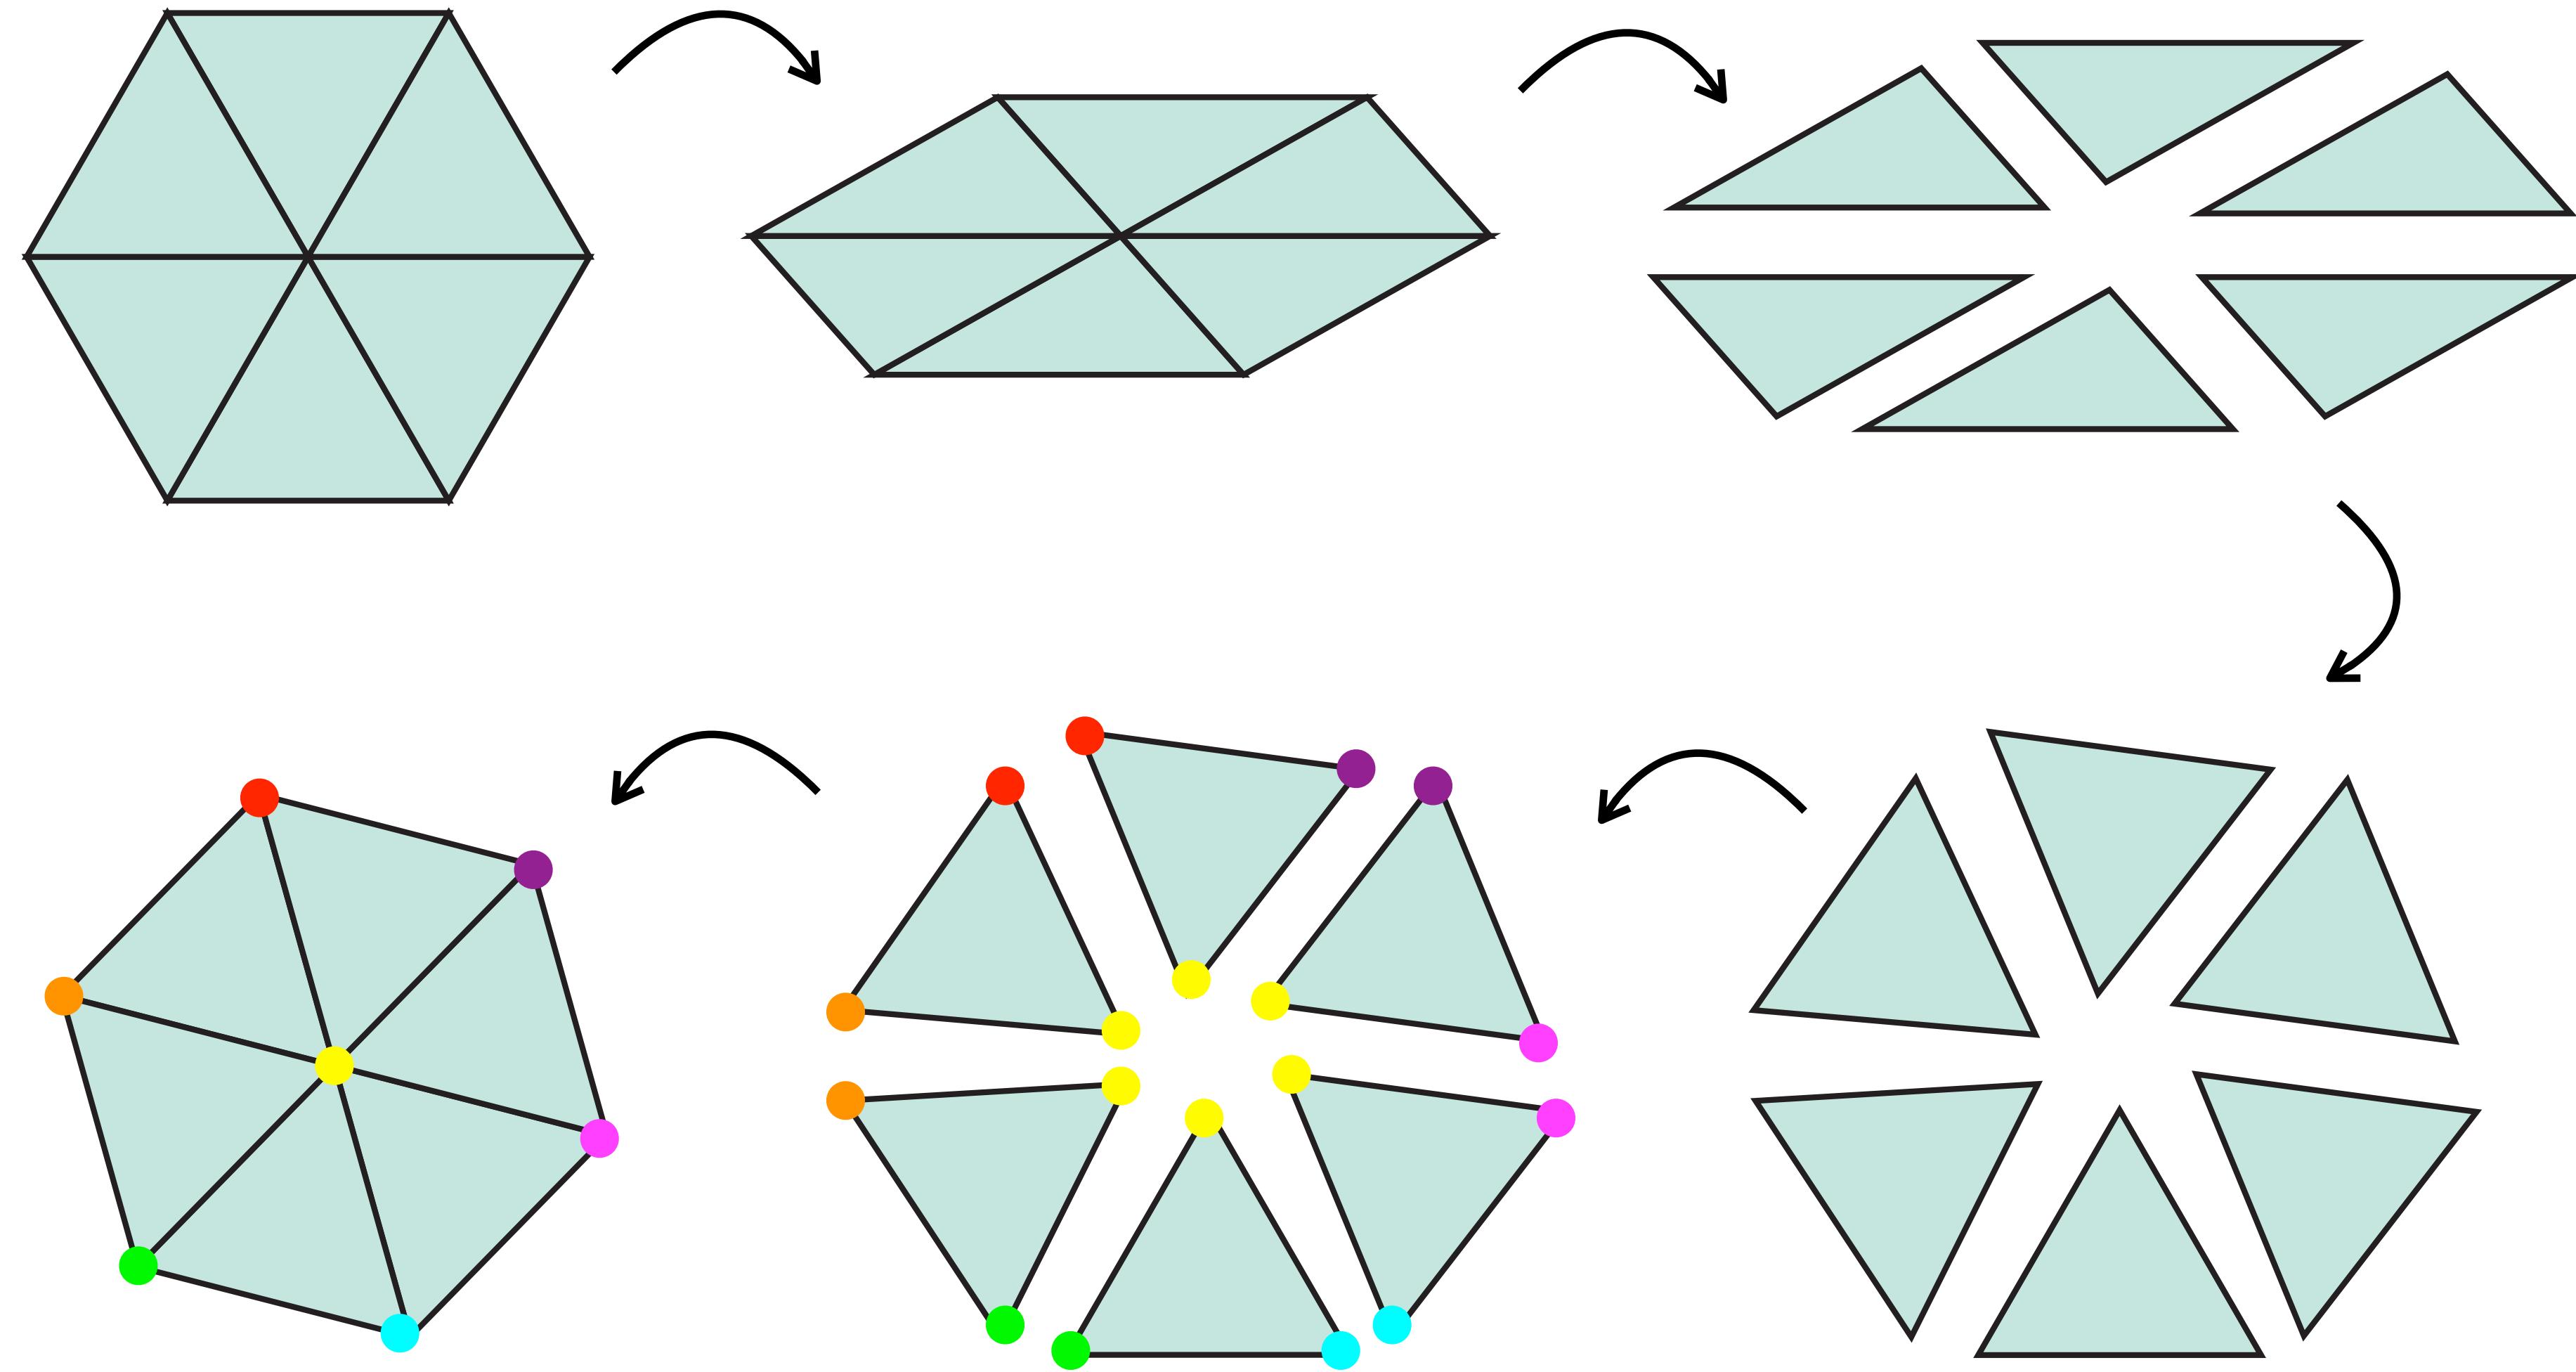
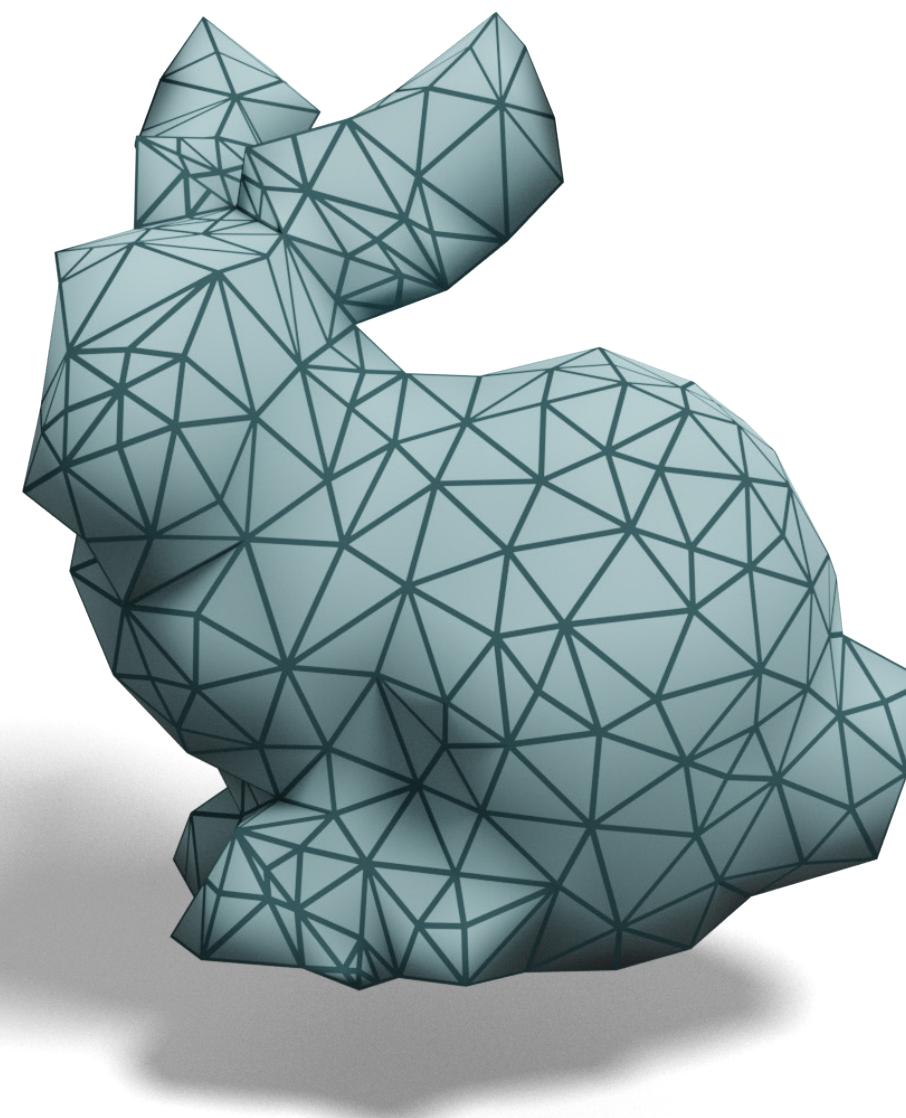
$$\mathbf{L}^{-1}$$

$$\mathbf{z}_k = (1 - \beta_k) \mathbf{y}_k + \beta_k \mathbf{L} \mathbf{s}_k$$

$$\mathbf{y}_k = \mathbf{g}_{k+1} - \mathbf{g}_k$$

$$\mathbf{s}_k = \mathbf{x}_{k+1} - \mathbf{x}_k$$

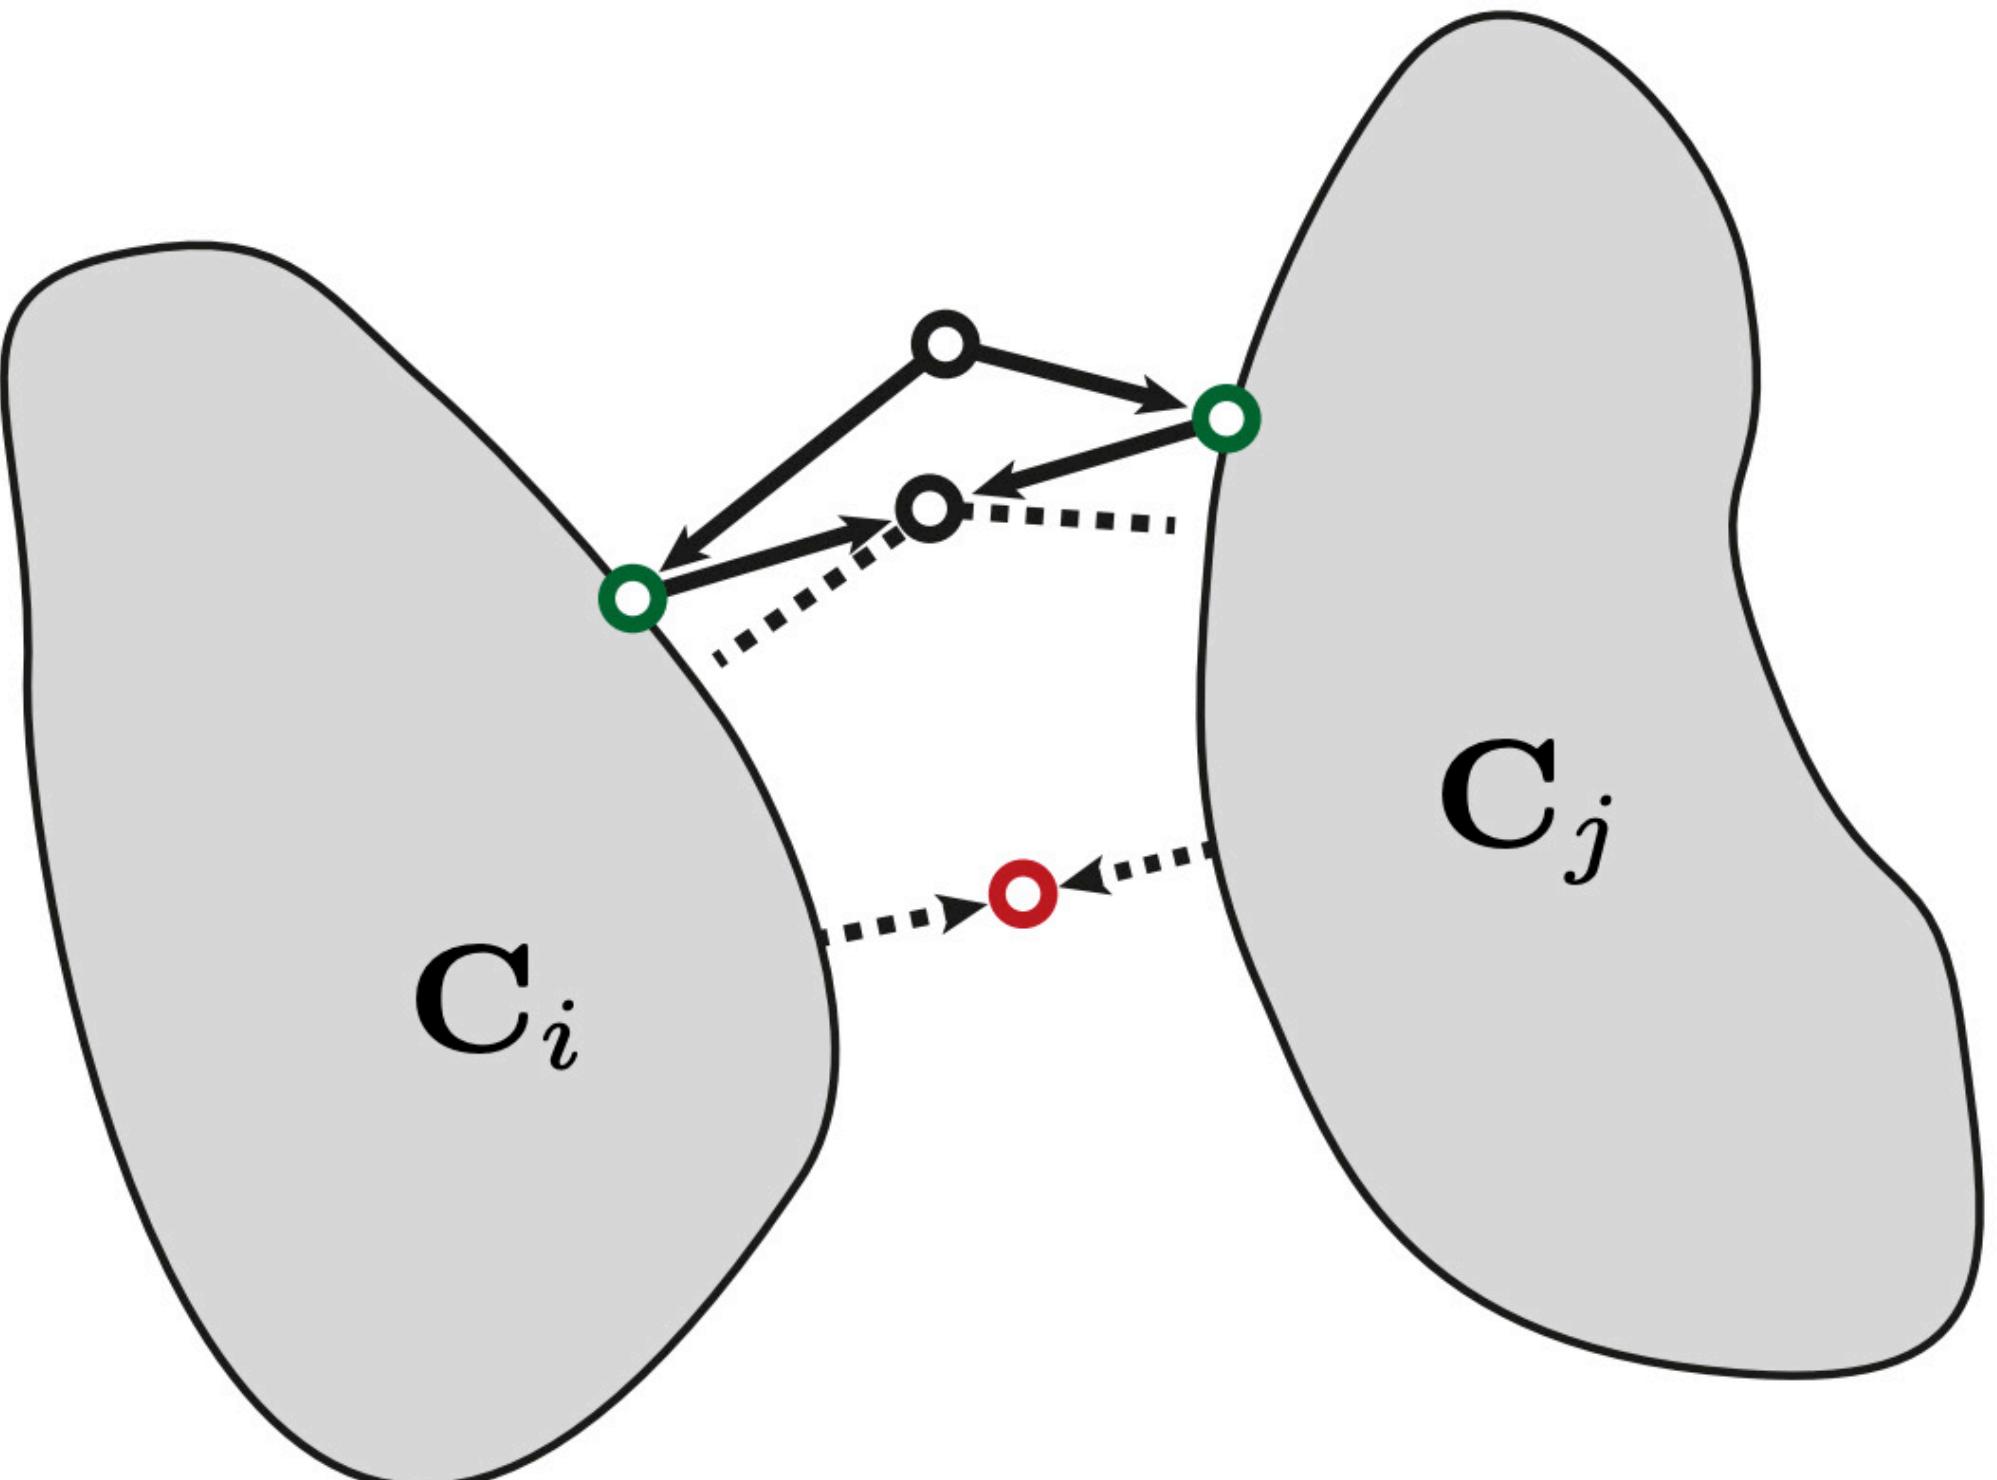
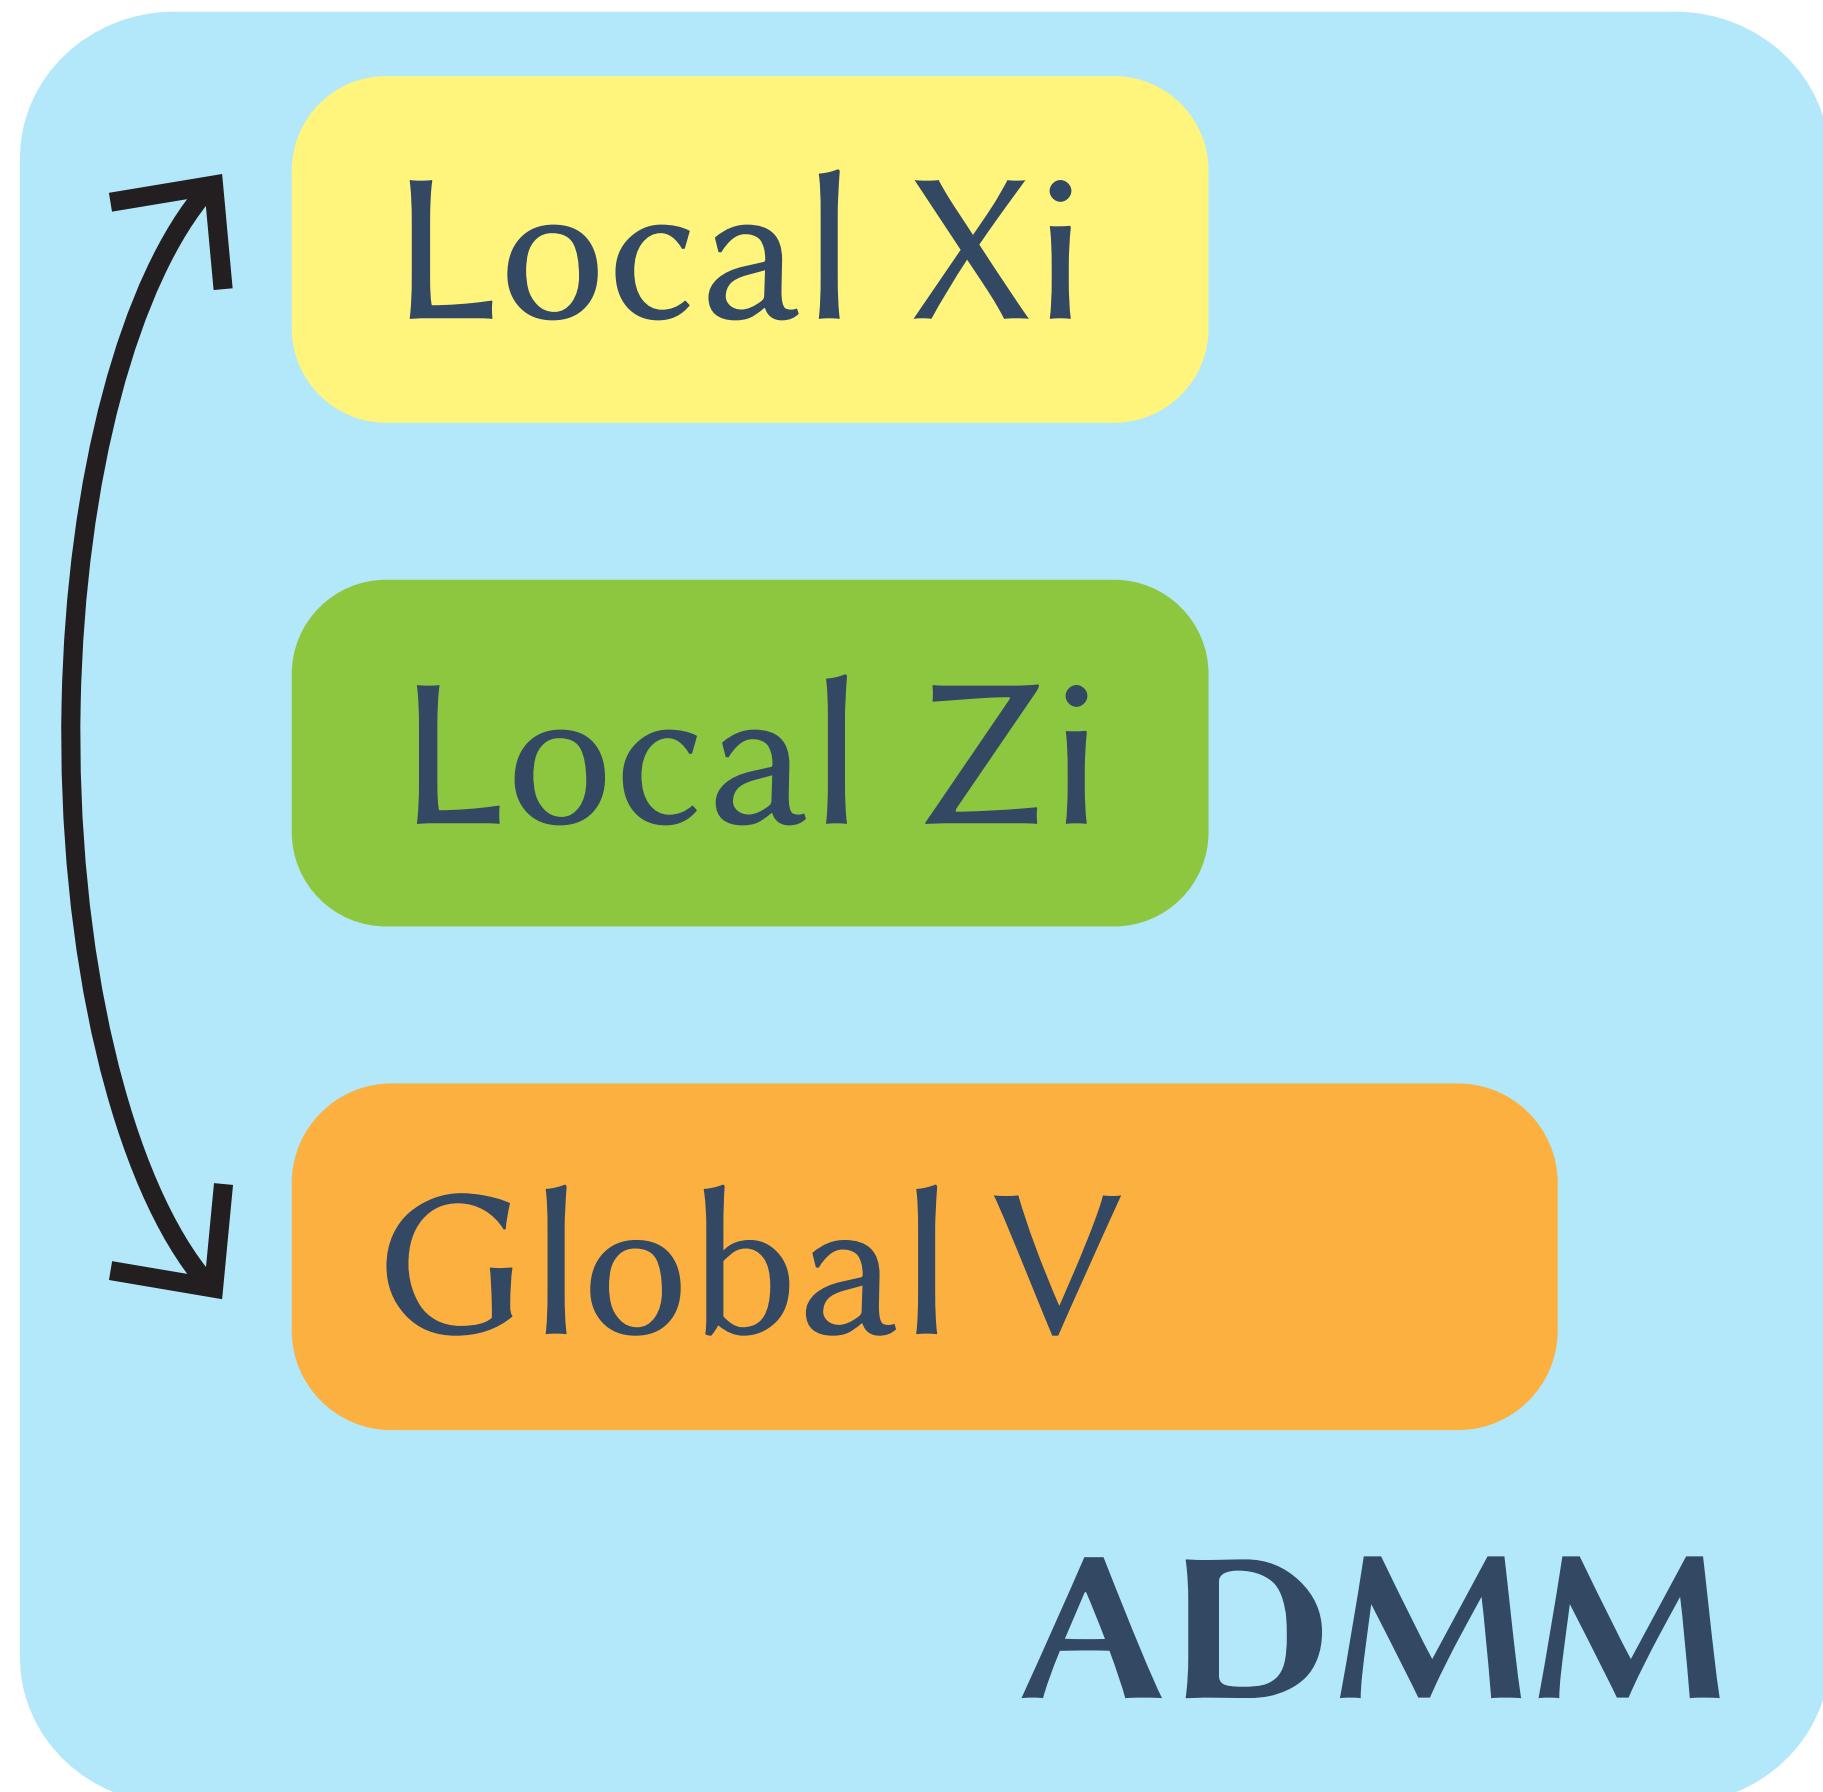
# Strategy 2: Exploit the Locality



# ADMM / Local-Global Solve / Projective Dynamics

$$\underset{\mathbf{V}, \{\mathbf{X}_i\}}{\text{minimize}} \quad \sum_{i \in V} \begin{array}{|c|} \hline \text{Elasticity} \\ \hline \end{array} + \begin{array}{|c|} \hline \text{(Optional) Other} \\ \hline \end{array}$$

# ADMM / Local-Global Solve / Projective Dynamics



# Block Coordinate Descent

arXiv:2403.06321v4 [cs.GR] 1 Jun 2024

## Vertex Block Descent

ANKA HE CHEN, University of Utah, USA

ZIHENG LIU, University of Utah, USA

YIN YANG, University of Utah, USA

CEM YUKSEL, University of Utah & Roblox, USA



Fig. 1. Example simulation results using our solver, both of those methods involve more than 100 million DoFs and 1 million active collisions.

We introduce vertex block descent, a block coordinate descent solution for the variational form of implicit Euler through vertex-level Gauss-Seidel iterations. It operates with local vertex position updates that achieve reductions in global variational energy with maximized parallelism. This forms a physics solver that can achieve numerical convergence with unconditional stability and exceptional computation performance. It can also fit in a given computation budget by simply limiting the iteration count while maintaining its stability and superior convergence rate.

We present and evaluate our method in the context of elastic body dynamics, providing details of all essential components and showing that it outperforms alternative techniques. In addition, we discuss and show examples of how our method can be used for other simulation systems, including particle-based simulations and rigid bodies.

CCS Concepts: • Computing methodologies → Physical simulation; Collision detection.

Additional Key Words and Phrases: physics-based simulation, elastic body, rigid body, time integration

## 1 INTRODUCTION

Physics-based simulation is the cornerstone of most graphics applications and the demands from simulation systems to deliver improved stability, accelerated computational performance, and enhanced visual realism are ever-growing. Particularly in real-time graphics applications, the stability and performance requirements are so strict that realism can sometimes be begrudgingly considered of secondary importance.

Notwithstanding the substantial amount of research and groundbreaking discoveries made on physics solvers over the past decades,

Authors' addresses: Anka He Chen, ankachan92@gmail.com, University of Utah, Salt Lake City, UT, USA; Ziheng Liu, NA, University of Utah, Salt Lake City, UT, USA; Yin Yang, NA, University of Utah, Salt Lake City, UT, USA; Cem Yuksel, cem@cemyuksel.com, University of Utah & Roblox, Salt Lake City, UT, USA.

existing methods still leave some things to be desired. They either deliver high-quality results, but fail to meet the computational demands of many applications or fit in a given computation time by sacrificing quality. Stability, on the other hand, is always a challenge, particularly with strict computation budgets.

In this paper, we introduce *vertex block descent (VBD)*, a physics solver that offers unconditional stability, superior computational performance than prior methods, and the ability to achieve numerical convergence to an implicit Euler integration. Though our method is a general solution that can be used for a variety of simulation problems, we present and evaluate it in the context of elastic body dynamics. Then, we briefly discuss how our method can be applied to some other simulation systems, including particle-based simulations and rigid bodies.

Our VBD method is based on block coordinate descent that performs vertex-based Gauss-Seidel iterations to solve the variational form of implicit Euler. For elastic body dynamics, each iteration runs a loop over the mesh vertices, adjusting the position of a single vertex at a time, temporarily fixing all others. This offers maximized parallelism when coupled with vertex-based mesh coloring, which can achieve an order of magnitude fewer colors (i.e. serialized workloads) as compared to element-based parallelization. Our local position-based updates can ensure that the variational energy is always reduced. Therefore, our method maintains its stability even with a single iteration per time step and large time steps, operating with un converged solutions containing a large amount of residual. With more iterations, it converges faster than its alternatives. Thus, it can more easily fit in a given computation budget, while maintaining stability with improved convergence.

We present all essential components of using VBD for elastic body dynamics, including formulations for damping, constraints,

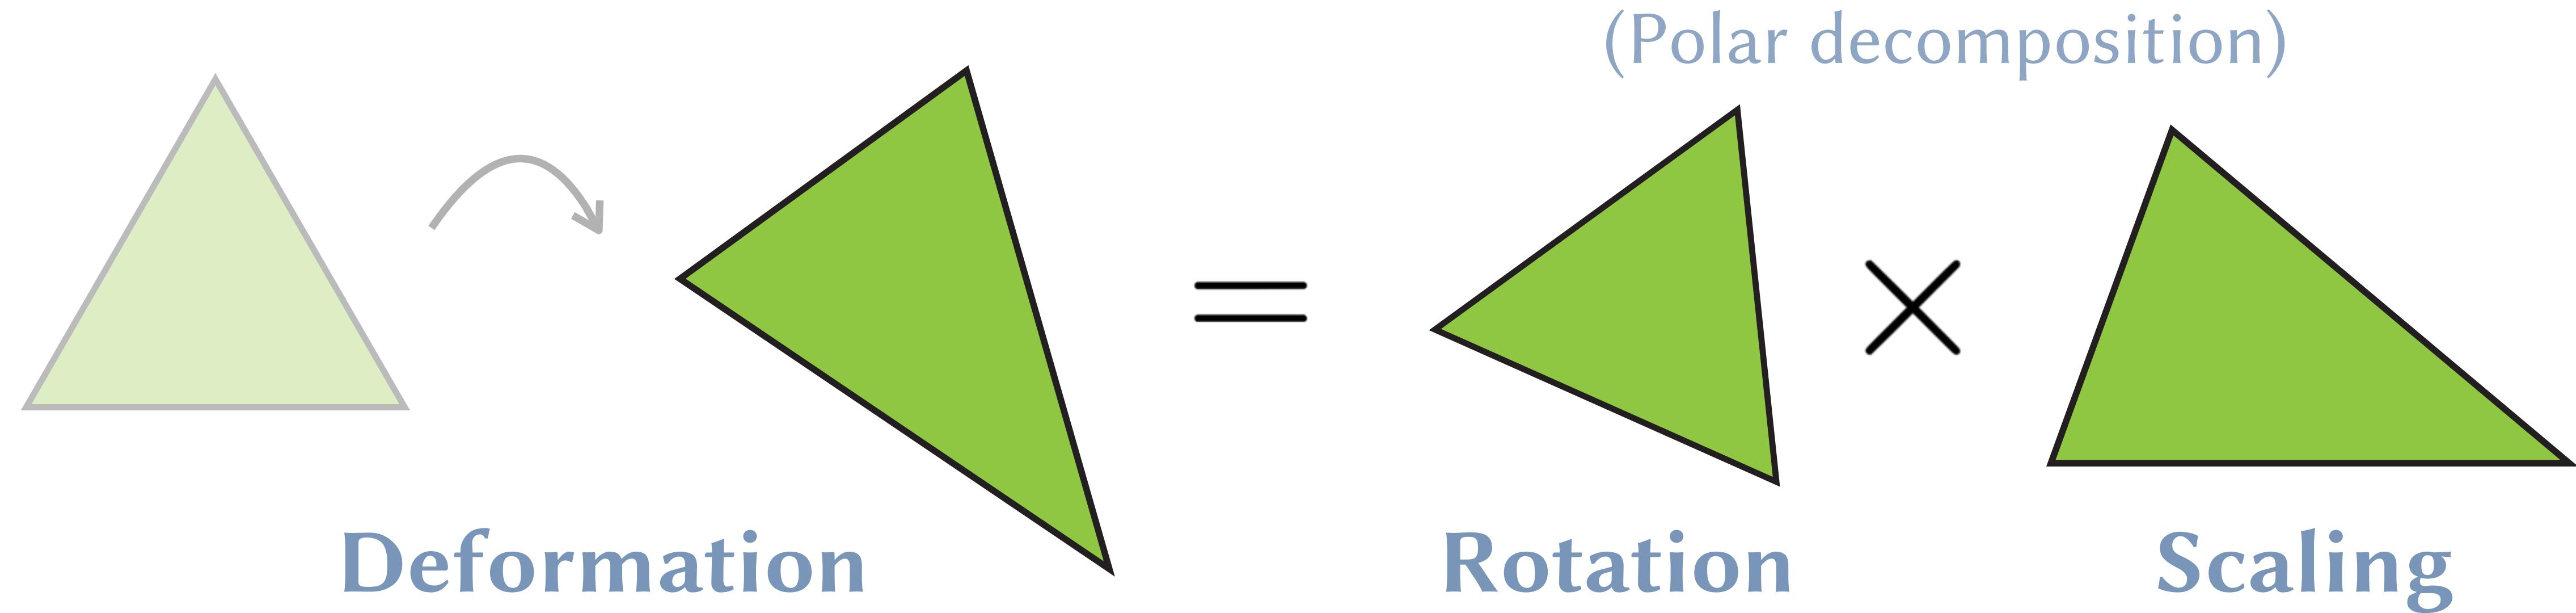
For each vertex (or element):

$$\mathbf{d}_i = -\mathbf{g}_i$$

or

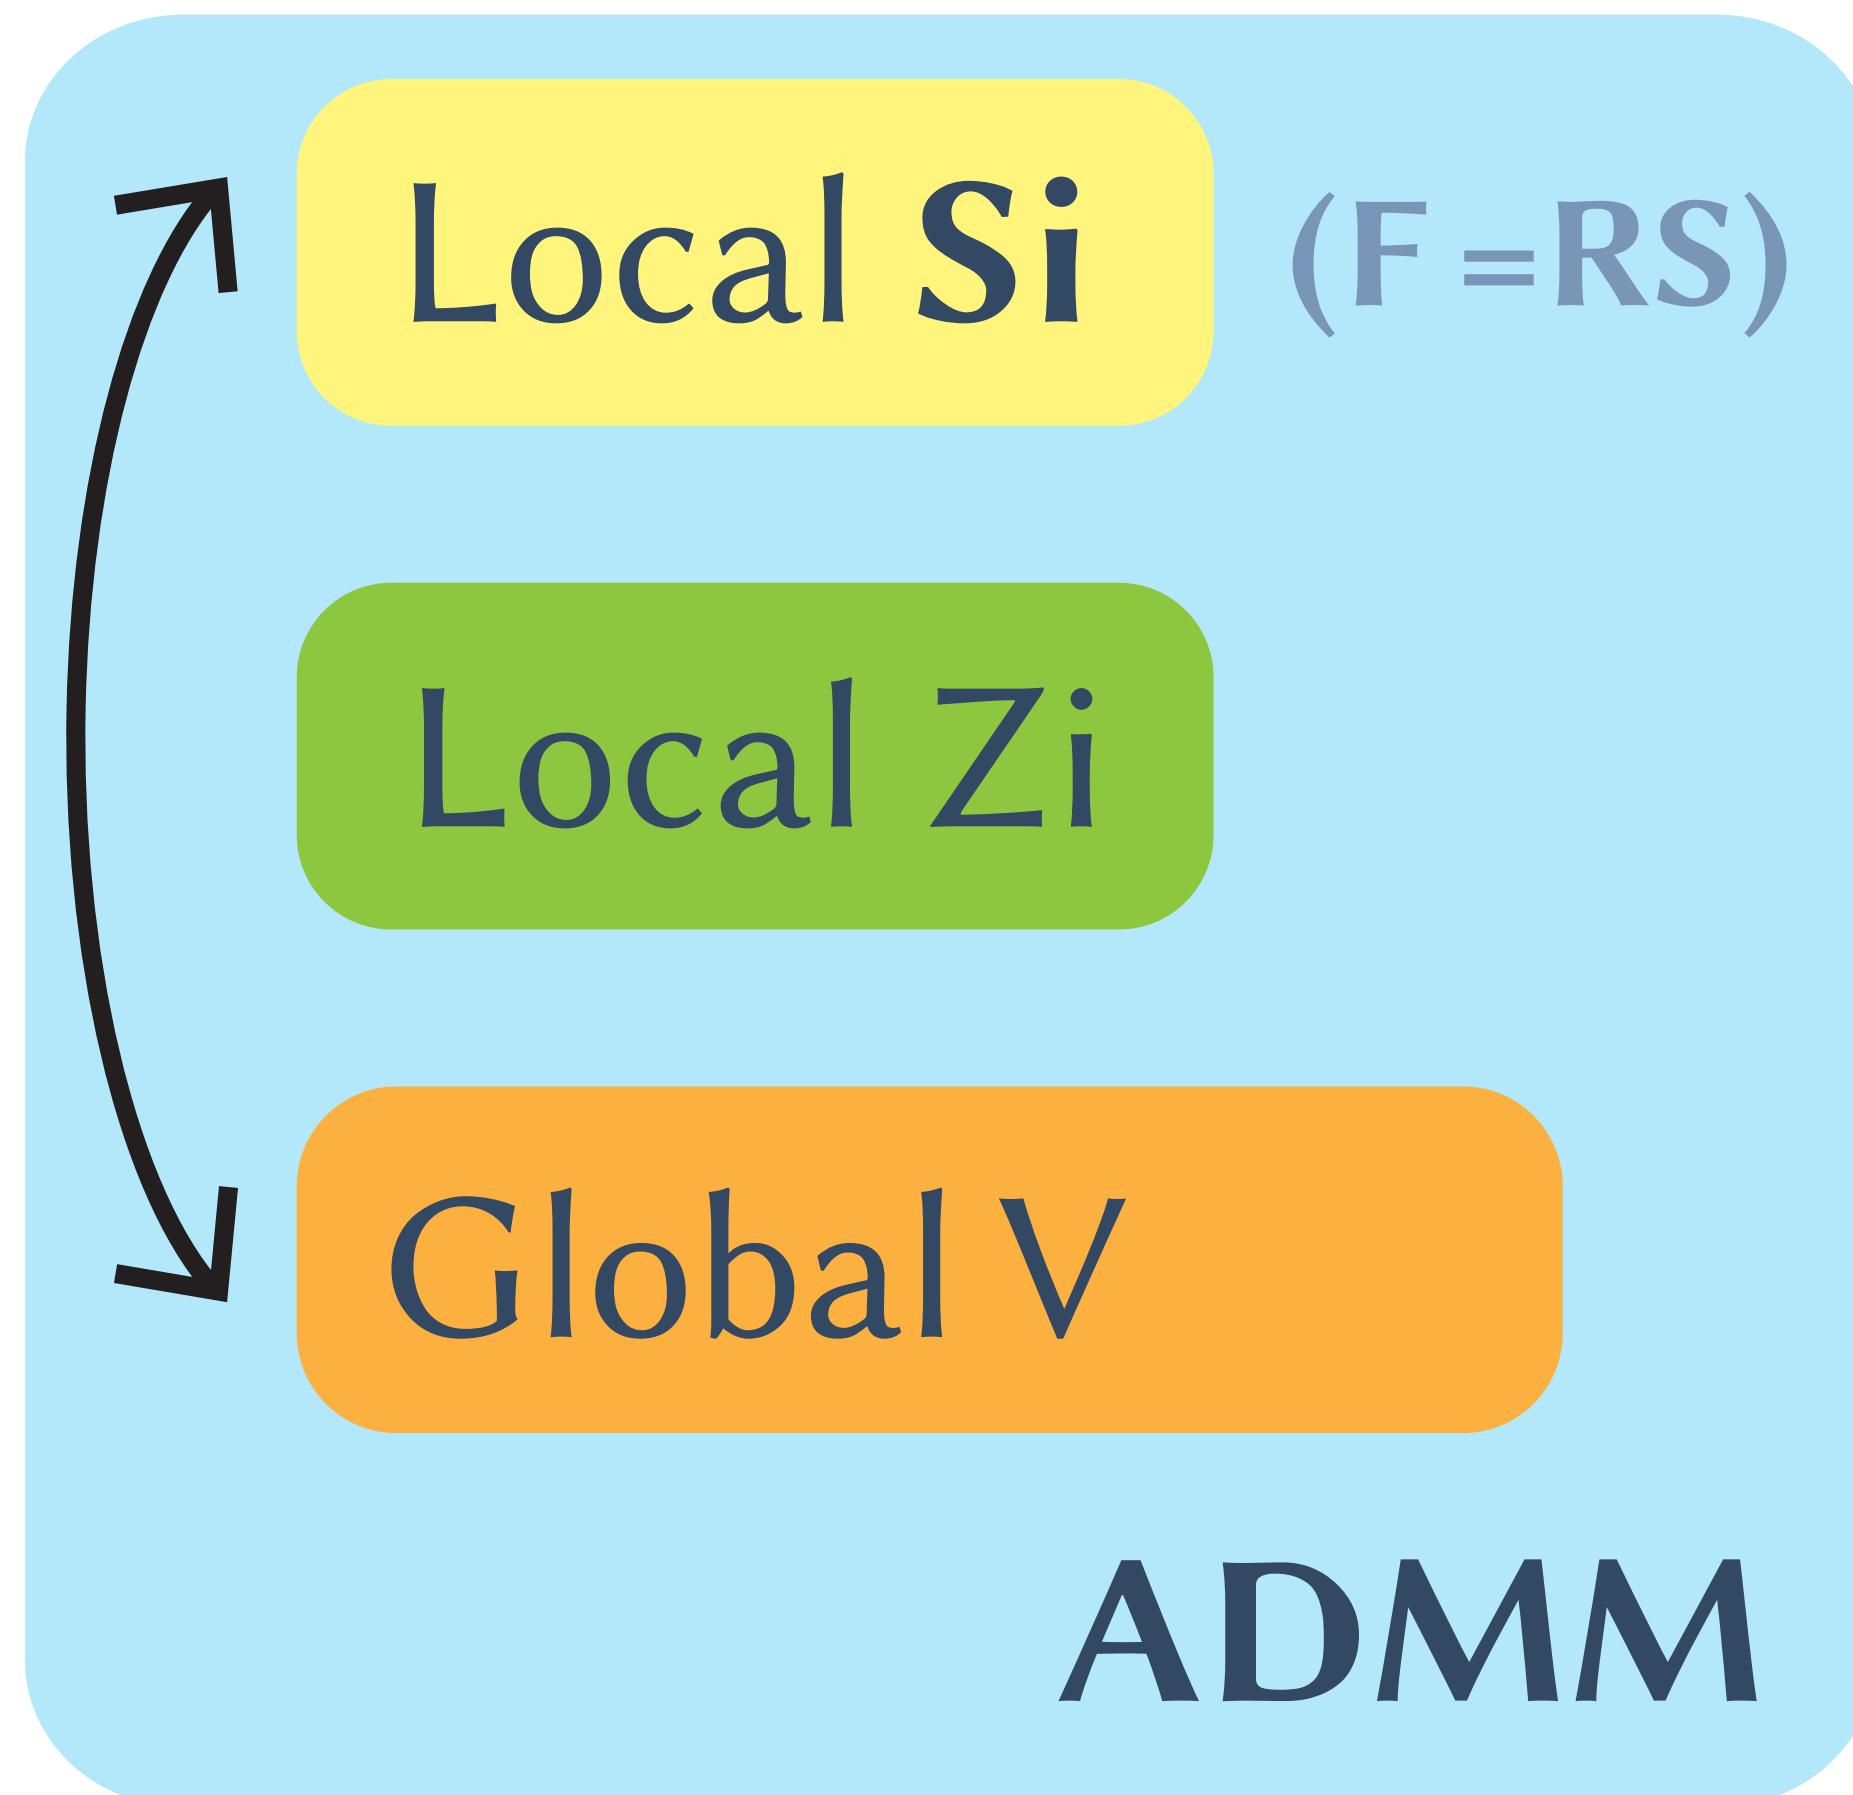
$$\mathbf{d}_i = -\mathbf{H}_i^{-1} \mathbf{g}_i$$

# Strategy 3: Make it Rotation-Aware



Rotation-invariant:  $f(\triangle) = f(\triangle)$

# Rotation-aware ADMM / Local-Global Solve



## WRAPD: Weighted Rotation-aware ADMM for Parameterization and Deformation

GEORGE E. BROWN, University of Minnesota  
RAHUL NARAIN, Indian Institute of Technology Delhi

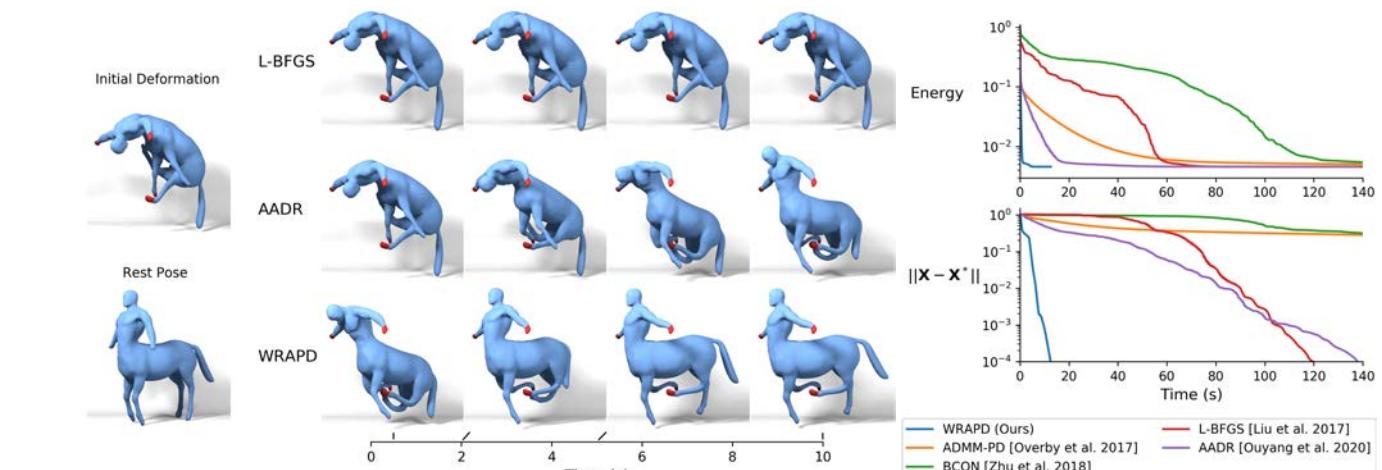


Fig. 1. Quasi-static simulation of a neo-Hookean centaur with 98K tetrahedra and 26K vertices, subject to pin constraints. *Left:* Initial and rest shapes. *Middle:* Selected frames from the first ten seconds of simulation comparing our algorithm to competitors. *Top right:* Objective value vs. time. *Bottom right:* Norm of position error relative to optimal state vs. time.

Local-global solvers such as ADMM for elastic simulation and geometry optimization struggle to resolve large rotations such as bending and twisting modes, and large distortions in the presence of barrier energies. We propose two improvements to address these challenges. First, we introduce a novel local-global splitting based on the polar decomposition that separates the geometric nonlinearity of rotations from the material nonlinearity of the deformation energy. The resulting ADMM-based algorithm is a combination of an L-BFGS solve in the global step and proximal updates of element stretches in the local step. We also introduce a novel method for dynamic reweighting that is used to adjust element weights at runtime for improved convergence. With both improved rotation handling and element weighting, our algorithm is considerably faster than state-of-the-art approaches for quasi-static simulations. It is also much faster at making early progress in parameterization problems, making it valuable as an initializer to jump-start second-order algorithms.

CCS Concepts: • Computing methodologies → Physical simulation.

Additional Key Words and Phrases: quasi-statics, parameterization, geometric nonlinearity, rotation-aware, reweighting, optimization, ADMM

Authors' addresses: George E. Brown, University of Minnesota, brow2327@umn.edu; Rahul Narain, Indian Institute of Technology Delhi, narain@cse.iitd.ac.in.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](https://permissions.acm.org). © 2021 Copyright held by the owner/author(s). Publication rights licensed to ACM. 0730-0301/2021/8-ART82 \$15.00 <https://doi.org/10.1145/3450626.3459942>

## ACM Reference Format:

George E. Brown and Rahul Narain. 2021. WRAPD: Weighted Rotation-aware ADMM for Parameterization and Deformation. *ACM Trans. Graph.* 40, 4, Article 82 (August 2021), 14 pages. <https://doi.org/10.1145/3450626.3459942>

## 1 INTRODUCTION

From quasi-static and dynamic simulation of elastic bodies in physics-based animation to interactive shape manipulation and mesh parameterization in geometry processing, recent work has converged on minimization of deformation energies as a central task in computer graphics. The deformation energy model takes many different forms in different applications, from physically-based hyperelastic models such as the neo-Hookean model, to geometrically motivated distortion metrics such as the symmetric Dirichlet energy. These energies can be highly nonlinear with respect to stretching deformations, but have the key property that they are invariant to rigid transformations.

Many popular techniques for minimizing such energies use a local-global approach [Sorkine and Alexa 2007; Bouaziz et al. 2012; Overby et al. 2017; Liu et al. 2017; Peng et al. 2018; Zhang et al. 2019; Ouyang et al. 2020], in which chosen per-element quantities are introduced as additional optimization variables. Optimization is then performed in an alternating fashion, with a local step acting on the per-element local variables and a global step updating the vertex positions. We focus on ADMM-PD [Overby et al. 2017], a local-global algorithm that has been shown to be effective for nonlinear energies

## Mixed Variational Finite Elements for Implicit Simulation of Deformables

TY TRUSTY, University of Toronto, Canada  
DANNY M. KAUFMAN, Adobe Research, USA  
DAVID I.W LEVIN, University of Toronto, Canada

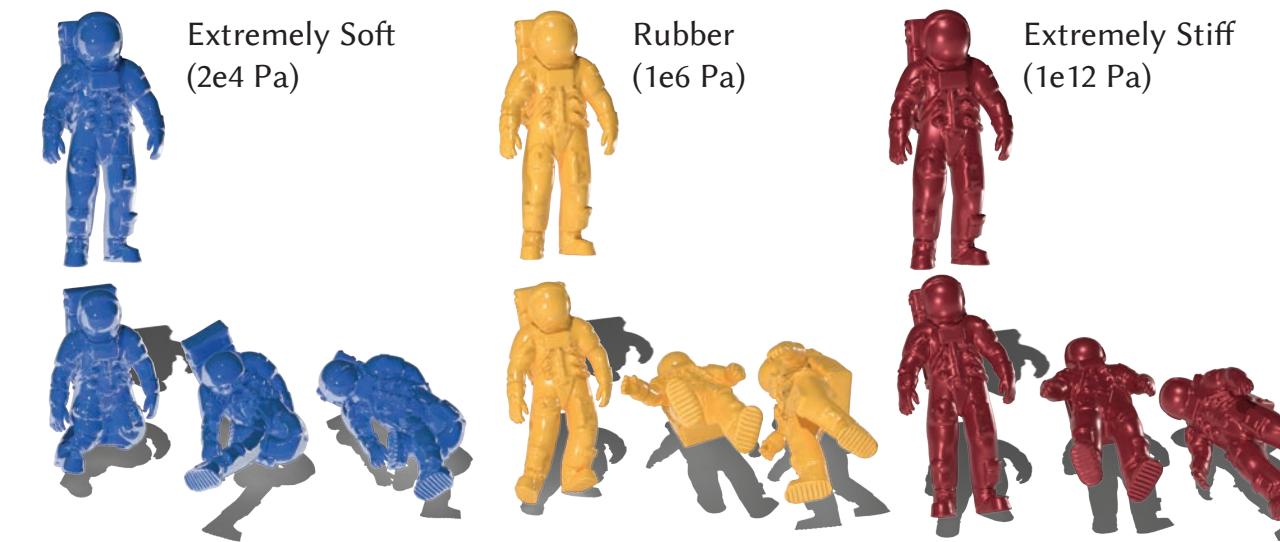


Fig. 1. Our new mixed finite element method (MFEM) can produce simulations of elastica with wildly different materials (including rigid) both accurately and quickly. Our key contribution is that our method is both capable of converging to an accurate solution, matching that of a Newton's method, as well as generate visually plausible results when stopped early. This makes it ideal for a plethora of engineering and graphics applications.

We propose and explore a new method for the implicit time integration of elastica. Key to our approach is the use of a mixed variational principle. In turn, its finite element discretization leads to an efficient and accurate sequential quadratic programming solver with a superset of the desirable properties of many previous integration strategies. This framework fits a range of elastic constitutive models and remains stable across a wide span of time step sizes and material parameters (including problems that are approximately rigid). Our method exhibits convergence on par with full Newton type solvers and also generates visually plausible results in just a few iterations comparable to recent fast simulation methods that do not converge. These properties make it suitable for both offline accurate simulation and performant applications with expressive physics. We demonstrate the efficacy of our approach on a number of simulated examples.

CCS Concepts: • Computing methodologies → Physical simulation.

Additional Key Words and Phrases: physics-based animation, physics simulation

### ACM Reference Format:

Ty Trusty, Danny M. Kaufman, and David I.W Levin. 2022. Mixed Variational Finite Elements for Implicit Simulation of Deformables. In *SIGGRAPH Asia 2022 Conference Papers (SA '22 Conference Papers)*, December 6–9, 2022, Daegu, Republic of Korea. ACM, New York, NY, USA, 8 pages. <https://doi.org/10.1145/3550469.3555418>

1

1

Republic of Korea. ACM, New York, NY, USA, 8 pages. <https://doi.org/10.1145/3550469.3555418>

### 1 INTRODUCTION

In this paper we explore the use of a mixed variational principle to build an efficient and general-purpose simulation algorithm for the physics-based animation of elastica.

Standard approaches for the implicit time integration of continua discretize with finite differences in time and finite elements in space. Recent methods often leverage the observation that, for these implicit time integration choices, each individual time step solve can then be cast as a minimization problem. In turn, the applied strategy for solving these optimization problems then leads to a wide range of well-known simulation algorithms [Li et al. 2019]. For example, a “standard” finite element approach involves minimizing an implicit integration energy via Newton’s method while solving the bottleneck of inner linear-systems solves either via direct or iterative methods. Extended Position-Based Dynamics replaces standard direct or iterative solvers with iterations (e.g., GS, Jacobi, and/or SOR) acting on the dual variables (constraint forces) while Projective Dynamics and its more recent generalizations apply various forms of ADMM-type solvers to split, augmented Lagrangian forms.

Despite their common variational origin, implicit solvers for elastica exhibit a wide range of features and limitations, and so tradeoffs.

## Subspace Mixed Finite Elements for Real-Time Heterogeneous Elastodynamics

Ty Trusty\*,  
University of Toronto  
Canada  
Otman Benchekroun\*  
University of Toronto  
Canada  
Eitan Grinspun  
University of Toronto  
Canada

Danny M. Kaufman  
University of Toronto and Adobe  
Research U.S.A.  
David I.W. Levin  
University of Toronto and NVIDIA  
Canada

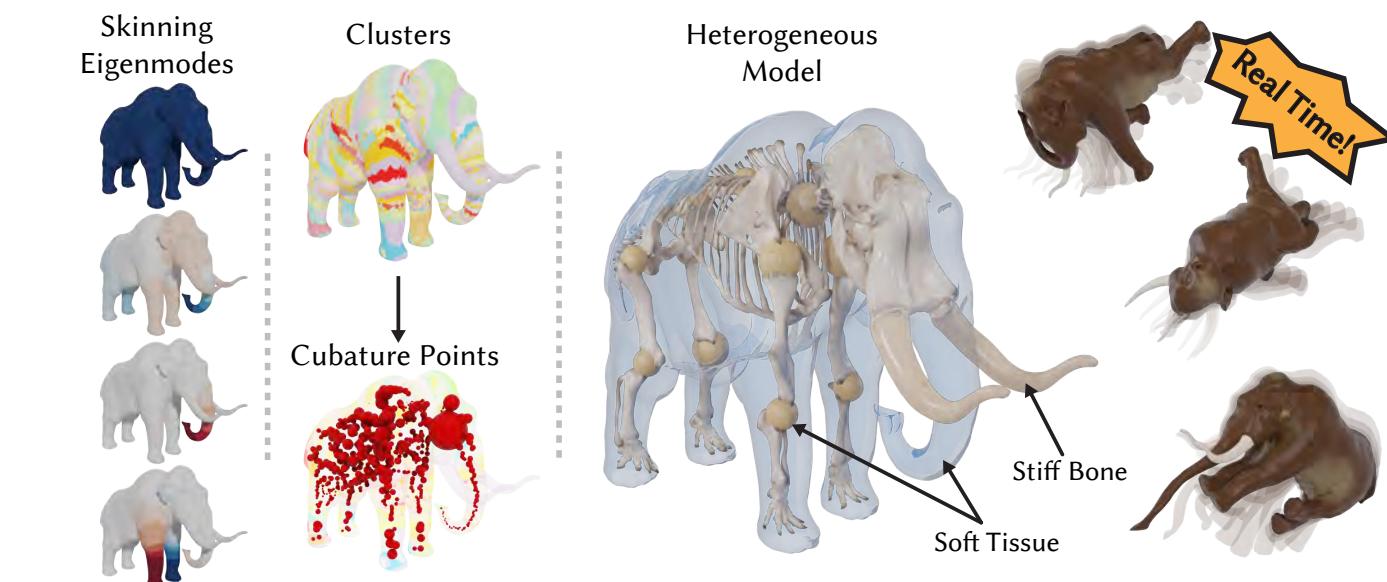


Figure 1: We propose a reduced space mixed finite element method (MFEM) built on a Skinning Eigenmode subspace and material-aware cubature scheme. Our solver is well-suited for simulating scenes with large material and geometric heterogeneities in real-time. This mammoth geometry is composed of 98,175 vertices and 531,565 tetrahedral elements and with a heterogeneous composition of widely varying materials of muscles ( $E = 5 \times 10^5$  Pa), joints ( $E = 1 \times 10^5$  Pa), and bone ( $E = 1 \times 10^{10}$  Pa). The resulting simulation runs at 120 frames per second (FPS).

### ABSTRACT

Real-time elastodynamic solvers are well-suited for the rapid simulation of homogeneous elastic materials, with high-rates generally enabled by aggressive early termination of timestep solves. Unfortunately, the introduction of strong domain heterogeneities can make these solvers slow to converge. Stopping the solve short creates visible damping artifacts and rotational errors. To address these challenges we develop a reduced mixed finite element solver that preserves rich rotational motion, even at low-iteration regimes. Specifically, this solver augments time-step solve optimizations with auxiliary stretch

\*Indicates joint first authors.

Authors’ addresses: Ty Trusty\*, University of Toronto, Canada; Otman Benchekroun\*, University of Toronto, Canada; Eitan Grinspun, University of Toronto, Canada; Danny M. Kaufman, University of Toronto and Adobe Research, U.S.A.; David I.W. Levin, University of Toronto and NVIDIA, Canada.

degrees of freedom at mesh elements, and maintains consistency with the primary positional degrees of freedoms at mesh nodes via explicit constraints. We make use of a Skinning Eigenmode subspace for our positional degrees of freedom. We accelerate integration of non-linear elastic energies with a cubature approximation, placing stretch degrees of freedom at cubature points. Across a wide range of examples we demonstrate that this subspace is particularly well suited for heterogeneous material simulation. Our resulting method is a subspace mixed finite element method completely decoupled from the resolution of the mesh that is well-suited for real-time simulation of heterogeneous domains.

### KEYWORDS

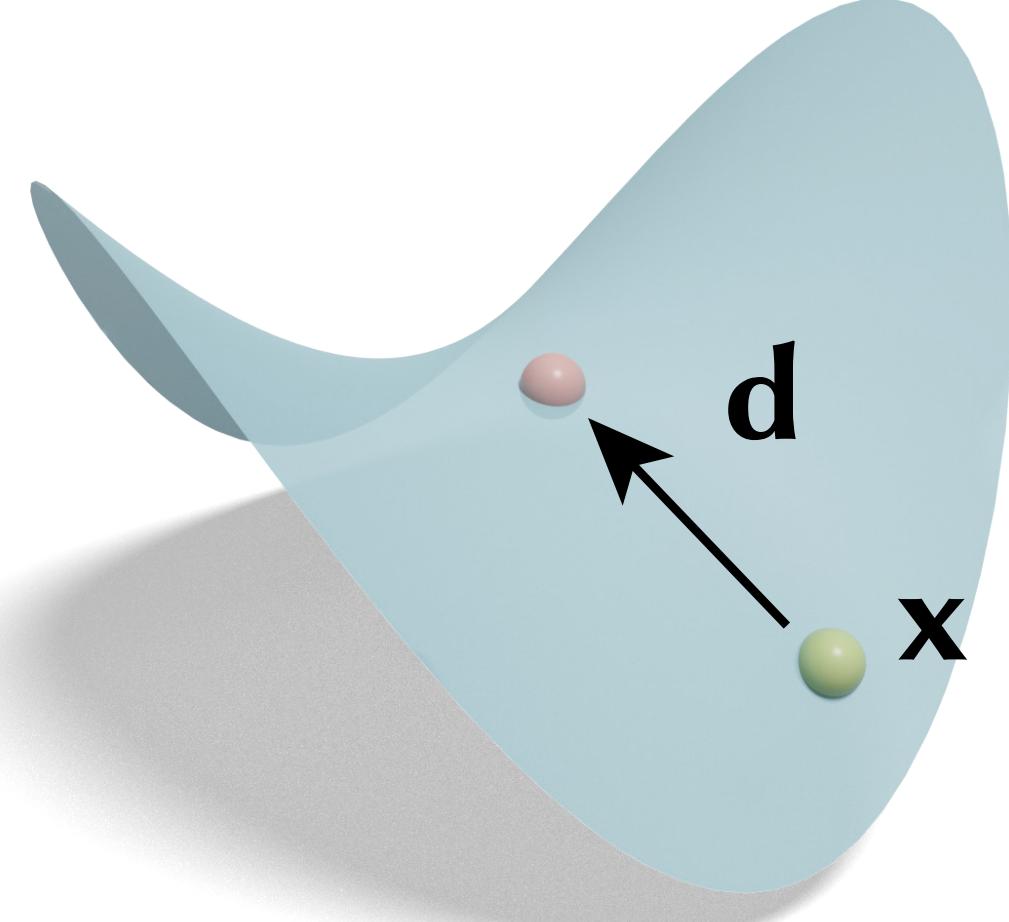
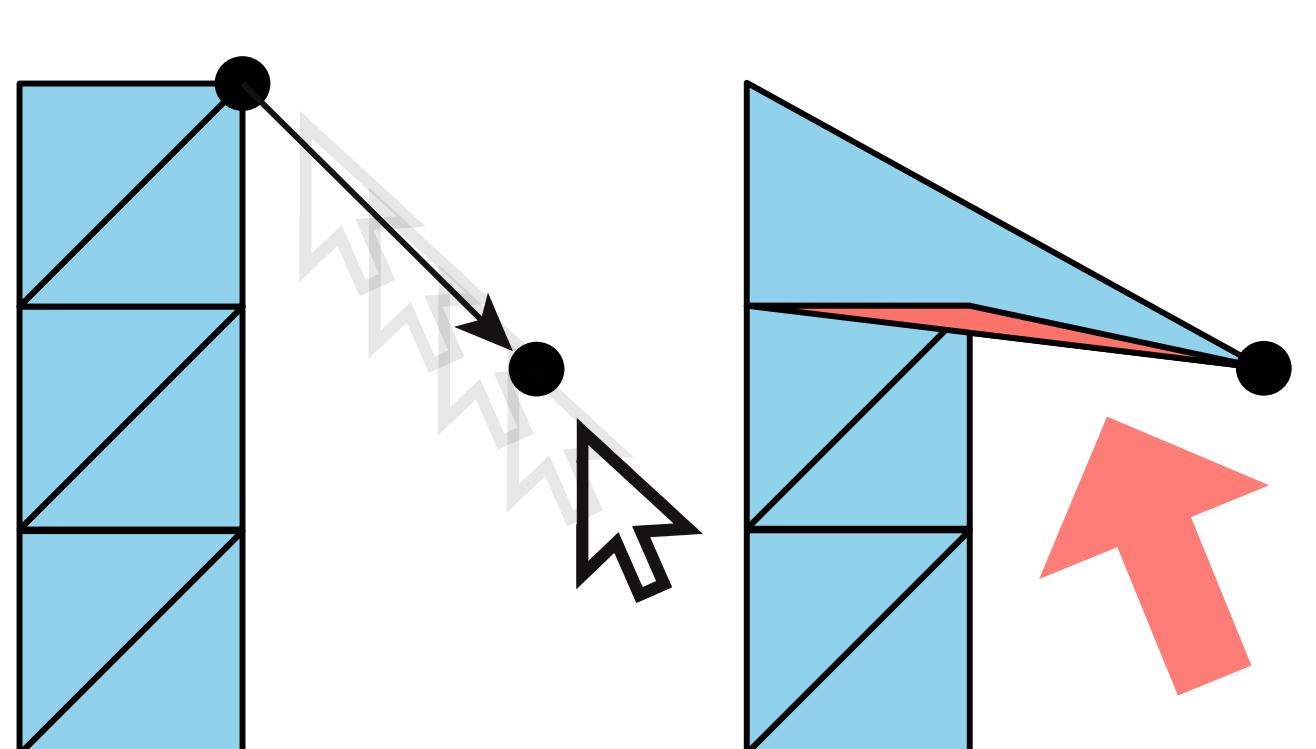
Mixed FEM, Heterogeneous Materials

## Rotation-aware Newton’s Method

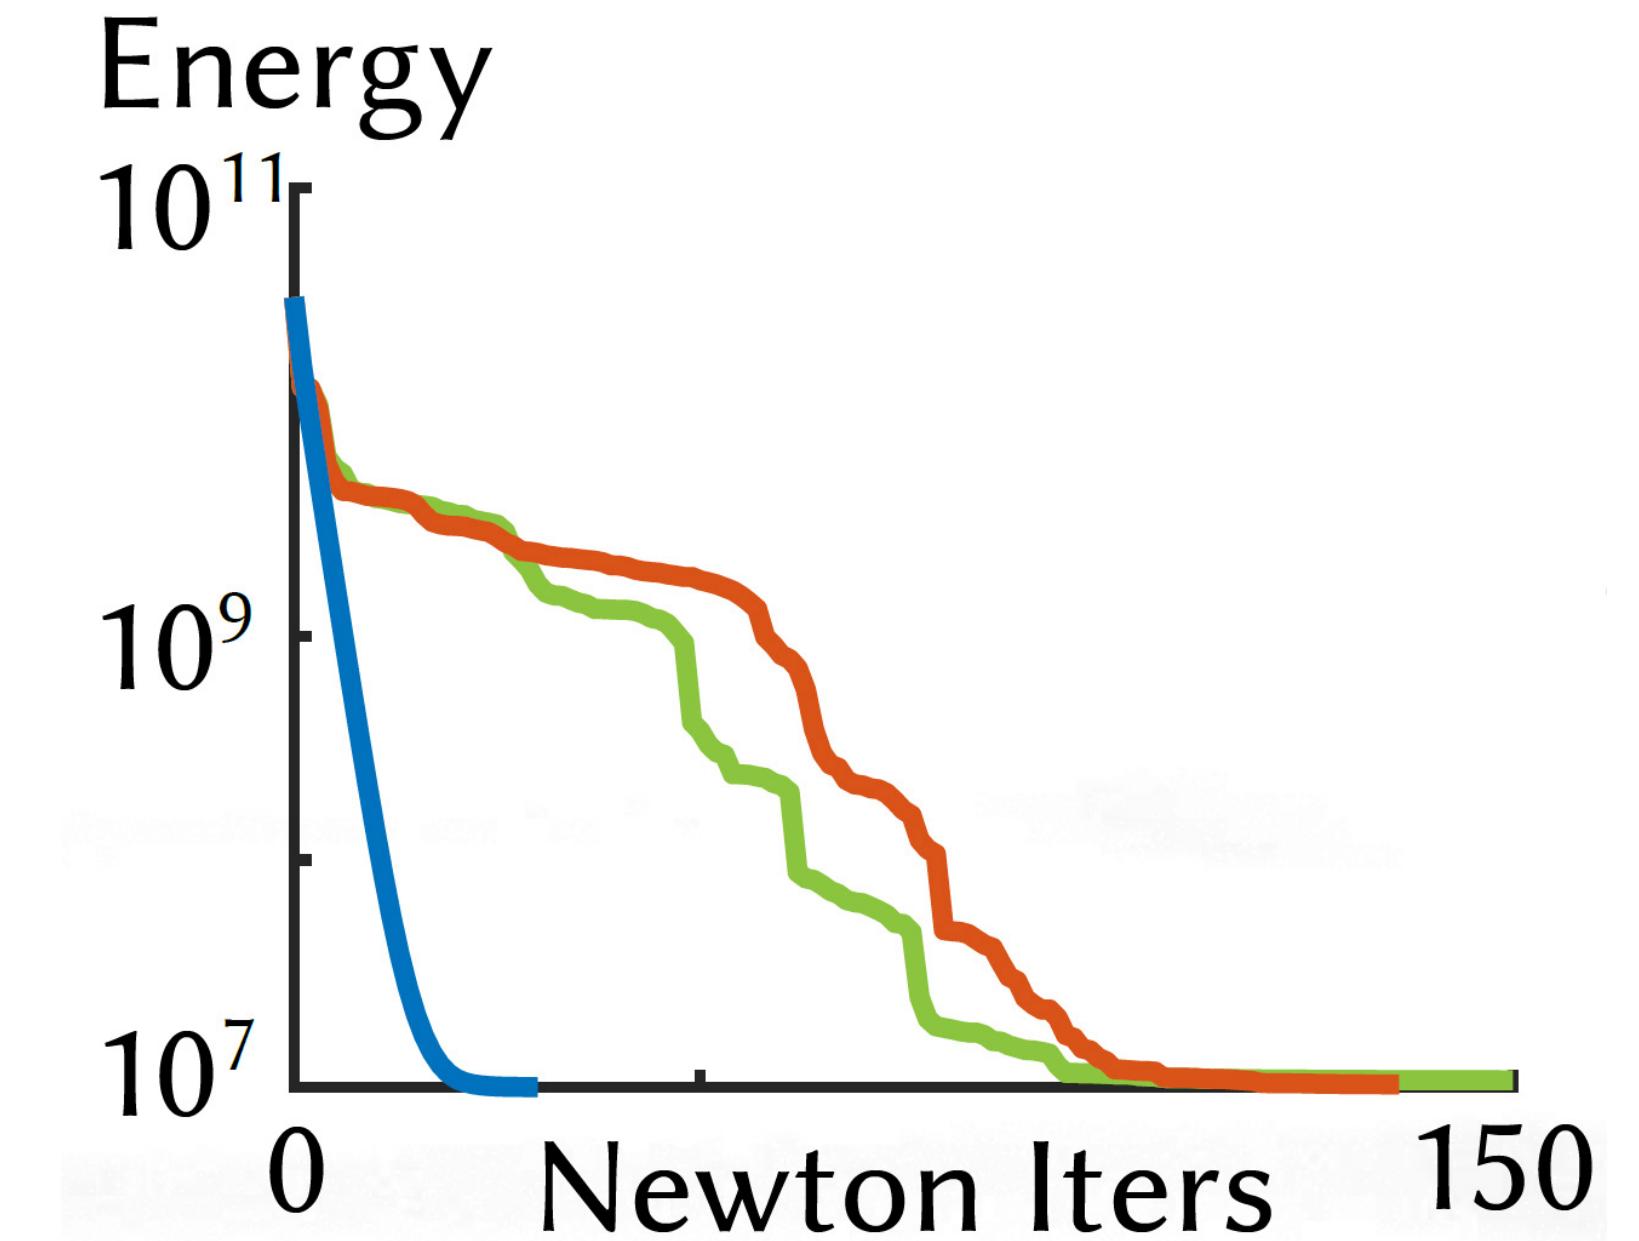
## Rotation-aware Subspace

# Recap

$$\mathbf{d} = -\mathbf{H}^{-1}\mathbf{g}$$

- Nonconvexity 
  - Infeasible points 
  - Large-scale 
- indefinite** 
- 
- 

# Thank you!



## Introduction to Optimization for Simulation

Honglin Chen

honglin.chen@columbia.edu

<https://www.cs.columbia.edu/~honglinchen/>